

CKernel(Beijing) 2017

eBPF In-kernel Virtual Machine & Cloud Computing

李枫

hkli2012@126.com

Oct 22, 2017

Agenda

I. Anatomy of eBPF

- eBPF
- LLVM
- Development

II. eBPF for Kernel Instrumentation

- BCC
- Applications
- Pros & Cons

III. Cloud Computing with eBPF

- Cilium
- Load Balance
- Security
- Go-based Cloud Ecosystem

IV. eBPF on ARM

- RPi3
- IO Visor

V. Wrap-up

I. Anatomy of eBPF

1) eBPF

- https://en.wikipedia.org/wiki/Berkeley_Packet_Filter

BPF (Berkeley Packet Filter, aka cBPF)

- Introduced in kernel 2.1.75 (1997)
- Originally designed for packet filtering (tcpdump...)
- Apply for seccomp filters, traffic control...
- <https://blog.cloudflare.com/bpf-the-forgotten-bytecode/>

| | |
|---|--|
| # tcpdump host 127.0.0.1 and port 22 -d | Optimizes packet filter performance |
| (000) ldh [12] | |
| (001) jeq #0x800 jt 2 jf 18 | |
| (002) ld [26] | |
| (003) jeq #0x7f000001 jt 6 jf 4 | |
| (004) ld [30] | |
| (005) jeq #0x7f000001 jt 6 jf 18 | 2 x 32-bit registers & scratch memory |
| (006) ldb [23] | |
| (007) jeq #0x84 jt 10 jf 8 | |
| (008) jeq #0x6 jt 10 jf 9 | |
| (009) jeq #0x11 jt 10 jf 18 | |
| (010) ldh [20] | User-defined bytecode executed by an in-kernel sandboxed virtual machine |
| (011) jset #0x1fff jt 18 jf 12 | |
| (012) ldxb 4*([14]&0xf) | |
| (013) ldh [x + 14] | |
| [...] | |

Steven McCanne and Van Jacobson, 1993

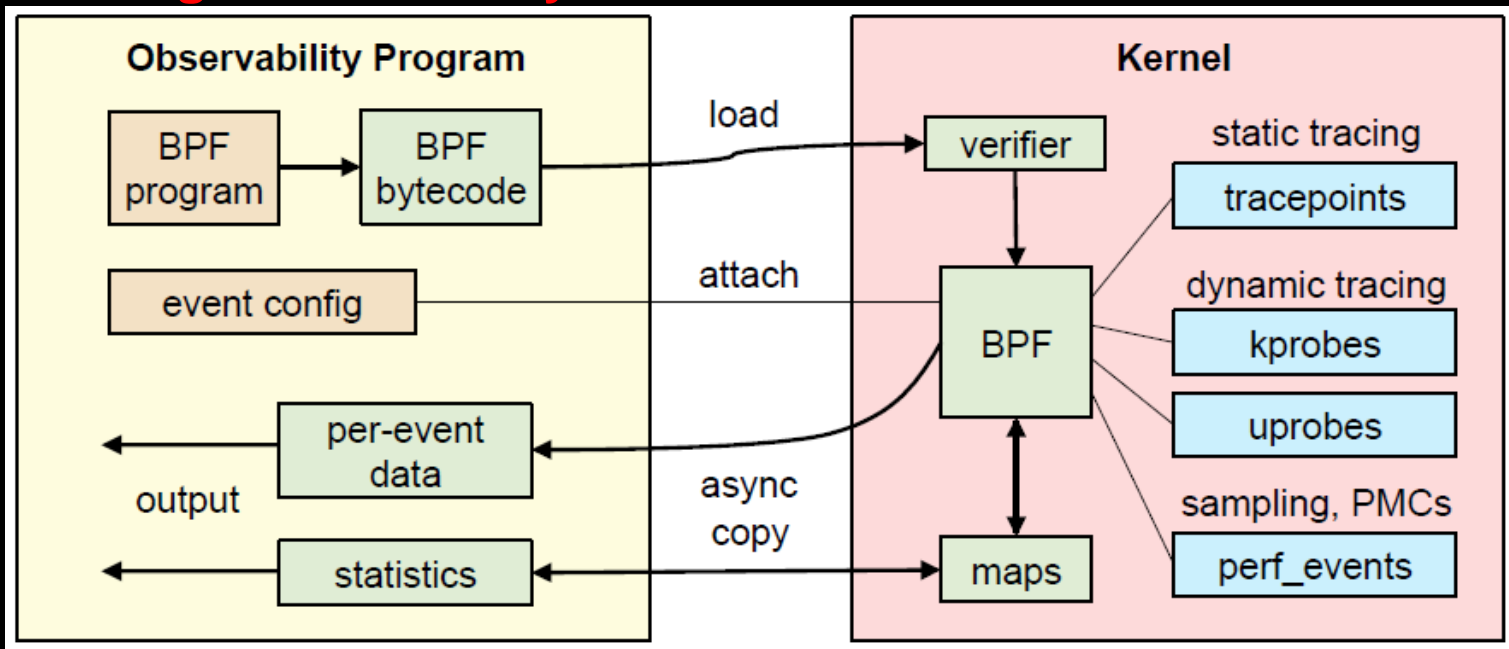
Source: <https://www.slideshare.net/brendangregg/kernel-recipes-2017-performance-analysis-with-bpf>

eBPF (extended BPF)

- Since Linux Kernel v3.15 and ongoing
- Aims at being a universal in-kernel virtual machine
- a simple way to extend the functionality of Kernel at runtime
- <https://lwn.net/Articles/655544>

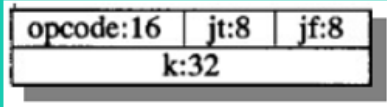
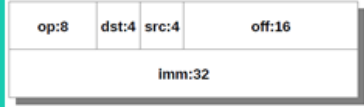
BPF for tracing is currently a hot area, Starovoitov said. It is a better alternative to [SystemTap](#) and runs two to three times faster than Oracle's [DTrace](#). Part of that speed comes from LLVM's optimizations plus the kernel's internal just-in-time compiler for BPF bytecode.

- it changes the old ways for Kernel instrumentation



Source: <https://www.slideshare.net/brendangregg/kernel-recipes-2017-performance-analysis-with-bpf>

Comparison

| | cBPF | eBPF |
|-------------|---|--|
| Register | Two 32 bit registers: A: accumulator X: indexing | Eleven 64 bit registers: R0: return value/exit value R1-R5: arguments R6-R9: callee saved registers R10: read-only frame pointer |
| Instruction | ~30  | ~90  |
| JIT | Support | Support (better mapping with newer architectures for JITing) |
| Toolchain | GCC, tools/net | LLVM eBPF backend |
| Platform | x86_64, ARM, ARM64, SPARC, PowerPC, MIPS and s390 | x86-64, aarch64, s390x |
| System Call | | <pre>#include <linux/bpf.h> int bpf(int cmd, union bpf_attr *attr, unsigned int size);</pre> <p>(CALL, MAP, LOAD...)</p> |

Internal

- \$KERNEL_SRC/Documentation/networking/filter.txt
- \$KERNEL_SRC/include/linux/filter.h

```
/* Helper macros for filter block array initializers. */

/* ALU ops on registers, bpf_add|sub|...: dst_reg += src_reg */

#define BPF_ALU64_REG(OP, DST, SRC) \
    ((struct bpf_insn) { \
        .code = BPF_ALU64 | BPF_OP(OP) | BPF_X, \
        .dst_reg = DST, \
        .src_reg = SRC, \
        .off = 0, \
        .imm = 0 })

#define BPF_ALU32_REG(OP, DST, SRC) \
    ((struct bpf_insn) { \
        .code = BPF_ALU | BPF_OP(OP) | BPF_X, \
        .dst_reg = DST, \
        .src_reg = SRC, \
        .off = 0, \
        .imm = 0 })
```

```
...
.insns = {
    BPF_MOV64_REG(BPF_REG_2, BPF_REG_10),
    BPF_ALU64_IMM(BPF_ADD, BPF_REG_2, -8),
    BPF_ST_MEM(BPF_DW, BPF_REG_2, 0, 0),
    BPF_LD_MAP_FD(BPF_REG_1, 0),
    BPF_EMIT_CALL(BPF_FUNC_map_lookup_elem),
    BPF_MOV64_REG(BPF_REG_1, BPF_REG_10),
    BPF_ALU64_IMM(BPF_ADD, BPF_REG_1, -152),
    BPF_STX_MEM(BPF_DW, BPF_REG_1, BPF_REG_0, 0),
    BPF_JMP_IMM(BPF_JEQ, BPF_REG_0, 0, 2),
    BPF_LDX_MEM(BPF_DW, BPF_REG_3, BPF_REG_1, 0),
    BPF_ST_MEM(BPF_DW, BPF_REG_3, 0, 42),
    BPF_EXIT_INSN(),
}
}
```

```
struct bpf_prog {
    u16 pages; /* Number of allocated pages */
    kmemcheck_bitfield_begin(meta);
    u16 jited:1, /* Is our filter JIT'ed? */
        locked:1, /* Program image locked? */
        gpl_compatible:1, /* Is filter GPL compatible? */
        cb_access:1, /* Is control block accessed? */
        dst_needed:1, /* Do we need dst entry? */
    kmemcheck_bitfield_end(meta);
    enum bpf_prog_type type; /* Type of BPF program */
    u32 len; /* Number of filter blocks */
    u32 jited_len; /* Size of jited insns in bytes */
    u8 tag[BPF_TAG_SIZE];
    struct bpf_prog_aux *aux; /* Auxiliary fields */
    struct sock_fprog_kern *orig_prog; /* Original BPF program */
    unsigned int (*bpf_func)(const void *ctx,
                             const struct bpf_insn *insns);
    /* Instructions for interpreter */
    union {
        struct sock_filter insns[0];
        struct bpf_insn insnsi[0];
    };
};
/* end bpf_prog */
```

```
...
bpf_prog_select_runtime
bpf_prog_free
bpf_prog_alloc
bpf_prog_realloc
__bpf_prog_free
bpf_prog_unlock_free
bpf_aux_classic_check_t
bpf_prog_create
bpf_prog_create_from_user
bpf_prog_destroy
sk_attach_filter
sk_attach_bpf
sk_reuseport_attach_filter
sk_reuseport_attach_bpf
sk_detach_filter
sk_get_filter
sk_filter_charge
sk_filter_uncharge
__bpf_call_base
bpf_int_jit_compile
bpf_jit_compile
bpf_helper_changes_pkt_data
bpf_patch_insn_single
xdp_do_generic_redirect
xdp_do_redirect
xdp_do_flush_map
bpf_warn_invalid_xdp_action
bpf_warn_invalid_xdp_redirect
do_sk_redirect_map
...
```

\$KERNEL_SRC/include/uapi/linux/bpf.h

```
enum bpf_prog_type {
    BPF_PROG_TYPE_UNSPEC,
    BPF_PROG_TYPE_SOCKET_FILTER,
    BPF_PROG_TYPE_KPROBE,
    BPF_PROG_TYPE_SCHED_CLS,
    BPF_PROG_TYPE_SCHED_ACT,
    BPF_PROG_TYPE_TRACEPOINT,
    BPF_PROG_TYPE_XDP,
    BPF_PROG_TYPE_PERF_EVENT,
    BPF_PROG_TYPE_CGROUP_SKB,
    BPF_PROG_TYPE_CGROUP_SOCK,
    BPF_PROG_TYPE_LWT_IN,
    BPF_PROG_TYPE_LWT_OUT,
    BPF_PROG_TYPE_LWT_XMIT,
    BPF_PROG_TYPE_SOCK_OPS,
    BPF_PROG_TYPE_SK_SKB,
};
```

```
/* BPF syscall commands, see bpf(2) man-page for details. */
enum bpf_cmd {
    BPF_MAP_CREATE,
    BPF_MAP_LOOKUP_ELEM,
    BPF_MAP_UPDATE_ELEM,
    BPF_MAP_DELETE_ELEM,
    BPF_MAP_GET_NEXT_KEY,
    BPF_PROG_LOAD,
    BPF_OBJ_PIN,
    BPF_OBJ_GET,
    BPF_PROG_ATTACH,
    BPF_PROG_DETACH,
    BPF_PROG_TEST_RUN,
    BPF_PROG_GET_NEXT_ID,
    BPF_MAP_GET_NEXT_ID,
    BPF_PROG_GET_FD_BY_ID,
    BPF_MAP_GET_FD_BY_ID,
    BPF_OBJ_GET_INFO_BY_FD,
};
```

```
enum bpf_map_type {
    BPF_MAP_TYPE_UNSPEC,
    BPF_MAP_TYPE_HASH,
    BPF_MAP_TYPE_ARRAY,
    BPF_MAP_TYPE_PROG_ARRAY,
    BPF_MAP_TYPE_PERF_EVENT_ARRAY,
    BPF_MAP_TYPE_PERCPU_HASH,
    BPF_MAP_TYPE_PERCPU_ARRAY,
    BPF_MAP_TYPE_STACK_TRACE,
    BPF_MAP_TYPE_CGROUP_ARRAY,
    BPF_MAP_TYPE_LRU_HASH,
    BPF_MAP_TYPE_LRU_PERCPU_HASH,
    BPF_MAP_TYPE_LPM_TRIE,
    BPF_MAP_TYPE_ARRAY_OF_MAPS,
    BPF_MAP_TYPE_HASH_OF_MAPS,
    BPF_MAP_TYPE_DEVMAP,
    BPF_MAP_TYPE_SOCKMAP,
};
```

```
struct bpf_insn {
    __u8 code; /* opcode */
    __u8 dst_reg:4; /* dest register */
    __u8 src_reg:4; /* source register */
    __s16 off; /* signed offset */
    __s32 imm; /* signed immediate constant */
};
```

```
/* User return codes for XDP prog type.
 * A valid XDP program must return one of these defined values. All other
 * return codes are reserved for future use. Unknown return codes will
 * result in packet drops and a warning via bpf_warn_invalid_xdp_action().
 */
```

```
enum xdp_action {
    XDP_ABORTED = 0,
    XDP_DROP,
    XDP_PASS,
    XDP_TX,
    XDP_REDIRECT,
};
```

```
/* user accessible metadata for XDP packet hook
 * new fields must be added to the end of this structure
 */
```

```
struct xdp_md {
    __u32 data;
    __u32 data_end;
};
```

```
enum sk_action {
    SK_ABORTED = 0,
    SK_DROP,
    SK_REDIRECT,
};
```

```
#define BPF_FUNC_MAPPER(FN) \
    FN(unspec), \
    FN(map_lookup_elem), \
    FN(map_update_elem), \
    FN(map_delete_elem), \
    FN(probe_read), \
    FN(ktime_get_ns), \
    FN(trace_printk), \
    FN(get_prandom_u32), \
    FN(get_smp_processor_id), \
    FN(skb_store_bytes), \
    FN(l3_csum_replace), \
    FN(l4_csum_replace), \
    FN(tail_call), \
    FN(clone_redirect), \
    FN(get_current_pid_tgid), \
    FN(get_current_uid_gid), \
    FN(get_current_comm), \
    FN(get_cgroup_classid), \
    FN(skb_vlan_push), \
    FN(skb_vlan_pop), \
    FN(skb_get_tunnel_key), \
    FN(skb_set_tunnel_key), \
    FN(perf_event_read), \
    ...
```


■ \$KERNEL_SRC/kernel/bpf

bpf
— arraymap.c
— bpf_lru_list.c
— bpf_lru_list.h
— cgroup.c
— core.c
— devmap.c
— hashtable.c
— helpers.c
— inode.c
— lpm_trie.c
— Makefile
— map_in_map.c
— map_in_map.h
— percpu_freelist.c
— percpu_freelist.h
— sockmap.c
— stackmap.c
— syscall.c
— tnum.c
— verifier.c

struct bpf_prog *bpf_prog_select_runtime(struct bpf_prog *fp, int *err)

static int bpf_prog_load(union bpf_attr *attr)

int bpf_check(struct bpf_prog **prog, union bpf_attr *attr)

■ \$KERNEL_SRC/arch/\$ARCH/net/bpf_jit_comp.c \$KERNEL_SRC/arch/\$ARCH/net/ebpf_jit.c

...

struct bpf_prog *bpf_int_jit_compile(struct bpf_prog *prog)

2) LLVM

- eBPF backend firstly introduced in LLVM 3.7 release
- <https://reviews.llvm.org/D6494>
- <http://llvm.org/docs/CodeGenerator.html#the-extended-berkeley-packet-filter-ebpf-backend>
- `$LLVM_SRC/lib/Target/BPF`

- Enabled by default with all major distributions
 - Registered targets: `llc --version`
 - `llc`'s BPF `-march` options: `bpf`, `bpfel`, `bpfef`
 - `llc`'s BPF `-mcpu` options: `generic`, `v1`, `v2`, `probe`
- Assembler output through `-S` supported
- `llvm-objdump` for disassembler and code annotations (via DWARF)
- Annotations correlate directly with kernel verifier log
- Outputs ELF file with maps as relocation entries
 - Processed by BPF loaders (e.g. `iproute2`) and pushed into kernel

Source: <https://ossna2017.sched.com/event/BCsg/making-the-kernels-networking-data-path-programmable-with-bpf-and-xdp-daniel-borkmann-covalent>

LLVM

- <https://en.wikipedia.org/wiki/LLVM>
- <http://clang.llvm.org/>



| GPL v3 | UIUC, MIT |
|----------------------|------------------|
| Front-end: CC1 / CPP | Front-end: Clang |
| ld.bfd / ld.gold | lld / mclinker |
| gdb | lldb |
| as / objdump | MC layer |
| libstdc++ | libc++ |
| libsupc++ | libc++abi |
| libgcc | libcompiler-rt |
| libgccjit | libLLVMMCJIT |

How is LLVM being used today?

XCode, Swift

FreeBSD, OpenMandriva Lx

Android

Debian experimenting with Clang as an additional compiler

...

Clang Goals

- GCC compatibility
- Fast compilation and low memory footprints
- Can reduce the linking time
- User friendly diagnostics
- Tooling
 - static analyzers
 - sanitizers



`$KERNEL_SRC/samples/bpf/Makefile`

3) Development

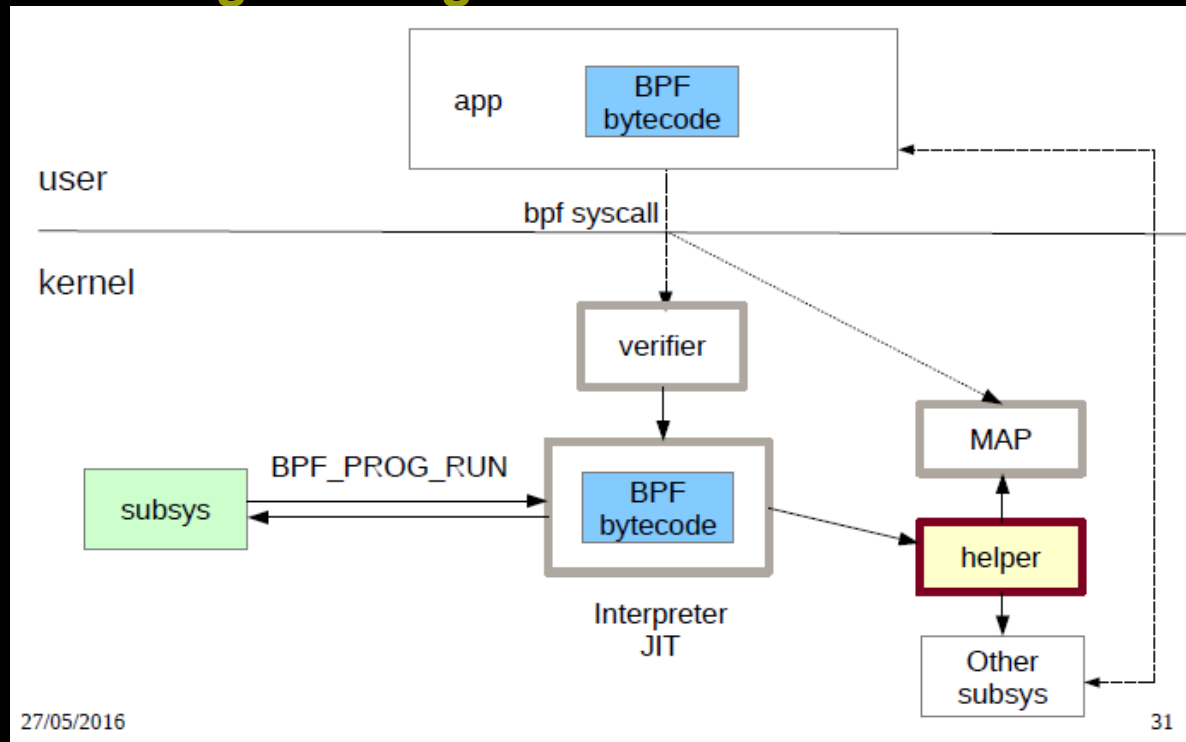
- Methods

- 1) eBPF assembly

- 2) **BCC**

...

- **BPF Programming Flow**

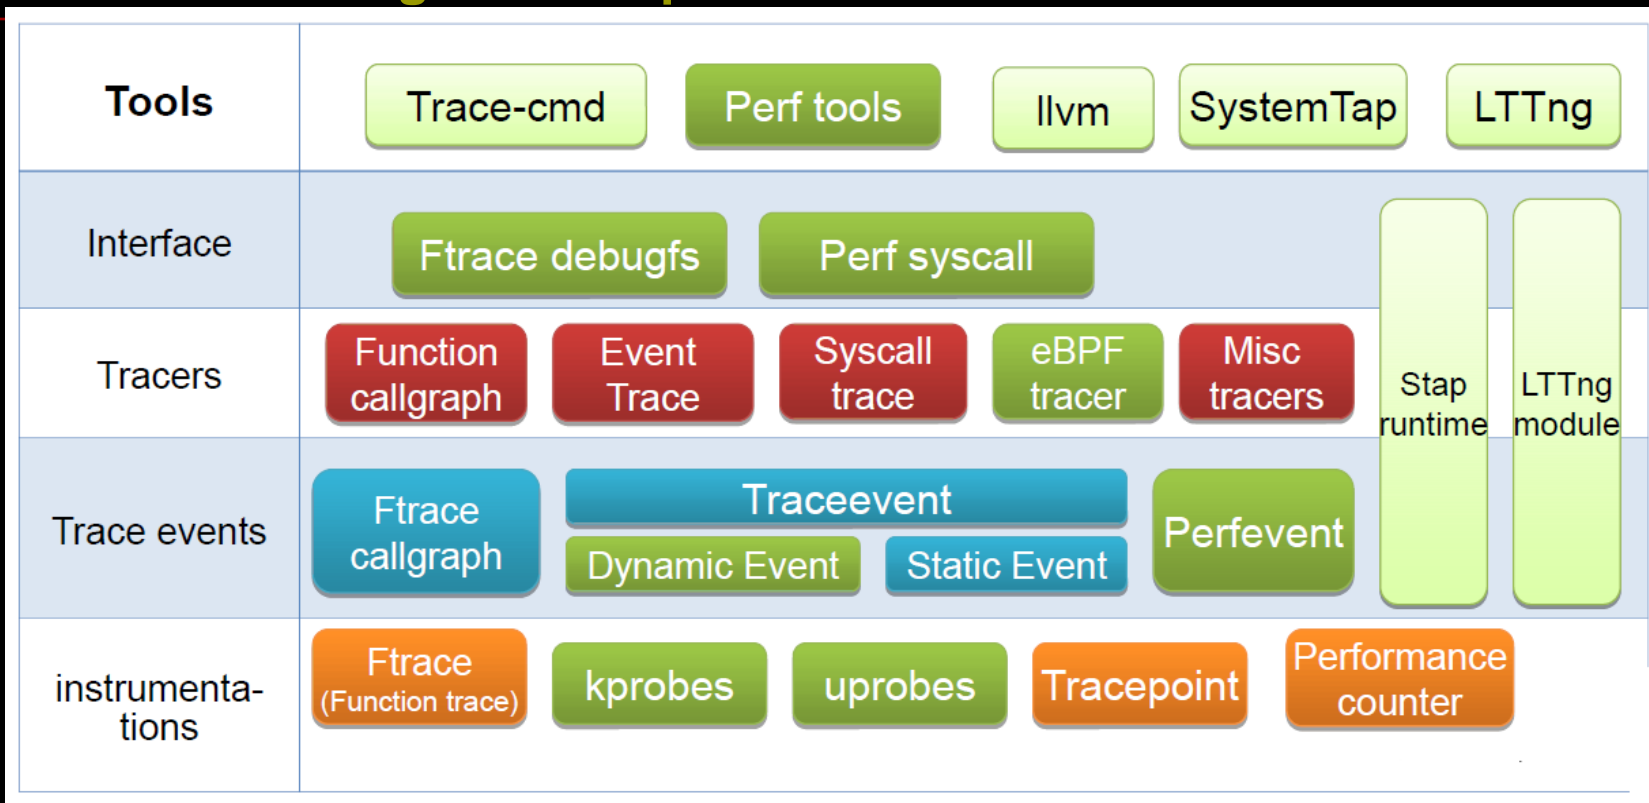


Source: <http://www.slideshare.net/vh21/meet-cutebetweenebpfandtracing>

2. eBPF for Kernel Instrumentation

Overview

■ The Tracing Landscape



Source: <http://tracingsummit.org/w/images/8/8c/TracingSummit2015-DynamicProbes.pdf>

1) BCC (BPF Compiler Collection)

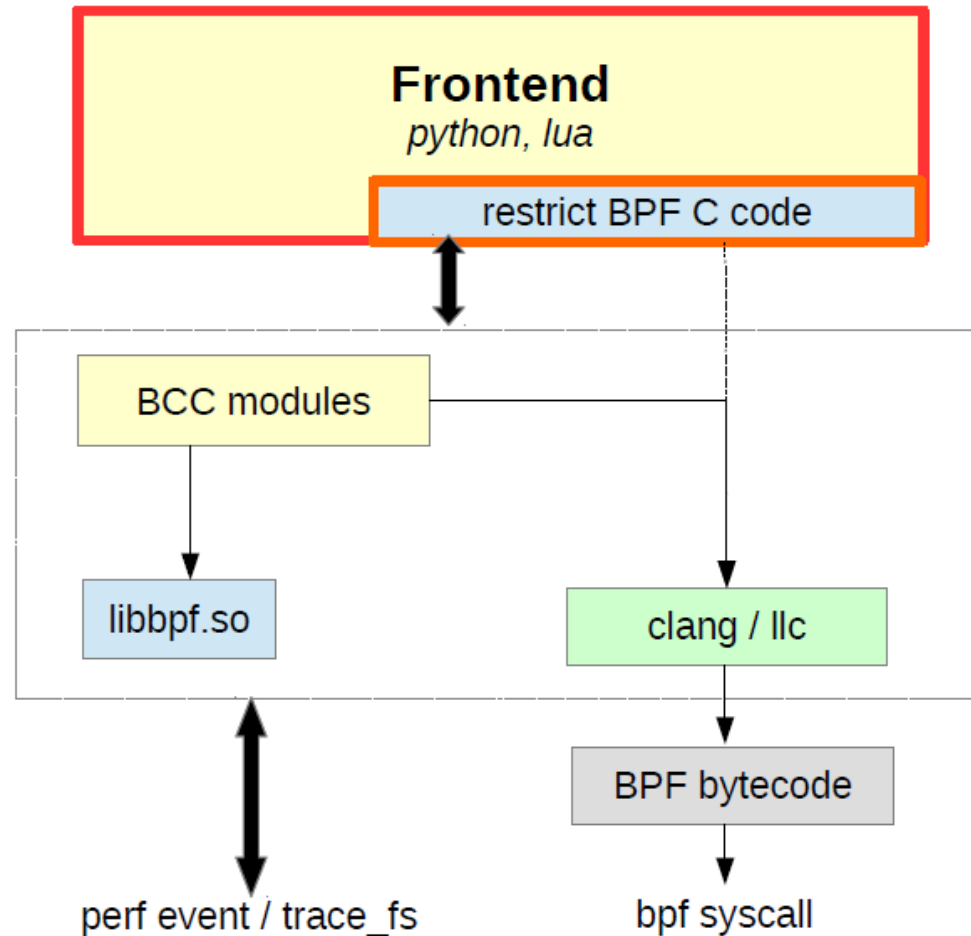
- <https://iovisor.github.io/bcc/>
- <https://github.com/iovisor/bcc.git>

A toolkit with Python/Lua frontend for compiling, loading, and executing BPF programs, which allows user-defined instrumentation on a live kernel image:

- **Compile BPF program from C source**
- **Attach BPF program to kprobe/uprobe/tracepoint/USDT/socket**
- **Poll data from BPF program**
- **Framework for building new tools or one-off scripts**
- **Contains a **P4** compiler for BPF targets**
- **Additional projects to support Go, Rust, and DTrace-style frontend**
- **...**

Arch

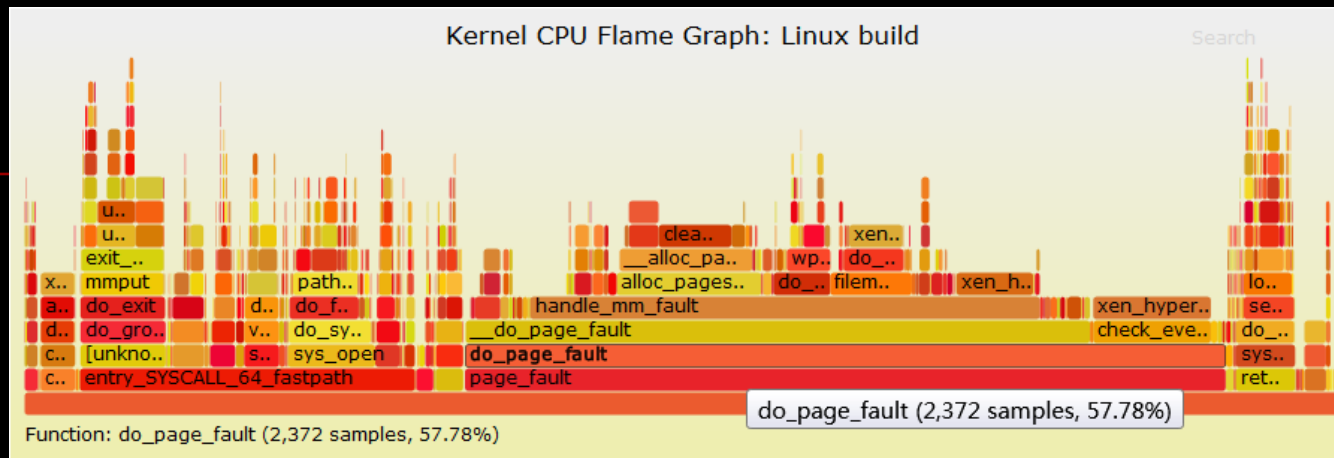
iovisor
BCC



27/05/2016

Source: <http://www.slideshare.net/vh21/meet-cutebetweenebpfandtracing>

Linux eBPF Flame Graph



Source: <http://www.brendangregg.com/blog/2016-01-20/ebpf-offcpu-flame-graph.html>

A Sample

bcc/examples/tracing/urandomread*.*

```
root@ubuntu: /opt/MyWorkspace/MyProjs/Open-Source/OS/In-Kernel-VM/eBPF/BCC/bcc/examples/tracing# ./urandomread.py
TIME(s)      COMM      PID      GOTBITS
3031.665037000  dd        6604     8192
3031.665365000  dd        6604     8192
3031.665642000  dd        6604     8192
3031.665924000  dd        6604     8192
3031.666202000  dd        6604     8192
3095.286445000  systemd   1        128
3095.286518000  systemd   1        128
3095.286582000  systemd   1        128
3095.286671000  systemd   1        128
```

```
nydev@ubuntu:/opt/Tmp$ dd if=/dev/urandom of=/dev/null bs=1k count=5
5+0 records in
5+0 records out
5120 bytes (5.1 kB, 5.0 KiB) copied, 0.00182226 s, 2.8 MB/s
```



```

1  #!/usr/bin/python
2  #
3  # urandomread-explicit  Example of instrumenting a kernel tracepoint.
4  #                        For Linux, uses BCC, BPF. Embedded C.
5  #
6  # This is an older example of instrumenting a tracepoint, which defines
7  # the argument struct and makes an explicit call to attach_tracepoint().
8  # See urandomread for a newer version that uses TRACEPOINT_PROBE().
9  #
10 # REQUIRES: Linux 4.7+ (BPF_PROG_TYPE_TRACEPOINT support).
11 #
12 # Test by running this, then in another shell, run:
13 #   dd if=/dev/urandom of=/dev/null bs=1k count=5
14 #
15 # Copyright 2016 Netflix, Inc.
16 # Licensed under the Apache License, Version 2.0 (the "License")
17
18 from __future__ import print_function
19 from bcc import BPF
20
21 # define BPF program
22 bpf_text = """
23 #include <uapi/linux/ptrace.h>
24
25 struct urandom_read_args {
26     // from /sys/kernel/debug/tracing/events/random/urandom_read/format
27     u64 __unused__;
28     u32 got_bits;
29     u32 pool_left;
30     u32 input_left;
31 };
32
33 int printarg(struct urandom_read_args *args) {
34     bpf_trace_printk("%d\\n", args->got_bits);
35     return 0;
36 };
37 """
38
39 # load BPF program
40 b = BPF(text=bpf_text)
41 b.attach_tracepoint("random:urandom_read", "printarg")
42
43 # header
44 print("%-18s %-16s %-6s %s" % ("TIME(s)", "COMM", "PID", "GOTBITS"))
45
46 # format output
47 while 1:
48     try:
49         (task, pid, cpu, flags, ts, msg) = b.trace_fields()
50     except ValueError:
51         continue
52     print("%-18.9f %-16s %-6d %s" % (ts, task, pid, msg))

```

include

/sys/kernel/debug/tracing/trace_pipe

```

290 TRACE_EVENT(urandom_read,
291 TP_PROTO(int got_bits, int pool_left, int input_left),
292 TP_ARGS(got_bits, pool_left, input_left),
293 TP_STRUCT__entry(
294     __field(int, got_bits)
295     __field(int, pool_left)
296     __field(int, input_left)
297 ),
298 TP_fast_assign(
299     __entry->got_bits = got_bits;
300     __entry->pool_left = pool_left;
301     __entry->input_left = input_left;
302 ),
303 TP_printk("got_bits %d nonblocking_pool_entropy_left %d "
304           "input_entropy_left %d", __entry->got_bits,
305           __entry->pool_left, __entry->input_left)
306 );

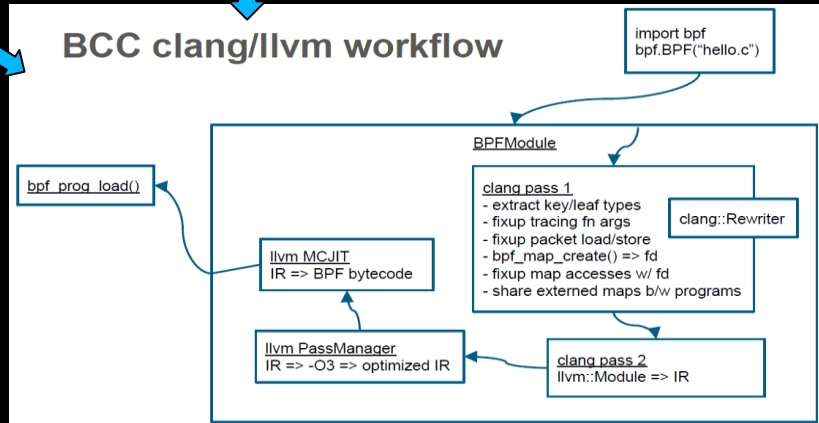
```

```

1  #!/usr/bin/python
2  #
3  # urandomread  Example of instrumenting a kernel tracepoint.
4  #              For Linux, uses BCC, BPF. Embedded C.
5  #
6  # REQUIRES: Linux 4.7+ (BPF_PROG_TYPE_TRACEPOINT support).
7  #
8  # Test by running this, then in another shell, run:
9  #   dd if=/dev/urandom of=/dev/null bs=1k count=5
10 #
11 # Copyright 2016 Netflix, Inc.
12 # Licensed under the Apache License, Version 2.0 (the "License")
13
14 from __future__ import print_function
15 from bcc import BPF
16
17 # load BPF program
18 b = BPF(text="""
19 TRACEPOINT_PROBE(random, urandom_read) {
20     // args is from /sys/kernel/debug/tracing/events/random/urandom_read/format
21     bpf_trace_printk("%d\\n", args->got_bits);
22     return 0;
23 };
24 """)
25
26 # header
27 print("%-18s %-16s %-6s %s" % ("TIME(s)", "COMM", "PID", "GOTBITS"))
28
29 # format output
30 while 1:
31     try:
32         (task, pid, cpu, flags, ts, msg) = b.trace_fields()
33     except ValueError:
34         continue
35     print("%-18.9f %-16s %-6d %s" % (ts, task, pid, msg))

```

BCC clang/llvm workflow

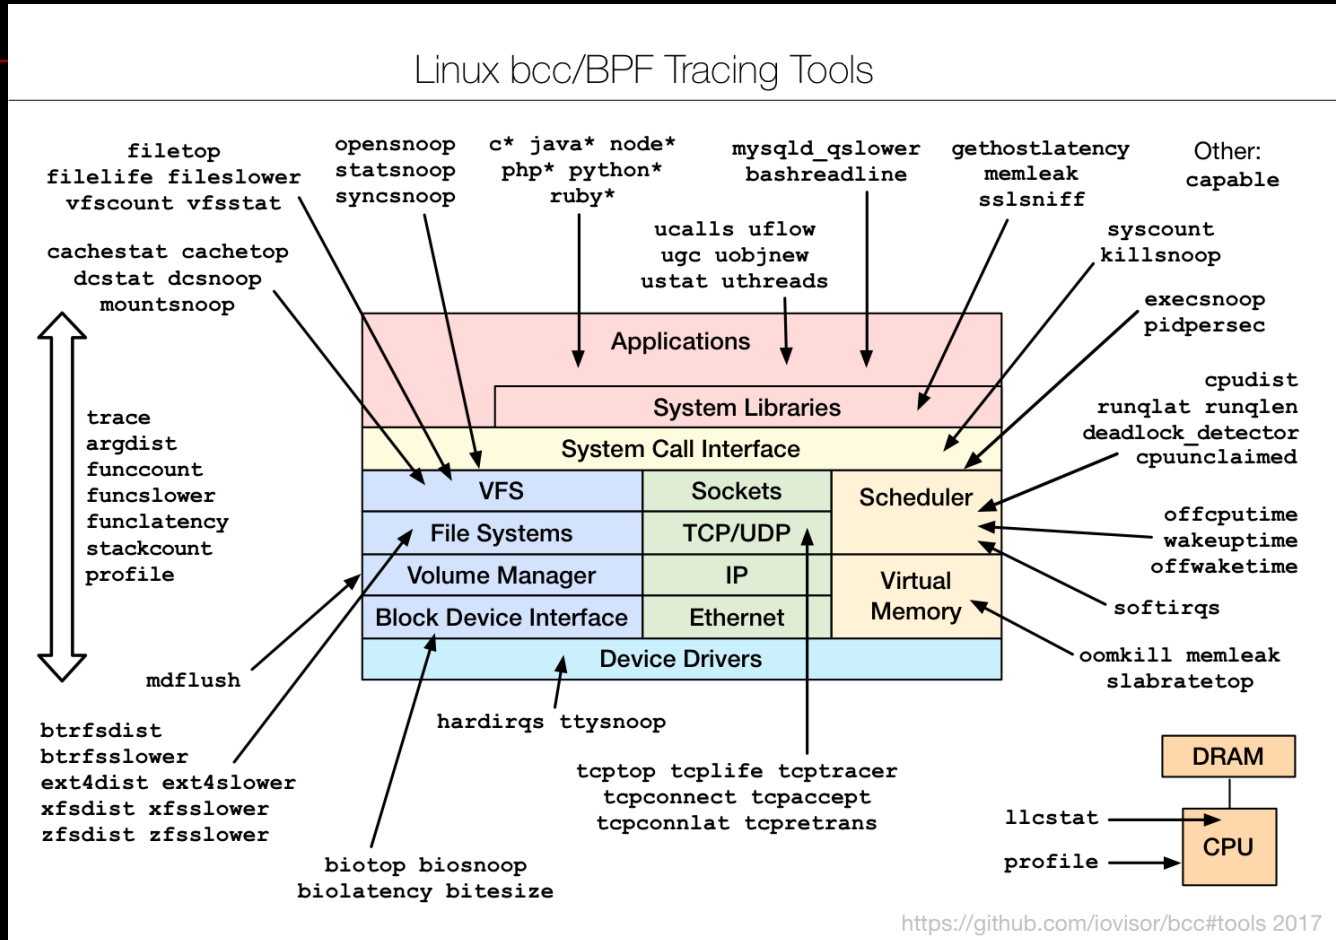


Source: http://linuxplumbersconf.org/2015/ocw/system/presentations/3249/original/bpf_llvm_2015aug19.pdf

3) Applications

Tuning

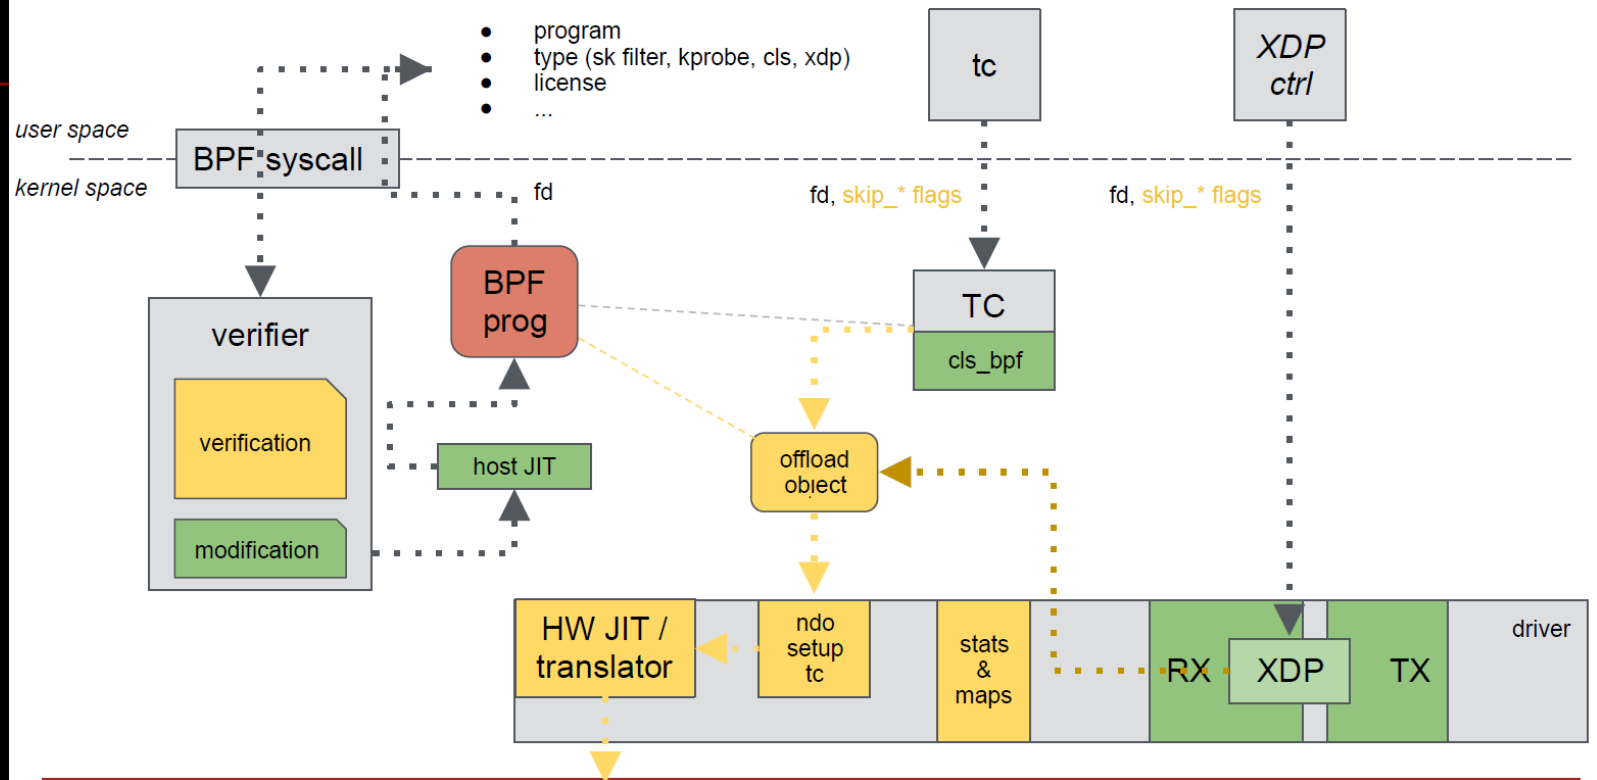
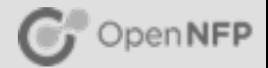
■ <http://www.brendangregg.com/blog/index.html>



Source: <https://github.com/iovisor/bcc/>

Offloading

Offload Architecture



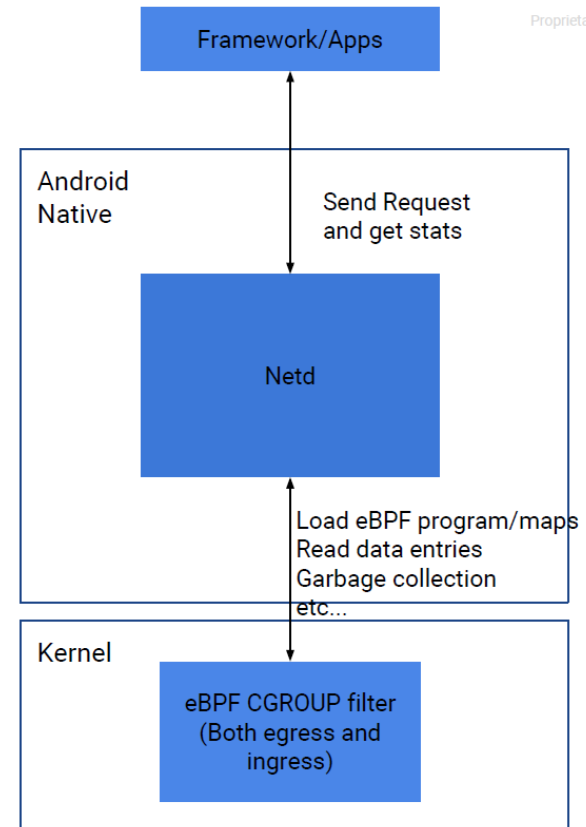
Source: <https://www.slideshare.net/Open-NFP/transparent-ebpf-offload-playing-nice-with-the-linux-kernel>

Netd on Android

- Old: xt_qtaguid module
- New: eBPF cgroup filters for data usage accounting

Basic Design

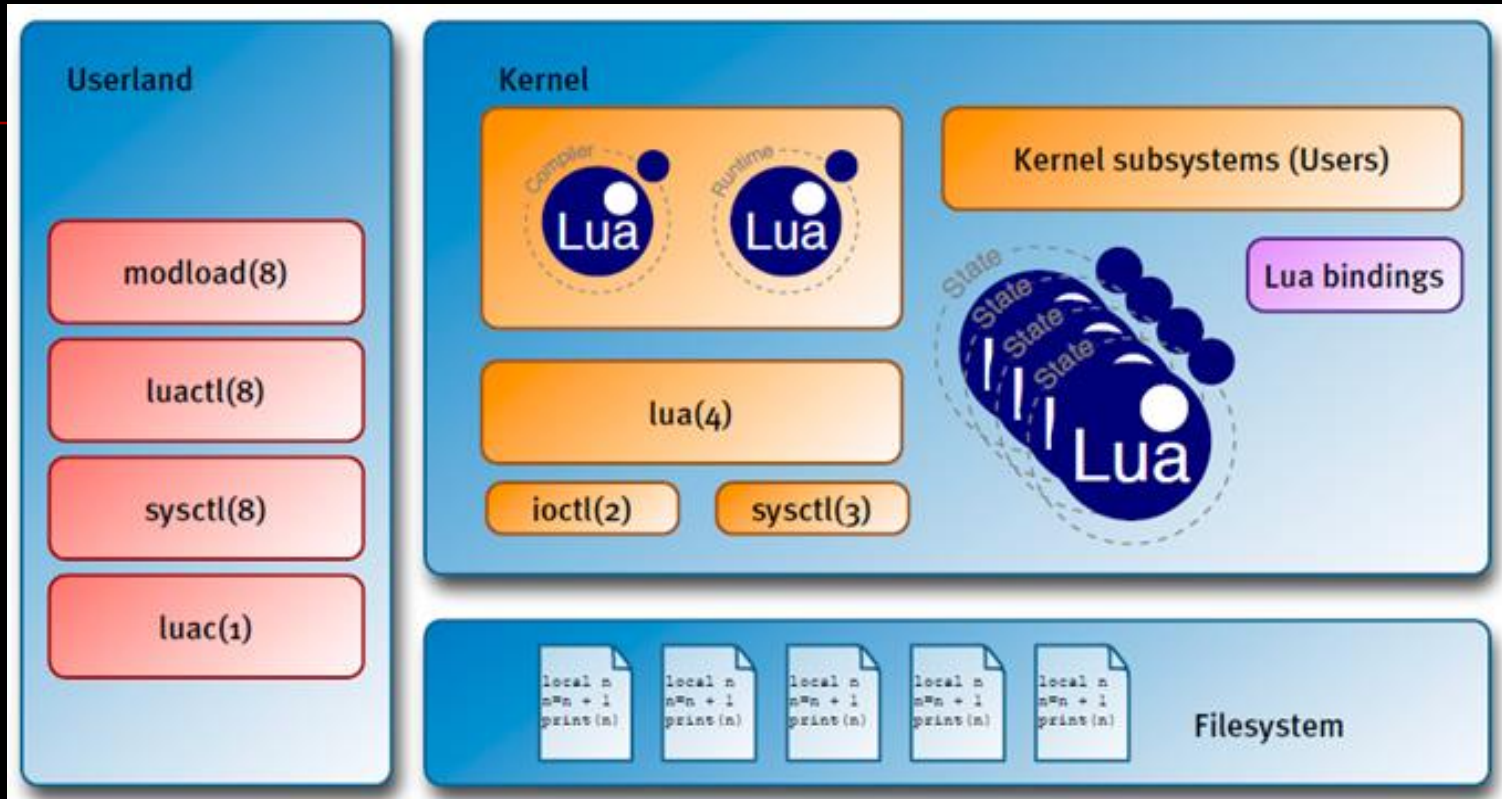
- Per-cgroup eBPF program to perform accounting
 - Ingress: Transport layer (e.g. tcp_v4_rcv), same as eBPF socket filter
 - Egress: Network layer (eg. ip_finish_output)
- Stats received are stored in eBPF maps.
- Stats periodically retrieved by privileged process from eBPF map
- Apps tag sockets by sending fd using binder call to privileged process



Source: <http://www.linuxplumbersconf.org/2017/ocw/proposals/4791>

Kernel Development

■ NetBSD Kernel scripting with Lua



Source: https://archive.fosdem.org/2013/schedule/event/lua_in_the_netbsd_kernel/

- deliver a higher-level programming environment to the Kernel
- great innovation in OS development

3) Pros & Cons

Pros

- Could replace lots of debugfs files
 - No need kernel debug symbols
 - Scalable for dynamic probing
 - Lower performance impact than even perf events
 - Security: sandboxing + verifier
 - On-the-fly program generation
 - ...
-

Cons

- Up to 512 bytes stack
- Max 4096 instructions per program
- No more than 64 maps
- ...

III. Cloud Computing with eBPF

1) Cilium Overview



■ <https://github.com/cilium/>

Cilium is open source software for providing and transparently securing network connectivity and loadbalancing between application containers and services deployed using Linux container management platforms like Docker and Kubernetes.

A new Linux kernel technology called **eBPF** is at the foundation of Cilium, which enables the dynamic insertion of BPF bytecode into the Linux kernel. Cilium generates eBPF programs for each individual application container to provide networking, security, loadbalancing and visibility.

Features Overview

- **Security Policies:** Enforcement of security policies at application (L7) and networking (L3-L4) layer. Application level policies include filtering of HTTP protocol properties such as method, path, host, and headers. Networking policies include container/pod/service interconnectivity rules based on labels, restriction of traffic to certain CIDR and/or port ranges for both ingress and egress.
- **Networking:** A simple flat Layer 3 network with the ability to span multiple clusters connects all application containers and services. Simple IP allocation using host scope allocators (dedicated /24 per cluster node for IPv4, dedicated /112 per cluster node for IPv6). Choice of either integrating with Linux routing to run a routing daemon or to create an overlay network using encapsulation (VXLAN/Geneve).
- **Load balancing:** Distributed load balancing for east-west traffic from application container to application container, e.g. implementation of Kubernetes services. North-south traffic to load balance external traffic, e.g. implementation of Kubernetes ingress. All load-balancing performed with direct server return (DSR) by default for improved performance.
- **Troubleshooting:** Built-in troubleshooting tools providing an alternative to traditional tcpdump troubleshooting techniques.
- **Integrations:**
 - Network plugin integrations: [CNI](#), [libnetwork](#)
 - Container runtime events: [containerd](#)
 - Kubernetes: [NetworkPolicy](#), [Labels](#), [Ingress](#), [Service](#)
 - Logging: [fluentd](#)

XDP

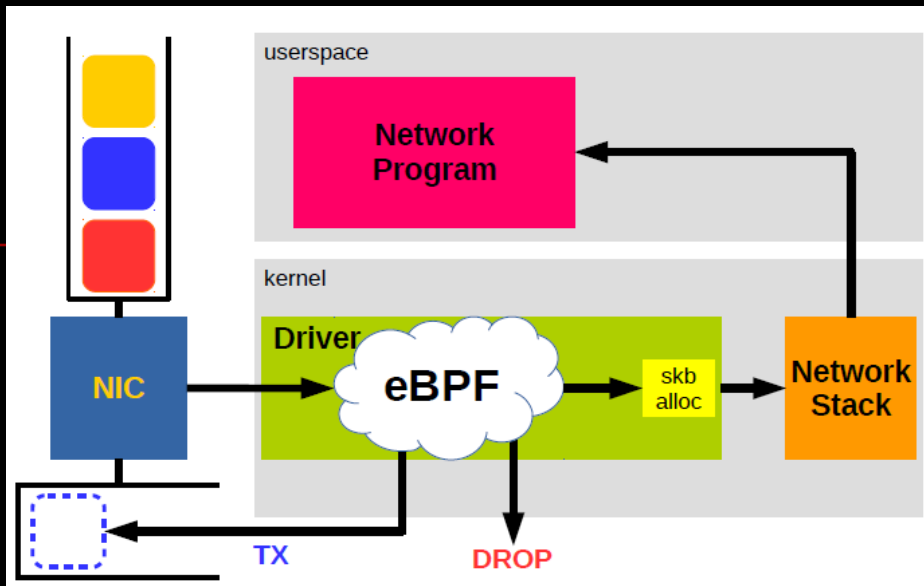
- <https://www.iovisor.org/technology/xdp>
- **eXpress Data Path**
- <https://lwn.net/Articles/708087/> //Debating the value of XDP
- **Generic hook**
- **eBPF-based “In-Kernel DPDK”**

XDP is a further step in evolution and enables to run a specific flavor of BPF programs from the network driver with direct access to the packet's DMA buffer. This is, by definition, the earliest possible point in the software stack, where programs can be attached to in order to allow for a programmable, high performance packet processor in the Linux kernel networking data path.

Source: <https://github.com/cilium/cilium>

- Works in concert with the kernel and its infrastructure (!)
- Advantages of XDP
 - Reuses upstream kernel drivers and tooling
 - Same security model as kernel for accessing hardware
 - Allows for flexible structuring of workloads
 - Punting to stable, efficient TCP/IP stack already available
 - No need for crossing boundaries when punting to sockets
 - No third party code/licensing required to use it
 - Shipped everywhere since kernel 4.8

Source: <https://ossna2017.sched.com/event/BCsg/making-the-kernels-networking-data-path-programmable-with-bpf-and-xdp-daniel-borkmann-covalent>

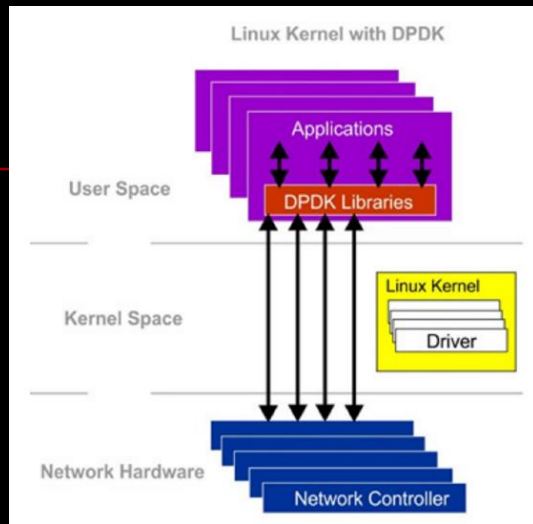


Source: <https://www.slideshare.net/lcplcp1/xdp-and-ebpfmaps>

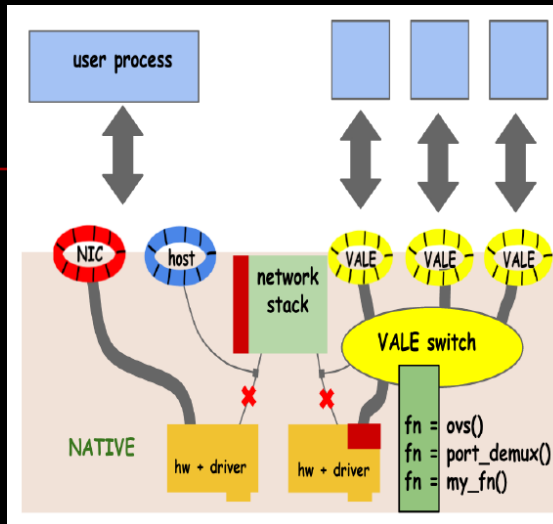
- eBPF trigger actions based on return codes
 - **XDP_DROP** - very fast drop by recycling
 - DDoS mitigation
 - **XDP_PASS** – pass possibly modified packet to network stack
 - Handle and pop new unknown encap protocols
 - **XDP_TX** – Transmit packet back out same interface
 - Facebook use it for load-balancing, and DDoS scrubber
 - **XDP_ABORTED** – also drop, but indicate error condition
 - Tracepoint: xdp_exception
 - **XDP_REDIRECT** – Transmit out other NICs
 - Very new (est.4.14), (plan also use for steering packets CPUs + sockets)

Source: http://people.netfilter.org/hawk/presentations/theCamp2017/theCamp2017_XDP_eBPF_technology_Jesper_Brouer.pdf

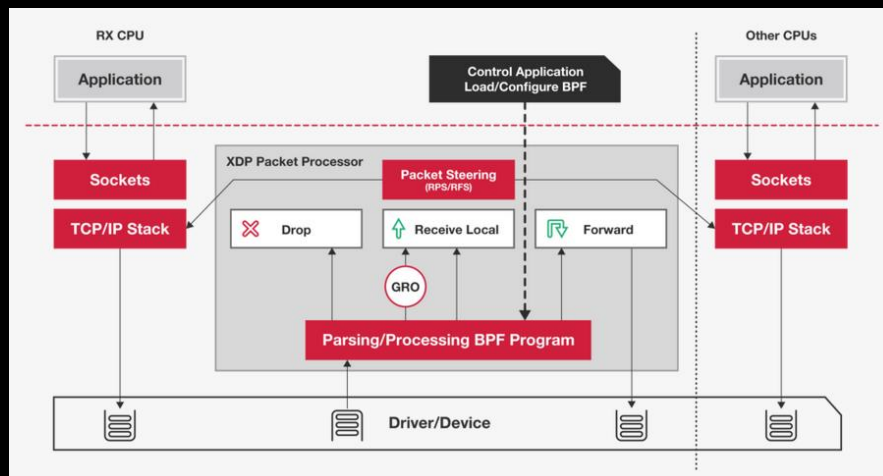
Comparison



DPDK



NetMap



XDP

eBPF Code Generation at Container Startup

- **Generate networking code at container startup, and tailored to each individual container**

On the fly BPF program generation means:

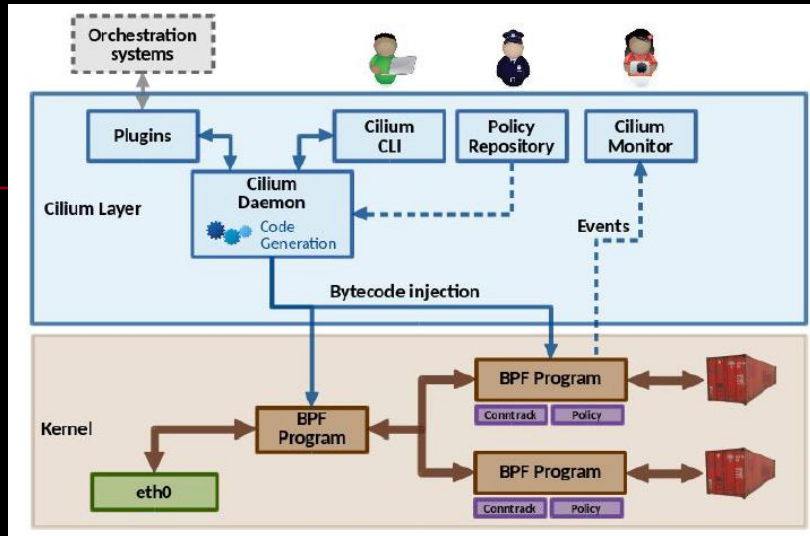
- Extensibility of userspace networking in the kernel
- MAC, IP, port number, ... all become constants
→ compiler can optimize heavily!
- BPF programs can be recompiled and replaced without interrupting the container and its connections
 - Features can be compiled in/out at runtime with container granularity
- Access to fast BPF maps and perf ring buffer to interact with userspace.
 - Drop monitor in $n * Mpps$ context
 - Use notifications for policy learning, IDS, logging, ...

Source: “Cilium: Fast IPv6 Container Networking with BPF and XDP” LinuxCon 2016, Toronto

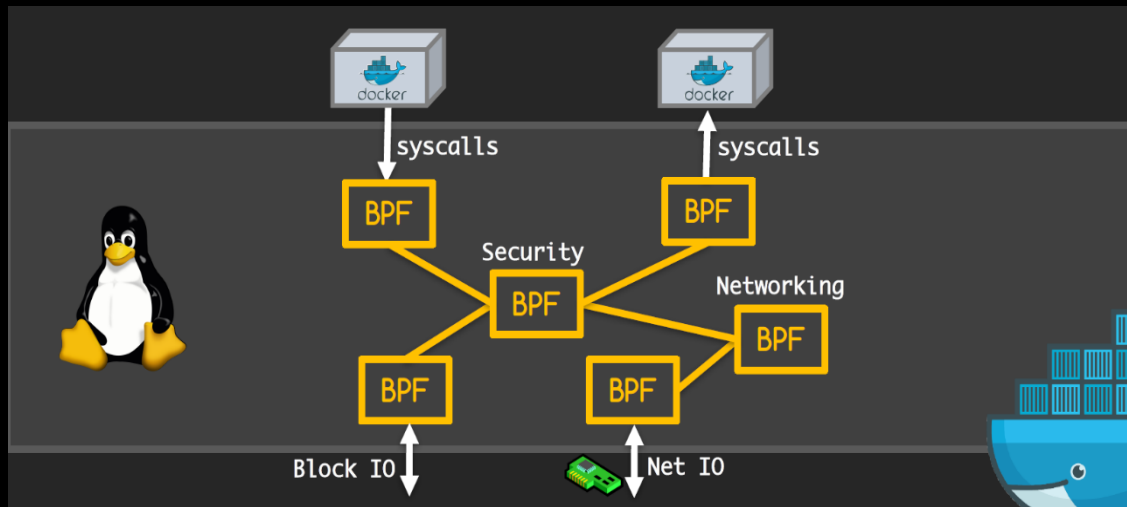
- **\$CILIUM_SRC/bpf**

```
bpf
├── bpf_features.h
├── bpf_lb.c
├── bpf_lxc.c
├── bpf_netdev.c
├── bpf_overlay.c
├── bpf_xdp.c
├── COPYING
├── filter_config.h
├── include
├── init.sh
├── join_ep.sh
├── lib
├── lxc_config.h
├── Makefile
├── netdev_config.h
├── node_config.h
├── probes
└── run_probes.sh
```

Arch



Source: <https://www.slideshare.net/ThomasGraf5/cilium-fast-ipv6-container-networking-with-bpf-and-xdp>

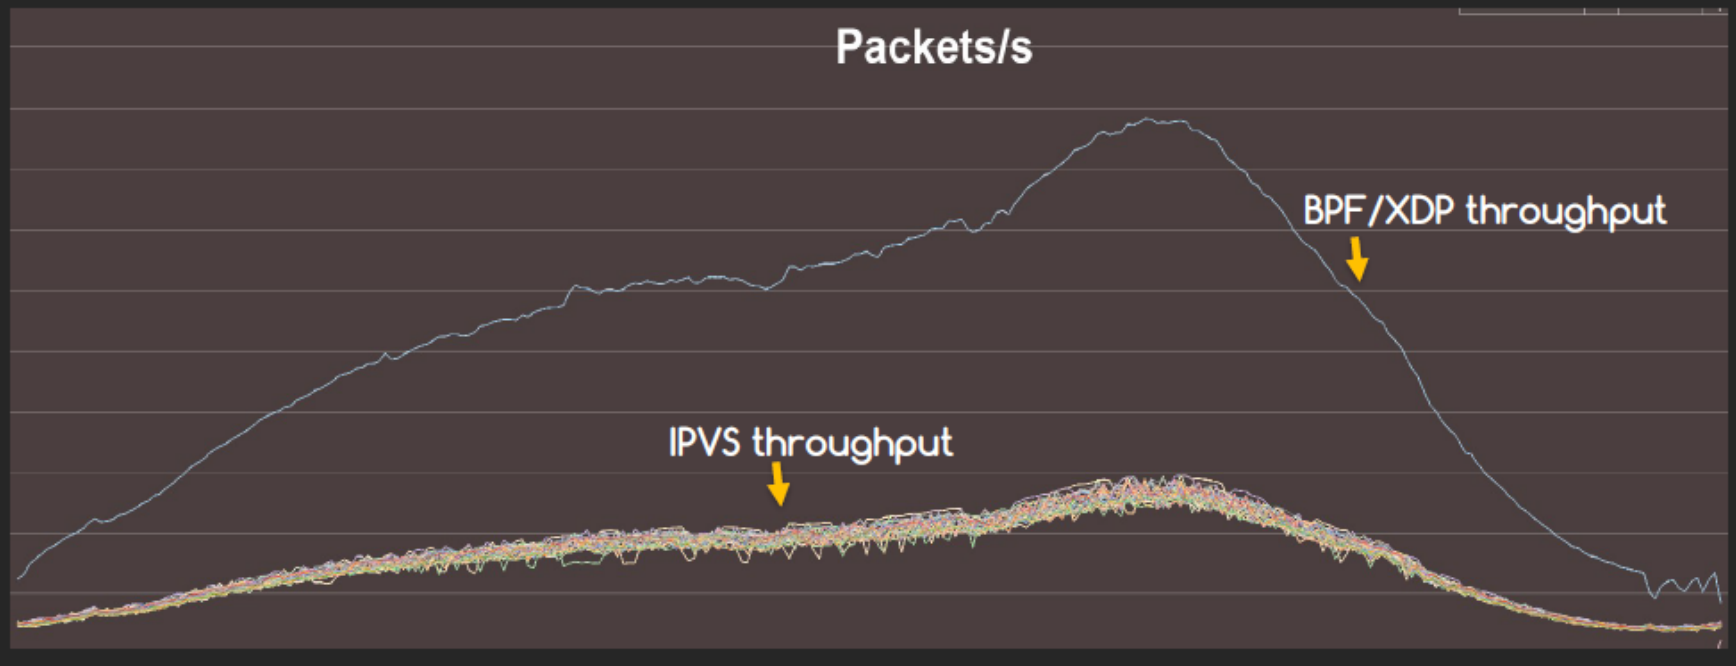


Source: <https://www.slideshare.net/ThomasGraf5/dockercon-2017-cilium-network-and-application-security-with-bpf-and-xdp>

2) Load Balance

■ <https://www.iovisor.org/>

Facebook published **BPF/XDP** numbers
for L3/L4 LB at Netdev 2.1

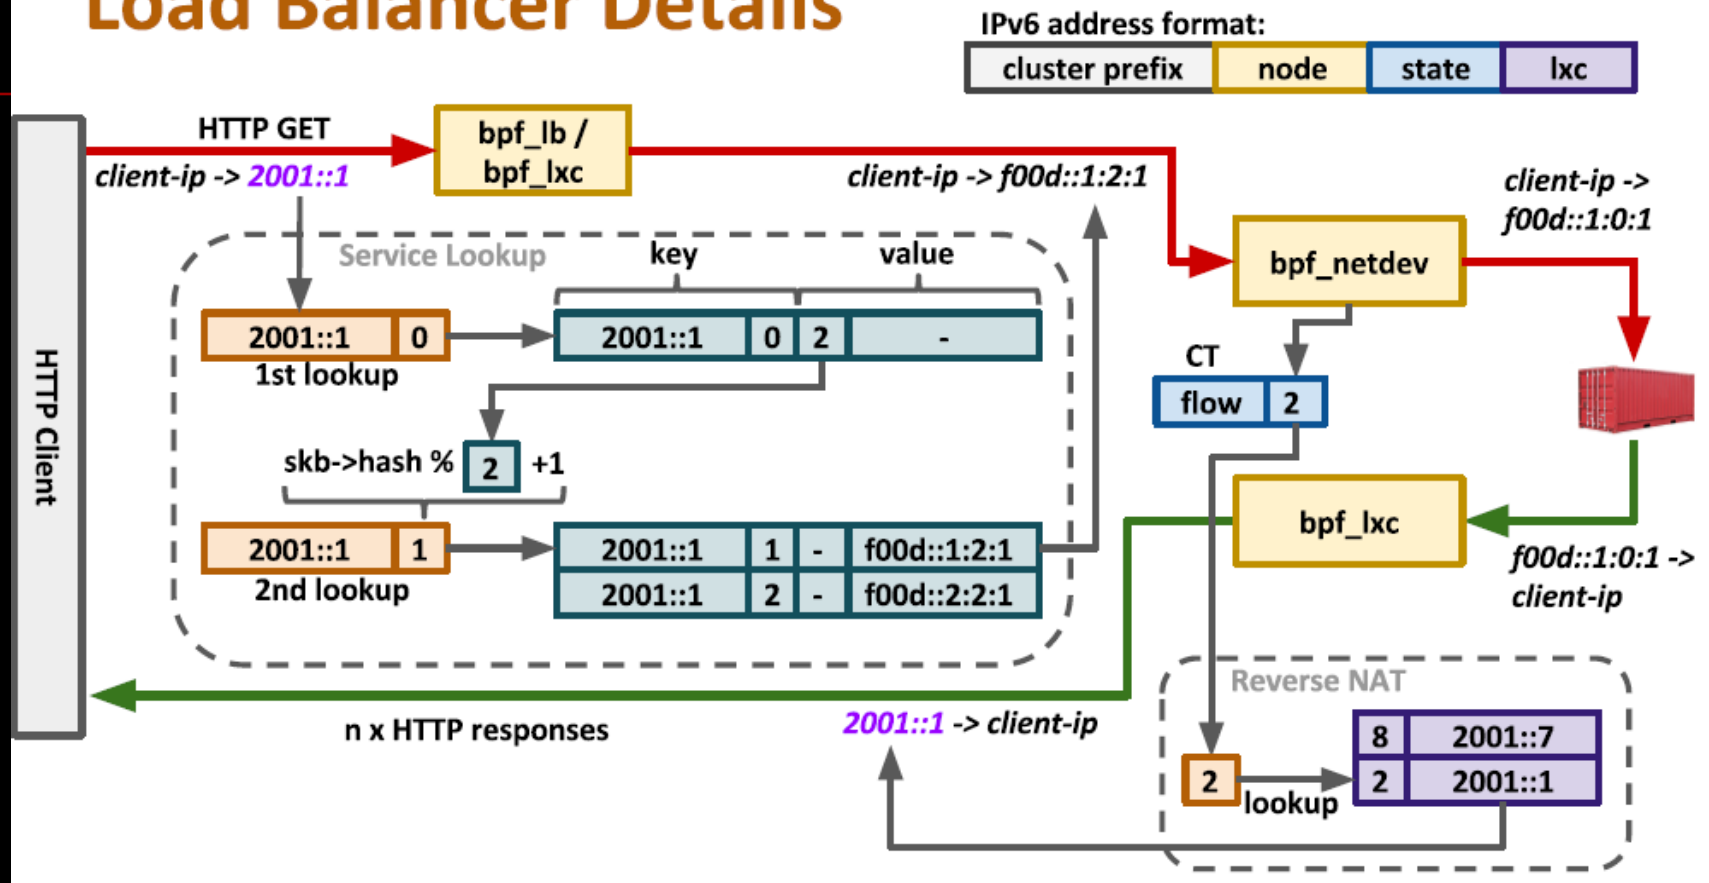


Source: <https://www.slideshare.net/ThomasGraf5/cilium-network-security-for-microservices>

LB in Cilium

■ scaling policy

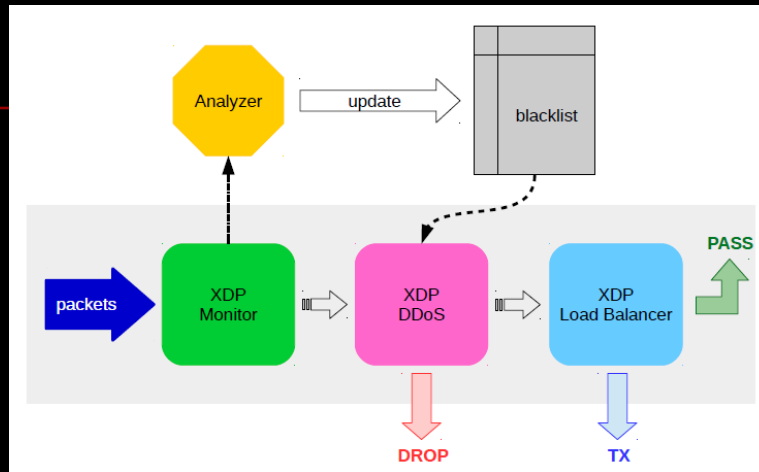
Load Balancer Details



Source: <https://www.slideshare.net/ThomasGraf5/cilium-container-networking-with-bpf-xdp>

3) Security

DDoS Protection



Source: <https://www.slideshare.net/lcplcp1/xdp-and-ebpfmaps>

| Metric | iptables / ipset | XDP |
|--------------------------------|------------------|--------|
| DDoS rate [packets/s] | 11.6M | 11.6M |
| Drop rate [packets/s] | 7.1M | 11.6M |
| Time to load rules [time] | 3 min 20 sec | 31 sec |
| Latency under load [ms] | 2.3ms | 0.1ms |
| Throughput under DDoS [Gbit/s] | 0.014 | 6.5 |
| Requests/s under DDoS [kReq/s] | 0.28 | 82.8 |

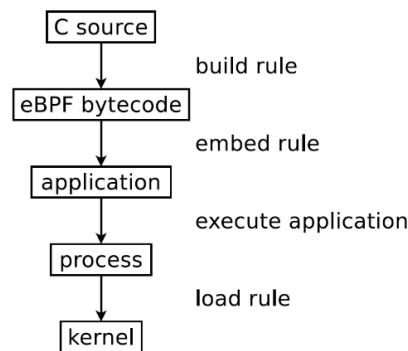
Source: <https://www.slideshare.net/ThomasGraf5/cilium-network-security-for-microservices>

LandLock

- **Linux Security Modules**
- https://en.wikipedia.org/wiki/Linux_Security_Modules
- <https://www.kernel.org/doc/Documentation/security/LSM.txt>
- <https://landlock.io/>
- <https://github.com/landlock-lsm/linux/commits/landlock-v7>
-

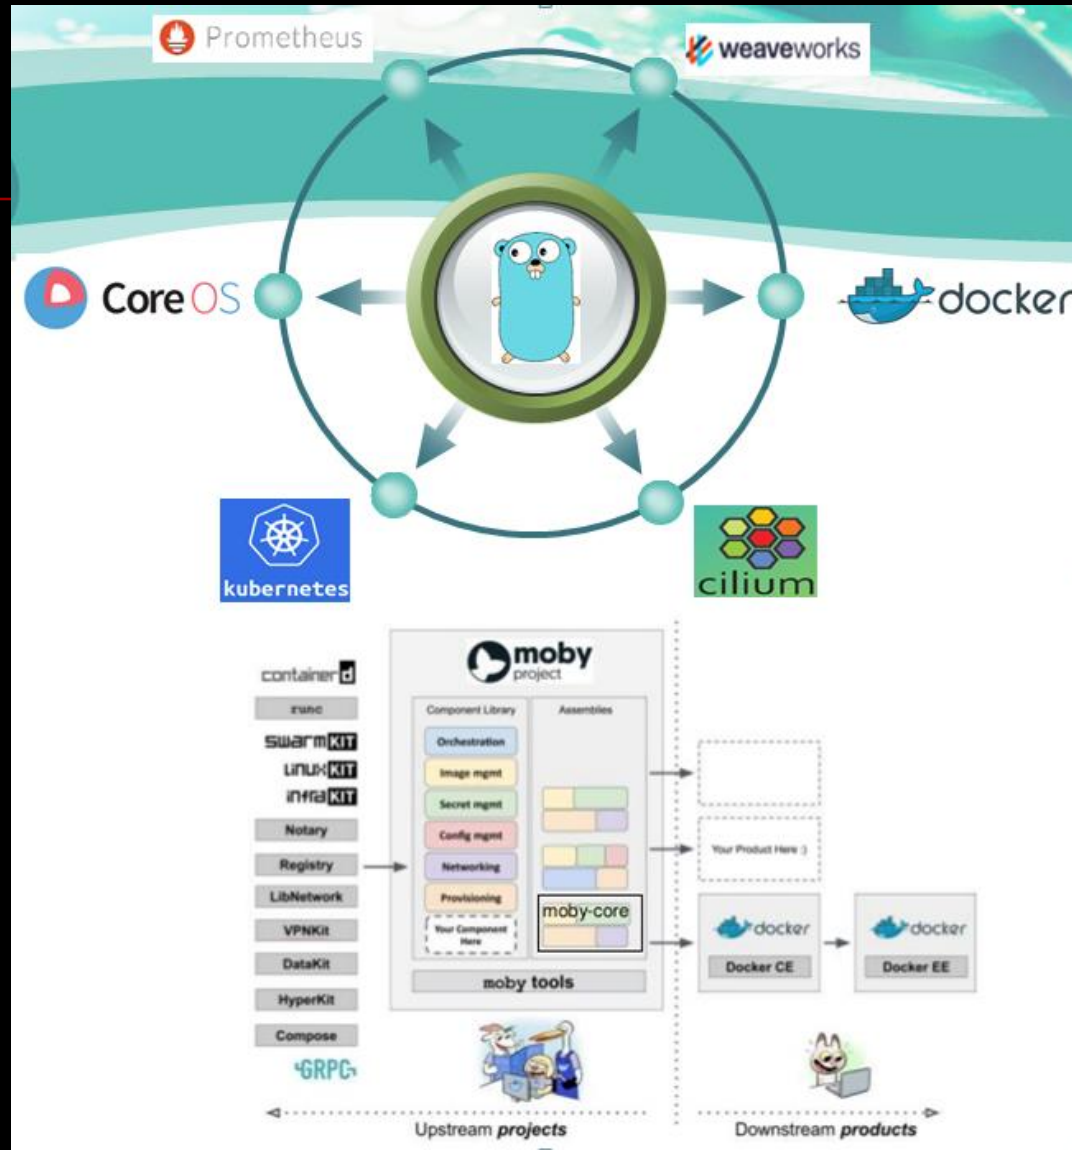
| | Fine-grained control | Embedded policy | Unprivileged use |
|--------------|----------------------|-----------------|------------------|
| SELinux. . . | ✓ | | |
| seccomp-bpf | | ✓ | ✓ |
| namespaces | | ✓ | ~ |
| Landlock | ✓ | ✓ | ✓ |

Life cycle of a Landlock rule



Source: http://events.linuxfoundation.org/sites/events/files/slides/2017-09-14_landlock-lss.pdf

4) Go-based Cloud Ecosystem



IV. eBPF on ARM



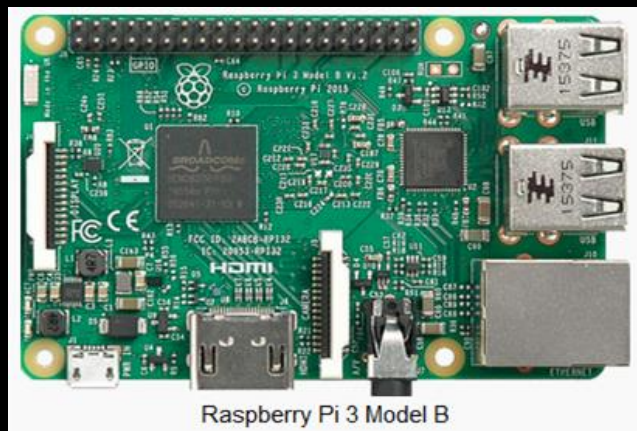
1) RPi3

- https://en.wikipedia.org/wiki/Raspberry_Pi

- <https://www.raspberrypi.org/>

RPi3 Model B

-

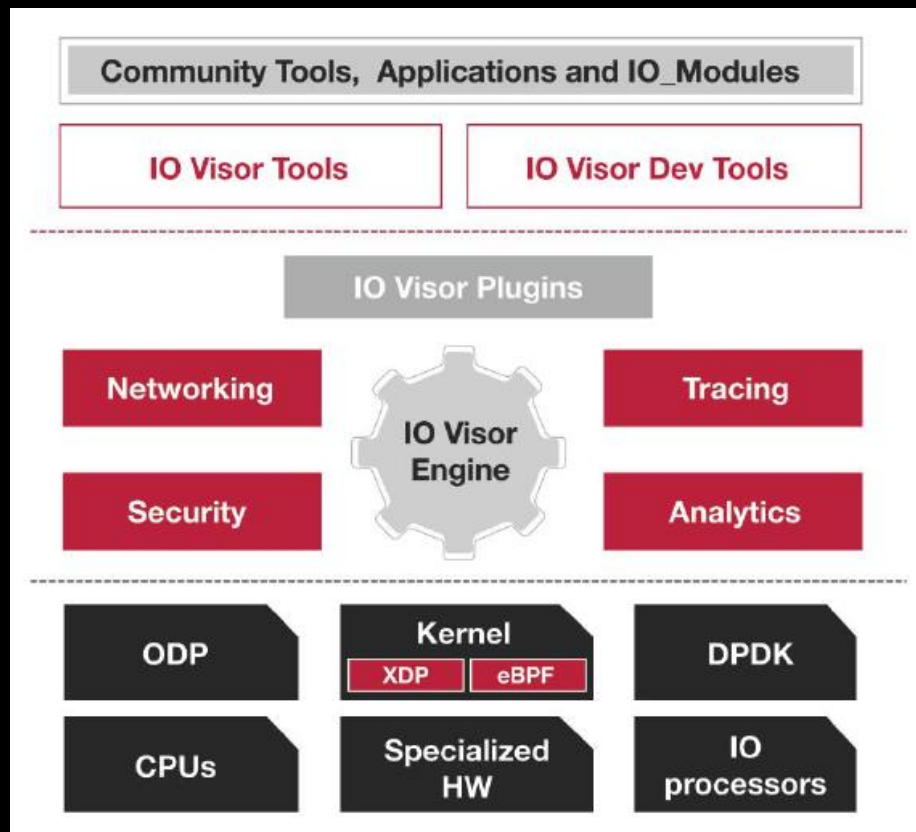


Limitations

- 1) 1.2 GHz 64-bit quad-core ARM Cortex-A53, 1GB LPDDR2 RAM @900MHz...
- 2) Official release (Raspbian with Linux Kernel 4.4 currently) does not support AArch64

2) IO Visor

- <https://www.iovisor.org/>
- Advancing In-Kernel IO Virtualization By Enabling Programmable Data Planes With Extensibility, Flexibility and High Performance
- eBPF-based



HypriotOS-RPi64

- <https://blog.hypriot.com/post/building-a-64bit-docker-os-for-rpi3/>
- <https://github.com/dieterreuter/workshop-raspberrypi-64bit-os>

HypriotOS/arm64: pirate@black-pearl in ~

\$ uname -a

Linux black-pearl 4.9.13-bee42-v8 #1 SMP PREEMPT Fri Mar 3 16:42:37 UTC 2017 aarch64 GNU/Linux

■ BCC on RPi3

HypriotOS/arm64: pirate@black-pearl in ~

\$ free -m

| | total | used | free | shared | buff/cache | available |
|-------|-------|------|-------|--------|------------|-----------|
| Mem: | 969 | 107 | 692 | 22 | 169 | 823 |
| Swap: | 16383 | 0 | 16383 | | | |

HypriotOS/arm64: pirate@black-pearl in ~

\$ cat /etc/ld.so.preload

/usr/local/jemalloc/lib/libjemalloc.so.2

HypriotOS/arm64: pirate@black-pearl in /usr/bin

\$ ll |grep ld

```
lrwxrwxrwx 1 root root      25 Oct 17 09:29 aarch64-linux-gnu-gold -> aarch64-linux-gnu-ld.gold*
lrwxrwxrwx 1 root root      24 Oct 17 09:29 aarch64-linux-gnu-ld -> aarch64-linux-gnu-ld.bfd*
-rwxr-xr-x 1 root root 1175520 Oct 17 09:29 aarch64-linux-gnu-ld.bfd*
-rwxr-xr-x 1 root root 5461968 Oct 17 09:29 aarch64-linux-gnu-ld.gold*
-rwxr-xr-x 1 root root 16567 Jul 14 01:20 build-rdeps*
-rwxr-xr-x 1 root root 6489 Jul 14 01:20 cvs-debuild*
-rwxr-xr-x 1 root root 37062 Jul 14 01:20 debuild*
-rwxr-xr-x 1 root root 1351 Oct 23 2016 dehtmldiff*
-rwxr-xr-x 1 root root 1335 Oct 14 07:19 dh_auto_build*
-rwxr-xr-x 1 root root 4583 Oct 14 07:19 dh_builddeb*
-rwxr-xr-x 1 root root 5868 Oct 14 07:19 dh_installdeb*
-rwxr-xr-x 1 root root 3274 Oct 14 07:19 dh_installdebconf*
-rwxr-xr-x 1 root root 1993 Oct 14 07:19 dh_installdirs*
-rwxr-xr-x 1 root root 14070 Oct 14 07:19 dh_installdocs*
-rwxr-xr-x 1 root root 7565 Oct 17 23:28 dpkg-buildflags*
-rwxr-xr-x 1 root root 29188 Oct 17 23:28 dpkg-buildpackage*
-rwxr-xr-x 1 root root 7503 Oct 17 23:28 dpkg-checkbuilddeps*
-rwxr-xr-x 1 root root 1005 Jul 14 01:20 dpkg-genbuilddeps*
-rwxr-xr-x 1 root root 16775 Oct 17 23:28 dpkg-genbuildinfo*
-rwxr-xr-x 1 root root 9921 Jul 28 22:54 equivs-build*
-rwxr-xr-x 1 root root 31400 Oct 2 17:51 fold*
-rwxr-xr-x 1 root root 4798 Jul 14 01:20 getbuildlog*
lrwxrwxrwx 1 root root      22 Oct 17 09:29 gold -> aarch64-linux-gnu-gold*
-rwxr-xr-x 1 root root 27704 Oct 16 12:26 gtk-builder-tool*
lrwxrwxrwx 1 root root      7 Oct 19 16:48 ld -> ld.gold*
lrwxrwxrwx 1 root root      24 Oct 17 09:29 ld.bfd -> aarch64-linux-gnu-ld.bfd*
-rwxr-xr-x 1 root root 5289 Aug 26 09:09 ldd*
lrwxrwxrwx 1 root root      25 Oct 17 09:29 ld.gold -> aarch64-linux-gnu-ld.gold*
lrwxrwxrwx 1 root root      36 Oct 13 19:20 lli-child-target-5.0 -> ../lib/llvm-5.0/bin/lli-child-target*
lrwxrwxrwx 1 root root      31 Oct 13 19:20 llvm-rtdyld-5.0 -> ../lib/llvm-5.0/bin/llvm-rtdyld*
```

repos

```
deb http://httpredir.debian.org/debian stretch main
deb-src http://httpredir.debian.org/debian stretch main

deb http://httpredir.debian.org/debian stretch-updates main
deb-src http://httpredir.debian.org/debian stretch-updates main

deb http://security.debian.org/ stretch/updates main
deb-src http://security.debian.org/ stretch/updates main

deb http://httpredir.debian.org/debian experimental main

deb http://deb.debian.org/debian sid main
deb-src http://deb.debian.org/debian sid main

deb http://deb.debian.org/debian sid-updates main
deb-src http://deb.debian.org/debian sid-updates main

deb http://security.debian.org/ sid/updates main
deb-src http://security.debian.org/ sid/updates main

deb http://deb.debian.org/debian stretch main contrib non-free
deb-src http://deb.debian.org/debian stretch main contrib non-free

deb http://deb.debian.org/debian stretch-updates main contrib non-free
deb-src http://deb.debian.org/debian stretch-updates main contrib non-free

deb http://security.debian.org/ stretch/updates main contrib non-free
deb-src http://security.debian.org/ stretch/updates main contrib non-free
```

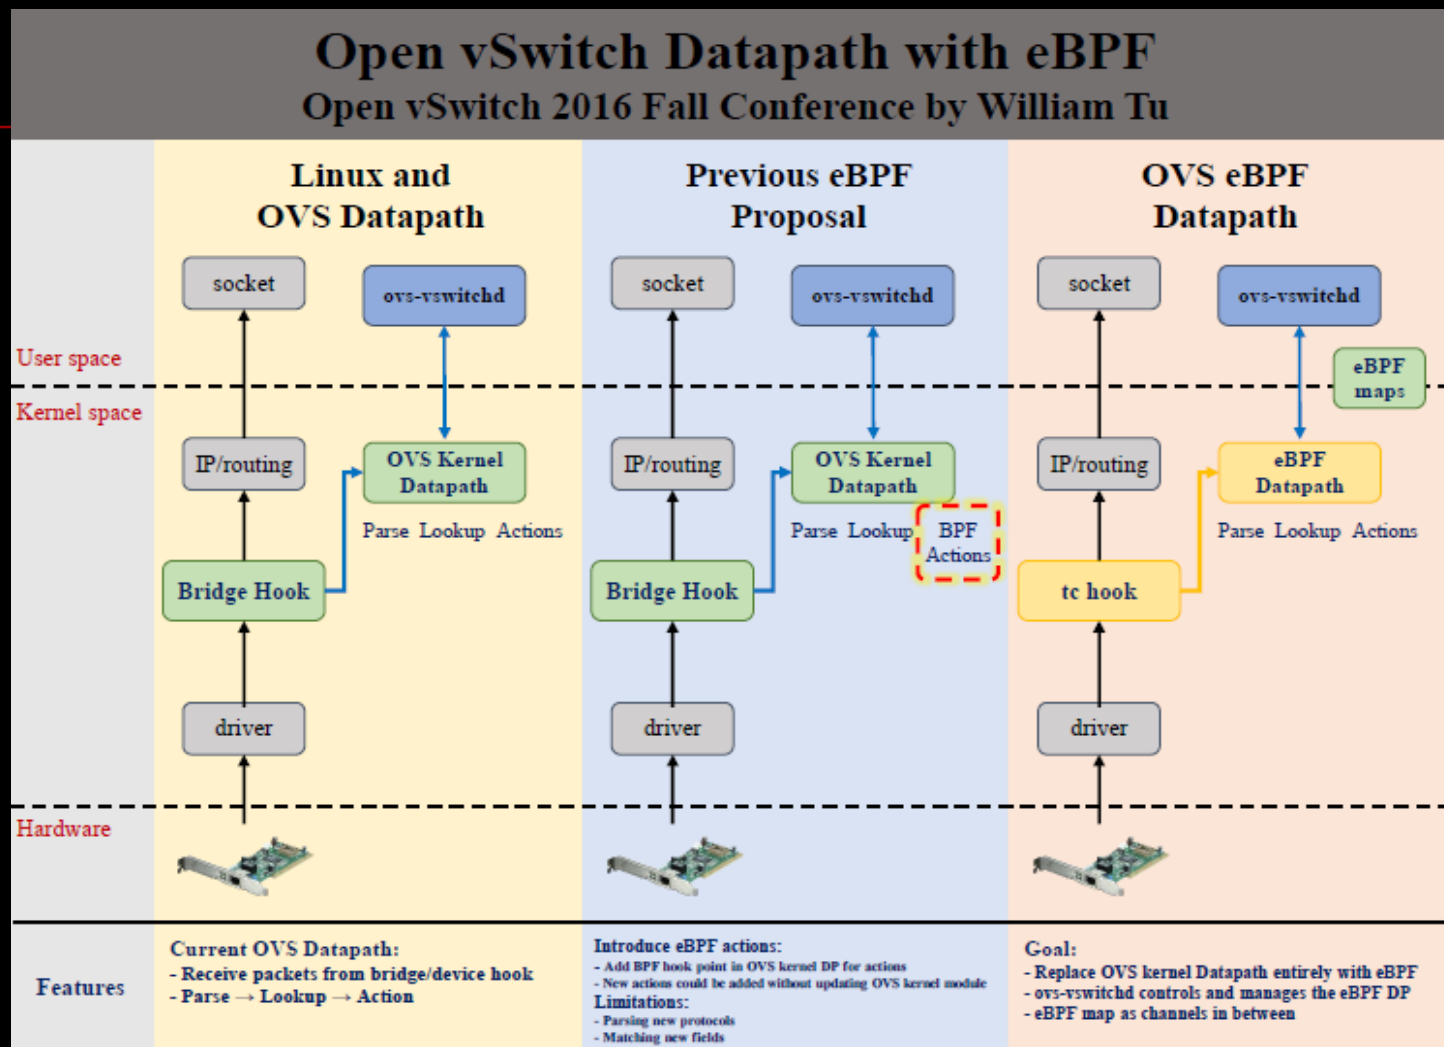
patch

```
option(ENABLE_CLANG_JIT "Enable Loading BPF through Clang Frontend" ON)
-option(ENABLE_USDT "Enable User-level Statically Defined Tracing" ON)
+option(ENABLE_USDT "Enable User-level Statically Defined Tracing" OFF)
CMAKE_DEPENDENT_OPTION(ENABLE_CPP_API "Enable C++ API" ON "ENABLE_USDT" OFF)
```

- **build BCC(master branch) by gcc 7.2.0 + jemalloc 5.0.1 + ld.gold ~ 43 minutes**

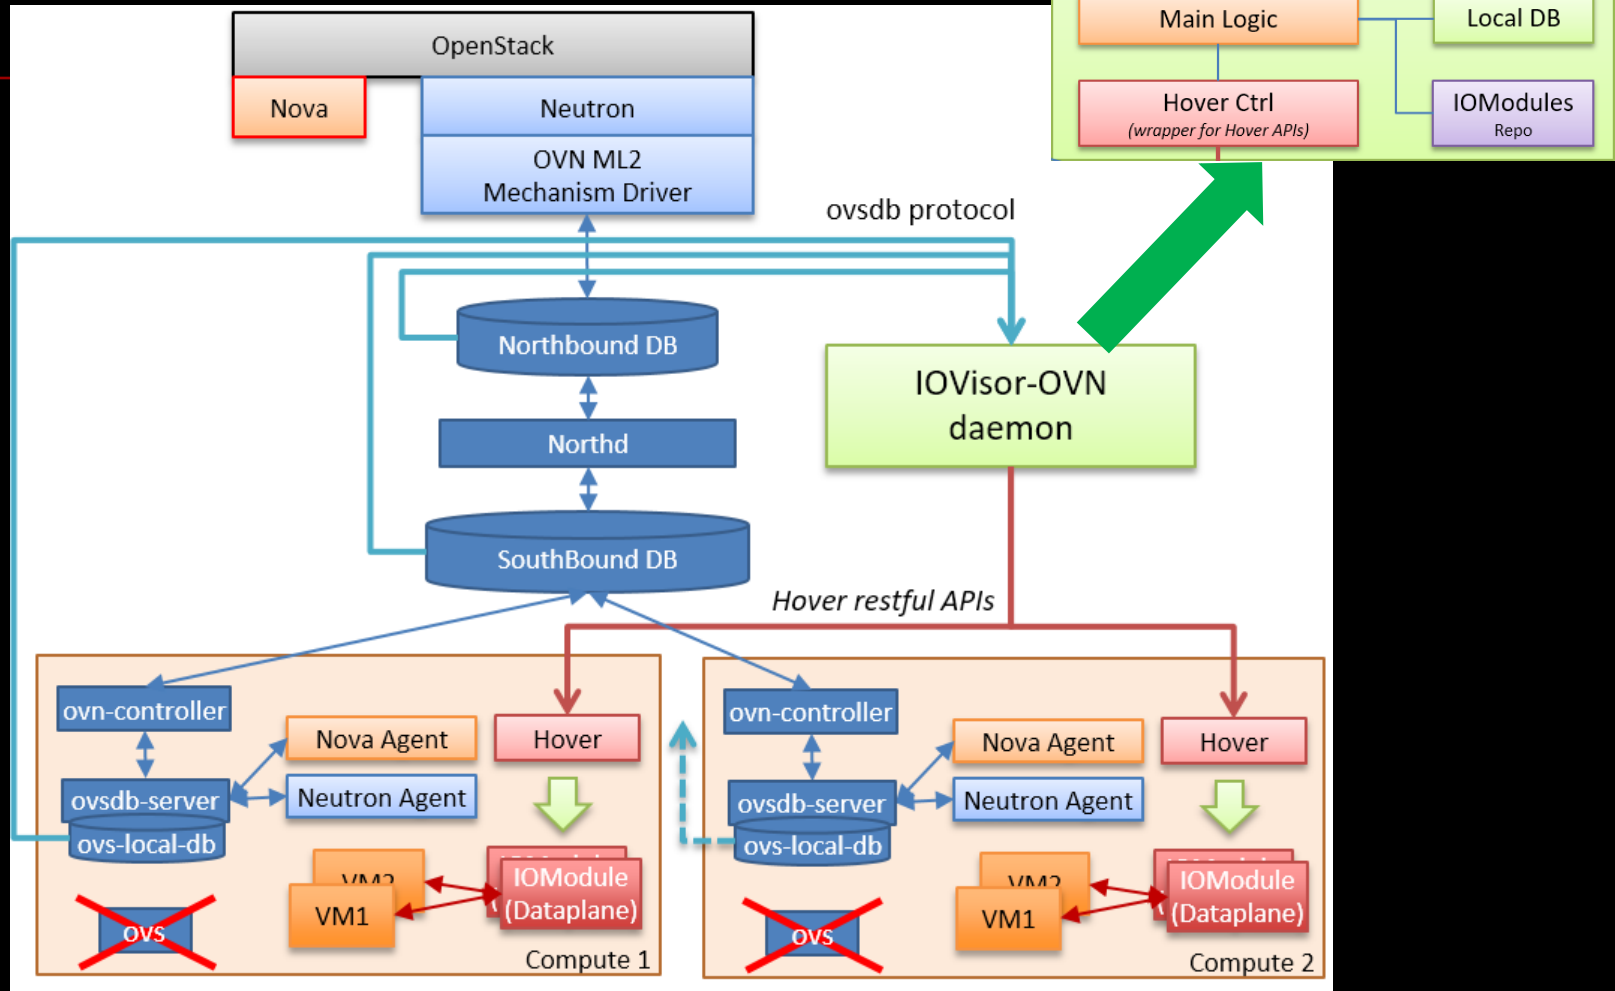
OVS

■ <http://openvswitch.org/>



iovisor-ovn

■ <https://github.com/iovisor/iovisor-ovn>



V. Wrap-up

■ A wide range of applications



Those who have publically stated they are using BPF or are planning to use BPF include

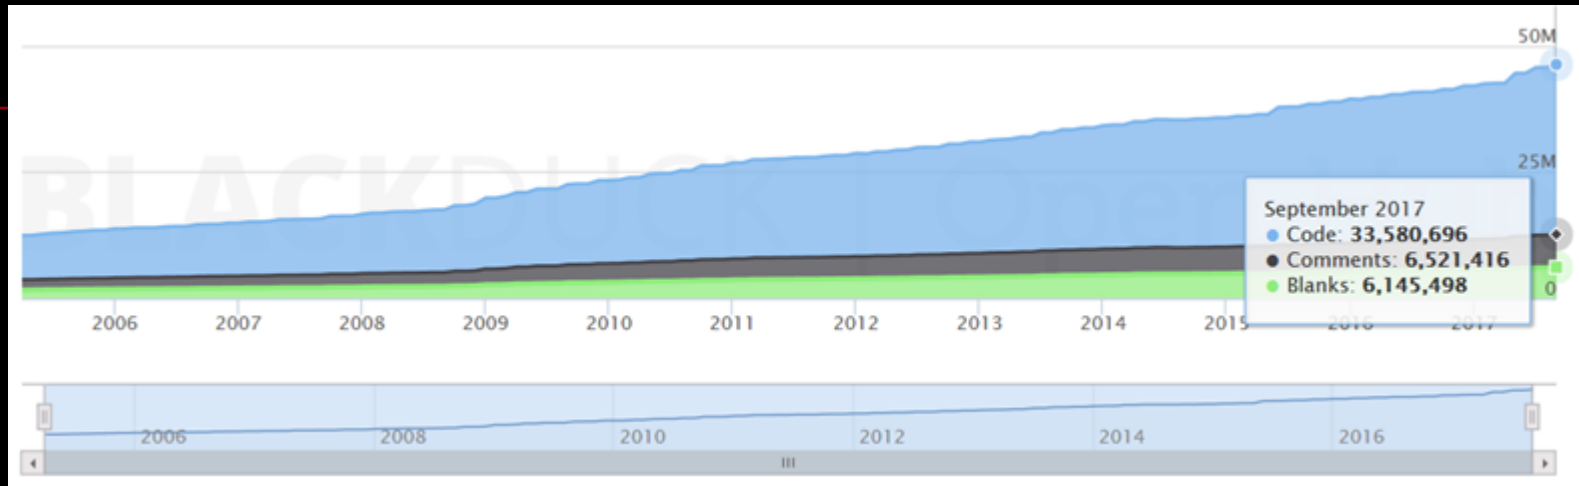
- Facebook-Load Balancing, Security
- Netflix-Network Monitoring
- Cilium Project
- Cloudflare-Security
- OVS-Virtual Switching

...

Source: <https://www.slideshare.net/Open-NFP/transparent-ebpf-offload-playing-nice-with-the-linux-kernel>

■ Polyglot VM

Changing the way you think about Linux Kernel development:



Source: https://www.openhub.net/p/linux/analyses/latest/languages_summary

■ User space/Kernel space Repartition & Unifying

eBPF is sure to play an important role in tomorrow's Linux!

Q & A

Thanks!



Reference

Slides/materials from many and varied sources:

- <http://en.wikipedia.org/wiki/>
- <http://www.slideshare.net/>
- <https://www.kernel.org/doc/Documentation/>
- <http://man7.org/linux/man-pages/man2/bpf.2.html>
- <https://www.python.org>
- <http://llvm.org>
- https://en.wikipedia.org/wiki/Just-in-time_compilation
- <http://dpdk.org/>
- <https://www.netbsd.org/gallery/presentations/>
- <https://www.opennetworking.org/>
- <https://www.opnfv.org/>
- ...