

Xen Summit

Nanjing 2018

Qubes in Action

李枫

hkli2013@126.com

Jun 20, 2018

Agenda

I. Overall Design

- Overview
- 4.0
- Security

II. Acceleration for Python-based ToolStack

- Overview
- SaltStack
- Python Runtimes
- GraalVM
- GraalPython

III. Future Evolution and Re-design


- Official
- Re-designing, Re-engineering, Re-inventing

I. Overall Design

1) Overview

■ <https://www.qubes-os.org/>



"If you're serious about security, Qubes OS is the best OS available today. It's what I use, and free." 

— Edward Snowden, *whistleblower and privacy advocate*

What's Inside of Qubes



Use Fedora, Debian, or even Windows



with  networking

SECURE COMPARTMENTALIZATION

Qubes brings to your personal computer the security of the Xen hypervisor, the same software relied on by many major hosting providers to isolate websites and services from each other. [Learn more](#)

OPERATING SYSTEM FREEDOM

Can't decide which Linux distribution you prefer? Still need that one Windows program for work? With Qubes, you're not limited to just one OS. [Learn more](#)

SERIOUS PRIVACY

With Whonix integrated into Qubes, using the Internet anonymously over the Tor network is safe and easy. [Learn more](#)



Joanna Rutkowska

rootkovska

Qubes OS and Invisible Things Lab

HCL

- <https://www.qubes-os.org/hcl/>
<https://www.qubes-os.org/doc/system-requirements/>

Qubes Release 4.x

Minimum

- 64-bit Intel or AMD processor (x86_64 aka x64 aka AMD64)
- Intel VT-x with EPT or AMD-V with RVI
- Intel VT-d or AMD-Vi (aka AMD IOMMU)
- 4 GB RAM
- 32 GB disk space

Qubes Release 3.x

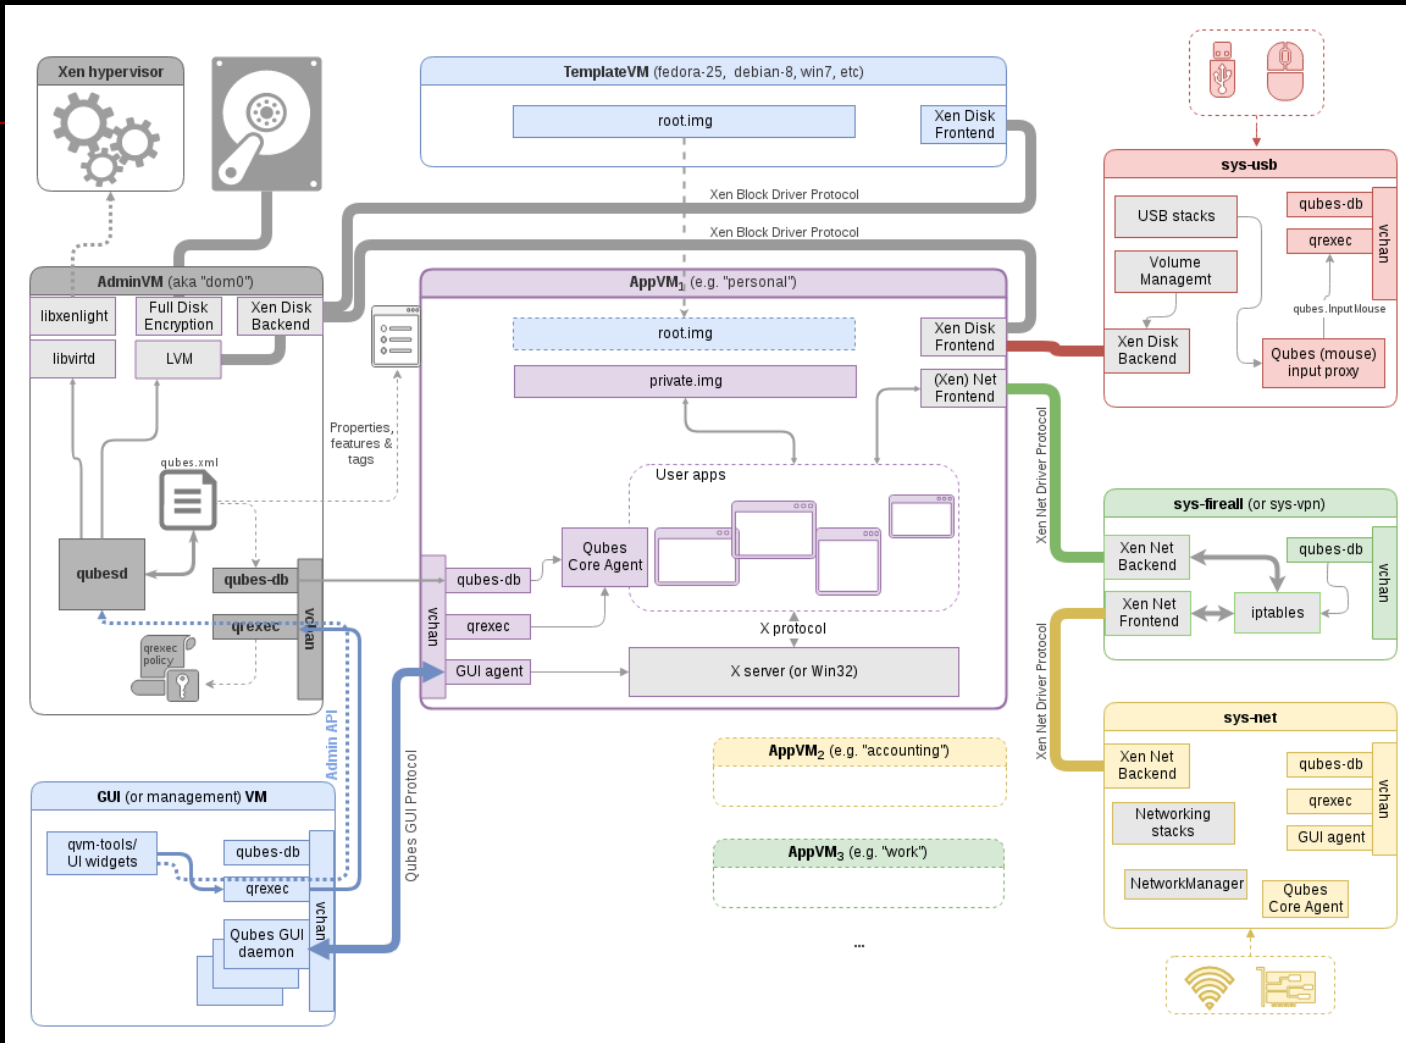
Minimum

- 64-bit Intel or AMD processor (x86_64 aka x64 aka AMD64)
- 4 GB RAM
- 32 GB disk space
- Legacy boot mode (required for R3.0 and earlier; UEFI is supported beginning with R3.1)



Core Stack

■ <https://www.qubes-os.org/news/2017/10/03/core3/>



2) 4.0

- <https://www.qubes-os.org/news/2018/03/28/qubes-40/>
- <https://www.qubes-os.org/doc/releases/4.0/release-notes/>

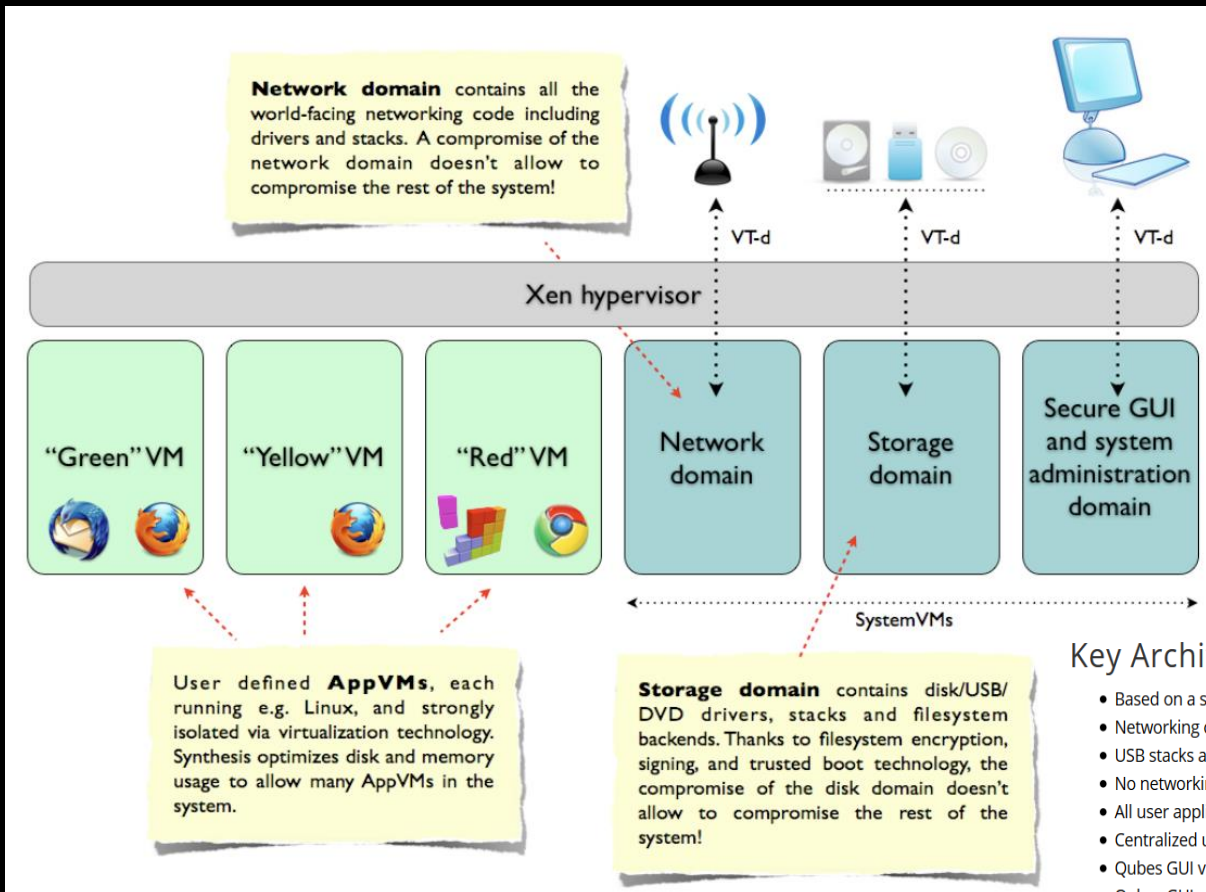
New features since 3.2

- Core management scripts rewrite with better structure and extensibility, [API documentation](#)
- [Admin API](#) allowing strictly controlled managing from non-dom0
- All `qvm-*` command-line tools rewritten, some options have changed
- Renaming VM directly is prohibited, there is GUI to clone under new name and remove old VM
- Use [PVH](#) and [HVM](#) by default to [mitigate Meltdown & Spectre](#) and lower the [attack surface on Xen](#)
- Create USB VM by default
- [Multiple Disposable VMs templates support](#)
- New [backup format](#) using scrypt key-derivation function
- Non-encrypted backups no longer supported
- [split VM packages](#), for better support minimal, specialized templates
- [Qubes Manager decomposition](#) - domains and devices widgets instead of full Qubes Manager; devices widget support also USB
- [More flexible firewall interface](#) for ease unikernel integration
- Template VMs do not have network interface by default, [qrexec-based updates proxy](#) is used instead
- More flexible IP addressing for VMs - [custom IP](#), [hidden from the IP](#)
- More flexible Qubes RPC policy - [related ticket](#), [documentation](#)
- [New Qubes RPC confirmation window](#), including option to specify destination VM
- [New storage subsystem design](#)
- Dom0 update to Fedora 25 for better hardware support
- Kernel 4.9.x



3) Security

■ <https://www.qubes-os.org/security>
Virtualization




Key Architecture features

- Based on a secure bare-metal hypervisor (Xen)
- Networking code sand-boxed in an unprivileged VM (using IOMMU/VT-d)
- USB stacks and drivers sand-boxed in an unprivileged VM (currently experimental feature)
- No networking code in the privileged domain (dom0)
- All user applications run in "AppVMs," lightweight VMs based on Linux
- Centralized updates of all AppVMs based on the same template
- Qubes GUI virtualization presents applications as if they were running locally
- Qubes GUI provides isolation between apps sharing the same desktop
- Secure system boot based (optional)

Source: <https://www.qubes-os.org/doc/architecture/>

Compartmentalization

- defects of monolithic systems
- separate your digital life into security domains



Source: <https://www.qubes-os.org/doc/architecture/>

■ Trust Level

- Untrusted appVM, red
 - Personal appVM, yellow
 - Work appVM, green
 - Vault appVM, black
-

...

■ VM Types

AppVM

ServiceVM

TemplateVM

DispVM

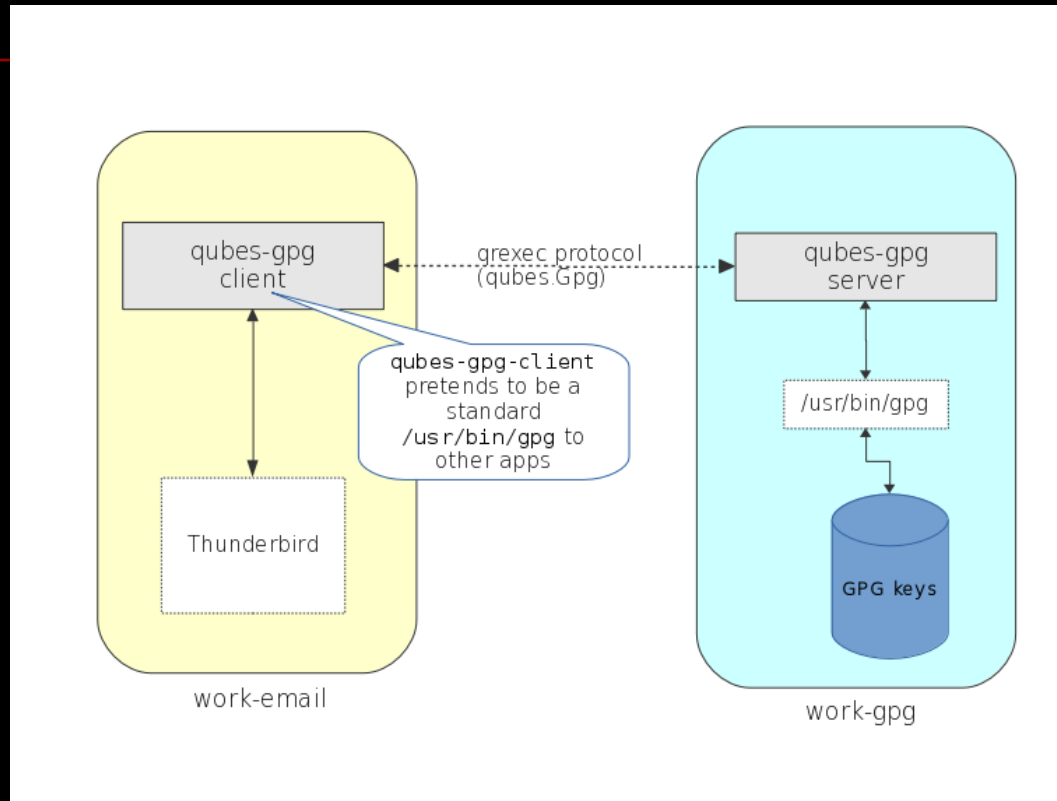
netVM proxyVM

...

■ VM Manager

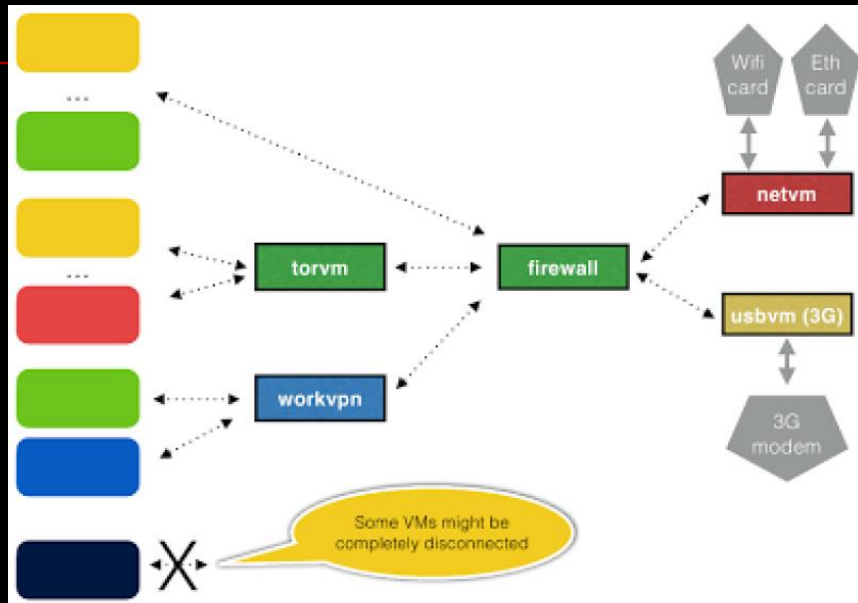
Split GPG

- <https://www.qubes-os.org/doc/split-gpg/>
- **Poor Man's Hardware Security Module (pmHSM)**



netVM

- <https://www.qubes-os.org/doc/networking/>
- <https://www.qubes-os.org/doc/firewall/>




Source: https://www.qubes-os.org/attachment/wiki/slides/RMLL_2016_Improving-client-systems-security.pdf

usbVM

- <https://www.qubes-os.org/doc/usb/#security-warning-about-usb-input-devices>
- **sys-usb**
- **qubes-usb-proxy**

Good Resource

- <https://www.qubes-os.org/doc/>
- <https://www.pearson.com/us/higher-education/program/Rankin-Linux-Hardening-in-Hostile-Networks-Server-Security-from-TLS-to-Tor/PGM137619.html>



Overview	Features	Contents	Resources
Chapter 2: Workstation Security 25			
Section 1: Security Fundamentals 25			
Section 2: Additional Workstation Hardening 33			
Section 3: Qubes 37			
Summary 52			
Chapter 3: Server Security 53			
Section 1: Server Security Fundamentals 53			
Section 2: Intermediate Server-Hardening Techniques 58			
Section 3: Advanced Server-Hardening Techniques 68			
Summary 74			
Chapter 4: Network 75			
Section 1: Essential Network Hardening 76			
Section 2: Encrypted Networks 87			
Section 3: Anonymous Networks 100			
Summary 107			

II. Acceleration for Python-based ToolStack

1) Overview







Why Python

■ <https://www.tiobe.com/tiobe-index/>


May 2018	May 2017	Change	Programming Language	Ratings	Change
1	1		Java	16.380%	+1.74%
2	2		C	14.000%	+7.00%
3	3		C++	7.668%	+2.92%
4	4		Python	5.192%	+1.64%
5	5		C#	4.402%	+0.95%

■ <http://pypl.github.io/PYPL.html>

■ <https://spectrum.ieee.org/computing/software/the-2017-top-programming-languages>

Language Rank	Types	Spectrum Ranking
1. Python		100.0
2. C		99.7
3. Java		99.5
4. C++		97.1
5. C#		87.7
6. R		87.7
7. JavaScript		85.6
8. PHP		81.2
9. Go		75.1
10. Swift		73.7

■ <https://www.kdnuggets.com/2018/05/poll-tools-analytics-data-science-machine-learning-results.html>



■ Other Python projects

Build: Meson, SCons... **DevOps:** Ansible, SaltStack...

Web: Django, web2py, Flask, Tornado, TurboGears...

AI: PyTorch, Theano... **Big Data:** PyData, PySpark...

Science: Scipy, Sage...

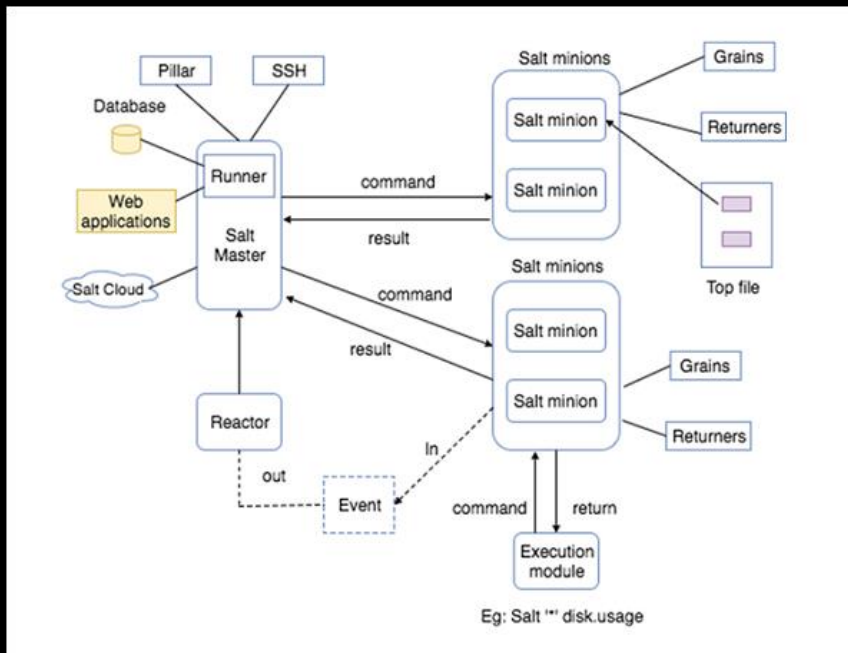
Cloud/DataCenter: OpenStack

Security: a swiss knife for hackers...

...

2) SaltStack

- <https://saltstack.com/>
Intelligent automation for a software-defined world
- <https://www.qubes-os.org/doc/salt/>
- default management engine in dom0 since Qubes 3.1
- **qubesctl** is inter-changeable and an alias for salt-call
- <https://docs.saltstack.com/en/latest/topics/virt/index.html>



Source: https://www.tutorialspoint.com/saltstack/saltstack_architecture.htm

3) Python Runtimes

Why Python is Slow

- dynamically typed
- no JIT support in the official CPython
- GIL (Global Interpreter Lock)

Python 3 programs versus Java

vs C vs C++ vs Go vs Java

by benchmark task performance

pidigits

source	secs	mem	gz	cpu	cpu load
Python 3	3.43	12,716	386	3.43	100% 1% 1% 0%
Java	3.13	36,984	938	3.36	4% 4% 99% 3%

regex-redux

source	secs	mem	gz	cpu	cpu load
Python 3	15.22	447,324	512	27.44	25% 33% 32% 91%
Java	10.51	573,972	929	31.30	70% 73% 70% 86%

n-body

source	secs	mem	gz	cpu	cpu load
Python 3	838.39	10,324	1196	838.20	95% 1% 5% 0%
Java	22.17	33,040	1489	22.27	100% 1% 0% 1%

spectral-norm

source	secs	mem	gz	cpu	cpu load
Python 3	180.97	15,876	443	720.51	100% 100% 100% 100%
Java	4.38	35,388	950	16.80	96% 96% 95% 97%

reverse-complement

source	secs	mem	gz	cpu	cpu load
Python 3	18.79	1,008,868	814	19.73	9% 69% 35% 30%
Java	3.15	680,424	2183	7.07	52% 70% 43% 63%

k-nucleotide

source	secs	mem	gz	cpu	cpu load
Python 3	77.65	182,700	1967	302.86	97% 99% 97% 98%
Java	8.75	385,056	1812	27.09	85% 72% 70% 85%

binary-trees

source	secs	mem	gz	cpu	cpu load
Python 3	93.55	280,624	589	337.74	92% 89% 87% 93%
Java	8.39	933,808	835	28.28	82% 86% 84% 88%

fasta

source	secs	mem	gz	cpu	cpu load
Python 3	59.47	15,996	1947	138.97	55% 55% 63% 66%
Java	2.27	43,628	2473	5.93	51% 75% 57% 81%

fannkuch-redux

source	secs	mem	gz	cpu	cpu load
Python 3	565.97	15,528	950	2,172.63	95% 94% 95% 100%
Java	18.27	31,820	1282	72.06	99% 99% 98% 98%

mandelbrot

source	secs	mem	gz	cpu	cpu load
Python 3	225.24	15,736	688	899.25	100% 100% 100% 100%
Java	6.10	76,520	796	23.59	97% 98% 98% 96%

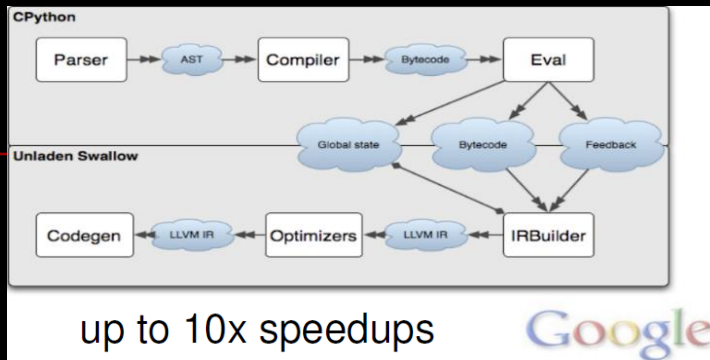
Python 3 Python 3.6.3

Java java 10 2018-03-20
Java(TM) SE Runtime Environment 18.3 (build 10+46)
Java HotSpot(TM) 64-Bit Server VM 18.3 (build 10+46, mixed mode)


<https://benchmarksgame-team.pages.debian.net/benchmarksgame/faster/python.html>

Runtimes

■ LLVM-based (VMKit, MCJIT...)




PySton



RPython Meta-tracing


...



Source: <http://speed.pypy.org/>

4) GraalVM

- <https://www.graalvm.org/>
- <http://www.oracle.com/technetwork/oracle-labs/program-languages/overview/index.html>
- <https://blogs.oracle.com/developers/announcing-graalvm>



Oracle Labs GraalVM

Thank you for downloading this release of the Oracle Labs GraalVM. With this release, one can execute Java applications with Graal, as well as applications written in JavaScript, Ruby, R, and Python, with our Polyglot language engines.

You must accept the OTN License Agreement to download this software.
☒ Accept License Agreement | ☐ Decline License Agreement

⚠ GraalVM based on JDK8, preview for Linux (1.0.0 RC2)
⚠ GraalVM based on JDK8, preview for macOS (1.0.0 RC2)

About this OTN Release
Oracle Labs GraalVM is a research artifact from Oracle Labs, whereas the current OTN release is a technology preview version of it. Henceforth, this release is intended for information purpose only, and may not be incorporated into any contract. This is not a commitment to deliver any material, code, or functionality to Oracle products, and thus should not be relied upon in making any purchase decisions. The development, release and timing of any features or functionality described for products of Oracle remains at the sole discretion of Oracle.

WARNING: This release contains older versions of the JRE and JDK that are provided to help developers debug issues in older systems. They are not updated with the latest security patches and are not recommended for use in production.

JVMCI JDK Downloads
To develop the Graal compiler, you need to accept the license above and download one of the JVMCI enabled JDK 8 binaries below:

⚠ labsjdk-8u172-jvmci-0.44-darwin-amd64.tar.gz
⚠ labsjdk-8u172-jvmci-0.44-solaris-sparcv9.tar.gz
⚠ labsjdk-8u172-jvmci-0.44-linux-amd64.tar.gz

Debug builds of the above JVMCI enabled JDK 8 binaries can be useful when diagnosing VM crashes. These binaries are provided below:

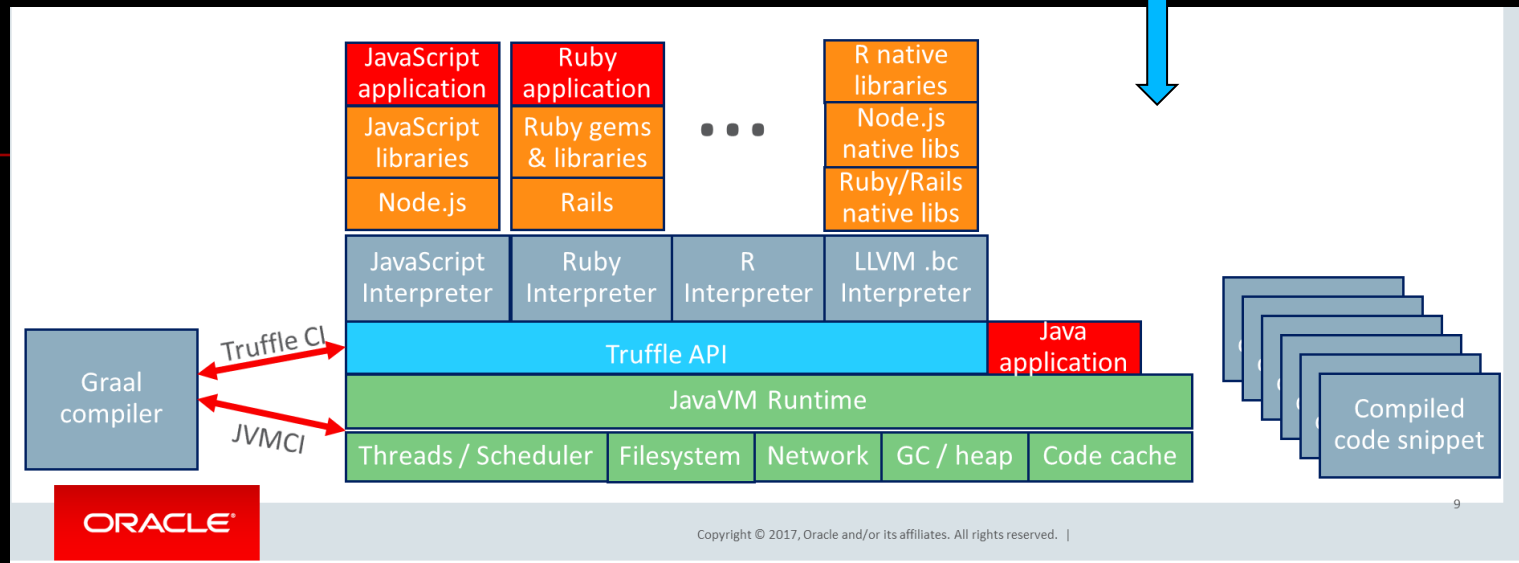
⚠ labsjdk-8u172-jvmci-0.44-fastdebug-darwin-amd64.tar.gz
⚠ labsjdk-8u172-jvmci-0.44-solaris-sparcv9.tar.gz
⚠ labsjdk-8u172-jvmci-0.44-fastdebug-linux-amd64.tar.gz

- **High-Performance Polyglot VM**
- **A meta-runtime for Language-Level Virtualization**
- **Currently based on Oracle Labs JDK 8 with JVMCI support**
- [http://openjdk.java.net/jeps/243\(JVMCI\)](http://openjdk.java.net/jeps/243(JVMCI)): **experimental in JDK 9**


Arch

■ A hybrid of static & dynamic runtimes

Substrate VM




Source: <https://ics.psu.edu/wp-content/uploads/2017/02/GraalVM-PSU.pptx>




Source: <https://www.slideshare.net/jyukuty/jvmgraalopenj9>

Performance



Source: http://lafo.ssw.uni-linz.ac.at/papers/2017_PLDI_GraalTutorial.pdf



Source: <http://dbpl2017.org/slides/DBPL-2017-s2.pdf>

■ but for GraalVM 1.0.0 RC1

	GRAALVM	ORACLE JDK 8	ORACLE JDK 9
AVERAGE OPS/S	6.795 ±(99.9%) 0.016	6.727 ±(99.9%) 0.017	7,136 ±(99.9%) 0,026
MIN	6.477	6.466	6,464
MAX	6.967	6.899	7,443
STD DEV	0.068	0.070	0,111
CI (99.9%) (ASSUMES NORMAL DISTRIBUTION)	[6.778, 6.811]	[6.710, 6.743]	[7,110, 7,162]

Source: <https://blog.frankel.ch/first-impressions-graalvm>

■ still have plenty of room for improvement!

5) GraalPython

Graal/Truffle-based implementation of Python

GraalVM provides an early-stage experimental implementation of Python. A primary goal is to support SciPy and its constituent libraries. This Python implementation currently aims to be compatible with Python 3.7, but it is a long way from there, and it is very likely that any Python program that requires any imports at all will hit something unsupported. At this point, the Python implementation is made available for experimentation and curious end-users.

- <https://github.com/graalvm/graalpython>
- <https://www.graalvm.org/docs/reference-manual/languages/python/>

	Java 10.0.1	CPython 3.6.5	GraalPython ee-1.0.0-rc2
n-body	9.676s	11m56.642s	15m57.543s

Test on Dell XPS 15z: i5-2410M@2.3Ghz, 6G RAM, Fedora 28 for X64 with Kernel 4.16.14

```
[mydev@myfedora Python]$ graalpython -V
Graal Python 3.7.0 (GraalVM 1.0.0-rc2)
[mydev@myfedora Python]$
[mydev@myfedora Python]$ graalpython knucleotide.py 0 < knucleotide-input1000.txt
Please note: This Python implementation is in the very early stages, and can run little more than basic benchmarks at this point.
Traceback (most recent call last):
  File "knucleotide.py", line 20, in <module>
    b'from os import cpu_count'
ImportError: cannot import name 'cpu_count'
```




ZipPy



ZipPy is a fast and lightweight Python 3 implementation built using the Truffle framework. ZipPy leverages the underlying Java JIT compiler and compiles Python programs to highly optimized machine code at runtime. [Repository on Bitbucket.](#)

- <http://thezhangwei.com/>
- <https://github.com/secreystems-lab/zippy>
- **Optimizations**
 - Numeric Types, Type Specializations, Efficient Data Representation
 - Control Flow Specializations, **Generator Peeling**, Optimizing Object Model and Calls

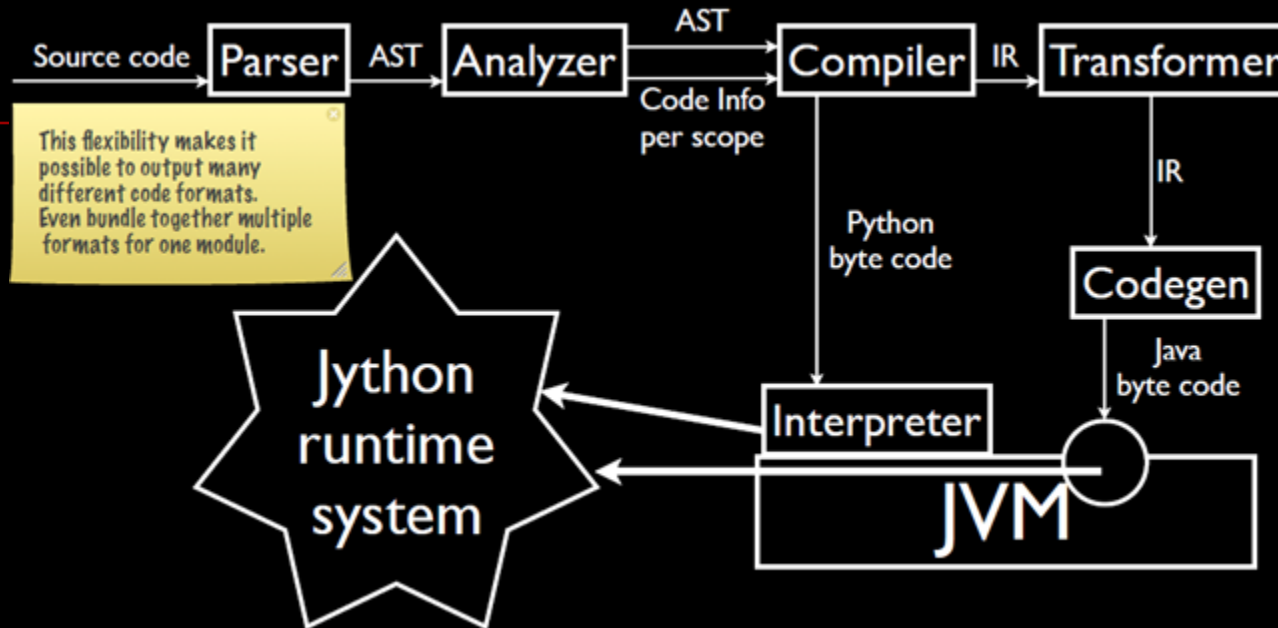
benchmark	CPython3	CPython	Jython	PyPy	PyPy3	ZipPy
binarytrees	1.00	0.94	1.99	2.60	2.70	7.31
fannkuchredux	1.00	0.97	0.51	44.53	47.29	87.50
fasta	1.00	1.04	1.55	11.73	11.24	15.57
mandelbrot	1.00	1.08	0.34	10.91	10.82	11.69
meteor	1.00	1.02	0.77	2.64	2.62	2.13
nbody	1.00	0.97	0.73	12.13	12.06	6.17
pidigits	1.00	1.00	0.62	0.98	0.95	0.60
spectralnorm	1.00	1.33	1.89	127.33	127.25	128.10
float	1.00	0.95	1.05	8.64	8.67	17.71
richards	1.00	0.94	1.21	29.53	29.25	50.13
chaos	1.00	1.17	1.55	40.88	25.69	68.28
deltablue	1.00	0.85	1.33	30.08	29.14	23.46
go	1.00	1.08	1.99	6.79	6.66	15.41
mean	1.00	1.02	1.05	12.15	11.68	15.34



Jython

■ <http://www.jython.org>

//No new release since 2015...



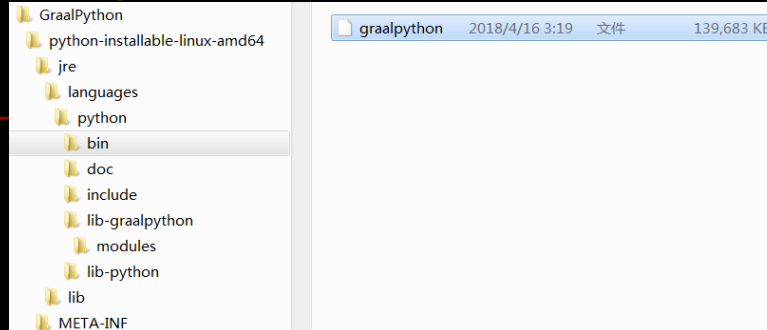
VOC

■ <https://github.com/pybee/voc/>

■ A transpiler that converts Python code into Java bytecode

Integration

- <https://github.com/graalvm/graalpython/releases/download/vm-1.0.0-rc2/python-installable-linux-amd64.jar>



■ GraalVM EE 1.0.0 RC2

graalvm-ee-1.0.0-rc2

```
3rd_party_licenses_graalpython.txt ->
3rd_party_licenses.txt
bin
COPYRIGHT
db
GRAALVM-README.md
include
javafx-src.zip
jre
lib
LICENSE
LICENSE_GRAALPYTHON -> jre/languages/
man
README.html
release
src.zip
THIRDPARTYLICENSEREADME-JAVAFX.txt
THIRDPARTYLICENSEREADME.txt
```

```
jre
bin
COPYRIGHT
languages
lib
LICENSE
plugin
README
THIRDPARTYLIC
THIRDPARTYLIC
tools
Welcome.html
```

```
bin
ControlPanel
gemasrv
graalpython -> ../languages/python/bin/graalpython
gu
java
javaws
jcontrol
jjs
js
keytool
lli
native-image
node
npm
orbd
pack200
policytool
polyglot
rmid
rmiregistry
servertool
tnameserv
unpack200
```

```
languages
js
asm-debug-all.jar
bin
graaljs.jar
icu4j
icu4j.jar
include
native-image.properties
NODE_README.md
npm
README.md
trufflenode.jar
llvm
bin
libsulong.bc
libsulong.so
native-image.properties
polyglot.h
README.md
sulong.jar
python
3rd_party_licenses_graalpython.txt
bin
doc
graalpython.jar
include
lib-graalpython
lib-python
LICENSE_GRAALPYTHON
native-image.properties
README_GRAALPYTHON.md
release
```

Practice

- <https://github.com/AdoptOpenJDK/openjdk-jdk> //OpenJDK11 src
- export **JDK_BOOT_DIR**=\$YOUR_OPENJDK10_HOME
- reserve at least 6GB disk space
- ~~on Laptop with Fedora 28 + Kernel 4.16.15 + GCC 8.1.1-1 + 8GB Memory (6GB DDR4 + 2GB Swap)~~
- cd \$YOUR_OPENJDK11_SRCHOME and run the commands:
bash configure **--disable-warnings-as-errors**
make JOBS=4 images

#build GraalPython & GraalVM

- setup mx
- patching for avoid **javaCompliance** limitation
- Fail to build GraalVM with previously built OpenJDK 11, something wrong in **javac**?

```
Compiling com.oracle.truffle.llvm.runtime with javac-daemon(JDK 11) failed
Shutting down
Shutting down
File "/opt/MyWorkSpace/DevSW/Tools/Build/mx/mx.py", line 17693, in <module>
    main()
File "/opt/MyWorkSpace/DevSW/Tools/Build/mx/mx.py", line 17674, in main
    retcode = c(command_args)
File "/opt/MyWorkSpace/DevSW/Tools/Build/mx/mx.py", line 11725, in build
    abort('{0} build tasks failed'.format(len(failed)))
File "/opt/MyWorkSpace/DevSW/Tools/Build/mx/mx.py", line 11251, in abort
    traceback.print_stack()
1 build tasks failed
```

- **Successfully build GraalPython via JDK 10**

challenges

- **prone to break build**
 - **deal with JDK, Truffle/Graal, LLVM...**
 - **customize GraalPython to meet our need**
 - **dynamically enable or reload Graal compiler at runtime**
 - **...**
-

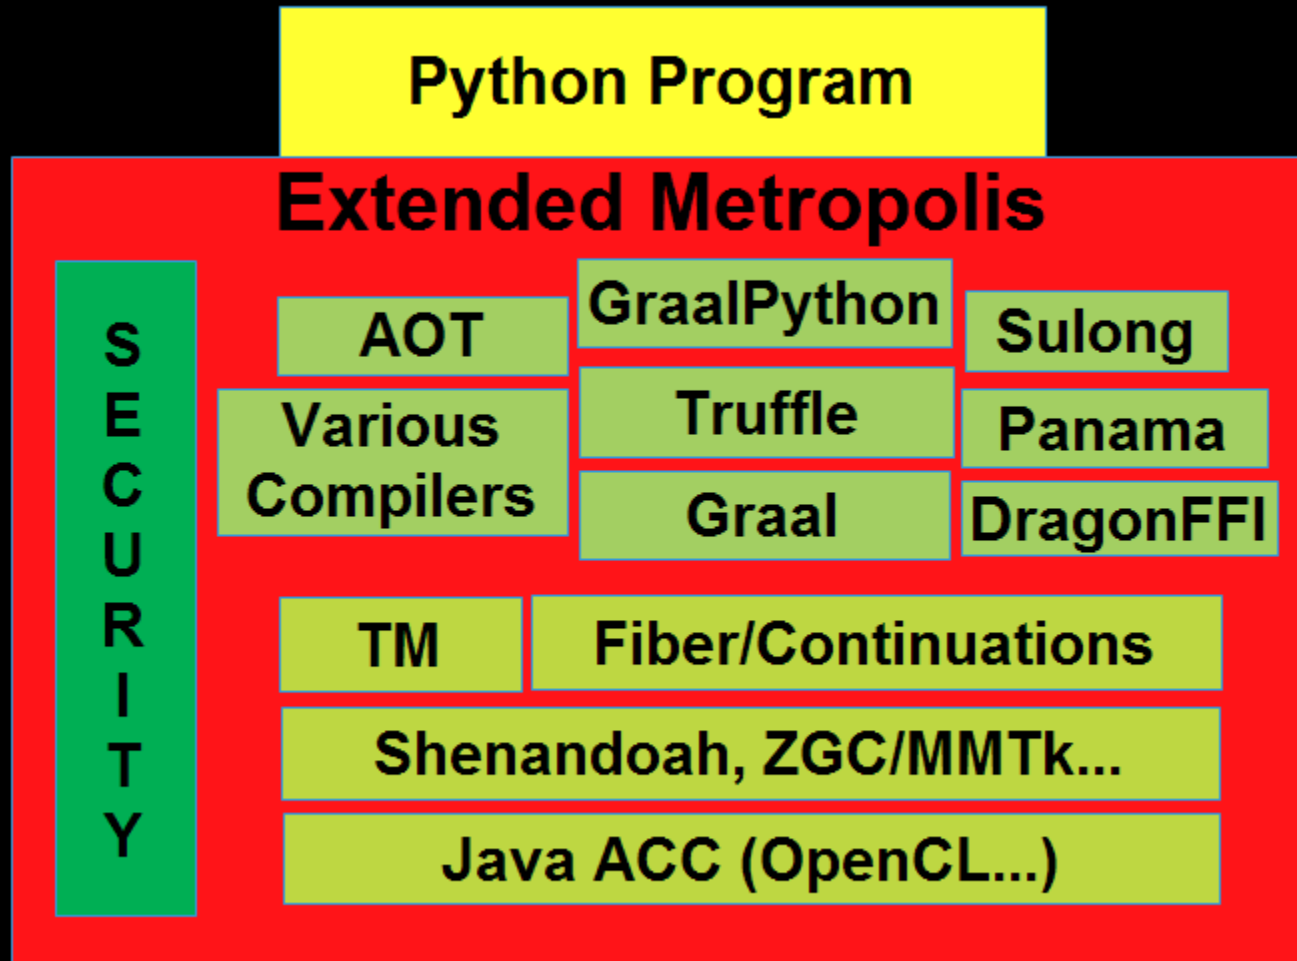
Rethinking of Python Runtime

- from my point of view, various Runtime Frameworks for Python implementation:

	OMR	LLVM	PyPy	GraalVM
Pros	easily leverage new hardware features low-maturity	high efficiency; high-maturity	productivity(RPython); high-maturity	combine continually improved JVM and LLVM techs; productivity(Java);
Cons	productivity (C++/C)?	death of VMKit...	mainly for dynamic language; PyPy3	low-maturity; memory footprint
Performance	experimental/not sure	not enough	not enough	not enough
Native		DragonFFI	CFFI, CPPYY	GNFI (Gaal Native Function Interface)
Related Projects	JBM J9/OpenJ9	Unladen Swallow, PySton	Psyco	ZipPy
License	EPL v2.0	LLVM	MIT	GPL v2

Future

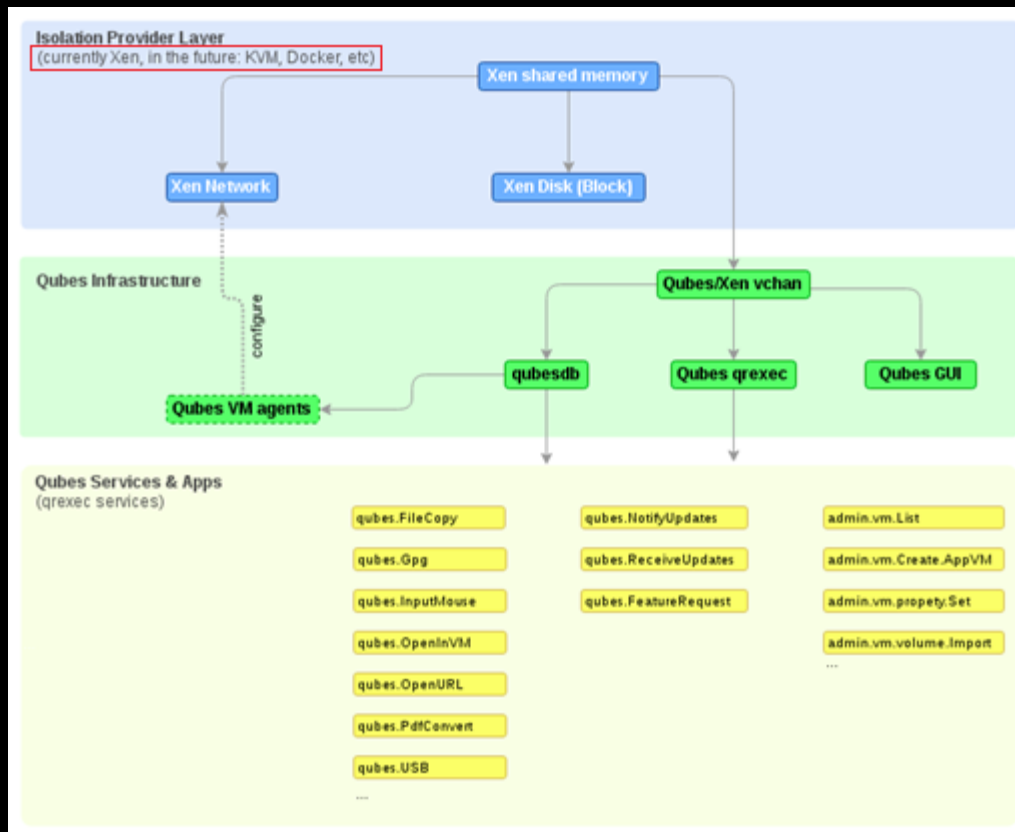
- <http://openjdk.java.net/projects/metropolis/>
- extend **Project Metropolis** and customize it for Python



III. Future Evolution & Re-design

1) Official Generalization

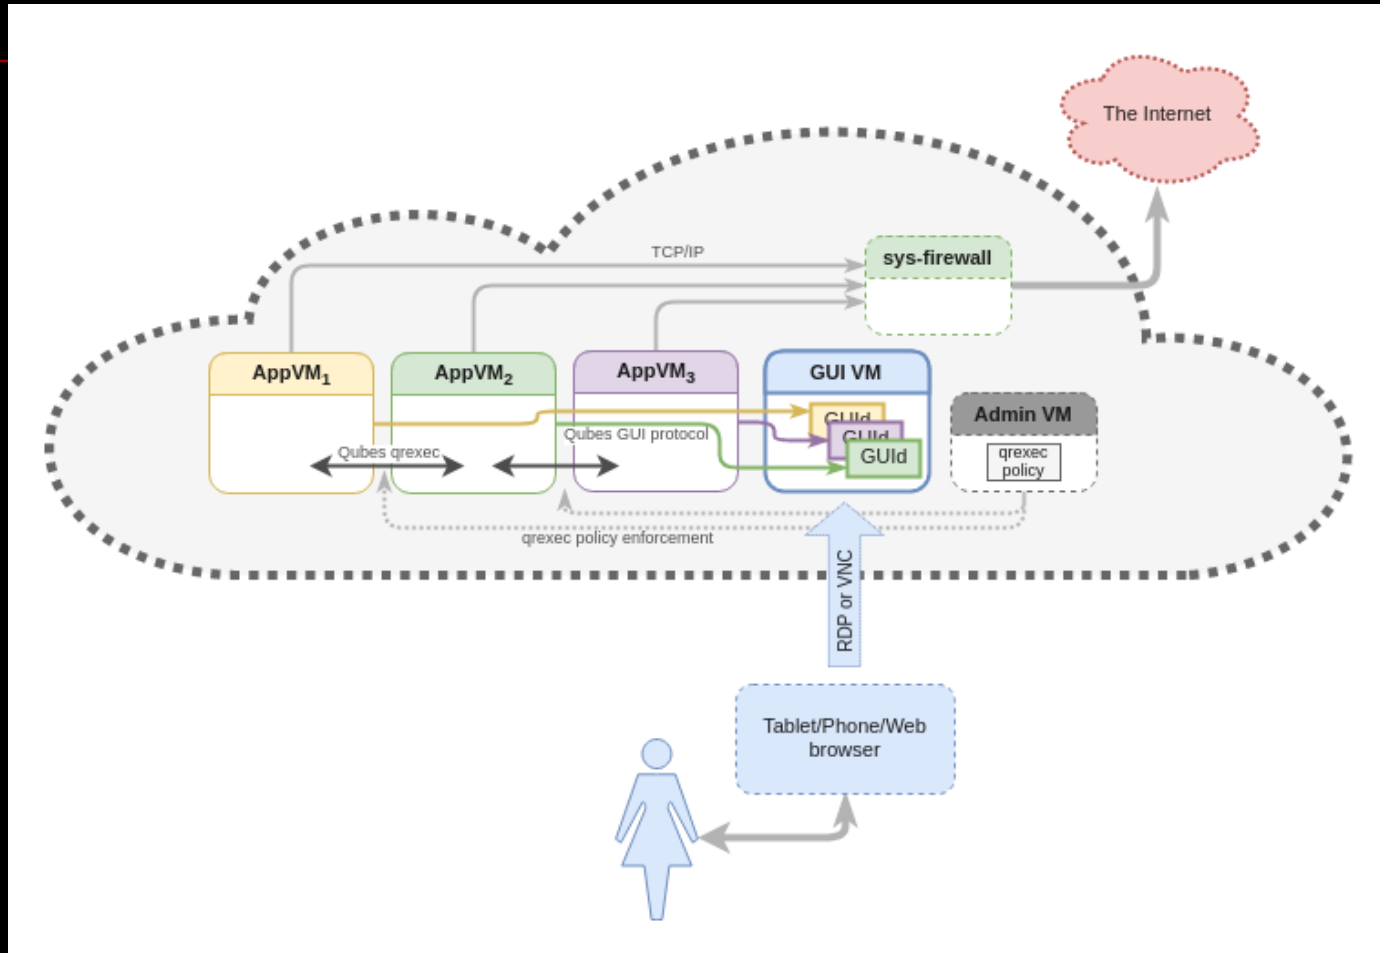
■ Generalizing the Qubes Architecture



Qubes Remote Execution (qrexec)

Qubes Air

- <https://www.qubes-os.org/news/2018/01/22/qubes-air/>
- **Qubes in the Cloud**



2) Re-designing, Re-engineering, Re-inventing



- A customized Linux distribution for Dom0 with various optimization
- Integrate GraalVM-based customized runtime for accelerating Toolstack & Applications in Dom0/DomU
- Support Wayland display server in TemplateVM
- ...

Q & A

Thanks!



Reference

Slides/materials from many and varied sources:

- <http://en.wikipedia.org/wiki/>
- https://en.wikipedia.org/wiki/Qubes_OS
- <https://www.python.org>
- <http://llvm.org>
- <https://en.wikipedia.org/wiki/CPython>
- https://en.wikipedia.org/wiki/Just-in-time_compilation
- <https://github.com/dropbox/pyston>
- ...