



PyCon China Hangzhou 2019

Python for Linux Kernel Debugging

Feng Li (李枫)

hkli2013@126.com

Oct 19, 2019



Agenda

I. Linux Kernel Debugging

- Overview
 - eBPF
 - Development
-

II. BCC

- Overview
- Development

III. LISA

- Overview
- Internals

IV. Drgn

- Overview

V. Practice on ARM

- Development Environment
- My Practice

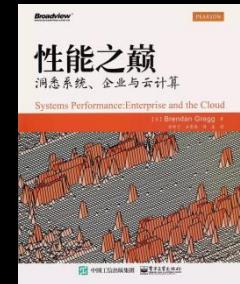
VI. Wrap-up

I. Linux Kernel Debugging

1) Overview

DTrace

- <https://en.wikipedia.org/wiki/DTrace>
a comprehensive dynamic tracing framework created by Sun Microsystems for troubleshooting kernel and application problems on production systems in real time. Originally developed for Solaris, it has since been released under the free CDDL in OpenSolaris and its descendant illumos...
- <https://www.oracle.com/technetwork/server-storage/solaris11/technologies/dtrace-1930301.html>
- <https://docs.oracle.com/en/operating-systems/oracle-linux/dtrace-guide/>
- ...
- <http://www.brendangregg.com/dtrace.html>



Linux Tracing Landscape

■ a

Sanitizers BCC

Tools	Trace-cmd	Perf tools	llvm	SystemTap	LTTng
Interface	Ftrace debugfs	Perf syscall			
Tracers	Function callgraph	Event Trace	Syscall trace	eBPF tracer	Misc tracers
Trace events	Ftrace callgraph	Traceevent Dynamic Event	Static Event	Perfevent	Stap runtime LTTng module
instrumentations	Ftrace (Function trace)	kprobes	uprobes	Tracepoint	Performance counter

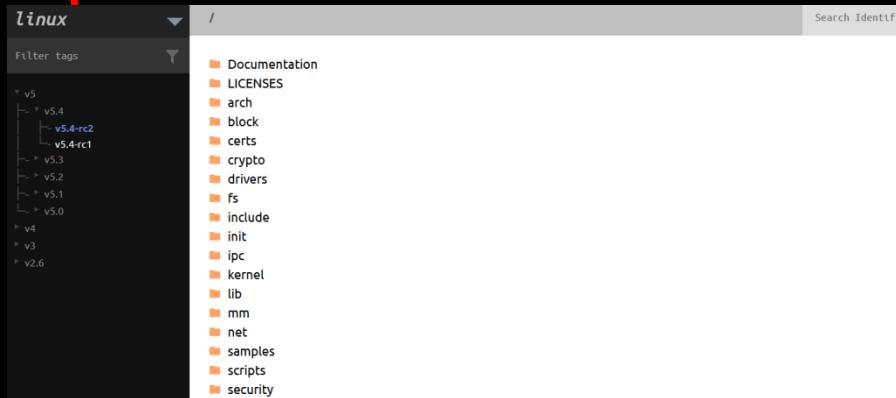
Source: “Dynamic Probes for Linux”, Masami Hiramatsu, Tracing Summit 2015

- **\$KERNEL_SRC/Documentation/kbuild/modules.rst**
- **\$KERNEL_SRC/Documentation/trace/events.rst**
- **\$KERNEL_SRC/Documentation/trace/ftrace*.rst**
- **\$KERNEL_SRC/Documentation/kprobes**

- **\$KERNEL_SRC/Documentation/trace/kprobetrace.rst**
- **\$KERNEL_SRC/Documentation/trace/uprobetracer.rst**
- **\$KERNEL_SRC/Documentation/trace/tracepoints*.rst**
- **\$KERNEL_SRC/Documentation/trace/features/*.***

Browsing Kernel Code Online

- <https://www.kernel.org/>
- <https://github.com/torvalds/linux>
- <https://elixir.bootlin.com/linux/latest/source>



2) eBPF

2.1 BPF (Berkeley Packet Filter, aka cBPF)

- https://en.wikipedia.org/wiki/Berkeley_Packet_Filter
- <http://www.tcpdump.org/papers/bpf-usenix93.pdf>
- History

- Before BPF, each OS (Sun, DEC, SGI etc) had its own packet filtering API
- In 1993: Steven McCanne & Van Jacobsen released a paper titled the *BSD Packet Filter (BPF)*
- Implemented as “Linux Socket Filter” in kernel 2.2
- While maintaining the BPF language (for describing filters), uses a different internal architecture

Source: [ebpfbasics-190611051559.pdf](#)

■ What is it

- Network packet filtering, Seccomp
- Filter Expressions → Bytecode → Interpret
- Small, in-kernel VM, Register based, switch dispatch interpreter, few instructions
- BPF uses a simple, non-shared buffer model made possible by today's larger address space

Source: [ebpfbasics-190611051559.pdf](#)

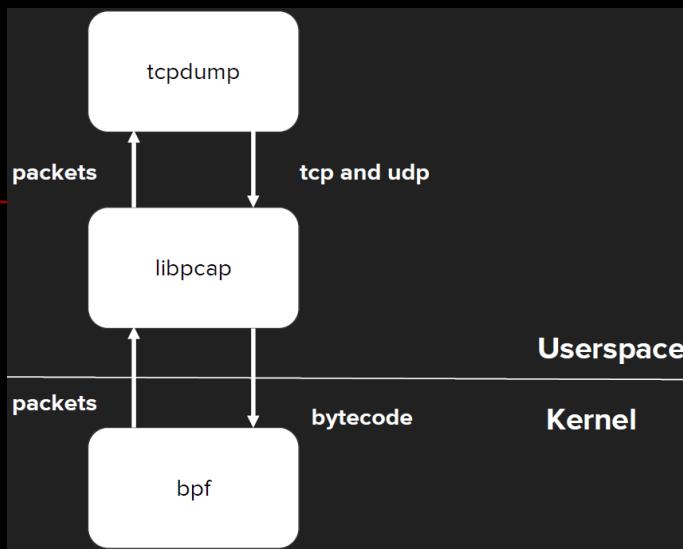
- Bytecode, register based VM, with a limited instruction set
- Runs in-kernel, designed for fast packet filtering
- 32-bit instructions (LOAD, STORE, ALU, BRANCH, RETURN)
- 2, 32-bit registers (A, X), hidden frame pointer

Source: [understandingebpfinahurry-190611040804.pdf](#)

<https://blog.cloudflare.com/bpf-the-forgotten-bytecode/>

...

■ tcpdump



Source: [understandingebpfinahurry-190611040804.pdf](https://www.slideshare.net/brendangregg/kernel-recipes-2017-performance-analysis-with-bpf)

```
# tcpdump host 127.0.0.1 and port 22 -d
(000) ldh      [12]
(001) jeq      #0x800          jt 2    jf 18
(002) ld      [26]
(003) jeq      #0x7f000001    jt 6    jf 4
(004) ld      [30]
(005) jeq      #0x7f000001    jt 6    jf 18
(006) ldb      [23]
(007) jeq      #0x84           jt 10   jf 8
(008) jeq      #0x6            jt 10   jf 9
(009) jeq      #0x11           jt 10   jf 18
(010) ldh      [20]
(011) jset     #0x1fff         jt 18   jf 12
(012) ldxb    4*([14]&0xf)
(013) ldh      [x + 14]
[...]
```

Optimizes packet filter performance

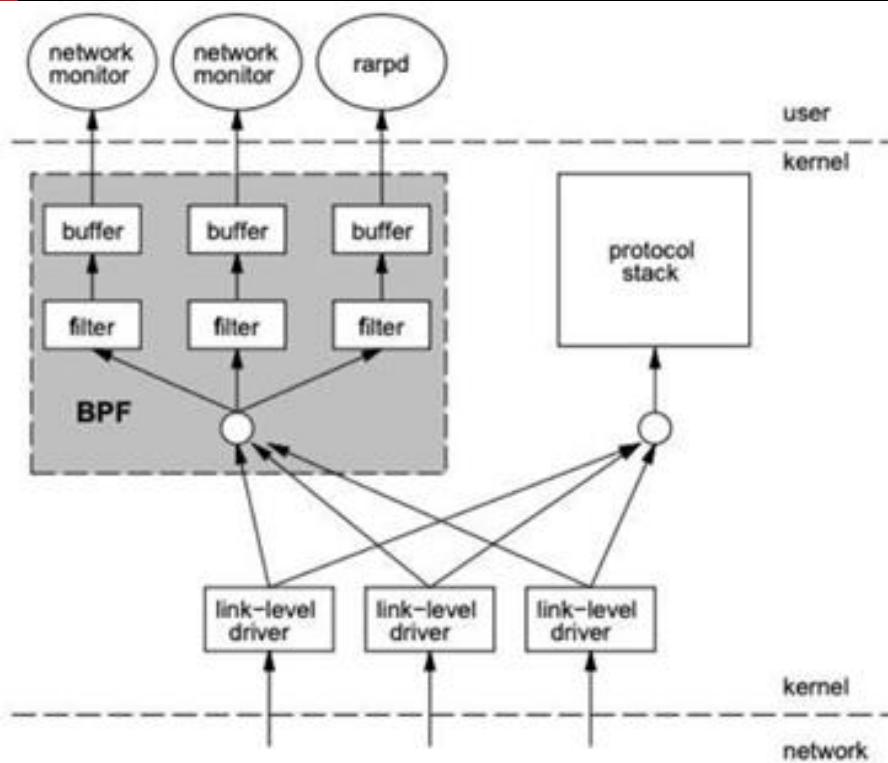
2 x 32-bit registers & scratch memory

User-defined bytecode executed by an in-kernel sandboxed virtual machine

Steven McCanne and Van Jacobson, 1993

Source: <https://www.slideshare.net/brendangregg/kernel-recipes-2017-performance-analysis-with-bpf>

Workflow

- 
- Utilizes sockets for passing/receiving packets to/from the kernel-space
- Filters are attached with the setsockopt(2) system call
- Create a special-purpose socket (i.e., PF_PACKET) 2
- Attach a BPF program to the socket using the setsockopt(2) system call
- Set the network interface to promiscuous mode with ioctl(2) (optionally)
- Read packets from the kernel, or send raw packets, by reading/writing to the corresponding file descriptor of /dev/bpf using read(2)/write(2) system calls
- Set various BPF parameters, (e.g. buffer size, attach some BPF filters) This is done using the ioctl(2) system call

Source: [ebpfbasics-190611051559.pdf](#)

2.2 eBPF (extended BPF)

- since Linux Kernel v3.15 and ongoing
 - aims at being a universal in-kernel virtual machine
 - a simple way to extend the functionality of Kernel at runtime
 - “dtrace for Linux”
-

BPF Features by Linux Kernel Version

- <https://github.com/iovisor/bcc/blob/master/docs/kernel-versions.md>

Instruction Set

- <https://github.com/iovisor/bpf-docs/blob/master/eBPF.md>

Projects using eBPF

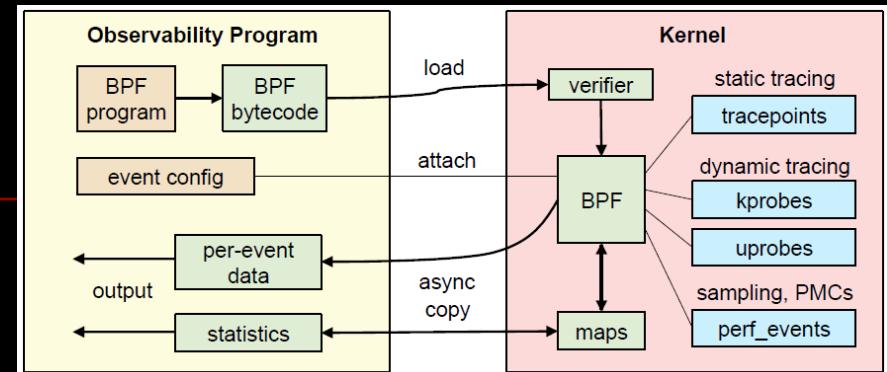
- <http://cilium.readthedocs.io/en/latest/bpf/#projects-using-bpf>

Good Resource

- <https://github.com/zoidbergwill/awesome-ebpf>

■ What is it

- User-defined, sandboxed bytecode executed by the kernel
- VM that implements a RISC-like assembly language in kernel space
- Similar to LSF, but with the following improvements:



- More registers, JIT compiler (flexible/ faster), verifier
- Attach on Tracepoint, Kprobe, Uprobe, USDT
- In-kernel trace aggregation & filtering
- Control via `bpf()`
- Designed for general event processing within the kernel
- All interactions between kernel/ user space are done through eBPF “maps”

Source: [ebpfbasics-190611051559.pdf](#)

Source: <https://kernel-recipes.org/en/2017/talks/performance-analysis-with-bpf/>

bpf() system call

- <http://www.man7.org/linux/man-pages/man2/bpf.2.html>

NAME [top](#)

bpf – perform a command on an extended BPF map or program

SYNOPSIS [top](#)

```
#include <linux/bpf.h>

int bpf(int cmd, union bpf_attr *attr, unsigned int size);
```

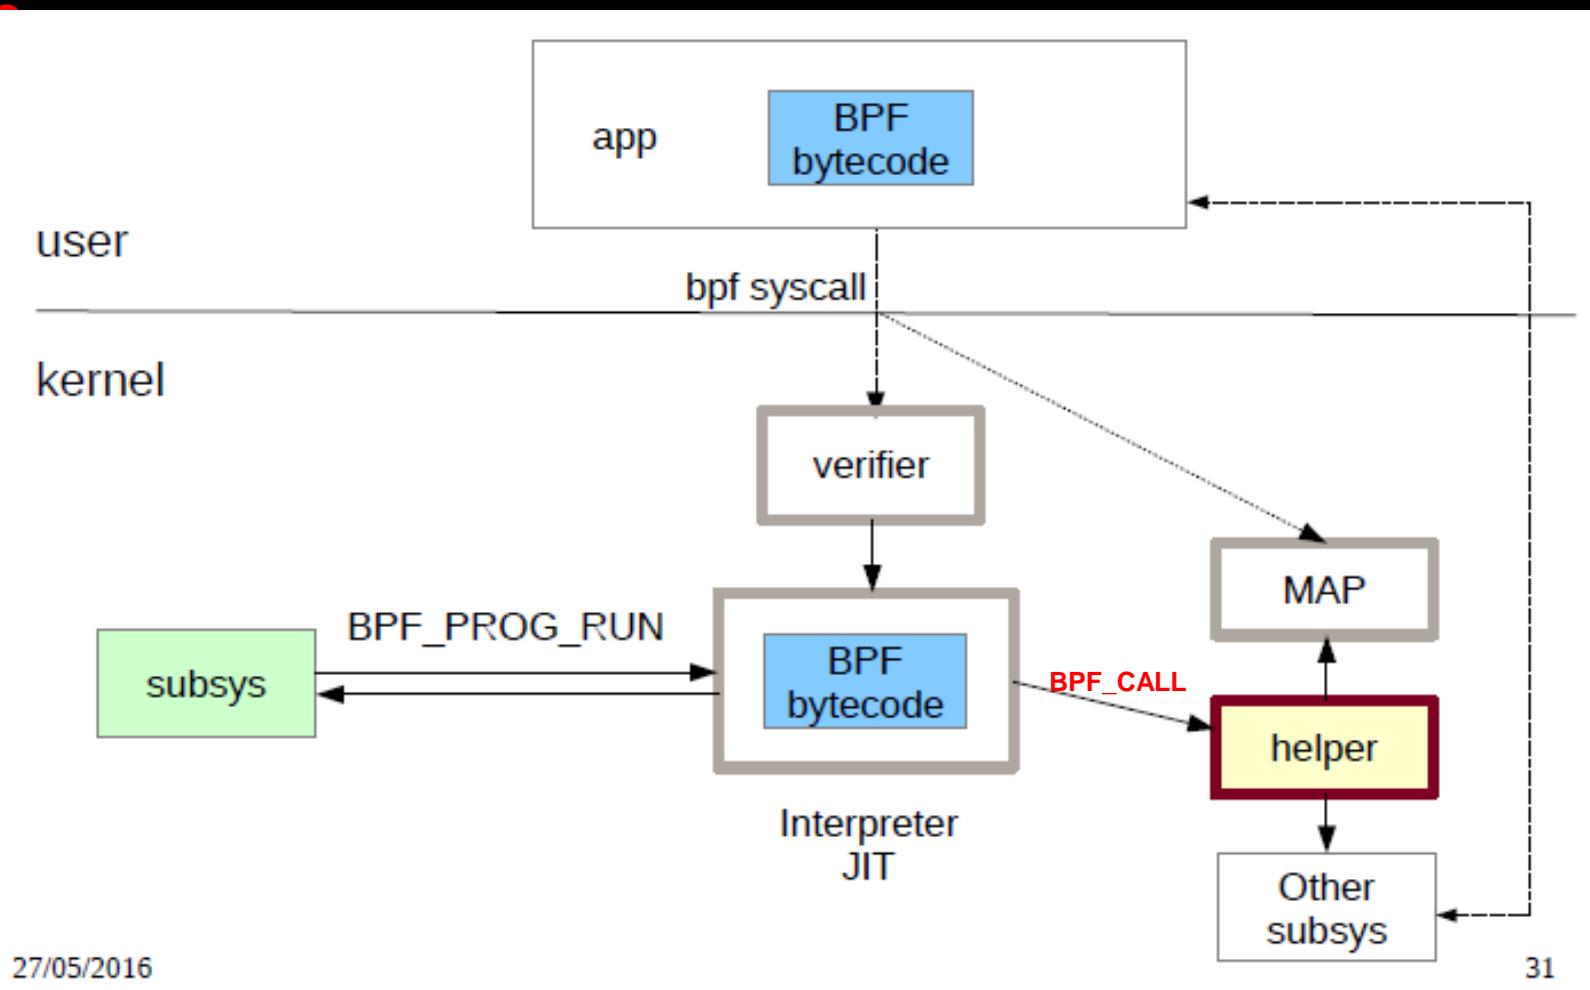
DESCRIPTION [top](#)

The **bpf()** system call performs a range of operations related to extended Berkeley Packet Filters. Extended BPF (or eBPF) is similar to the original ("classic") BPF (cBPF) used to filter network packets. For both cBPF and eBPF programs, the kernel statically analyzes the programs before loading them, in order to ensure that they cannot harm the running system.

eBPF extends cBPF in multiple ways, including the ability to call a fixed set of in-kernel helper functions (via the **BPF_CALL** opcode extension provided by eBPF) and access shared data structures such as eBPF maps.

...

Workflow



Source: <http://www.slideshare.net/vh21/meet-cutebetweenebpffandtracing>

Internals

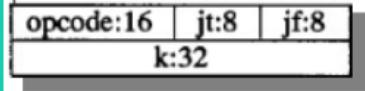
- Pls refer to my presentation "eBPF in Action" at LC3 Beijing (on Jun 25, 2018)
- **\$KERNEL_SRC/Documentation/networking/filter.txt**
- **https://kernelnewbies.org/Linux_5.3**

6. Tracing, perf and BPF

- BPF
 - libbpf: Add BTF-to-C dumping support, allowing to output a subset of BTF types as a compilable C type definitions. This is useful by itself, as raw BTF output is not easy to inspect and comprehend. But it's also a big part of BPF CO-RE (compile once - run everywhere) initiative aimed at allowing to write relocatable BPF programs, that won't require on-the-host kernel headers (and would be able to inspect internal kernel structures, not exposed through kernel headers) [commit](#), [commit](#)
 - Implements initial version (as discussed at LSF/MM2019 conference) of a new way to specify BPF maps, relying on BTF type information, which allows for easy extensibility, preserving forward and backward compatibility [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#)
 - Adds support for propagating congestion notifications to TCP from cgroup inet skb egress BPF programs [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#)
 - Add `SO_DETACH_REUSEPORT_BPF` to detach BPF prog from reuseport sk [commit](#), [commit](#)
 - Add a `sock_ops` callback that can be selectively enabled on a socket by socket basis and is executed for every RTT. BPF program frequency can be further controlled by calling `bpf_ktime_get_ns` and bailing out early [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#)
 - Allow CGROUP_SKB programs to use `bpf_skb_cgroup_id()` helper [commit](#)
 - Eliminate zero extensions for sub-register writes [commit](#), [commit](#)
 - Export `bpf_sock` for `BPF_PROG_TYPE_CGROUP_SOCK_ADDR` prog type [commit](#) and for `BPF_PROG_TYPE SOCK_OPS` prog type [commit](#)
 - allow wide (u64) aligned stores for some fields of `bpf_sock_addr` [commit](#), [commit](#), [commit](#)
 - Adds support for fq's Earliest Departure Time to HBM (Host Bandwidth Manager) [commit](#)
 - Introduces verifier support for bounded loops and other improvements [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#)
 - bpf: getsockopt and setsockopt hooks [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#)
 - libbpf: add `bpf_link` and tracing attach APIs [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#), [commit](#)

...

Comparison

	cBPF	eBPF
Register	Two 32 bit registers: A: accumulator X: indexing	Eleven 64 bit registers: R0: return value/exit value R1-R5: arguments R6-R9: callee saved registers R10: read-only frame pointer
Instruction	~30 	~90 
JIT	Support	Support (better mapping with newer architectures for JITing)
Toolchain	GCC, tools/net	LLVM eBPF backend
Platform	x86_64, ARM, ARM64, SPARC, PowerPC, MIPS and s390	x86-64, aarch64, s390x...
System Call		#include <linux/bpf.h> int bpf(int cmd, union bpf_attr *attr, unsigned int size); (CALL, MAP, LOAD...)
Application	tcpdump... apply for seccomp filters, traffic control...	DDoS Mitigation, Intrusion Detection, Container Security, SDN Configuration, Observability...

3) Development

3.1 Toolchain

LLVM

- eBPF backend firstly introduced in LLVM 3.7 release
- <http://llvm.org/docs/CodeGenerator.html#the-extended-berkeley-packet-filter-ebpf-backend>

- Enabled by default with all major distributions
 - Registered targets: llc --version
 - llc's BPF -march options: bpf, bpfeb, bpfel
 - llc's BPF -mcpu options: generic, v1, v2, probe
- Assembler output through -S supported
- llvm-objdump for disassembler and code annotations (via DWARF)
- Annotations correlate directly with kernel verifier log
- Outputs ELF file with maps as relocation entries
 - Processed by BPF loaders (e.g. iproute2) and pushed into kernel

Source: <https://ossna2017.sched.com/event/BCsg/making-the-kernels-networking-data-path-programmable-with-bpf-and-xdp-daniel-borkmann-coalent>

- **\$LLVM_SRC/lib/Target/BPF**
- <http://cilium.readthedocs.io/en/latest/bpf/>
- **LLVM 9.0 Released With Ability To Build The Linux x86_64**

GCC (GCC support for eBPF is on the way)

- https://www.phoronix.com/scan.php?page=news_item&px=GNU-Binutils-eBPF-Support
//GNU Binutils Begins Landing eBPF Support
- https://www.phoronix.com/scan.php?page=news_item&px=GCC-10-eBPF-Backend-Plans
//Oracle Is Aiming To Contribute An eBPF Backend To The GCC 10 Compiler
- https://www.phoronix.com/scan.php?page=news_item&px=Oracle-More-DTrace-Linux-eBPF
//Oracle Is Working To Upstream More Of DTrace To The Linux Kernel & eBPF Implementation
- https://www.phoronix.com/scan.php?page=news_item&px=Oracle-GCC-10-eBPF-V2
//2019-8-17::Oracle Continues Working On eBPF Support For GCC 10
- https://www.phoronix.com/scan.php?page=news_item&px=GCC-10-eBPF-Port-Lands
//GCC 10 Lands The eBPF Port For Targeting The Linux In-Kernel VM

Status

- - Phase 1: add eBPF target to the toolchain
 - bpf-unknown-none
 - binutils (upstream since May 2019)
 - GCC (upstream since September 2019)
 - Phase 2: make the generated programs palatable for the kernel loaders and verifier, and **keep it that way**.
 - Phase 3: provide development goodies for eBPF developers
 - GNU simulator
 - GDB
 - ...

Source: “eBPF support in the GNU Toolchain”, Jose E. Marchesi (Oracle),
Linux Plumbers Conference 2019

3.2 Debugging

- https://www.netronome.com/documents/143/UG_Getting_Started_with_eBPF_Offload.pdf
- `int bpf(int cmd, union bpf_attr *attr, unsigned int size);`

log_level

- 0: No debug output.
- 1: Debug information from the verifier (all instructions).
- 2: More information: add all register states after each instruction.

- `llvm-objdump, llvm-mc...`

`bpftool`

- `$KERNEL_SRC/tools/bpf /bpftool`

```
bpftool
  bash-completion
  btf.c
  btf_dumper.c
  cfg.c
  cfg.h
  cgroup.c
  common.c
  Documentation
  feature.c
  jit_disasm.c
  json_writer.c
  json_writer.h
  main.c
  main.h
  Makefile
  map.c
  map_perf_ring.c
  net.c
  netlink_dumper.c
  netlink_dumper.h
  perf.c
  prog.c
  tracelog.c
  xlated_dumper.c
  xlated_dumper.h
```

```
# bpftool prog show
1337: sched_cls name cls_entry tag e202124da7c84e89
          loaded_at Mar 08/19:53 uid 0
          xlated 304B not jited memlock 4096B
```

```
# bpftool prog dump xlated id 1337
  0: (71) r6 = *(u8 *)(r1 +142)
  1: (54) (u32) r6 &= (u32) 1
  2: (15) if r6 == 0x0 goto pc+7
  3: (bf) r6 = r1
  [...]
  37: (95) exit
```

```
# bpftool map
1234: array name ch_rings flags 0x0
          key 4B value 4B max_entries 7860 memlock 65536B

# bpftool map dump id 1234
key: 00 00 00 00 value: 00 00 00 00
key: 01 00 00 00 value: 00 00 00 00
key: 02 00 00 00 value: 00 00 00 00
key: 03 00 00 00 value: 00 00 00 00
[...]
Found 7860 elements
```

4) Summary

Evaluation



“eBPF is Linux’s new superpower”

Gaurav Gupta



“eBPF does to Linux what JavaScript does to HTML”

Brendan Gregg



“Run code in the kernel without having to write a kernel module”

Liz Rice

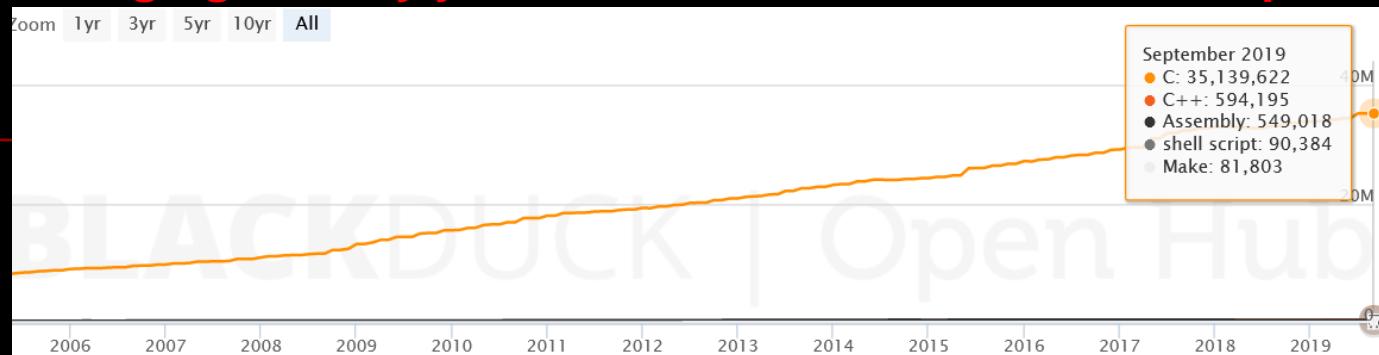
Source: [ebpfbasics-190611051559.pdf](#)

Kernel Space & User Space Instrumentation



■ Polyglot VM

Changing the way you think about Linux Kernel development:



Language	Code Lines	Comment Lines	Comment Ratio	Blank Lines	Total Lines	Total Percentage
C	35,139,622	6,241,390	15.1%	6,333,153	47,714,165	<div style="width: 95.7%; background-color: orange;"></div> 95.7%
C++	594,195	262,892	30.7%	116,246	973,333	<div style="width: 2.0%; background-color: orange;"></div> 2.0%
Assembly	549,018	124,794	18.5%	96,032	769,844	<div style="width: 1.5%; background-color: grey;"></div> 1.5%

Source: https://www.openhub.net/p/linux/analyses/latest/languages_summary

■ The Next Linux Superpower:

User Space/Kernel Space Repartition & Unifying
Reconstructing nearly every aspect of Linux Networking and Security subsystem

II. BCC(BPF Compiler Collection)

1) Overview

- <https://github.com/iovisor/bcc/>

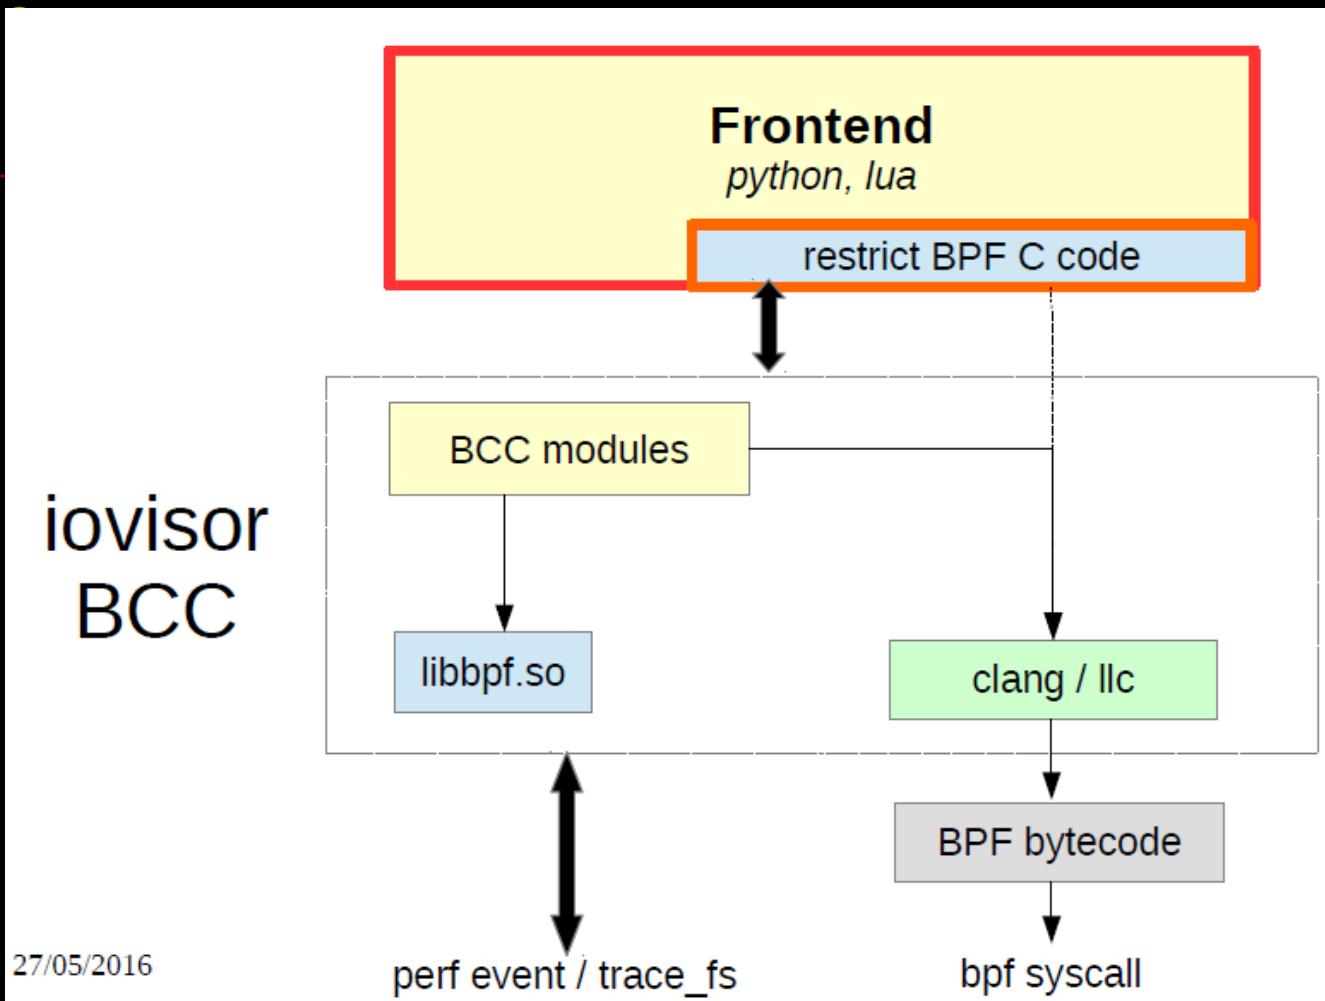


A toolkit with Python/Lua frontend for compiling, loading, and executing BPF programs, which allows user-defined instrumentation on a live kernel image:

- compile BPF program from C source
- attach BPF program to kprobe/uprobe/tracepoint/USDT/socket
- poll data from BPF program
- framework for building new tools or one-off scripts
- additional projects to support Go, Rust, and DTrace-style frontend
- ...

Architecture

iovisor
BCC



27/05/2016

Source: <http://www.slideshare.net/vh21/meet-cutebetweenbpffandtracing>

A Sample

- [https://lwn.net/Articles/747640/ //Some advanced BCC topics](https://lwn.net/Articles/747640/)

```
#!/usr/bin/env python

from bcc import BPF
from time import sleep

program = """
BPF_HASH(callers, u64, unsigned long);

TRACEPOINT_PROBE(kmem, kmalloc) {
    u64 ip = args->call_site;
    unsigned long *count;
    unsigned long c = 1;

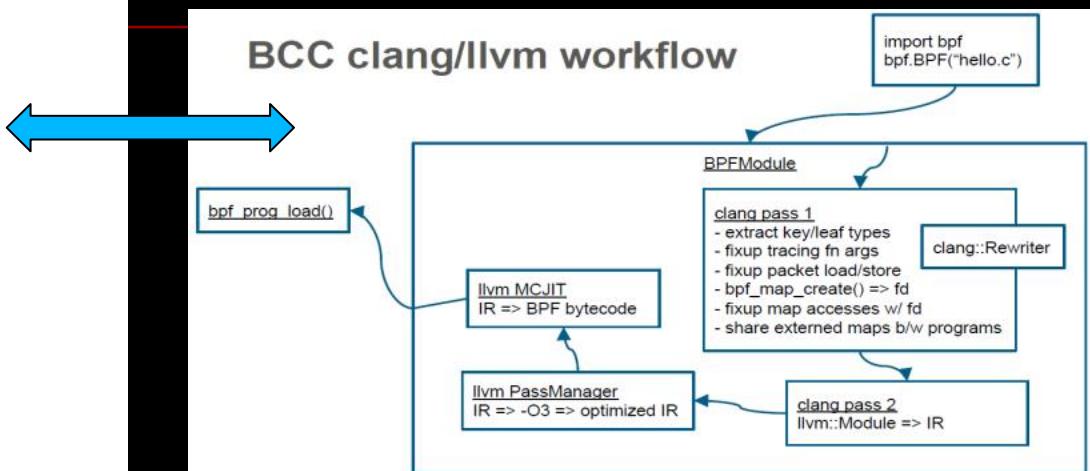
    count = callers.lookup((u64 *)&ip);
    if (count != 0)
        c += *count;

    callers.update(&ip, &c);

    return 0;
}
"""

b = BPF(text=program)

while True:
    try:
        sleep(1)
        for k, v in sorted(b["callers"].items()):
            print ("%s %u" % (b.ksym(k.value), v.value))
        print
    except KeyboardInterrupt:
        exit()
```



Source: http://linuxplumbersconf.org/2015/ocw/system/presentations/3249/original/bpf_llvm_2015aug19.pdf

The output from this little program looks like:

```
# ./example.py
i915_sw_fence_await_dma_fence 4
intel_crtc_duplicate_state 4
Sys_memfd_create 1
drm_atomic_state_init 4
sg_kmalloc 7
intel_atomic_state_alloc 4
seq_open 504
Sys_bpf 22
```

2) Development

- <https://github.com/iovisor/bcc/blob/master/docs/tutorial.md>
- https://github.com/iovisor/bcc/blob/master/docs/reference_guide.md
- https://github.com/iovisor/bcc/blob/master/docs/tutorial_bcc_python_developer.md
- ...

2.1 New Ideas

py2bpf

- <https://github.com/facebookresearch/py2bpf>
- translates functions from Python to BPF

```
q = py2bpf.datastructures.BpfQueue(ctypes.c_int)

@py2bpf.kprobe.probe('sys_close')
def on_sys_close(pt_regs):
    pid = py2bpf.funcs.get_current_pid_tgid() & 0xffffffff
    ptr = py2bpf.funcs.addrof(pid)
    cpuid = py2bpf.funcs.get_smp_processor_id()
    py2bpf.funcs.perf_event_output(pt_regs, q, cpuid, ptr)
    return 0

with on_sys_close():
    for pid in q:
        print('pid={}'.format(pid))
```

gobpf

- <https://github.com/iovisor/gobpf>
- provides go bindings for the BCC framework as well as low-level routines to load and use eBPF programs from .elf files



- requirement: install the latest released version of libbcc
- still is in early stage, but usable

3) Summary

- <https://www.infoworld.com/article/3444198/the-best-open-source-software-of-2019.html>

The diagram illustrates the Linux bcc/BPF Tracing Tools architecture. It shows a central vertical stack of system components: Applications, Runtimes, System Libraries, System Call Interface, VFS, Sockets, File Systems, TCP/UDP, Volume Manager, IP, Block Device, Net Device, and Device Drivers. Arrows point from various tools on the left to these components. A large box on the right lists specific tools grouped by category.

Linux bcc/BPF Tracing Tools

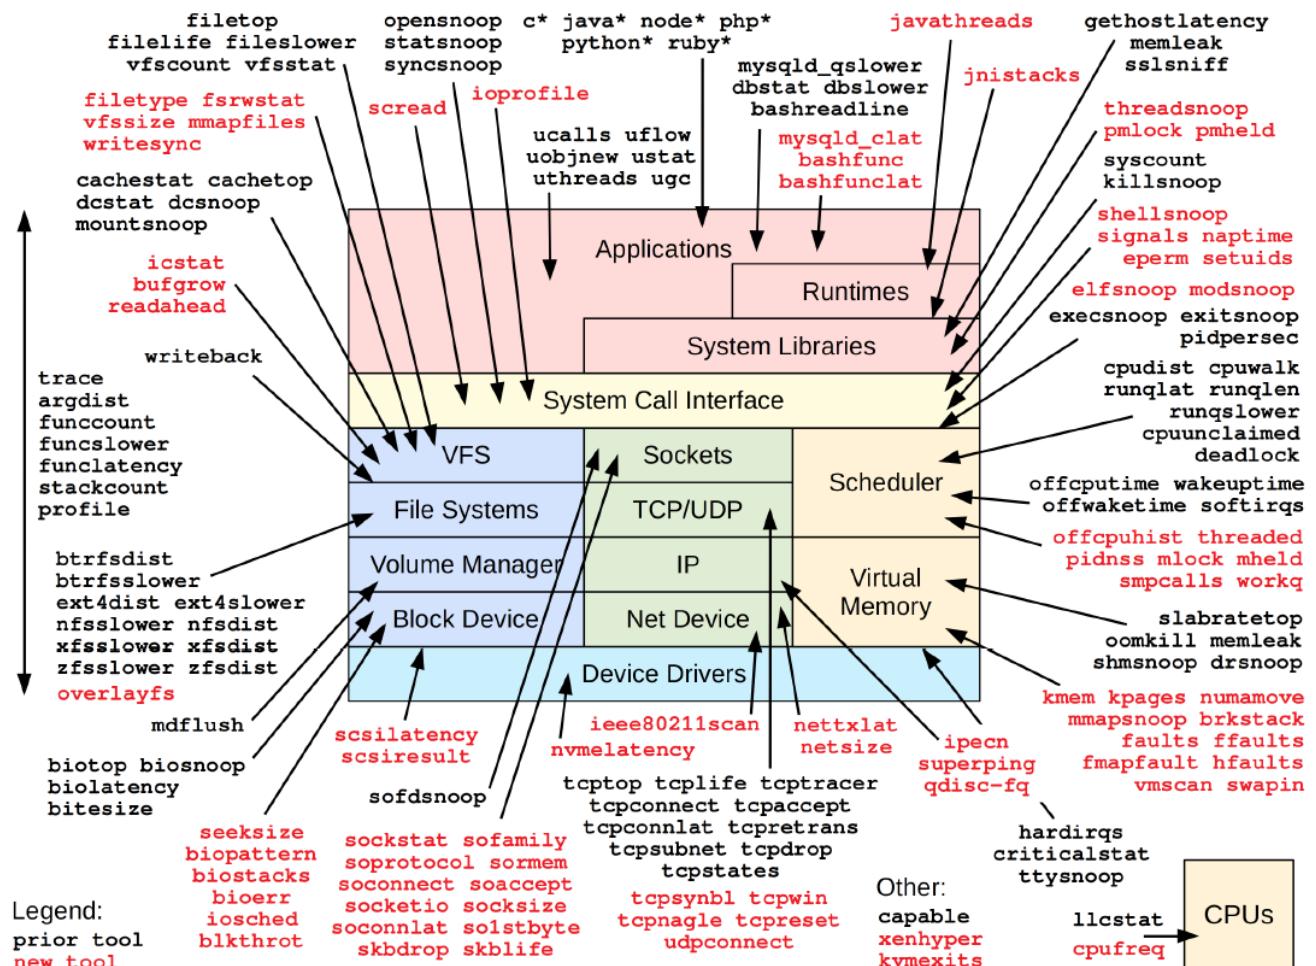
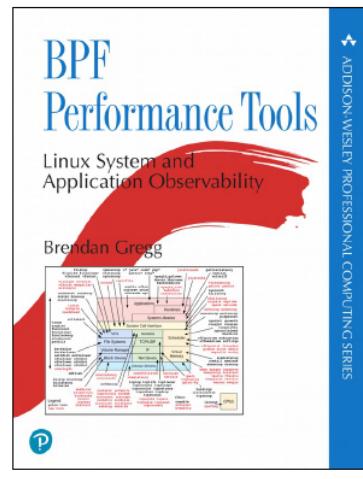
- Applications:** ucalls, uflow, uobjnew, ustat, uthreads, ugc, * java*, node*, php*, Python*, ruby*
- Runtimes:** mysqld_qslower, dbstat, dbslower, bashreadline
- System Libraries:** gethostlatency, memleak, ssleniff
- System Call Interface:** syscount, killnsnoop, execsnoop, exitsnoop, pidpersec
- VFS:** cpudist, cpuwalk, runqlat, runqlen, rungslower, cpounclaimed, deadlock
- Sockets:** offcpertime, offwaketime, softirqs
- File Systems:** slabratoep, comkill, memleak, shmsnoop, drsnoop
- TCP/UDP:** hardirqs, criticalstat, tsysnoop
- Volume Manager:** mdflush
- IP:** l1cstat
- Block Device:** biosnoop, biolatency, bitesize
- Net Device:** teptop, tcplife, tcptracer, tpcconnect, tpcaccept, tpcconnlat, tpcpretrans, tcpsubnet, tcpdrop, tcpstates
- Device Drivers:** sofdsnoop
- Other:** capable
- CPU:** CPUS

<https://github.com/levisor/bcc/tools/2019>

BOSSIE
2019 AWARDS

InfoWorld

BPF Perf Tools



Source: “BPF Tracing Tools”, Brendan Gregg, Linux Plumbers Conference 2019

III. LISA

1) Overview

■ **Linux Integrated System Analysis**

■ <https://github.com/ARM-software/lisa>

- The LISA project provides a toolkit that supports regression testing and interactive analysis of Linux kernel behavior. LISA stands for Linux Integrated/Interactive System Analysis. LISA's goal is to help Linux kernel developers to measure the impact of modifications in core parts of the kernel. The focus is on the scheduler (e.g. EAS), power management and thermal frameworks. However LISA is generic and can be used for other purposes too.

LISA has a *host/target* model. LISA itself runs on a *host* machine, and uses the [devlib](#) toolkit to interact with the *target* via SSH, ADB or telnet. LISA is flexible with regard to the target OS; its only expectation is a Linux kernel-based system. Android, GNU/Linux and busybox style systems have all been used.

LISA provides features to describe workloads (notably using [rt-app](#)) and run them on targets. It can collect trace files from the target OS (e.g. strace and ftrace traces), parse them via the [TRAPPy](#) framework. These traces can then be parsed and analysed in order to examine detailed target behaviour during the workload's execution.

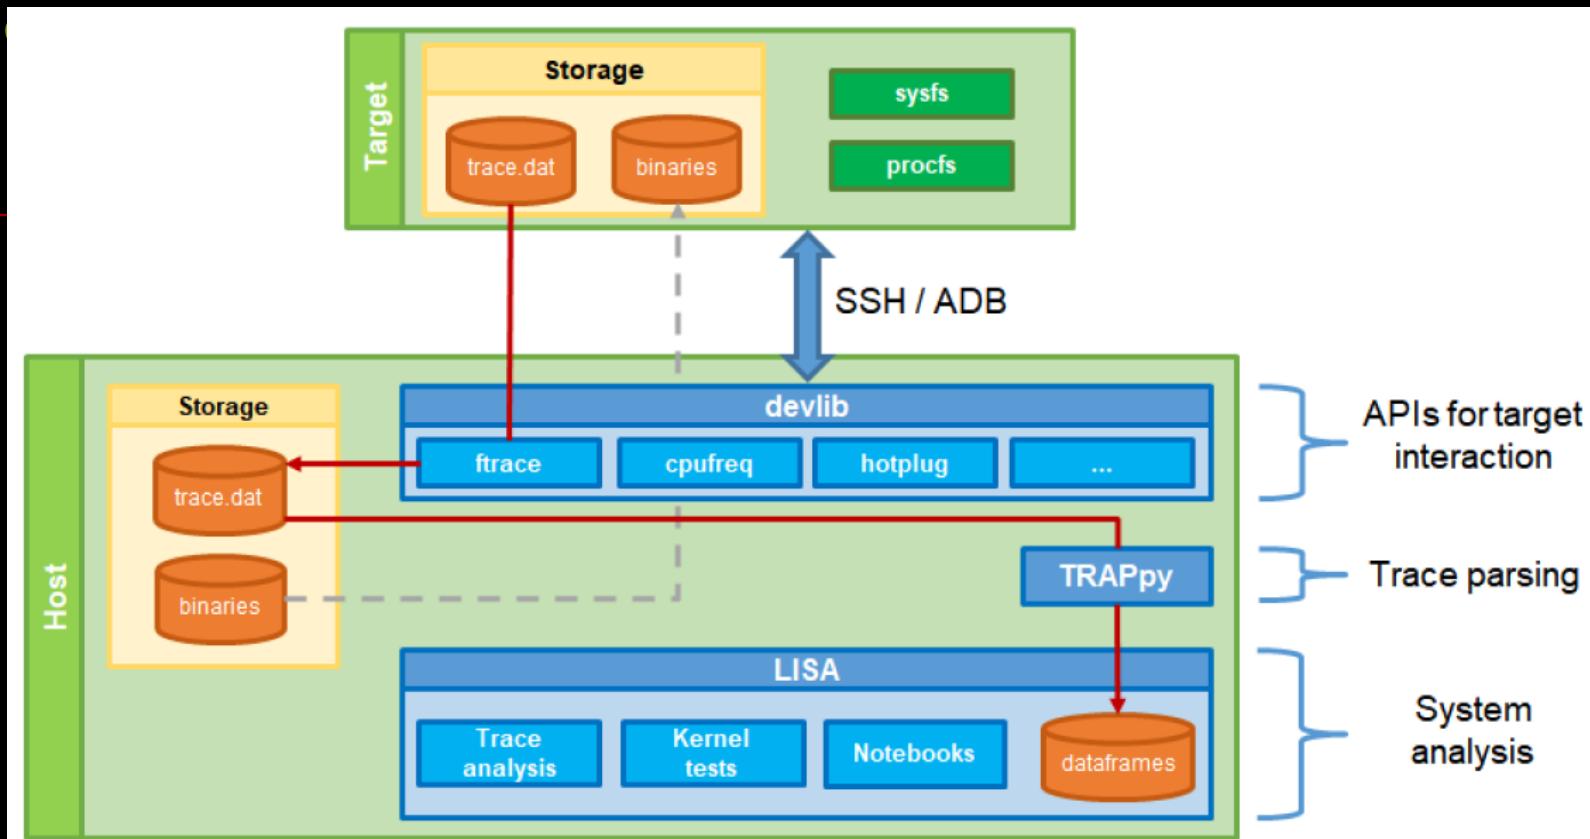
Some LISA features may require modifying the target OS. For example, in order to collect ftrace files the target kernel must have [CONFIG_DYNAMIC_FTRACE](#) enabled.

There are two "entry points" for running LISA:

- Via the [Jupyter/IPython notebook framework](#). This allows LISA to be used interactively and supports visualisation of trace data. Some notebooks are provided with example and ready-made LISA use-cases.
- Via the automated test framework. This framework allows the development of automated pass/fail regression tests for kernel behaviour. The [BART](#) toolkit provides additional domain-specific test assertions for this use-case. LISA provides some ready-made automated tests under the `lisa/tests/` directory.

- Unlike **devlib** or **trappy**, **LISA** is now **Python 3 only**.
 - A big refactoring effort was started in mid 2018, which produced a lot of (much needed) changes.
-

Architecture



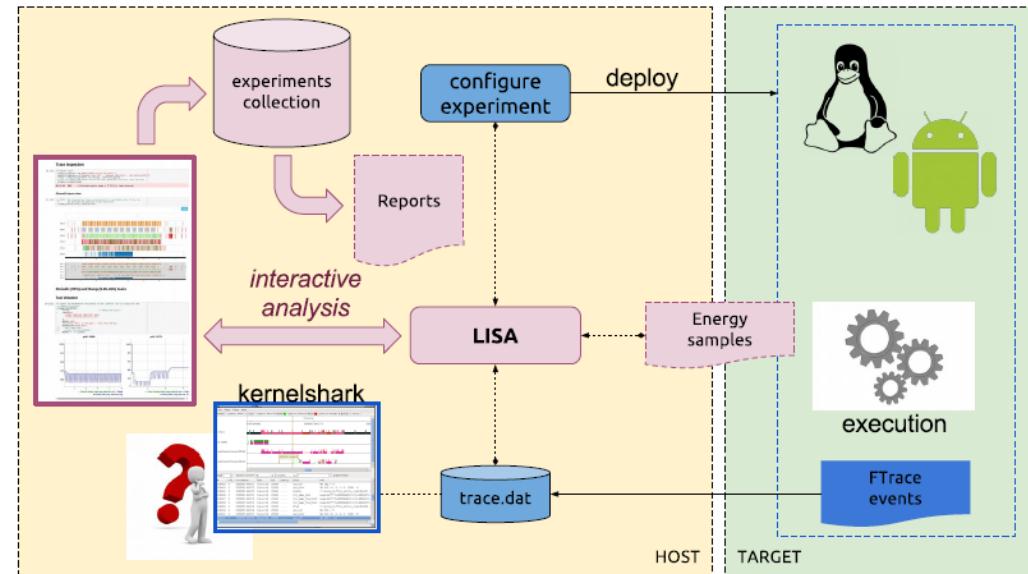
Source: <https://lisa-linux-integrated-system-analysis.readthedocs.io/en/master/>

2) Internals

Workflow

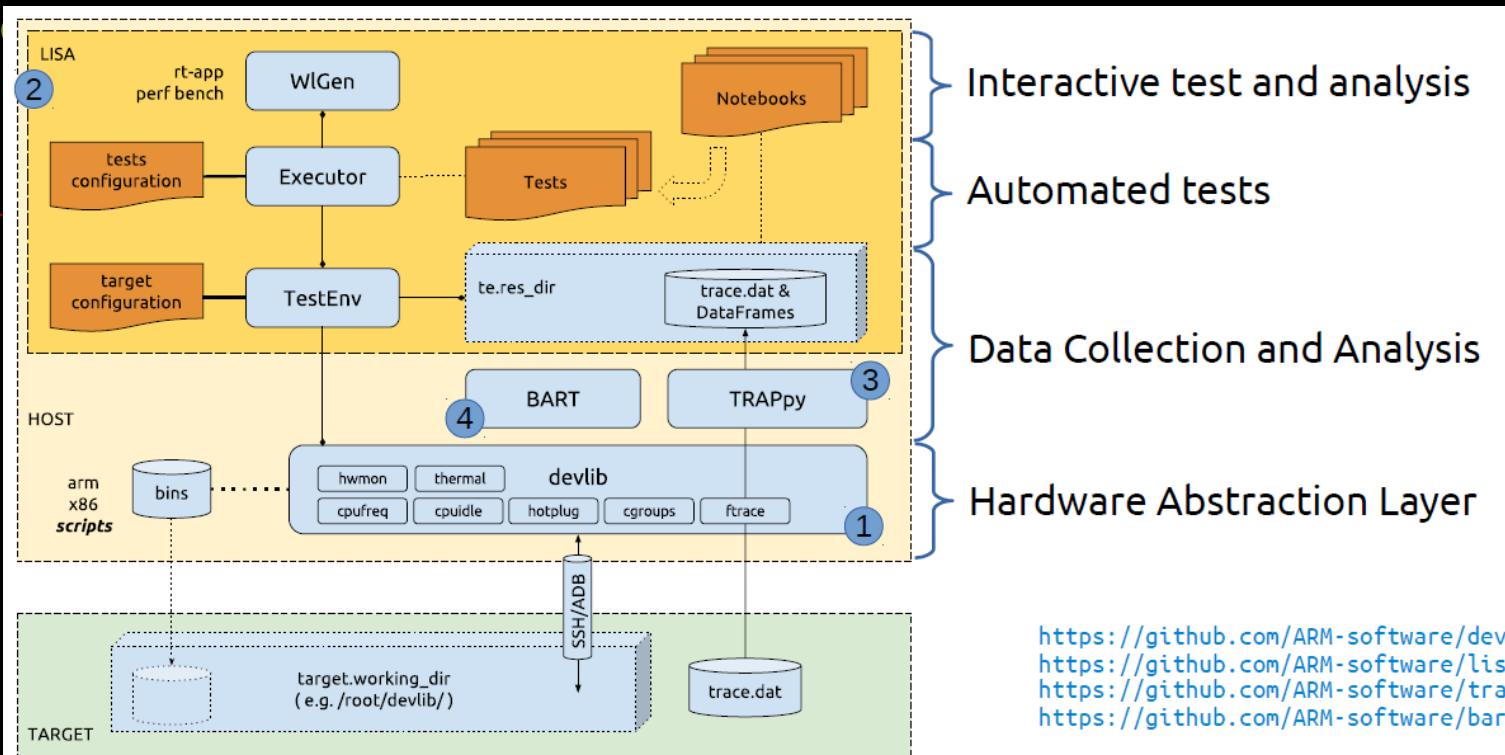
- Experimenting using an “interactive environment”
- Data analysis and post-processing
- Tests definitions to support regression analysis

*Evaluate trade-offs
on Power/Performances*



Source: “Linux Integrated System Analysis (LISA) & Friends”, Patrick Bellasi, ELC 2016

Toolkit

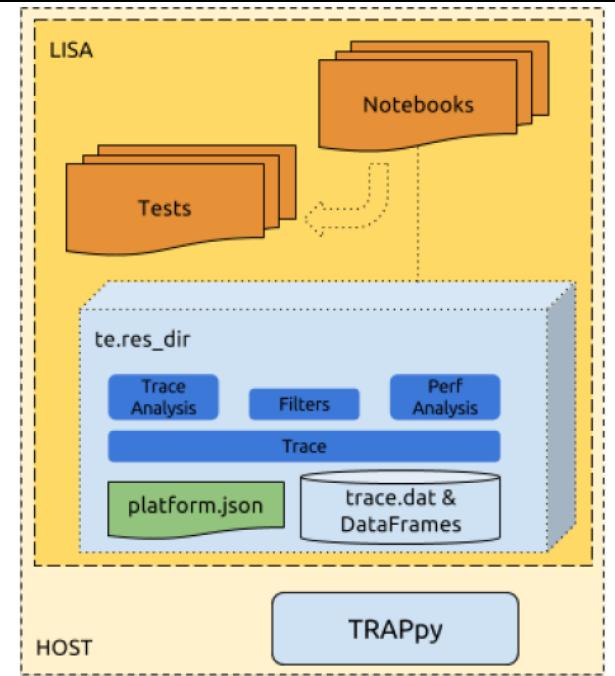


<https://github.com/ARM-software/devlib> [1]
<https://github.com/ARM-software/lisa> [2]
<https://github.com/ARM-software/trappy> [3]
<https://github.com/ARM-software/bart> [4]

Source: “Linux Integrated System Analysis (LISA) & Friends ”, Patrick Bellasi, ELC 2016

Data Analysis

- LISA::Trace glues platform data with TRAPPy DataFrames
more complete analysis dataset
- LISA::Filters
commonly used events filtering functions
- LISA::TraceAnalysis
commonly used trace events plots
- LISA::PerfAnalysis
commonly used performance plots



Source: "Linux Integrated System Analysis (LISA) & Friends ", Patrick Bellasi, ELC 2016

trappy

- <https://github.com/ARM-software/trappy>
 - Under the hood, it is used to convert the trace events into **pandas.DataFrame** which are suited to handling large data sets
-

rt-app

- <https://github.com/scheduler-tools/rt-app>
- emulates typical mobile and real-time systems use cases and gives runtime information

workload-automation

- <https://github.com/ARM-software/workload-automation>
- a framework for executing workloads and collecting measurements on Android and Linux devices. It includes automation for nearly 40 workloads and supports some common instrumentation (ftrace, hwmon) along with a number of output formats

2.1 Automated Testing

- **exekall** is the test runner of LISA. **lisa-test** command is a thin wrapper on top of exekall.
- use `exekall run *subcommand*` to start a test session
e.g., `exekall run lisa/tests/ --conf target_conf.yml`
- **bisector** allows setting up the steps of a test iteration, repeating them an infinite number of times (by default).

e.g.,

```
# As a convenience, myreport.yml.gz.log will also be created, with a
# behaviour similar to: bisector run ... 2>&1 | tee myreport.yml.gz.log
bisector run --steps steps.yml --report myreport.yml.gz
```

- **check-test-fix.py** tool can be used to check that a fix to a test resolved errors or a regression.

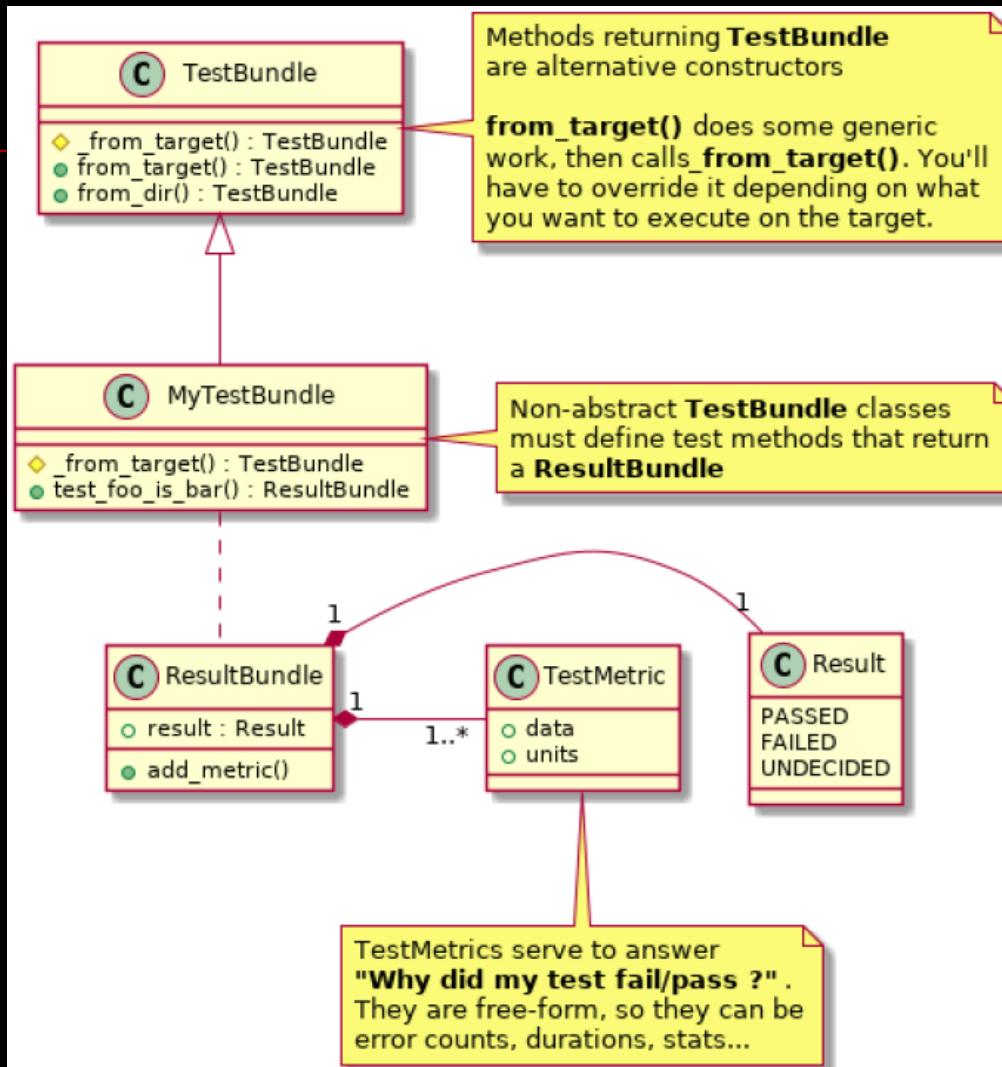
e.g.,

```
# The test to check is selected using --select in the same way as for `exekall run`.
# hikey960.report.yml.gz is a bisector report generated using `bisector run`
# All options coming after the report are passed to `bisector report` to
# control what artifacts are downloaded and what TestBundle are used.
check-test-fix.py --select 'OneSmallTask:test_task_placement' hikey960.report.yml.gz -
→oiterations=1-20
```

- **Wltest** enables the comparison of power and performance impacts of kernel changes on Android devices.
- ...

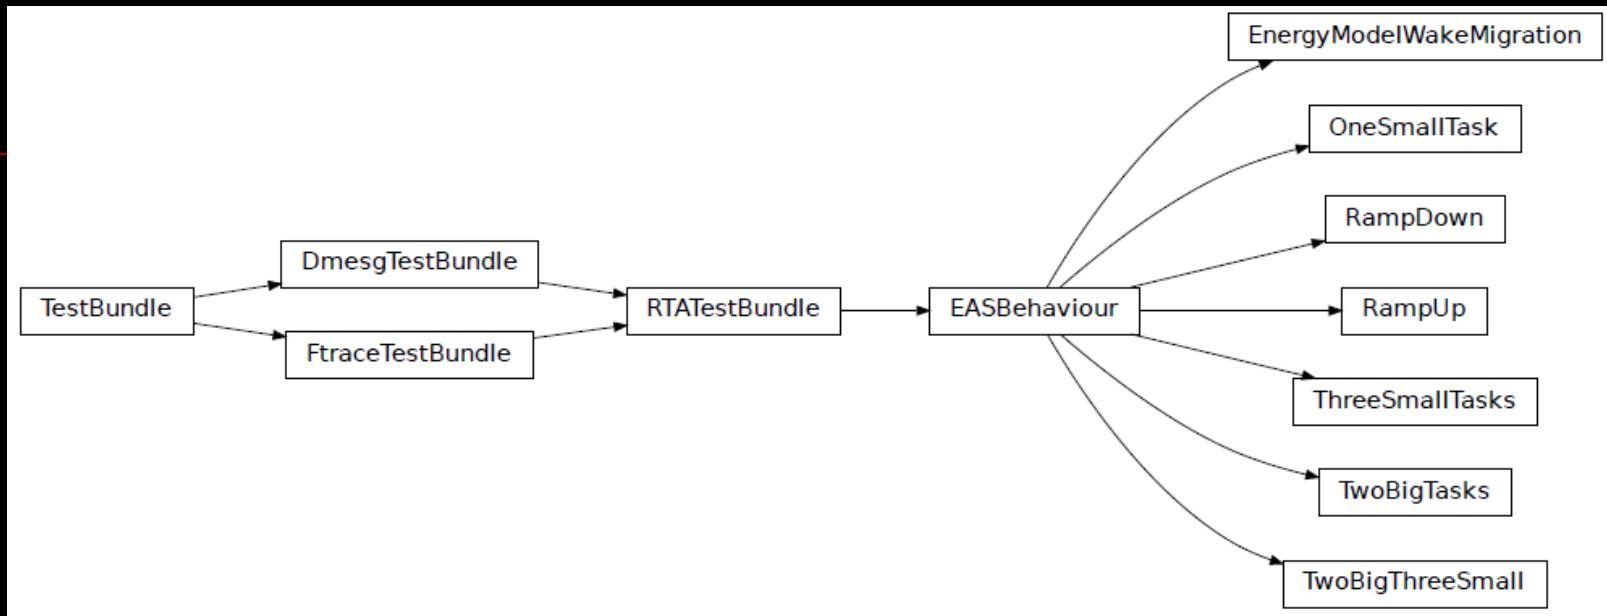
Writing Tests

- the main class of the kernel tests is **TestBundle**



Source: <https://lisa-linux-integrated-system-analysis.readthedocs.io/en/master/>

- E.g., scheduler tests
EAS tests



Source: <https://lisa-linux-integrated-system-analysis.readthedocs.io/en/master/>

- **HWMON: ARM Energy Probe (AEP)**

ARM Energy Probes are **lightweight power measurement tools** for software developers. They can monitor up to three voltage rails simultaneously.



Source: <https://lisa-linux-integrated-system-analysis.readthedocs.io/en/master/>

Caiman

<https://github.com/ARM-software/caiman>

Caiman can usually auto-detect the Energy Probe device. But on older versions of Linux which do not have `libudev.so.0`, such as Red Hat, it cannot be auto-detected. In that case the device name will need to be provided, usually `/dev/ttyACM0`.

setting up the ARM Energy Probe or National Instruments DAQ and running caiman...

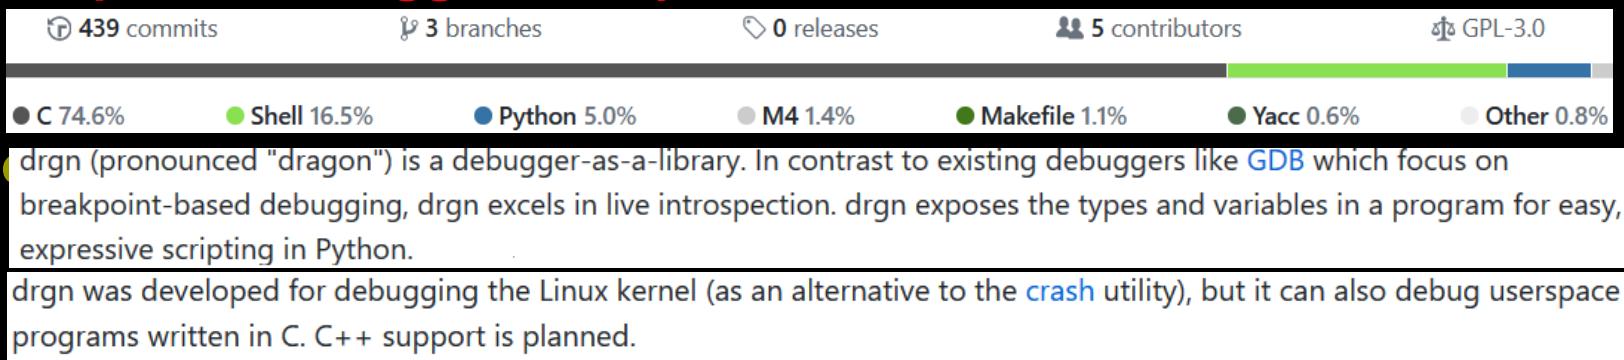
3) Development

- <https://github.com/ARM-software/lisa/tree/master/doc>
 - <https://lisa-linux-integrated-system-analysis.readthedocs.io/en/master/>
 - ...
-

IV. Drgn

1) Overview

- <https://github.com/osandov/drgn>
scriptable debugger library



- <https://lwn.net/Articles/789641/> //A kernel debugger in Python
- <https://drgn.readthedocs.io/en/latest/>

drgn debugs the running kernel by default; run `sudo drgn`. To debug a running program, run `sudo drgn -p $PID`. To debug a core dump (either a kernel vmcore or a userspace core dump), run `drgn -c $PATH`. The program must have debugging symbols available.

Installation
User Guide
▪ Quick Start
▪ Core Concepts
▪ Programs
▪ Objects
▪ Helpers
▪ Other Concepts
▪ Symbols
▪ Stack Traces
▪ Types
▪ Platforms
▪ Command Line Interface
▪ Script Mode
▪ Interactive Mode
▪ Next Steps
Advanced Usage
API Reference
Helpers

Script Mode

Script mode is useful for reusable scripts. Simply pass the path to the script along with any arguments:

```
$ cat script.py
import sys
from drgn.helpers.linux import find_task

pid = int(sys.argv[1])
uid = find_task(prog, pid).cred.uid.val.value_()
print(f'PID {pid} is being run by UID {uid}')
$ sudo drgn script.py 601
PID 601 is being run by UID 1000
```

It's even possible to run drgn scripts directly with the proper [shebang](#):

```
$ cat script2.py
#!/usr/bin/env drgn

mounts = prog['init_task'].nsproxy.mnt_ns.mounts.value_()
print(f'You have {mounts} filesystems mounted')
$ sudo ./script2.py
You have 36 filesystems mounted
```

Future

- - Combining with BPF tracing
 - Better debug information
 - Better vmcores

Source: “[drgn: Programmable Debugging](#)”, Omar Sandoval ([Facebook](#)), [Linux Plumbers Conference 2019](#)

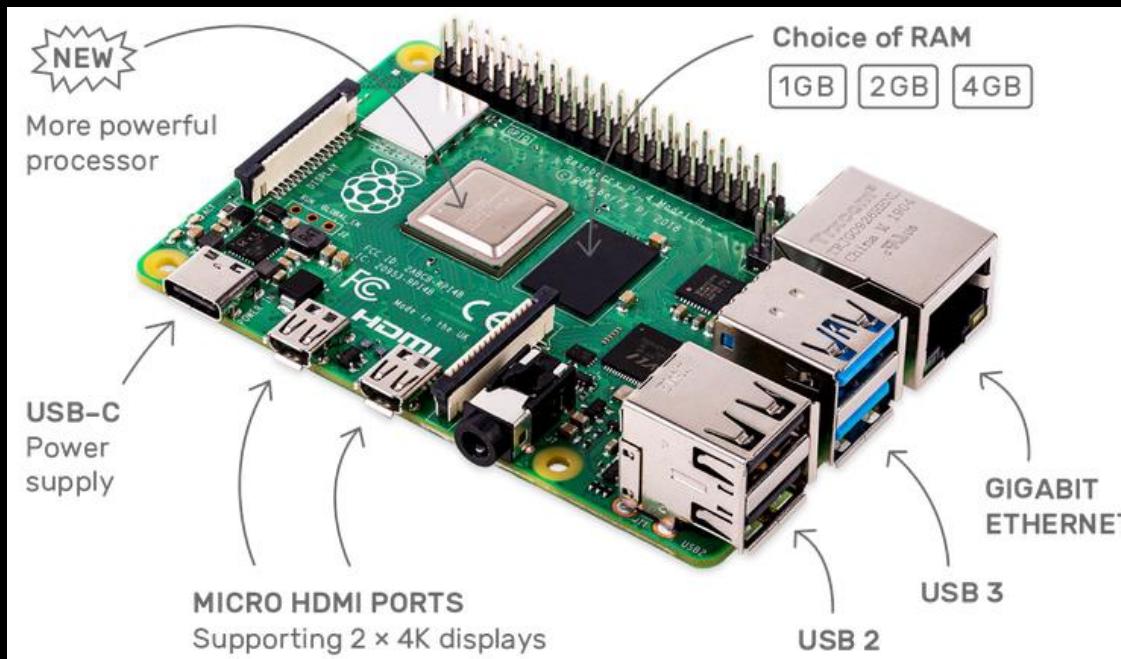
V. Practice on ARM

1) Development Environment

1.1 HW/SW Environment

RPi 4B

- <https://www.raspberrypi.org/products/raspberry-pi-4-model-b/>



■ Model 4B vs 3B+

Features/Specs	Raspberry Pi 4B	Raspberry Pi 3 B+
Release date	24th June 2019	14th March 2018
SoC	Broadcom BCM2711 quad-core Cortex-A72 @ 1.5 GHz	Broadcom BCM2837B0 quad-core Cortex-A53 @ 1.4 GHz
GPU	VideoCore VI with OpenGL ES 1.1, 2.0, 3.0	VideoCore IV with OpenGL ES 1.1, 2.0
Video Decode	H.265 4Kp60, H.264 1080p60	H.264 & MPEG-4 1080p30
Video Encode		H.264 1080p30
Memory	1GB, 2GB, or 4GB LPDDR4	1GB LPDDR2
Storage		microSD card
Video & Audio Output	2x micro HDMI ports up to 4Kp60 3.5mm AV port (composite + audio) MIPI DSI connector	1x HDMI 1.4 port up to 1080p60 3.5mm AV port (composite + audio) MIPI DSI connector
Camera		MIPI CSI connector
Ethernet	Native Gigabit Ethernet	Gigabit Ethernet over USB (300 Mbps max.)
WiFi		Dual band 802.11 b/g/n/ac
Bluetooth	Bluetooth 5.0 + BLE	Bluetooth 4.2 + BLE
USB	2x USB 3.0 + 2x USB 2.0	4x USB 2.0
Expansion		40-pin GPIO header
Power Supply	5V via USB type-C up to 3A 5V via GPIO header up to 3A Power over Ethernet via PoE HAT	5V via micro USB up to 2.5A 5V via GPIO header up to 3A Power over Ethernet via PoE HAT
Dimensions		85x56 mm
Default OS	Raspbian (after June 24, 2019)	Raspbian (after March 2018)
Price	\$35 (1GB RAM), \$45 (2GB RAM), \$55 (4GB RAM)	\$35 (1GB RAM)

Source: <https://www.cnx-software.com/2019/06/24/raspberry-pi-4-vs-pi-3-what-are-the-differences/>

Manjaro

- an open-source Linux distribution based on the Arch Linux operating system
- https://en.wikipedia.org/wiki/Manjaro_Linux
- <https://manjaro.org>

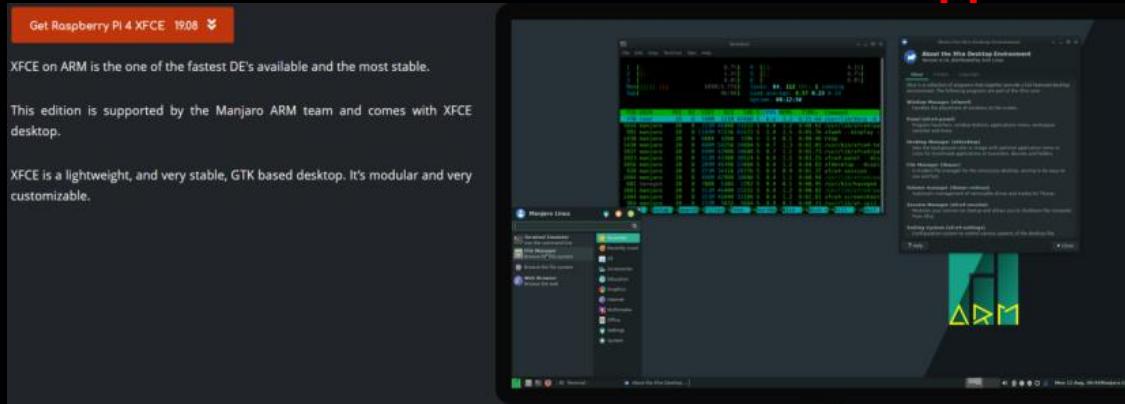
Rank	Distribution	H.PD*
1	Mint	2880▼
2	Debian	1668-
3	Ubuntu	1347-
4	openSUSE	1234▲
5	elementary	1085-
6	Manjaro	1039▼
7	Fedora	982▼
8	Zorin	914-
9	CentOS	773-
10	deepin	736▲

2016-->2019

Rank	Distribution	HPD*
1	MX Linux	4895
2	Manjaro	2597
3	Mint	2048
4	Debian	1543
5	Ubuntu	1398
6	elementary	1302
7	Solus	1087
8	Fedora	992
9	deepin	847
10	Zorin	830

<https://www.distrowatch.com>

- the first ARM64 Linux distribution that support RPi 4



■ Advantages (from my point of view)

- update packages quickly

- group install

<https://www.archlinux.org/groups/>

- Python 3 is the default Python version now

<https://hg.python.org/peps/rev/76d43e52d978>

- much more stable than expected

■ Preparation

```
sudo pacman -Syu                                //update system
pacman -Ss packagename                         //search package

sudo pacman -S --needed base base-devel
sudo pacman -S --needed linux-tools

sudo pacman -S --needed vim colordiff axel      //utilities
sudo pacman -S --needed cmake ninja maven scons   //Build tools
sudo pacman -S --needed valgrind c-ares protobuf  //Dev tools

sudo pacman -S --needed llvm clang lld llc       //LLVM
sudo pacman -S --needed libxml2 doxygen swig subversion //for build LLVM
sudo pacman -S --needed python-pip cython python-netaddr python-pyroute2 //Python packages
sudo pacman -S --needed docker qemu libvirt podman skopeo buildah    //Virtualization

sudo pacman -S --needed iperf ethtool            //Network tools
sudo pacman -S --needed bc rsync                 //for build Kernel

sudo pacman -S --needed rust
sudo pacman -S --needed lua luajit luarocks      //Rust
                                                //Lua
```

<https://wiki.archlinux.org/index.php/Pacman>

■ <https://www.howtogeek.com/430556/why-i-switched-from-ubuntu-to-manjaro-linux/>

1.2 rebuild the Linux Kernel with desired configurations

get Kernel configuration

- sudo modprobe configs

```
zcat /proc/config.gz > old.config //currently v4.19.79-2
```

edit new Kernel configuration

- e.g. edit “**rpi4kernel-5.3_defconfig**” (use **old.config** as the base) by adding following configures (remove the old one if it was configured twice)

```
#  
# eBPF  
#  
CONFIG_CGROUP_BPF=y  
CONFIG_BPF=y  
CONFIG_BPF_SYSCALL=y  
CONFIG_BPF_JIT_ALWAYS_ON=y  
CONFIG_IPV6_SEG6_BPF=y  
# CONFIG_BPFILTER is not set  
CONFIG_NET_CLS_BPF=m  
CONFIG_NET_ACT_BPF=m  
CONFIG_BPF_JIT=y  
CONFIG_BPF_STREAM_PARSER=y  
CONFIG_LWTUNNEL_BPF=y  
CONFIG_HAVE_EBPF_JIT=y  
CONFIG_BPF_LIRC_MODE2=y  
CONFIG_BPF_EVENTS=y  
# CONFIG_TEST_BPF is not set  
# XDP  
CONFIG_XDP_SOCKETS=y  
CONFIG_XDP_SOCKETS_DIAG=m
```

```
#  
# Miscs  
#  
CONFIG_KSM=y  
CONFIG_PROC_EVENTS=y  
#CONFIG_IKHEADERS=y  
CONFIG_REMOTEPROC=m  
CONFIG_REISERFS_PROC_INFO=y  
CONFIG_PROC_CHILDREN=y  
...  
...
```

```
#  
# Debug  
#  
CONFIG_ARCH_HAS_DEBUG_VIRTUAL=y  
CONFIG_STRIP_ASM_SYMS=y  
CONFIG_UNUSED_SYMBOLS=y  
CONFIG_PM_DEBUG=y  
CONFIG_CMA_DEBUGFS=y  
CONFIG_L2TP_DEBUGFS=m  
CONFIG_6LOWPAN_DEBUGFS=y  
CONFIG_CFG80211_DEBUGFS=y  
CONFIG_MAC80211_DEBUGFS=y  
CONFIG_DEBUG_DEVRES=y  
CONFIG_SCSI_DEBUG=m  
CONFIG_DM_DEBUG=y  
CONFIG_DM_DEBUG_BLOCK_MANAGER_LOCKING=y  
CONFIG_ATH9K_DEBUGFS=y  
CONFIG_ATH6KL_DEBUG=y  
CONFIG_SUNRPC_DEBUG=y  
CONFIG_DLM_DEBUG=y  
CONFIG_DYNAMIC_DEBUG=y  
CONFIG_DEBUG_INFO=y  
#CONFIG_DEBUG_INFO_DWARF4=y  
CONFIG_DEBUG_SECTION_MISMATCH=y  
CONFIG_DEBUG_RODATA_TEST=y  
CONFIG_DEBUG_VM=y  
CONFIG_DEBUG_SHIRQ=y  
CONFIG_DEBUG_LIST=y  
CONFIG_ARM64_PTDUMP_DEBUGFS=y  
CONFIG_DEBUG_WX=y  
CONFIG_PROC_KCORE=y  
CONFIG_PROC_VMCORE=y  
CONFIG_PROC_VMCORE_DEVICE_DUMP=y  
CONFIG_FTRACE_SYSCALLS=y  
CONFIG_HWLAT_TRACER=y
```

setup

- `git clone https://github.com/raspberrypi/linux`
//currently, master is branch rpi-4.19.y
`git checkout rpi-5.3.y`
~~`cp rpi4kernel-5.3_defconfig $RPILINUX_SRC/arch/arm64/configs`~~

build and install the new Kernel

- `cd $RPILINUX_SRC`
`make ARCH=arm64 rpi4kernel-5.3_defconfig //generate .config`
`make -j4`
//on RPi 4B with Manjaro-ARM-xfce-rpi4-19.08 + GCC 9.1.0 +
gnu ld 2.32 + 4GB LPDDR4 +  , and JOBS=4:
~1h50m for a full build
`sudo make headers_install modules_install install`
- `cd /boot`
`sudo cp vmlinux-5.3.5-MANJARO-ARM+ kernel8.img`
//same as `$RPILINUX_SRC/arch/arm64/boot/Image`
`sudo reboot`

■ [myrpi4@myrpi4h1 ~]\$ `uname -a`
Linux myrpi4h1 5.3.5-MANJARO-ARM+ #1 SMP PREEMPT Sat Oct 12 14:08:59 UTC 2019 aarch64 GNU/Linux

2) Experiments

2.1 BCC

build

- <https://github.com/iovisor/bcc/blob/master/INSTALL.md>

Required

```
CONFIG_BPF=y
CONFIG_BPF_SYSCALL=y
# [optional, for tc filters]
CONFIG_NET_CLS_BPF=m
# [optional, for tc actions]
CONFIG_NET_ACT_BPF=m
CONFIG_BPF_JIT=y
# [for Linux kernel versions 4.1 through 4.6]
CONFIG_HAVE_BPF_JIT=y
# [for Linux kernel versions 4.7 and later]
CONFIG_HAVE_EBPF_JIT=y
# [optional, for kprobes]
CONFIG_BPF_EVENTS=y
```

Optional

```
CONFIG_NET_SCH_SFQ=m
CONFIG_NET_ACT_POLICE=m
CONFIG_NET_ACT_GACT=m
CONFIG_DUMMY=m
CONFIG_VXLAN=m
```

- **git clone https://github.com/iovisor/bcc
cd \$BCC_SRC
git submodule update --init --recursive
//<https://github.com/iovisor/bcc/tree/master/src/cc>**

The libbpf directory is a git submodule for repository
<https://github.com/libbpf/libbpf>

If you have any change in libbpf directory, please upstream to linux
first as libbpf repo is a mirror of [linux/tools/lib/bpf](#) directory.

- ```
mkdir $BCC_SRC/build; cd $BCC_SRC/build
cmake .. -DCMAKE_INSTALL_PREFIX=/usr
make -j4
sudo make install
```

~~/on RPi 4B with Manjaro-ARM-xfce-rpi4-19.08 + Kernel 5.3.5 +~~  
~~GCC 9.1.0 + gnu ld 2.32 + 4GB LPDDR4 +~~  ~~U3 A2~~, and JOBS=4:  
~11m for a full build

- run an example  

```
cd /usr/share/bcc/examples/tracing
```

```
[myrpi4@myrpi4h1 tracing]$ sudo ./bitehist.py
[sudo] password for myrpi4:
Tracing... Hit Ctrl-C to end.
^C
log2 histogram

 kbytes : count distribution
 0 -> 1 : 0
 2 -> 3 : 0
 4 -> 7 : 35
 8 -> 15 : 16
 16 -> 31 : 0
 32 -> 63 : 2

linear histogram

 kbytes : count distribution
 0 : 0
 1 : 0
 2 : 0
 3 : 0
 4 : 35
 5 : 0
 6 : 0
 7 : 0
 8 : 11
 9 : 0
 10 : 0
 11 : 0
 12 : 5
 13 : 0
 14 : 0
 15 : 0
 16 : 0
```

## 2.2 LISA overview

### ■ \$LISA\_SRC

```
-rw-r--r-- 1 myrp14 myrp14 2275 Sep 15 16:21 CONTRIBUTING.md
-rw-r--r-- 1 myrp14 myrp14 407 Oct 3 12:41 devmode_requirements.txt
drwxr-xr-x 7 myrp14 myrp14 4096 Oct 9 10:22 doc/
drwxr-xr-x 5 myrp14 myrp14 4096 Oct 3 12:41 external/
drwxr-xr-x 8 myrp14 myrp14 4096 Oct 16 13:22 .git/
drwxr-xr-x 3 myrp14 myrp14 4096 Sep 15 16:21 .github/
-rw-r--r-- 1 myrp14 myrp14 221 Sep 15 16:21 .gitignore
-rw-r--r-- 1 myrp14 myrp14 2114 Sep 15 16:21 init_env
-rwrxr-xr-x 1 myrp14 myrp14 11245 Oct 9 10:22 install_base.sh*
lrwxrwxrwx 1 myrp14 myrp14 15 Sep 15 16:21 install_base_ubuntu.sh -> install_base.sh*
drwxr-xr-x 8 myrp14 myrp14 4096 Sep 15 16:21 ipynb/
-rw-r--r-- 1 myrp14 myrp14 11358 Sep 15 16:21 LICENSE.txt
drwxr-xr-x 8 myrp14 myrp14 4096 Oct 9 10:22 lisa/
-rw-r--r-- 1 myrp14 myrp14 3063 Sep 15 16:21 logging.conf
-rw-r--r-- 1 myrp14 myrp14 17727 Sep 15 16:21 .pylintrc
-rw-r--r-- 1 myrp14 myrp14 4951 Sep 15 16:21 README.rst
-rw-r--r-- 1 myrp14 myrp14 577 Oct 3 12:41 .readthedocs.yml
-rw-r--r-- 1 myrp14 myrp14 55 Sep 15 16:21 requirements.txt
-rw-r--r-- 1 myrp14 myrp14 111 Sep 15 16:21 setup.cfg
-rwrxr-xr-x 1 myrp14 myrp14 3763 Oct 9 10:22 setup.py*
drwxr-xr-x 2 myrp14 myrp14 4096 Sep 15 16:21 shell/
-rw-r--r-- 1 myrp14 myrp14 2258 Sep 15 16:21 target_conf.yml
drwxr-xr-x 3 myrp14 myrp14 4096 Oct 2 12:11 tests/
drwxr-xr-x 7 myrp14 myrp14 4096 Oct 14 13:53 tools/
-rw-r--r-- 1 myrp14 myrp14 1265 Sep 15 16:21 .travis.yml
-rw-r--r-- 1 myrp14 myrp14 2440 Sep 15 16:21 Vagrantfile
```

### patching and build

- <https://lisa-linux-integrated-system-analysis.readthedocs.io/en/master/setup.html>  
`sudo ./install_base.sh --install-all`  
//on **RPi 4B with Manjaro-ARM-xfce-rpi4-19.08 + Kernel 5.3.5 + GCC 9.1.0 + gnu ld 2.32 + 4GB LPDDR4 + ~4.5h for our patched build**



# ■ disable Android & remove limitations for some Python packages

```
diff --git a/install_base.sh b/install_base.sh
index fc4b8478..127e7af2 100755
--- a/install_base.sh
+++ b/install_base.sh
@@ -153,15 +153,15 @@ apt_packages=(

pacman-based distributions like Archlinux or its derivatives
pacman_packages=(
- coreutils
- git
- openssh
- base-devel
- wget
- unzip
- python
- python-pip
- python-setuptools
+ #coreutils
+ #git
+ #openssh
+ #base-devel
+ #wget
+ #unzip
+ #python
+ #python-pip
+ #python-setuptools

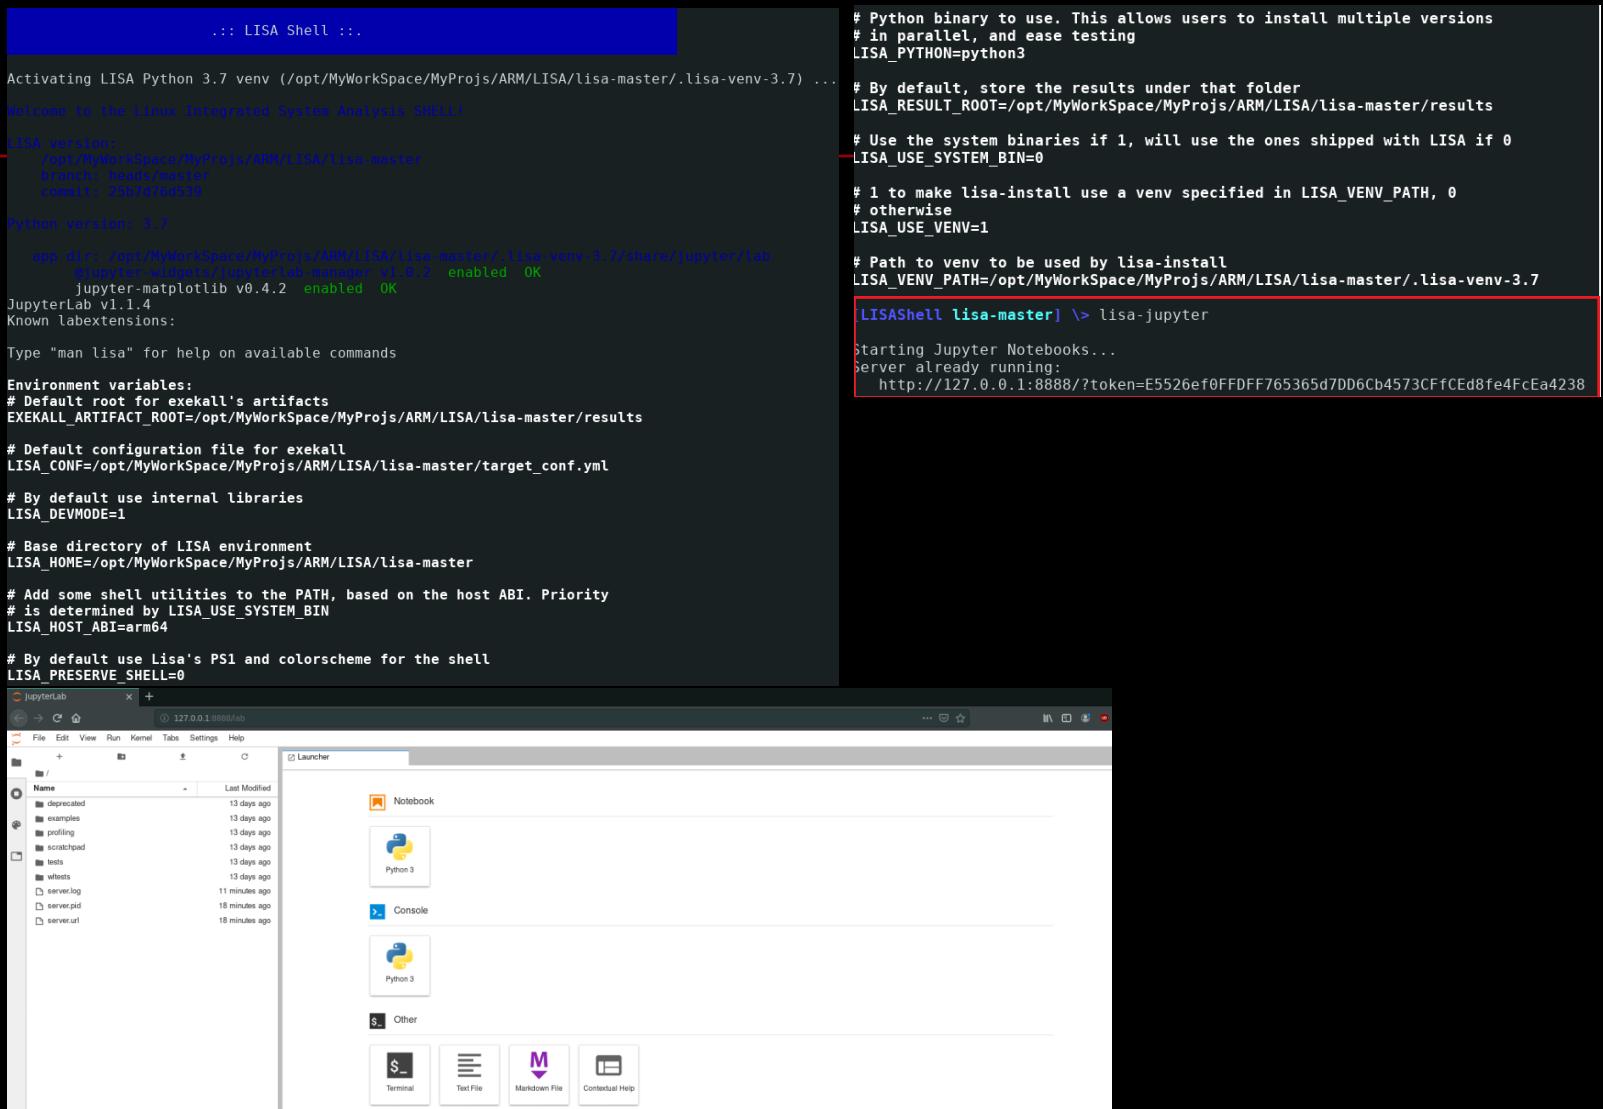
These packages can be installed from AUR
kernelshark
@@ -233,13 +233,13 @@ for arg in "$@"; do

TODO: remove --install-android-sdk, since it is only temporarily there to
give some time to migrate CI scripts
- "--install-android-sdk" | "--install-android-tools" | "--install-all")
+ "--install-android-sdk" | "--install-android-tools")
 install_functions+=(
 install_android_sdk_manager # Needed by install_android_build_tools
 install_android_tools
)
 apt_packages+=($openjdk-$ANDROID_SDK_JAVA_VERSION-jre $openjdk-$ANDROID_SDK_JAVA_VERSION-jdk)
- pacman_packages+=(jreANDROID_SDK_JAVA_VERSION-$openjdk jdkANDROID_SDK_JAVA_VERSION-$openjdk)
+ pacman_packages+=(jreANDROID_SDK_JAVA_VERSION-$openjdk jdkANDROID_SDK_JAVA_VERSION-$openjdk)
```

```
diff --git a/external/devlib/setup.py b/external/devlib/setup.py
index df3377d0..b6af03df 100644
--- a/external/devlib/setup.py
+++ b/external/devlib/setup.py
@@ -84,11 +84,9 @@ params = dict(
 'pyserial', # Serial port interface
 'wrapt', # Basic for construction of decorator functions
 'future', # Python 2-3 compatibility
- 'enum34;python_version<"3.4"', # Enums for Python < 3.4
- 'contextlib2;python_version<"3.0"', # Python 3 contextlib backport for Python 2
- 'numpy<=1.16.4; python_version<"3"',
+ 'enum34', # Enums for Python < 3.4
+ 'contextlib2', # Python 3 contextlib backport for Python 2
- 'numpy; python_version>="3"',
- 'pandas<=0.24.2; python_version<"3"',
+ 'pandas; python_version>="3"',
),
extras_require={

diff --git a/external/trappy/setup.py b/external/trappy/setup.py
index b750aae4..bf0566a1 100644
--- a/external/trappy/setup.py
+++ b/external/trappy/setup.py
@@ -66,7 +66,6 @@ setup(name='TRAPPy',
 "Environment :: Console",
 "License :: OSI Approved :: Apache Software License",
 "Operating System :: POSIX :: Linux",
- "Programming Language :: Python :: 2.7",
- "Programming Language :: Python :: 3",
 "# As we depend on trace data from the Linux Kernel/FTrace
 "Topic :: System :: Operating System Kernels :: Linux",
diff --git a/external/workload-automation/setup.py b/external/workload-automation/setup.py
index 36c697c6..45380032 100755
--- a/external/workload-automation/setup.py
+++ b/external/workload-automation/setup.py
@@ -80,7 +80,6 @@ params = dict(
 maintainer='ARM Architecture & Technology Device Lab',
 maintainer_email='workload-automation@arm.com',
 setup_requires=[
- 'numpy<=1.16.4; python_version<"3"',
- 'numpy; python_version>="3"',
],
 install_requires=[
@@ -93,7 +92,6 @@ params = dict(
 'devlib>={}'.format(devlib_version), # Interacting with devices
 'louie-latest', # callbacks dispatch
 'wrapt', # better decorators
- 'pandas<=0.23.0,<>0.24.2; python_version<="3.5.3"', # Data analysis and manipulation
- 'pandas>=0.23.0; python_version>="3.5.3"', # Data analysis and manipulation
 'future', # Python 2-3 compatibility
),
```

## ■ start the Jupyter Notebook server cd \$LISA\_SRC; source init\_env



Activating LISA Python 3.7 venv (/opt/MyWorkSpace/MyProjs/ARM/LISA/lisa-master/.lisa-venv-3.7) ...  
Welcome to the Linux Integrated System Analysis SHELL!

LISA version:  
/opt/MyWorkSpace/MyProjs/ARM/LISA/lisa-master  
branch: heads/master  
commit: 25b7d75d519

Python version: 3.7

app dir: /opt/MyWorkSpace/MyProjs/ARM/LISA/lisa-master/.lisa-venv-3.7/share/jupyter/lab  
jupyter-widgets/jupyterlab-manager v1.0.2 enabled OK  
jupyter-matplotlib v0.4.2 enabled OK

JupyterLab v1.1.4  
Known labextensions:

Type "man lisa" for help on available commands

Environment variables:  
# Default root for exekall's artifacts  
EXEKALL\_ARTIFACT\_ROOT=/opt/MyWorkSpace/MyProjs/ARM/LISA/lisa-master/results

# Default configuration file for exekall  
LISA\_CONF=/opt/MyWorkSpace/MyProjs/ARM/LISA/lisa-master/target\_conf.yml

# By default use internal libraries  
LISA\_DEVMODE=1

# Base directory of LISA environment  
LISA\_HOME=/opt/MyWorkSpace/MyProjs/ARM/LISA/lisa-master

# Add some shell utilities to the PATH, based on the host ABI. Priority  
# is determined by LISA\_USE\_SYSTEM\_BIN  
LISA\_HOST\_ABI=arm64

# By default use Lisa's PS1 and colorscheme for the shell  
LISA\_PRESERVE\_SHELL=0

```
Python binary to use. This allows users to install multiple versions
in parallel, and ease testing
LISA_PYTHON=python3

By default, store the results under that folder
LISA_RESULT_ROOT=/opt/MyWorkSpace/MyProjs/ARM/LISA/lisa-master/results

Use the system binaries if 1, will use the ones shipped with LISA if 0
LISA_USE_SYSTEM_BIN=0

1 to make lisa-install use a venv specified in LISA_VENV_PATH, 0
otherwise
LISA_USE_VENV=1

Path to venv to be used by lisa-install
LISA_VENV_PATH=/opt/MyWorkSpace/MyProjs/ARM/LISA/lisa-master/.lisa-venv-3.7

[LISAShell lisa-master] >> lisa-jupyter

Starting Jupyter Notebooks...
Server already running:
http://127.0.0.1:8888/?token=E5526ef0FFDF765365d7DD6Cb4573CFFCEd8fe4FcEa4238
```

The terminal shows the activation of the LISA Python 3.7 virtual environment and the configuration of the LISA shell. It also lists installed JupyterLab extensions and environment variables. A red box highlights the command 'lisa-jupyter' and its output, which shows that the Jupyter Notebook server is already running on port 8888.

The JupyterLab interface is shown in a separate window, displaying a file tree with notebooks and a launcher panel containing two Python 3 notebooks and other options like Terminal, Text File, Markdown File, and Contextual Help.

## 2.3 Drgn

- <https://drgn.readthedocs.io/en/latest/installation.html>

### build & install

- build **libkdumpfile** if you want support for kdump-compressed kernel core dumps //Optional
  - cd \$DRGN\_SRC
  - python setup.py build
  - sudo python setup.py install

### run an example

- ```
[myrpi4@myrpi4h1 Drgn]$ sudo drgn
could not get debugging information for:
/boot/vmlinuz-5.3.5-MANJARO-ARM+: not an ELF file
drgn 0.0.1 (using Python 3.7.4, with libkdmpfile)
For help, type help(drgn).
>>> import drgn
>>> from drgn import cast, container_of, execscript, NULL, Object, reinterpret
>>> from drgn.helpers.linux import *
>>> prog['init_task'].comm
Traceback (most recent call last):
  File "<console>", line 1, in <module>
KeyError: 'init_task'
>>> 
```

Solution

```
dwarf_index.c (libdrgn) dwarf_index.h (libdrgn) linux_kernel.c (libdrgn) *
} /* end report_kernel_modules */

static struct drgn_error *
report_default_vmlinux(struct drgn_program *prog,
                      struct drgn_dwarf_index *dindex,
                      bool *vmlinux_is_pending)
{
    static const char * const vmlinux_paths[] = {
        /*
         * The files under /usr/lib/debug should always have debug
         * information, so check for those first.
         */
        "/usr/lib/debug/boot/vmlinux-%s",
        "/usr/lib/debug/lib/modules/%s/vmlinux",
        "/boot/vmlinux-%s",
        "/lib/modules/%s/build/vmlinux",
        "/lib/modules/%s/vmlinux",
        NULL,
    };
}
```

\$RPILINUX_SRC

```
drwxr-xr-x 26 myrpi4 myrpi4 4096 Oct 15 23:59 sound/
-rw-r--r-- 1 myrpi4 myrpi4 3657102 Oct 16 01:12 System.map
-rw-r--r-- 1 myrpi4 myrpi4 1458280 Oct 16 01:11 .tmp_kallsyms1.S
-rw-r--r-- 1 myrpi4 myrpi4 8560443 Oct 16 01:11 .tmp_kallsyms1.S
-rw-r--r-- 1 myrpi4 myrpi4 1458280 Oct 16 01:12 .tmp_kallsyms2.o
-rw-r--r-- 1 myrpi4 myrpi4 8560443 Oct 16 01:12 .tmp_kallsyms2.S
-rw-r--r-- 1 myrpi4 myrpi4 3657102 Oct 16 01:12 .tmp_System.map
-rwxr-xr-x 1 myrpi4 myrpi4 222986776 Oct 16 01:11 .tmp_vmlinux1*
-rwxr-xr-x 1 myrpi4 myrpi4 224428928 Oct 16 01:12 .tmp_vmlinux2*
drwxr-xr-x 36 myrpi4 myrpi4 4096 Oct 15 23:26 tools/
drwxr-xr-x 3 myrpi4 myrpi4 4096 Oct 15 23:34 usr/
-rw-r--r-- 1 myrpi4 myrpi4 2 Oct 16 01:10 .version
drwxr-xr-x 4 myrpi4 myrpi4 4096 Oct 16 00:07 virt/
-rwxr-xr-x 1 myrpi4 myrpi4 224428928 Oct 16 01:12 vmlinux*
-rw-r--r-- 1 myrpi4 myrpi4 134 Oct 16 01:12 vmlinux.cmd
-rw-r--r-- 1 myrpi4 myrpi4 582146608 Oct 16 01:11 vmlinux.o
```

**sudo mkdir -p /usr/lib/debug/boot
sudo cp \$RPILINUX_SRC/linux/vmlinux /usr/lib/debug/boot/
vmlinux-5.3.5-MANJARO-ARM+**

```
[myrpi4@myrpi4h1 /]$ sudo drgn
[sudo] password for myrpi4:
drgn 0.0.1 (using Python 3.7.4, with libkdmpfile)
For help, type help(drgn).
>>> import drgn
>>> from drgn import cast, container_of, execscript, NULL, Object, reinterpret
>>> from drgn.helpers.linux import *
>>> prog['init_task'].comm
(char [16])"swapper/0"
```

```
>>> from drgn.helpers.linux import list_for_each_entry
>>> for mod in list_for_each_entry('struct module',prog['modules'].address_of_(),'list'):
...     if mod.refcnt.counter > 10:
...         print(mod.name)
...
(char [56])"rfcomm"
(char [56])"bluetooth"
(char [56])"ipv6"
>>> 
```

2.4 Rethink extending LISA

Limitations of LISA (from my point of view)

- Bias towards automated testing, which has overlap in functionality with **LAVA**(<https://www.lavasoftware.org/>) from Linaro in some extent
- **Android**-oriented, while it is required to adapt to a much more generalized Linux system
- Lack of plug-in mechanism to integrated existed Kernel analysis and test projects

...

LISA has provided a good base framework and starting point for integrated Linux Kernel analysis.

Try to extend LISA in the following ways

- In addition to Ftrace, combining with eBPF support
- **Architecture redesign (for better support existing Kernel analysis and test utilities)**
e.g. flexible plug-in mechanism and well-designed interfaces for system extension...

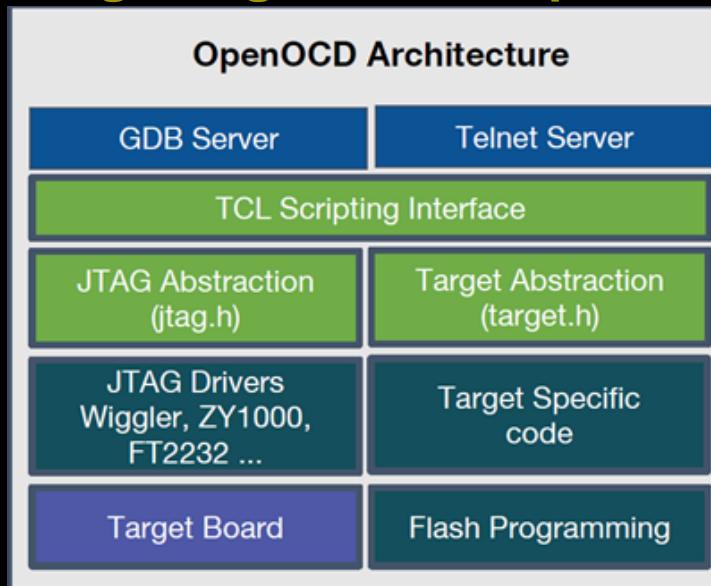
■ **Hardware-assisted debugging**

1. better support hardware debug interface like **JTAG** and **SWD**

<https://en.wikipedia.org/wiki/JTAG>

https://www.arm.com/files/pdf/Serial_Wire_Debug.pdf

2. integrate good on-chip debugger like **OpenOCD**



Source: “OpenOCD support for AArch64 targets”, Omair Javaid, Linaro Connect HK 2018

3. support more mechanism like like **Semihosting**

<https://developer.arm.com/architectures/learn-the-architecture/debugger-usage/semitesting>

■ Full CoreSight support

<https://developer.arm.com/architectures/cpu-architecture/debug-visibility-and-trace/coresight-architecture>

The screenshot displays a grid of six items, each with an icon, a title, a brief description, and a 'Learn more' button.

- Architecture Specifications**
 - Embedded Trace Macrocell Architecture Specification
 - CoreSight Program Flow Trace Architecture Specification
 - CoreSight Architecture Specification

[Learn more](#)
- Arm Debug Interface (ADI) architecture v6**
 - Access to embedded debug functionality

[Learn more](#)
- Arm Debug Interface (ADI) architecture v5**
 - Access to embedded debug functionality

[Learn more](#)
- High Speed Serial Trace Port (HSSTP)**
 - Lower ASIC pin count
 - Increase possible bandwidth
 - Reduce silicon area

[Learn more](#)
- Serial Wire Debug**
 - 2-pin debug port
 - Low pin count
 - High-performance alternative to JTAG

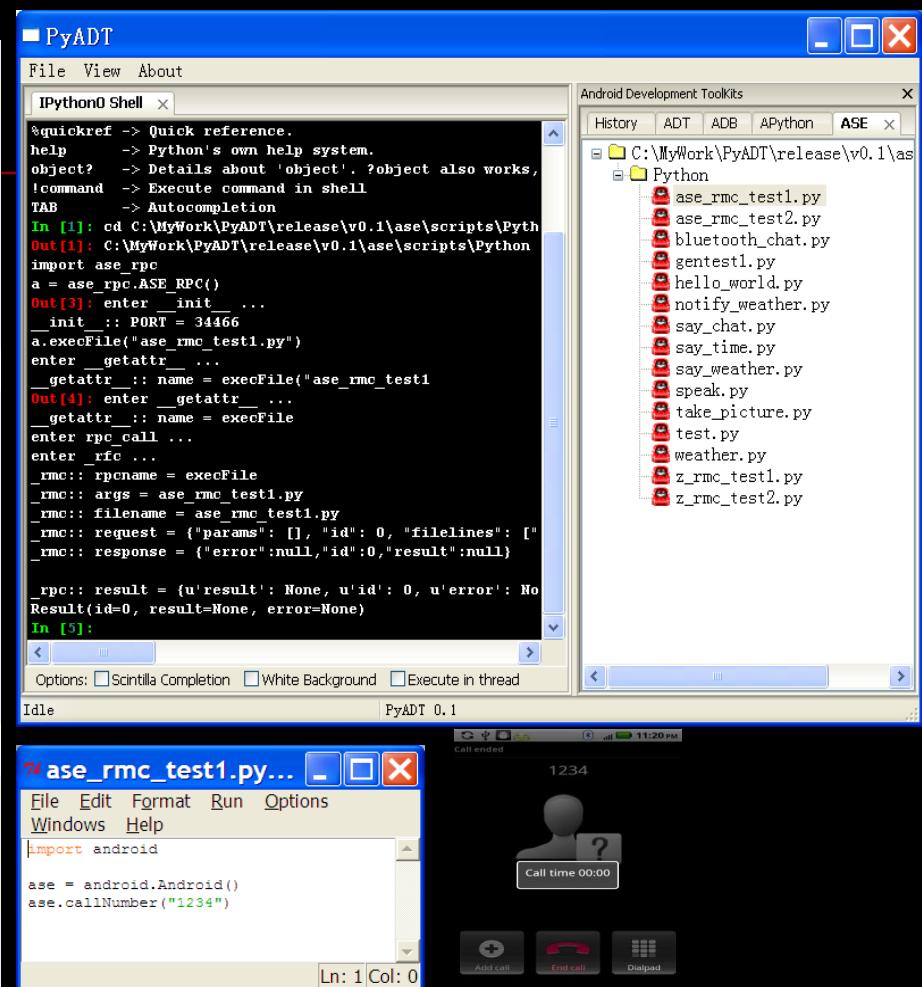
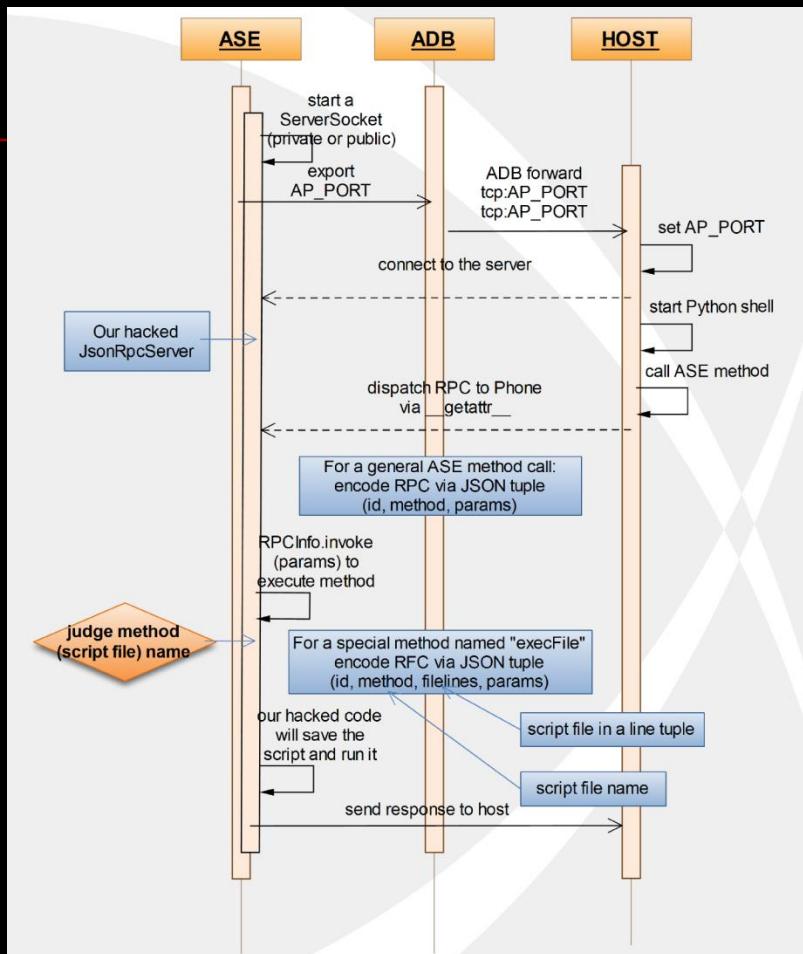
[Learn more](#)
- CoreSight-BSA**
 - Ensures that debug and performance optimization software works seamlessly on all devices.

[Learn more](#)

■ The new project **LISA2** is hosted at:

<https://github.com/Yuwengfeng2019/LISA2>

■ new interactive method for eBPF-based tracing with BCC my old toy “PyADT”

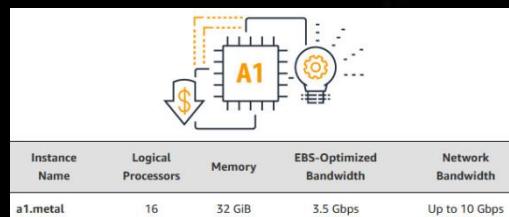


base on ASE, a.k.a SL4A (Scripting Layer for Android)

https://github.com/aosp-caf-upstream/platform_external_sl4a

V. Wrap-up

- Hacking everything with  and eBPF!



Q & A

Thanks!



Reference

Slides/materials from many and varied sources:

- <http://en.wikipedia.org/wiki/>
 - <http://www.slideshare.net/>
 - <https://www.python.org>
 - <http://llvm.org>
 - <http://www.brendangregg.com/>
 - https://en.wikipedia.org/wiki/Just-in-time_compilation
 - <http://dwarfstd.org/>
 - ...
-