

# Kata Containers Meetup

---

## Shanghai 2021

### Kata Containers for Edge Computing

Feng Li (李枫)

hkli2013@126.com

May 22, 2021

# Agenda

## I. Overview

- Edge Computing
  - Kata Containers
  - Testbed
- 

## II. LF Edge

- Overview
- EVE

## III. Lightweight Kubernetes

- K3S
- Arkade
- Sandboxed Containers Operator

## IV. Project ACRN

- Overview
- X86 only

## V. Rethinking Infrastructure for Edge Cloud

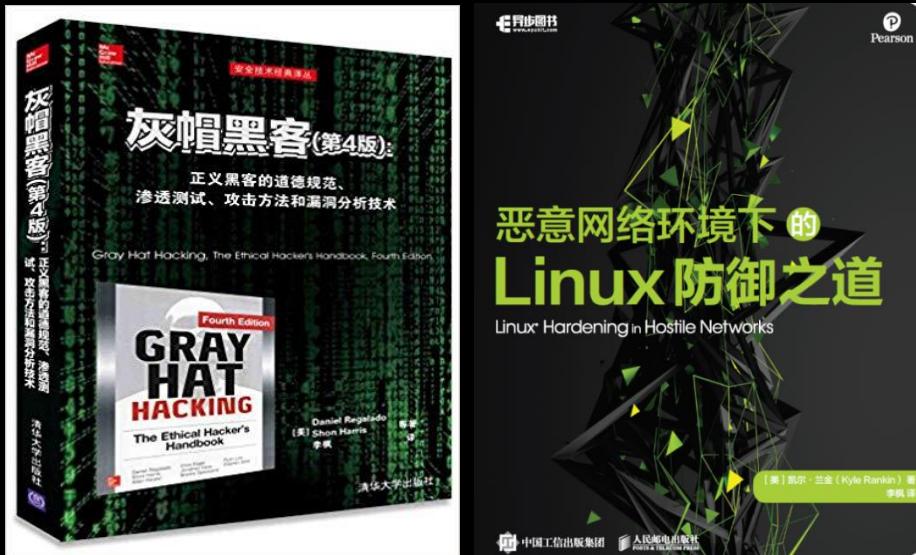
- Linux

- eBPF-centric
  - Rust
  - WebAssembly
  - Edge AI
  - Serverless
  - Unified Runtime
  - Autonomous Vehicles
  - RISC-V
  - HW-SW Co-designed System Security
- 

## VI. Wrap-up

## Who Am I

- The main translator of the book «Gray Hat Hacking The Ethical Hacker's Handbook, Fourth Edition» (ISBN: 9787302428671) & «Linux Hardening in Hostile Networks, First Edition» (ISBN: 9787115544384)

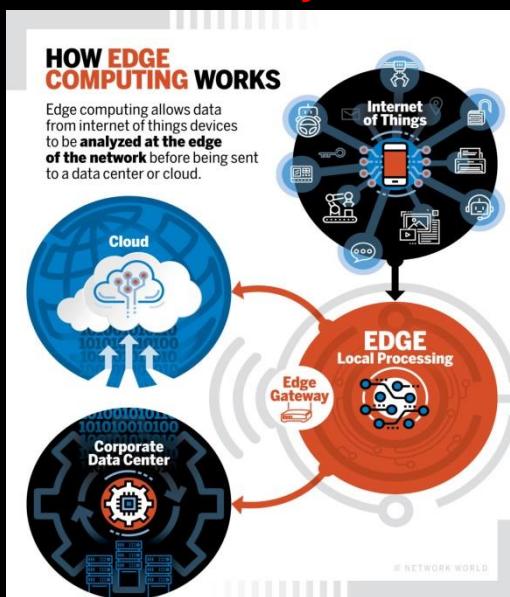


- Actively participate in various activities of the open source community
- <https://github.com/XianBeiTuoBaFeng2015/MySlides>
- Recently, focus on infrastructure of Edge Computing, AI, RISC-V, EDA...

# I. Overview

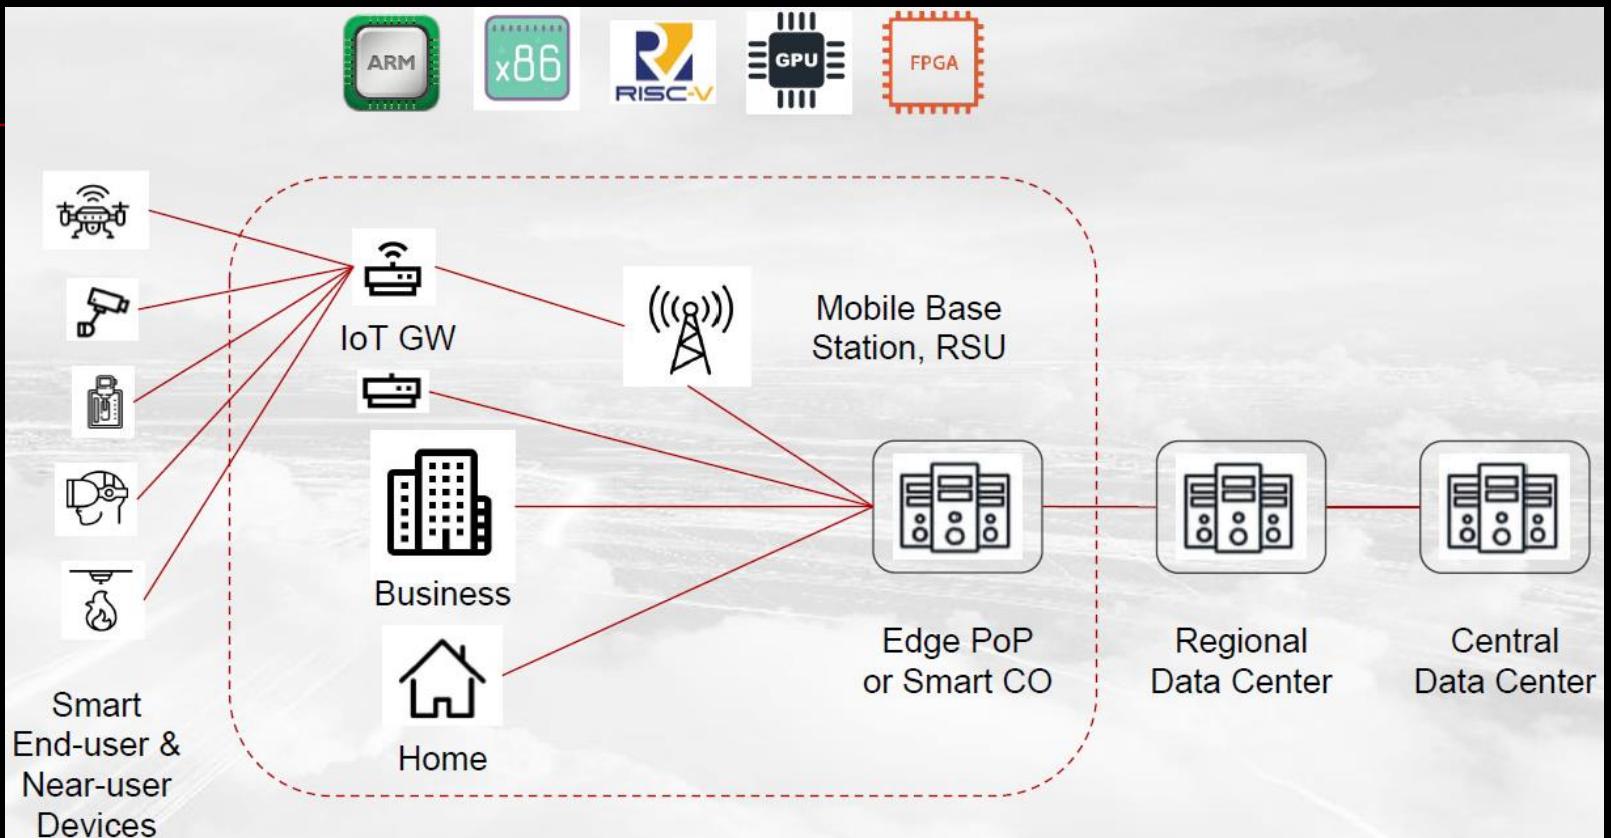
## 1) Edge Communicating

- [https://en.wikipedia.org/wiki/Edge\\_computing](https://en.wikipedia.org/wiki/Edge_computing)  
a distributed computing paradigm which brings computation and data storage closer to the location where it is needed, to improve response times and save bandwidth....
- <https://www.networkworld.com/article/3224893/what-is-edge-computing-and-how-it-s-changing-the-network.html>  
a way to streamline the flow of traffic from IoT devices and provide real-time local data analysis



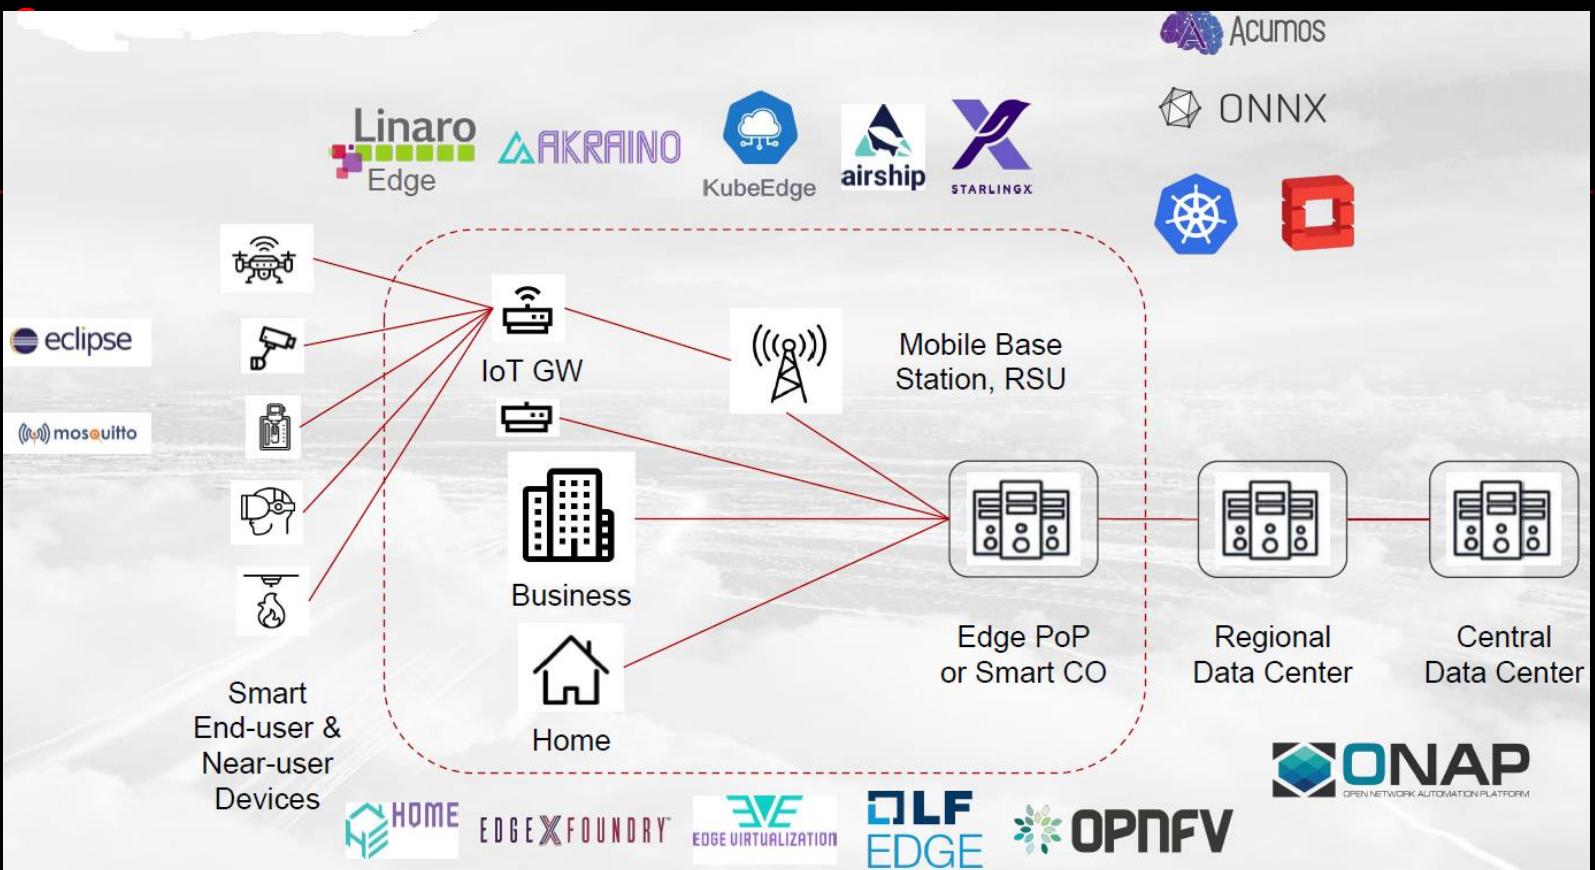
# 1.1 Intelligent Edge with Hardware, Software and Data

## Hardware



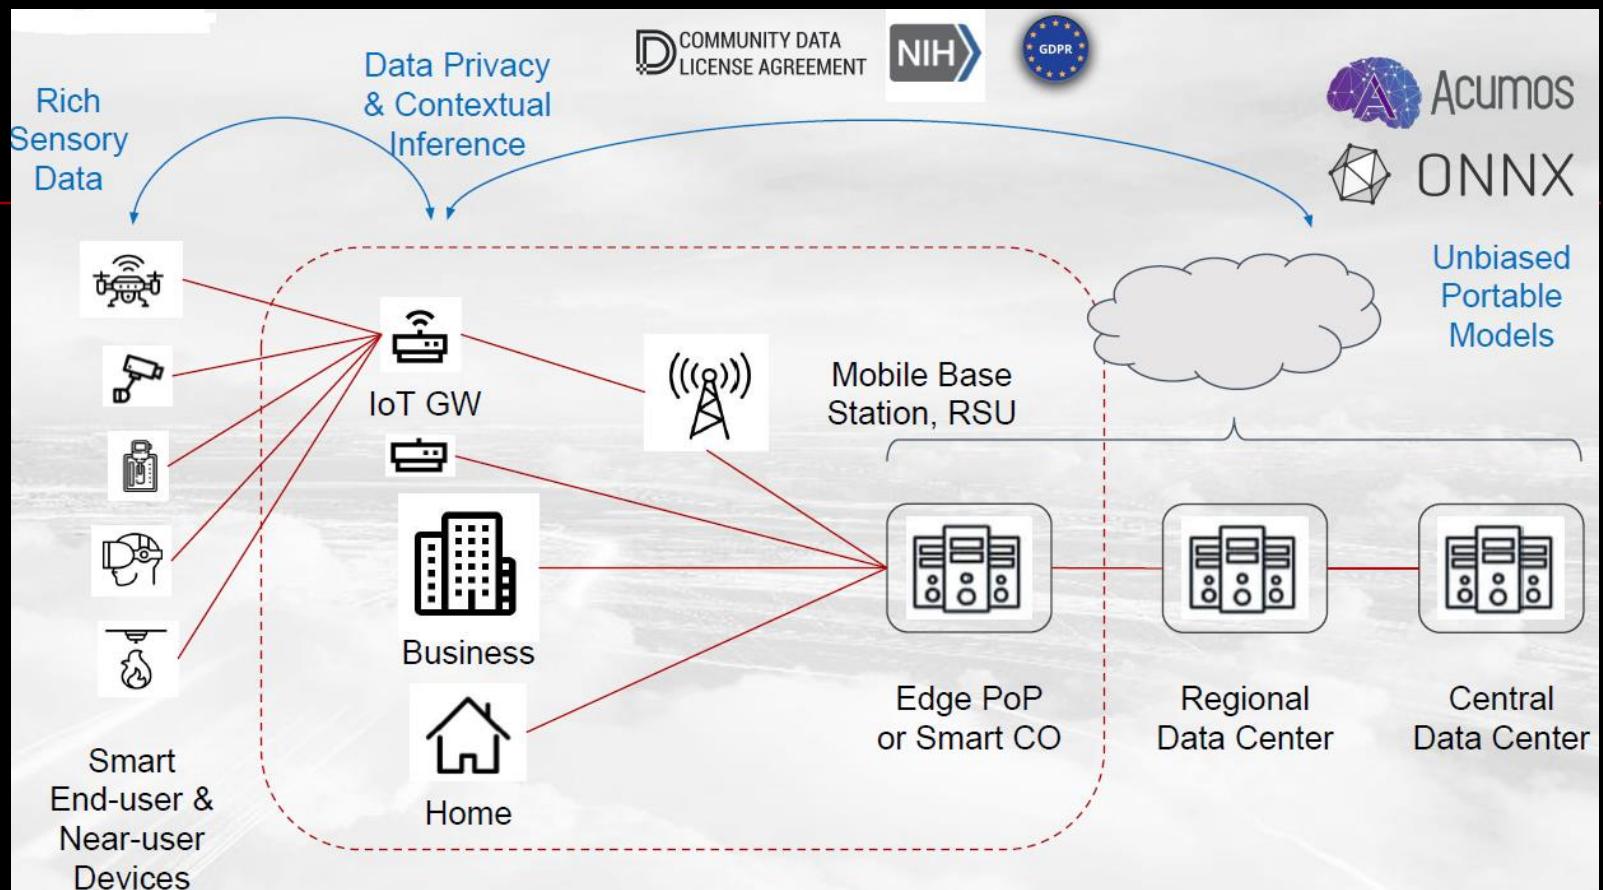
Source: “Deploying the New Intelligent Edge with Open Hardware, Software and Data”,  
Wenjing Chu, ONS North America 2019

## Software



Source: “Deploying the New Intelligent Edge with Open Hardware, Software and Data”,  
Wenjing Chu, ONS North America 2019

# Data

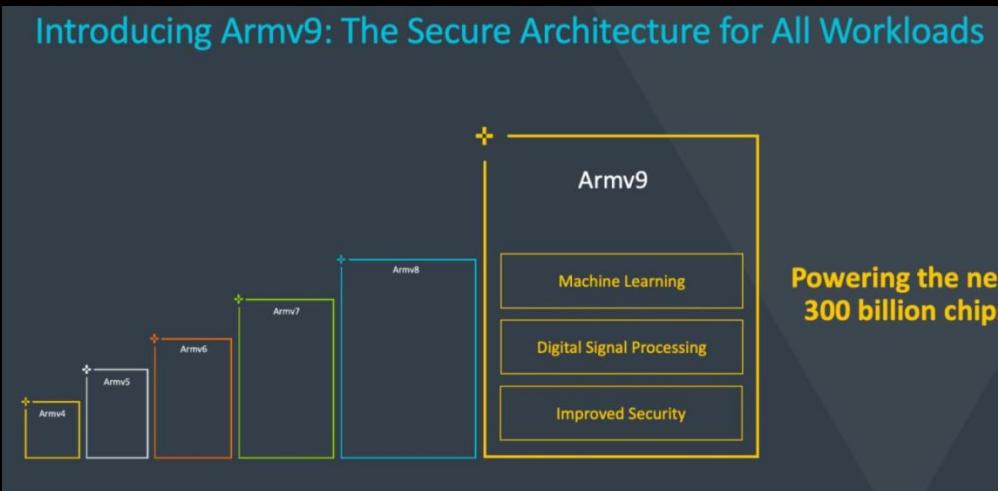
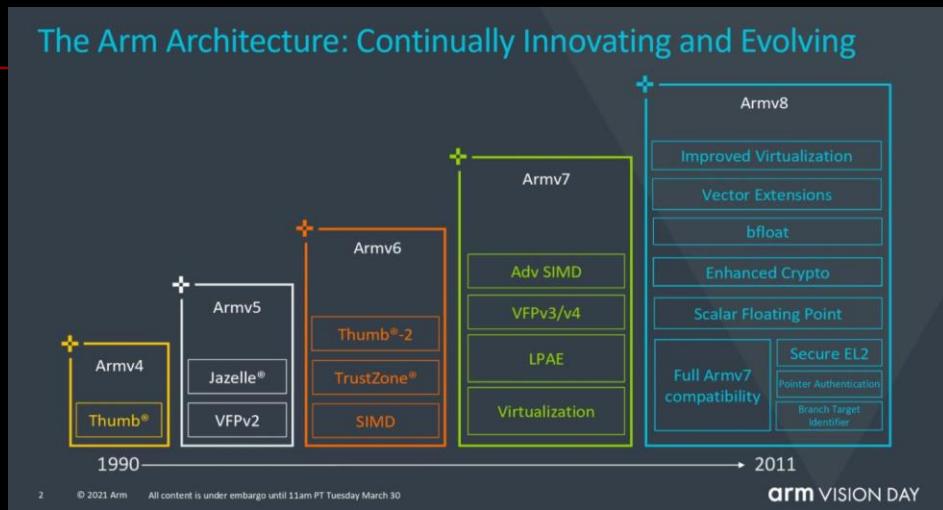


Source: “Deploying the New Intelligent Edge with Open Hardware, Software and Data”,  
Wenjing Chu, ONS North America 2019

# 1.2 ARM Ecosystem for Edge Computing in 2021

## ARM v9

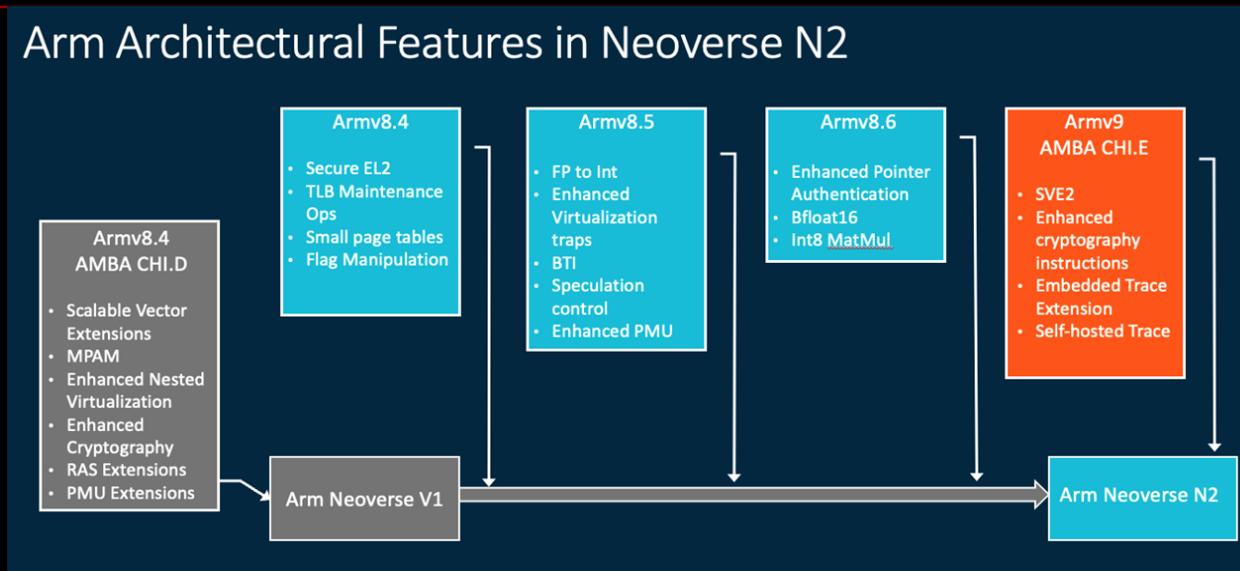
- <https://www.arm.com/campaigns/arm-vision>



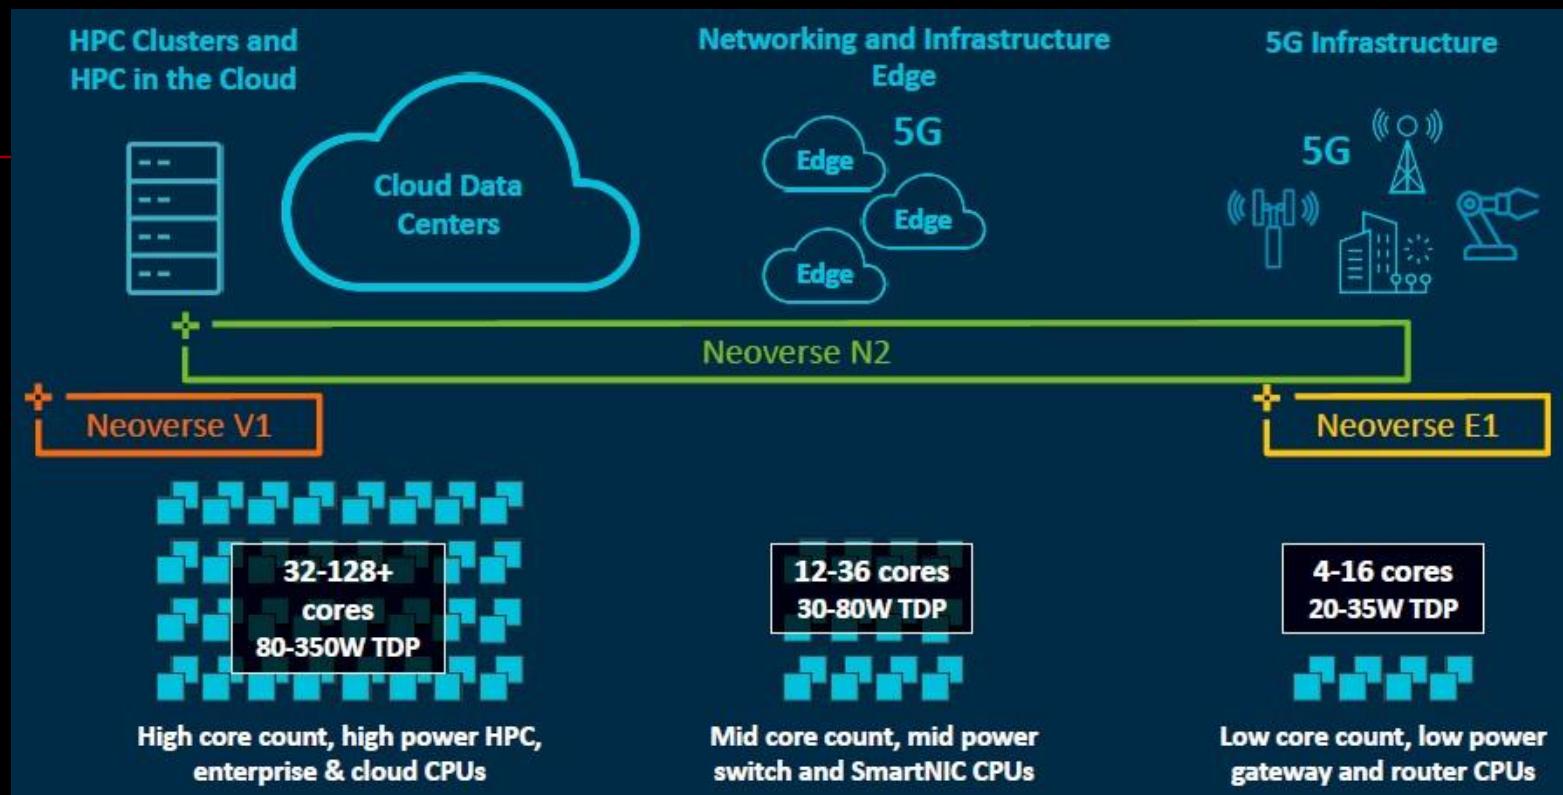
## ***Neoverse N2***

- <https://community.arm.com/developer/ip-products/processors/b/processors-ip-blog/posts/arm-neoverse-n2-industry-leading-performance-efficiency>

### Arm Architectural Features in Neoverse N2

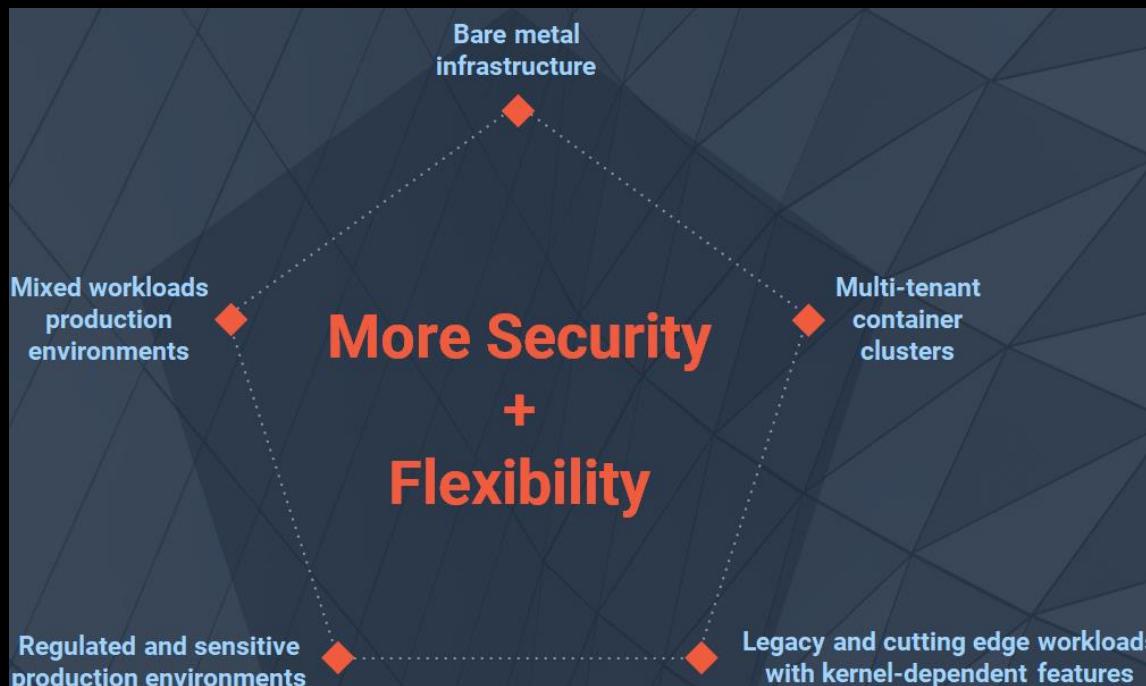


- <https://www.nextplatform.com/2021/04/27/arm-puts-some-muscle-into-future-neoverse-server-cpu-designs/>



## 2) Kata Containers

- <https://katacontainers.io/>
- a standard implementation of lightweight Virtual Machines (VMs) that feel and perform like containers, but provide the workload isolation and security advantages of VMs.

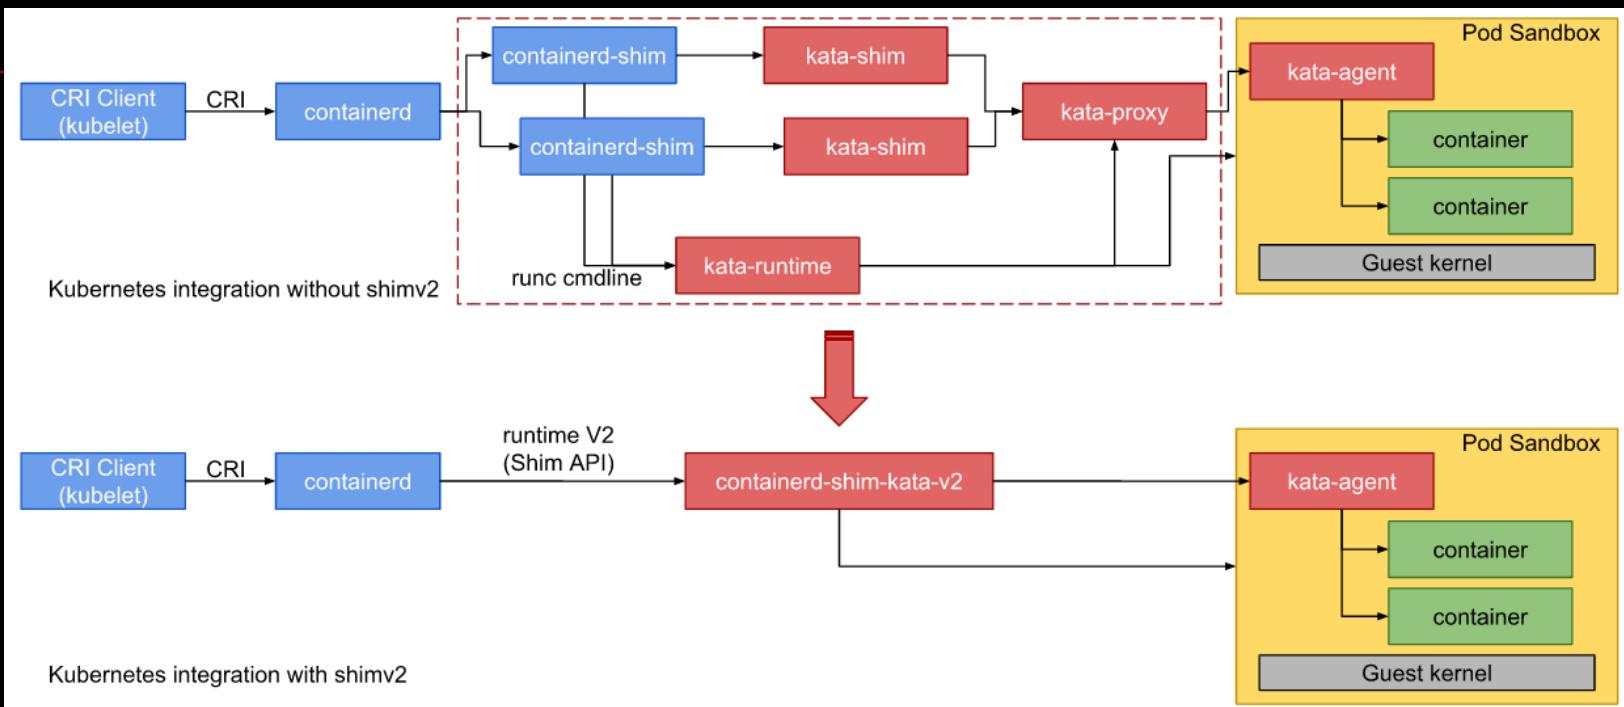


Source: <https://katacontainers.io/collateral/kata-containers-onboarding-deck.pptx>

- <https://medium.com/kata-containers/kata-containers-2020-annual-report-c9ad7f44e0d4>

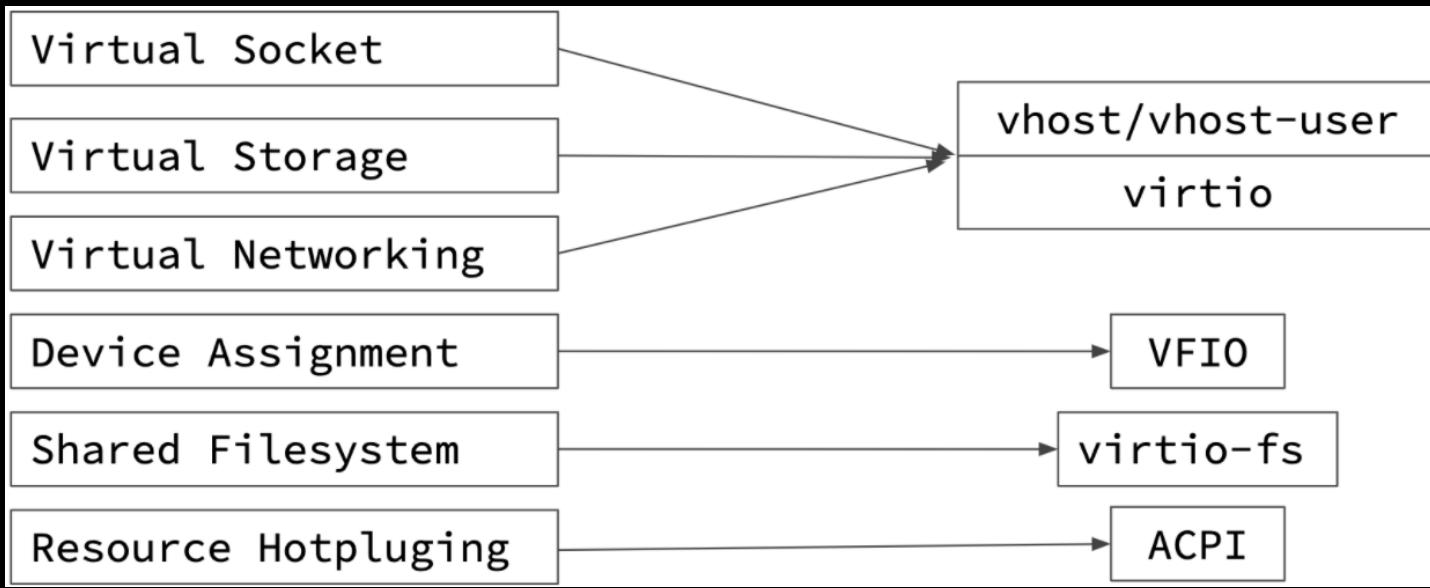
## Architecture

- <https://github.com/kata-containers/documentation/blob/master/design/architecture.md>



## Virtualization

- <https://github.com/kata-containers/documentation/blob/master/design/virtualization.md>
- Hypervisor/VMM support:
  - ACRN hypervisor
  - Cloud Hypervisor/KVM
  - Firecracker/KVM
  - QEMU/KVM
- **Each hypervisor/VMM varies on how or if it handles each of the specific para-virtualized devices or virtualization technologies as below:**



## Cloud Native

- Kata Containers runtime is **OCI-compatible**, **CRI-O-compatible**, and **Containerd-compatible(CNI, CSI...)**, allowing it to work seamlessly with both Docker and Kubernetes respectively
- <https://github.com/kata-containers/documentation/blob/master/how-to/how-to-use-k8s-with-cri-containerd-and-kata.md>
-

## 2.x

- <https://github.com/kata-containers/kata-containers/releases/tag/2.0.0>
- <https://medium.com/kata-containers/kata-containers-version-2-0-e45df4dd328>

### **What's New**

- Kata-agent has been rewritten in rust to significantly reduce memory overhead overall attack surface.
- Agent protocol has been simplified to use ttRPC from grpc.
- Component called Kata-monitor has been introduced to improve observability and manageability
- New component called agent-ctl has been introduced to help validate agent API
- We have moved to a mono-repo model, all code and document repositories consolidated into one single repo.
- Virtio-fs is now the default shared file system type which mean better POSIX compliance compared to virtio-9p
- Latest Cloud Hypervisor support on par with QEMU
- Move to support solely shimv2 api to simplify code and reduce attack surface further. This also meanskata-shim and kata-proxy used in 1.x are no longer required.
- Guest kernel updated to v5.4.71
- QEMU updated to v5.0.0

Source: <https://openanolis-downloads.oss-cn-beijing.aliyuncs.com/meetup/kata-update-2020-11.pdf>

- <https://www.openstack.org/videos/summits/virtual/Observability-in-Kata-containers-2.0>

## Security

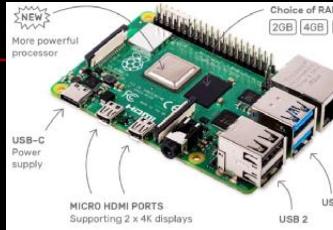
- <https://github.com/kata-containers/documentation/blob/master/Limitations.md>
- <https://containerjournal.com/topics/container-security/kata-container-security-is-good-but-theres-an-achilles-heel/>
- “Escaping Virtualized Containers”, Yuval Avrahami, BlackHat USA 2020
- <https://www.bugcrowd.com/blog/big-bugs-cve-2020-28914/>
- <https://github.com/kata-containers/community/tree/master/VMT/KCSA>

## Getting Started

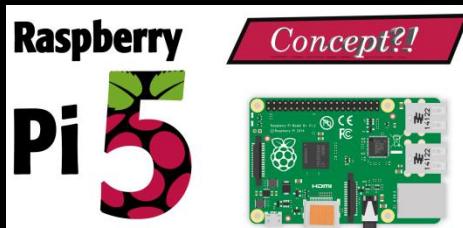
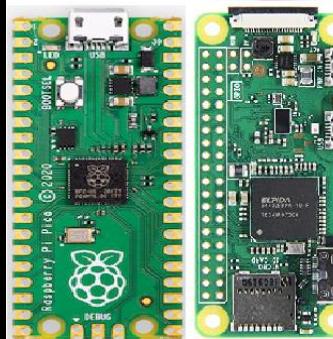
- <https://github.com/kata-containers/kata-containers#getting-started>
- <https://github.com/kata-containers/kata-containers/tree/main/docs>
- <https://github.com/kata-containers/kata-containers/tree/main/docs/install>
- <https://github.com/kata-containers/kata-containers/blob/main/docs/install/README.md#packaged-installation-methods>
- <https://docs.01.org/clearlinux/latest/tutorials/kata.html>
- <https://github.com/kata-containers/kata-containers/tree/main/docs/install#build-from-source-installation>
- <https://github.com/kata-containers/documentation/blob/master/how-to/containerd-kata.md>

### 3) Testbed Raspberry Pi

- <https://www.raspberrypi.org/products/raspberry-pi-4-model-b/>



Features/Specs	Raspberry Pi 4B	Raspberry Pi 3 B+
Release date	24th June 2019	14th March 2018
SoC	Broadcom BCM2711 quad-core Cortex-A72 @ 1.5 GHz	Broadcom BCM2837B0 quad-core Cortex-A53 @ 1.4 GHz
GPU	VideoCore VI with OpenGL ES 1.1, 2.0, 3.0	VideoCore IV with OpenGL ES 1.1, 2.0
Video Decode	H.265 4Kp60, H.264 1080p60	H.264 & MPEG-4 1080p30
Video Encode		H.264 1080p30
Memory	1GB, 2GB, or 4GB LPDDR4	1GB LPDDR2
Storage		microSD card
Video & Audio Output	2x micro HDMI ports up to 4Kp60 3.5mm AV port (composite + audio) MIPI DSI connector	1x HDMI 1.4 port up to 1080p60 3.5mm AV port (composite + audio) MIPI DSI connector
Camera		MIPI CSI connector
Ethernet	Native Gigabit Ethernet	Gigabit Ethernet over USB (300 Mbps max.)
WiFi		Dual band 802.11 b/g/n/ac
Bluetooth	Bluetooth 5.0 + BLE	Bluetooth 4.2 + BLE
USB	2x USB 3.0 + 2x USB 2.0	4x USB 2.0
Expansion		40-pin GPIO header
Power Supply	5V via USB type-C up to 3A 5V via GPIO header up to 3A Power over Ethernet via PoE HAT	5V via micro USB up to 2.5A 5V via GPIO header up to 3A Power over Ethernet via PoE HAT
Dimensions		85x56 mm
Default OS	Raspbian (after June 24, 2019)	Raspbian (after March 2018)
Price	\$35 (1GB RAM), \$45 (2GB RAM), \$55 (4GB RAM)	\$35 (1GB RAM)

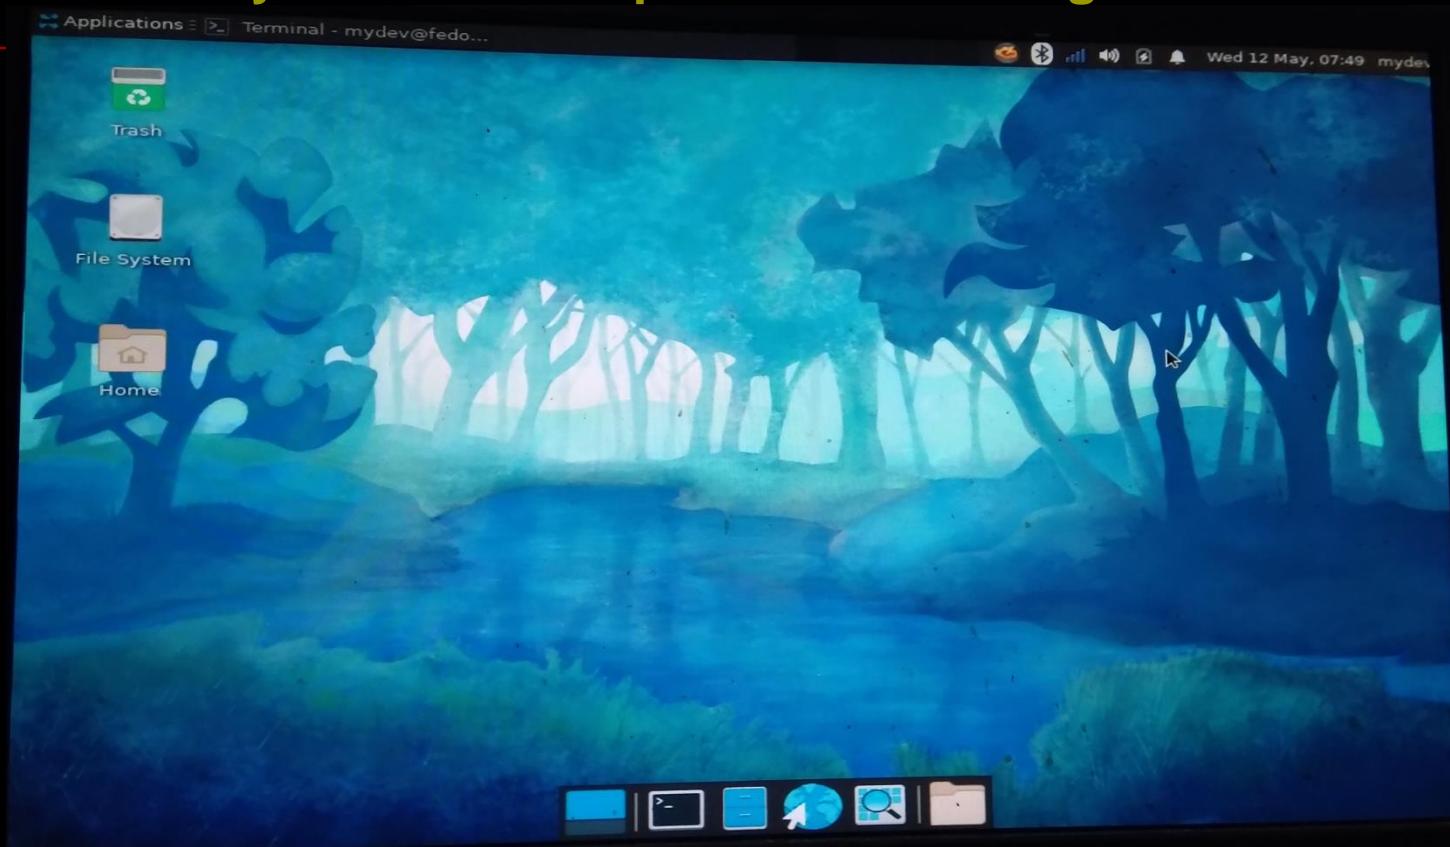


### 3.1 Fedora

- a Linux distribution developed by the community-supported Fedora Project which is sponsored primarily by Red Hat
- [https://en.wikipedia.org/wiki/Fedora\\_\(operating\\_system\)](https://en.wikipedia.org/wiki/Fedora_(operating_system))
- <https://getfedora.org/>
- <https://alt.fedoraproject.org/alt/>
- <https://spins.fedoraproject.org/>
- <https://fedoraproject.org/wiki/Architectures/ARM>
- <https://fedoramagazine.org/>
- <https://silverblue.fedoraproject.org/>
- Developer friendly!

## Fedora 34 AArch64 XFCE

- <https://fedoramagazine.org/announcing-fedora-34/>
- <https://archives.fedoraproject.org/pub/fedora-secondary/releases/34/Spins/aarch64/images/>



- [mydev@fedora ~]\$ uname -a  
Linux fedora 5.11.20-300.fc34.aarch64 #1 SMP Wed May 12 12:27:11 UTC 2021 aarch64 aarch64 aarch64 GNU/Linux

```
[mydev@fedora ~]$ cat /boot/config-5.11.20-300.fc34.aarch64 |grep -i kvm
CONFIG_KVM=y
CONFIG_HAVE_KVM_IRQCHIP=y
CONFIG_HAVE_KVM_IRQFD=y
CONFIG_HAVE_KVM_IRQ_ROUTING=y
CONFIG_HAVE_KVM_EVENTFD=y
CONFIG_KVM_MMIO=y
CONFIG_HAVE_KVM_MSIX=y
CONFIG_HAVE_KVM_CPU_RELAX_INTERCEPT=y
CONFIG_KVM_VFI0=y
CONFIG_HAVE_KVM_ARCH_TLB_FLUSH_ALL=y
CONFIG_KVM_GENERIC_DIRTYLOG_READ_PROTECT=y
CONFIG_HAVE_KVM_IRQ_BYPASS=y
CONFIG_HAVE_KVM_VCPU_RUN_PID_CHANGE=y
[mydev@fedora ~]$
[mydev@fedora ~]$ cat /boot/config-5.11.20-300.fc34.aarch64 |grep -i virtual
CONFIG_VIRTUALIZATION=y
# Virtual GPIO drivers
# end of Virtual GPIO drivers
CONFIG_REGULATOR_VIRTUAL_CONSUMER=m
CONFIG_FB_VIRTUAL=m
CONFIG_DMA_VIRTUAL_CHANNELS=y
CONFIG_ARCH_HAS_DEBUG_VIRTUAL=y
# CONFIG_DEBUG_VIRTUAL is not set
[mydev@fedora ~]$
```

```
[mydev@fedora ~]$ cat /boot/config-5.11.20-300.fc34.aarch64 |grep -i bpt
CONFIG_CGROUP_BPF=y
CONFIG_BPF=y
CONFIG_BPF_LSM=y
CONFIG_BPF_SYSCALL=y
CONFIG_ARCH_WANT_DEFAULT_BPF_JIT=y
CONFIG_BPF_JIT_ALWAYS_ON=y
CONFIG_BPF_JIT_DEFAULT_ON=y
CONFIG_BPF_PRELOAD=y
CONFIG_BPF_PRELOAD_UMD=m
CONFIG_IPV6_SEG6_BPF=y
CONFIG_NETFILTER_XT_MATCH_BPF=m
# CONFIG_BPFILTER is not set
CONFIG_NET_CLS_BPF=m
CONFIG_NET_ACT_BPF=m
CONFIG_BPF_JIT=y
CONFIG_BPF_STREAM_PARSER=y
CONFIG_LWTUNNEL_BPF=y
CONFIG_HAVE_EBPF_JIT=y
CONFIG_BPF_LIRC_MODE2=y
CONFIG_BPF_EVENTS=y
# CONFIG_BPF_KPROBE_OVERRIDE is not set
# CONFIG_TEST_BPF is not set
[mydev@fedora ~]$
[mydev@fedora ~]$ cat /boot/config-5.11.20-300.fc34.aarch64 |grep -i cgroup
CONFIG_CGROUPS=y
CONFIG_BLK_CGROUP=y
CONFIG_CGROUP_WRITEBACK=y
CONFIG_CGROUP_SCHED=y
CONFIG_CGROUP_PIDS=y
# CONFIG_CGROUP_RDMA is not set
CONFIG_CGROUP_FREEZER=y
# CONFIG_CGROUP_HUGETLB is not set
CONFIG_CGROUP_DEVICE=y
CONFIG_CGROUP_CPUACCT=y
CONFIG_CGROUP_PERF=y
CONFIG_CGROUP_BPF=y
# CONFIG_CGROUP_DEBUG is not set
CONFIG_SOCK_CGROUP_DATA=y
CONFIG_BLK_CGROUP_RWSTAT=y
CONFIG_BLK_CGROUP_IOLATENCY=y
CONFIG_BLK_CGROUP_IOCOST=y
# CONFIG_BFQ_CGROUP_DEBUG is not set
CONFIG_NETFILTER_XT_MATCH_CGROUP=m
CONFIG_NET_CLS_CGROUP=y
CONFIG_CGROUP_NET_PRIO=y
CONFIG_CGROUP_NET_CLASSID=y
[mydev@fedora ~]$
```

## ■ Three years ago build Kata Containers 1.0.0 from source on Raspberry Pi 3B+ with Fedora Minimal 28 AArch64 & Go 1.10.3

```
2018-06-19 00:36:39  CLEAN    clean
2018-06-19 00:36:42  GENERATE cli/config-generated.go
2018-06-19 00:36:42  CONFIG   data/kata-collect-data.sh
2018-06-19 00:36:42  kata-runtime - version 1.0.0 (commit 42821b7c0a572bb6b1497e6a2b9a3ad6301c09bb)
2018-06-19 00:36:42
2018-06-19 00:36:42  • architecture:
2018-06-19 00:36:42    Host: aarch64
2018-06-19 00:36:42    golang: arm64
2018-06-19 00:36:42    Build: arm64
2018-06-19 00:36:42
2018-06-19 00:36:42  • golang:
2018-06-19 00:36:42    go version gol.10.3 linux/arm64
2018-06-19 00:36:42
2018-06-19 00:36:42  • Summary:
2018-06-19 00:36:42
2018-06-19 00:36:42    binary install path (DESTTARGET)      : /usr/local/bin/kata-runtime
2018-06-19 00:36:42    config install path (DESTCONFIG)       : /usr/share/defaults/kata-containers/configuration.toml
2018-06-19 00:36:42    alternate config path (DESTSYSCONFIG) : /etc/kata-containers/configuration.toml
2018-06-19 00:36:42    hypervisor path (QEMUPATH)          : /usr/bin/qemu-system-aarch64
2018-06-19 00:36:42    assets path (PKGDATA DIR)        : /usr/share/kata-containers
2018-06-19 00:36:42    proxy+shim path (PKGLIBEXEC DIR) : /usr/libexec/kata-containers
2018-06-19 00:36:42
2018-06-19 00:36:42  BUILD   /opt/MyWorkSpace/MyProjs/Virtual/VM-Container/Kata/src/github.com/kata-containers/runtime
/kata-runtime
2018-06-19 00:38:11  CONFIG   cli/config/configuration.toml
```

Source: “OpenStack on ARM”, Feng Li, OpenInfra Days 2018(Beijing)

## ■ Nowadays

```
[mydev@fedora ~]$ dnf search kata*
Last metadata expiration check: 21:12:50 ago on Sun 09 May 2021 09:49:21 AM PDT.
=====
Name & Summary Matched: kata* =====
kata-agent.aarch64 : Kata guest agent
kata-containers.aarch64 : Kata Containers version 2.x repository
kata-ksm-throttler.aarch64 : Kata KSM throttling daemon
kata-osbuilder.aarch64 : Kata guest initrd and image build scripts
kata-runtime.aarch64 : Kata runtime to run containers in virtual machines
=====
Name Matched: kata* =====
kata-proxy.aarch64 : Proxy for Kata Containers
kata-shim.aarch64 : A shim running in the host for the Kata Containers project
[mydev@fedora ~]$
```

## Kata Containers 2.x

- **sudo dnf install libvirt qemu-kvm oci-kvm-hook virt-install genisoimage libvirt-client libvirt-python3 libvirt-daemon-kvm libvirt-daemon-config-network fakeroot policycoreutils-python-utils libguestfs qemu-user-static kpartx jq podman skopeo buildah wine...**
- **<https://docs.docker.com/engine/install/fedora/>**

```
[mydev@fedora ~]$ which kata-runtime
/usr/bin/kata-runtime
[mydev@fedora ~]$
[mydev@fedora ~]$ kata-runtime check
/usr/share/kata-containers/defaults/configuration.toml: file /usr/bin/qemu-kvm does not exist
[mydev@fedora ~]$
[mydev@fedora ~]$ cd /usr/bin
[mydev@fedora bin]$ sudo ln -sf qemu-system-aarch64 qemu-kvm
[mydev@fedora bin]$
[mydev@fedora bin]$ ls -l qemu-kvm
lrwxrwxrwx. 1 root root 19 May 19 02:15 qemu-kvm -> qemu-system-aarch64
[mydev@fedora bin]$
[mydev@fedora bin]$ kata-runtime check
Invalid command "check"
[mydev@fedora bin]$ kata-runtime kata-check
Newer minor release available: 2.1.0 (url: https://github.com/kata-containers/kata-containers/releases/download/2.1.0/kata-static-2.1.0-x86\_64.tar.xz, date: 2021-05-14T18:37:59Z)
WARN[0001] modprobe insert module failed: modprobe: ERROR: could not insert 'vhost_net': Operation not permitted arch=arm64 error="exit status 1" module=vhost_net name= pid=19893 source=runtime
ERROR[0001] kernel property not found arch=arm64 description="Host kernel accelerator for virtio network" name=vhost_net pid=19893 source=runtime type=module
WARN[0001] modprobe insert module failed: modprobe: ERROR: could not insert 'vhost_vsock': Operation not permitted arch=arm64 error="exit status 1" module=vhost_vsock name= pid=19893 source=runtime
ERROR[0001] kernel property not found arch=arm64 description="Host Support for Linux VM Sockets" name=vhost_vsock pid=19893 source=runtime type=module
WARN[0001] modprobe insert module failed: modprobe: ERROR: could not insert 'vhost': Operation not permitted arch=arm64 error="exit status 1" module=vhost name= pid=19893 source=runtime
ERROR[0001] kernel property not found arch=arm64 description="Host kernel accelerator for virtio" name=vhost pid=19893 source=runtime type=module
ERROR[0001] System is not capable of running Kata Containers arch=arm64 name= pid=19893 source=runtime
ERROR: System is not capable of running Kata Containers
[mydev@fedora bin]$
[mydev@fedora bin]$ sudo kata-runtime kata-check
WARN[0000] Not running network checks as super user
System is capable of running Kata Containers
System can currently create Kata Containers
[mydev@fedora bin]$
```

## ■ Env

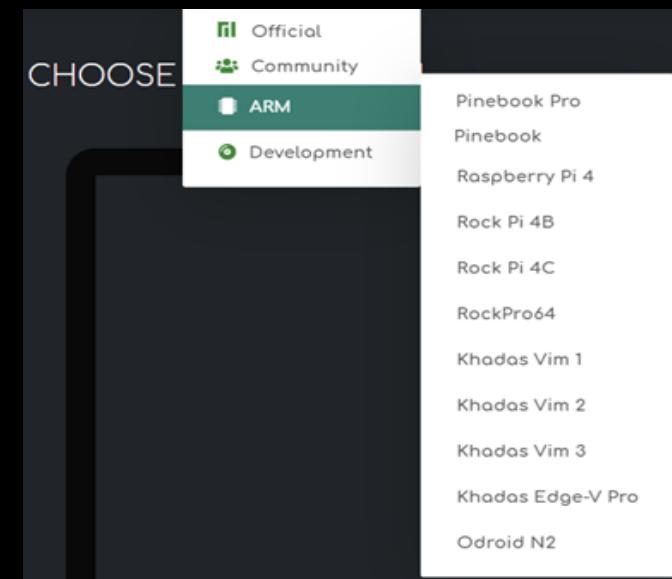
```
[mydev@fedora bin]$ sudo kata-runtime kata-env
[Meta]
  Version = "1.0.25"

[Runtime]
  Debug = false
  Trace = false
  DisableGuestSeccomp = true
  DisableNewNetNs = false
  SandboxCgroupOnly = true
  Path = "/usr/bin/kata-runtime"
[Runtime.Version]
  OCI = "1.0.1-dev"
[Runtime.Version.Version]
  Semver = "2.0.3"
  Major = 2
  Minor = 0
  Patch = 3
  Commit = "77e069db5e352a4044554ab8059b70e16b51875a"
[Runtime.Config]
  Path = "/usr/share/kata-containers/defaults/configuration.toml"

[Hypervisor]
  MachineType = "virt"
  Version = "QEMU emulator version 5.2.0 (qemu-5.2.0-5.fc34.1)\nCopyright (c) 2003-2020 Fabrice Bellard and the QEMU Project developers"
  Path = "/usr/bin/qemu-system-aarch64"
  BlockDeviceDriver = "virtio-scsi"
  EntropySource = "/dev/urandom"
  SharedFS = "virtio-fs"
  VirtioFSDaemon = "/usr/libexec/virtiofsd"
  Msize9p = 8192
  MemorySlots = 10
  PCIeRootPort = 0
  HotplugVFIOnRootBus = false
  Debug = false
```

## 3.2 Manjaro

- an open-source Linux distribution based on the Arch Linux operating system
- [https://en.wikipedia.org/wiki/Manjaro\\_Linux](https://en.wikipedia.org/wiki/Manjaro_Linux)
- <https://distrowatch.com/>
- <https://manjaro.org>



```
[mydev@MyRPi4Manjaro2 ~]$ uname -a
Linux MyRPi4Manjaro2 5.12.3-1-MANJARO-ARM #1 SMP PREEMPT Thu May 13 12:58:54 CDT 2021 aarch64 GNU/Linux
[mydev@MyRPi4Manjaro2 ~]$
[mydev@MyRPi4Manjaro2 ~]$ dmesg |grep -i kvm
[    0.739737] kvm [1]: IPA Size Limit: 44 bits
[    0.742958] kvm [1]: vgic interrupt IRQ9
[    0.745142] kvm [1]: Hyp mode initialized successfully
[mydev@MyRPi4Manjaro2 ~]$
[mydev@MyRPi4Manjaro2 ~]$ cat /lib/modules/5.12.3-1-MANJARO-ARM/build/.config |grep -i kvm
CONFIG_KVM=y
CONFIG_HAVE_KVM_IRQCHIP=y
CONFIG_HAVE_KVM_IRQFD=y
CONFIG_HAVE_KVM_IRQ_ROUTING=y
CONFIG_HAVE_KVM_EVENTFD=y
CONFIG_KVM_MMIO=y
CONFIG_HAVE_KVM_MSII=y
CONFIG_HAVE_KVM_CPU_RELAX_INTERCEPT=y
CONFIG_KVM_VFI0=y
CONFIG_HAVE_KVM_ARCH_TLB_FLUSH_ALL=y
CONFIG_KVM_GENERIC_DIRTYLOG_READ_PROTECT=y
CONFIG_HAVE_KVM_IRQ_BYPASS=y
CONFIG_HAVE_KVM_VCPU_RUN_PID_CHANGE=y
[mydev@MyRPi4Manjaro2 ~]$
[mydev@MyRPi4Manjaro2 ~]$ cat /lib/modules/5.12.3-1-MANJARO-ARM/build/.config |grep -i virtual
CONFIG_VIRTUALIZATION=y
# Virtual GPIO drivers
# end of Virtual GPIO drivers
# CONFIG_REGULATOR_VIRTUAL_CONSUMER is not set
# CONFIG_FB_VIRTUAL is not set
CONFIG_DMA_VIRTUAL_CHANNELS=y
CONFIG_ARCH_HAS_DEBUG_VIRTUAL=y
# CONFIG_DEBUG_VIRTUAL is not set
[mydev@MyRPi4Manjaro2 ~]$
```

- my presentation "Python for Linux Kernel Debugging" at PyCon China Hangzhou 2019 contains instructions to build Linux Kernel directly on RPi4 Manjaro

```
[mydev@MyRPi4Manjaro2 ~]$ cat /lib/modules/5.12.3-1-MANJARO-ARM/build/.config |grep -i bpf
CONFIG_CGROUP_BPF=y
CONFIG_BPF=y
CONFIG_BPF_SYSCALL=y
CONFIG_ARCH_WANT_DEFAULT_BPF_JIT=y
# CONFIG_BPF_PRELOAD is not set
CONFIG_NETFILTER_XT_MATCH_BPF=m
# CONFIG_BPILTER is not set
# CONFIG_NET_CLS_BPF is not set
# CONFIG_NET_ACT_BPF is not set
# CONFIG_BPF_JIT is not set
# CONFIG_BPF_STREAM_PARSER is not set
CONFIG_LWTUNNEL_BPF=y
CONFIG_HAVE_EBPF_JIT=y
CONFIG_BPF_LIRC_MODE2=y
CONFIG_BPF_EVENTS=y
# CONFIG_BPF_KPROBE_OVERRIDE is not set
# CONFIG_TEST_BPF is not set
[mydev@MyRPi4Manjaro2 ~]$
```

```
[mydev@MyRPi4Manjaro2 ~]$ cat /lib/modules/5.12.3-1-MANJARO-ARM/build/.config |grep -i cgroup
CONFIG_CGROUPS=y
CONFIG_BLK_CGROUP=y
CONFIG_CGROUP_WRITEBACK=y
CONFIG_CGROUP_SCHED=y
CONFIG_CGROUP_PIDS=y
CONFIG_CGROUP_RDMA=y
CONFIG_CGROUP_FREEZER=y
CONFIG_CGROUP_DEVICE=y
CONFIG_CGROUP_CPUACCT=y
CONFIG_CGROUP_PERF=y
CONFIG_CGROUP_BPF=y
# CONFIG_CGROUP_DEBUG is not set
CONFIG_SOCK_CGROUP_DATA=y
CONFIG_BLK_CGROUP_RWSTAT=y
# CONFIG_BLK_CGROUP_IOLATENCY is not set
# CONFIG_BLK_CGROUP_IOCOST is not set
# CONFIG_BFQ_CGROUP_DEBUG is not set
CONFIG_NETFILTER_XT_MATCH_CGROUP=m
CONFIG_NET_CLS_CGROUP=m
CONFIG_CGROUP_NET_PRIO=y
CONFIG_CGROUP_NET_CLASSID=y
[mydev@MyRPi4Manjaro2 ~]$
```

## Build Kata Containers

- [https://wiki.archlinux.org/title/Kata\\_Containers](https://wiki.archlinux.org/title/Kata_Containers)
- No available distro packages for Manjaro RPi4, just DIY it
- QEMU 6.x is available on Manjaro

```
[mydev@MyRPi4Manjaro2 ~]$ which qemu-system-aarch64
/usr/bin/qemu-system-aarch64
[mydev@MyRPi4Manjaro2 ~]$
[mydev@MyRPi4Manjaro2 ~]$ /usr/bin/qemu-system-aarch64 -version
QEMU emulator version 6.0.0
Copyright (c) 2003-2021 Fabrice Bellard and the QEMU Project developers
[mydev@MyRPi4Manjaro2 ~]$
```

- <https://github.com/kata-containers/kata-containers/tree/main/docs/install#build-from-source-installation>
- <https://github.com/kata-containers/kata-containers/blob/main/versions.yaml>
- <https://github.com/kata-containers/kata-containers/blob/main/docs/Developer-Guide.md>

```
$ go get -d -u github.com/kata-containers/kata-containers
$ cd $GOPATH/src/github.com/kata-containers/kata-containers/src/runtime
$ make && sudo -E PATH=$PATH make install
```

**~150 seconds on my Manjaro RPi4 with the latest Go 1.16.4**

## ■ build artifacts

```
[mydev@MyRPi4Manjaro2 ~]$ ll /usr/local/bin
total 90196
drwxr-xr-x 3 root root 4096 May 21 20:06 .
drwxr-xr-x 11 root root 4096 Sep 22 2020 ../
-rwxr-xr-x 1 root root 33401728 May 21 20:06 containerd-shim-kata-v2*
-rwxr-xr-x 1 root root 16678 May 21 20:06 kata-collect-data.sh*
-rwxr-xr-x 1 root root 24705280 May 21 20:06 kata-monitor*
-rwxr-xr-x 1 root root 34208992 May 21 20:06 kata-runtime*
-rwxr-xr-x 1 root root 26 Jan 27 06:49 stop-dmesg*
drwxr-xr-x 2 root root 4096 Apr 18 14:20 tflite/
[mydev@MyRPi4Manjaro2 ~]$
[mydev@MyRPi4Manjaro2 ~]$ file /usr/local/bin/kata-runtime
/usr/local/bin/kata-runtime: ELF 64-bit LSB pie executable, ARM aarch64, version 1 (SYSV), dynamically linked, interpreter /lib/ld-linux-aarch64.so.1, BuildID[sha1]=a4b985d98f2e24f90fe1c56cc32212bea7ecde2a, for GNU/Linux 3.7.0, not stripped
[mydev@MyRPi4Manjaro2 ~]$
[mydev@MyRPi4Manjaro2 ~]$ ll /usr/share/defaults/kata-containers
total 60
drwxr-xr-x 2 root root 4096 May 21 20:06 .
drwxr-xr-x 4 root root 4096 May 21 19:26 ../
-rw-r--r-- 1 root root 11106 May 21 20:06 configuration-clh.toml
-rw-r--r-- 1 root root 15614 May 21 20:06 configuration-fc.toml
-rw-r--r-- 1 root root 23174 May 21 20:06 configuration-qemu.toml
lrwxrwxrwx 1 root root 23 May 21 20:06 configuration.toml -> configuration-qemu.toml
[mydev@MyRPi4Manjaro2 ~]$
```

## ■ Kata-check failed

```
[mydev@MyRPi4Manjaro2 ~]$ kata-runtime kata-check
ERROR[0000] /usr/share/defaults/kata-containers/configuration-qemu.toml: file /usr/share/kata-containers/vmlinux.container does not exist
ERROR[0000] arch=arm64 name=kata-runtime pid=8982 source=runtime
/usr/share/defaults/kata-containers/configuration-qemu.toml: file /usr/share/kata-containers/vmlinux.container does not exist
[mydev@MyRPi4Manjaro2 ~]$
[mydev@MyRPi4Manjaro2 ~]$ vim /usr/share/defaults/kata-containers/configuration-qemu.toml
```

```
[hypervisor.qemu]
path = "/usr/bin/qemu-system-aarch64"
kernel = "/usr/share/kata-containers/vmlinux.container"
image = "/usr/share/kata-containers/kata-containers.img"
machine_type = "virt"
```

No Kernel and Rootfs for container...

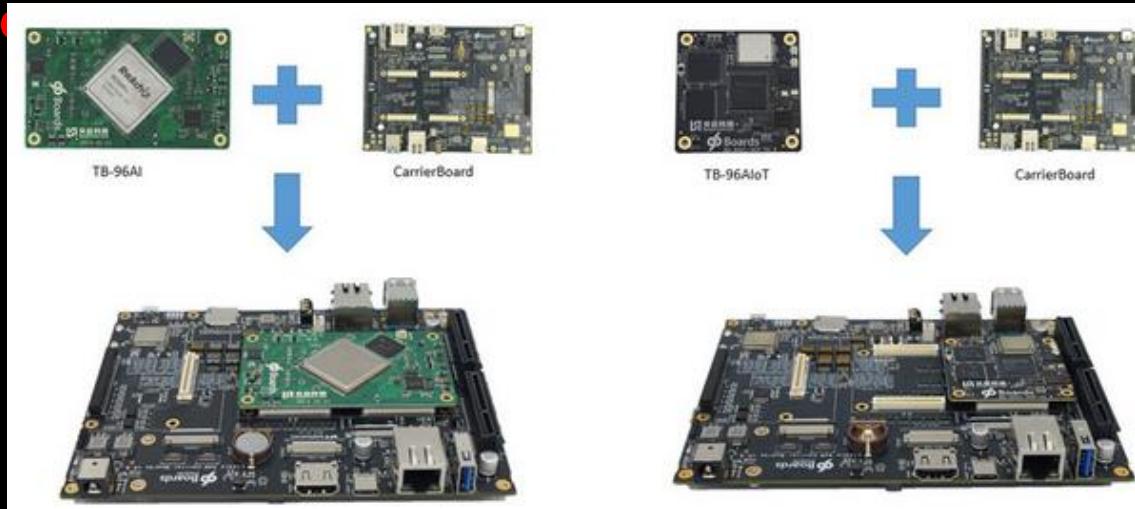
- <https://github.com/kata-containers/packaging/tree/master/kernel#build-kata-containers-kernel>  
<https://github.com/kata-containers/documentation/blob/master/Developer-Guide.md#create-a-rootfs-image>

---

<https://github.com/kata-containers/documentation/blob/master/Developer-Guide.md#create-an-initrd-image---optional>
-

### 3.3 Clustering SOM/COM

- <https://www.linaro.org/news/linaro-announces-launch-of-96boards-system-on-module-som-specification/>
- <http://linuxgizmos.com/linaro-launches-two-96boards-som-specifications/>
- <http://static.linaro.org/assets/specifications/96BoardsComputeSoMSpecificationV1.0.pdf>
- <https://static.linaro.org/assets/specifications/96BoardsWirelessSoMSpecificationV1.0.pdf>
- 



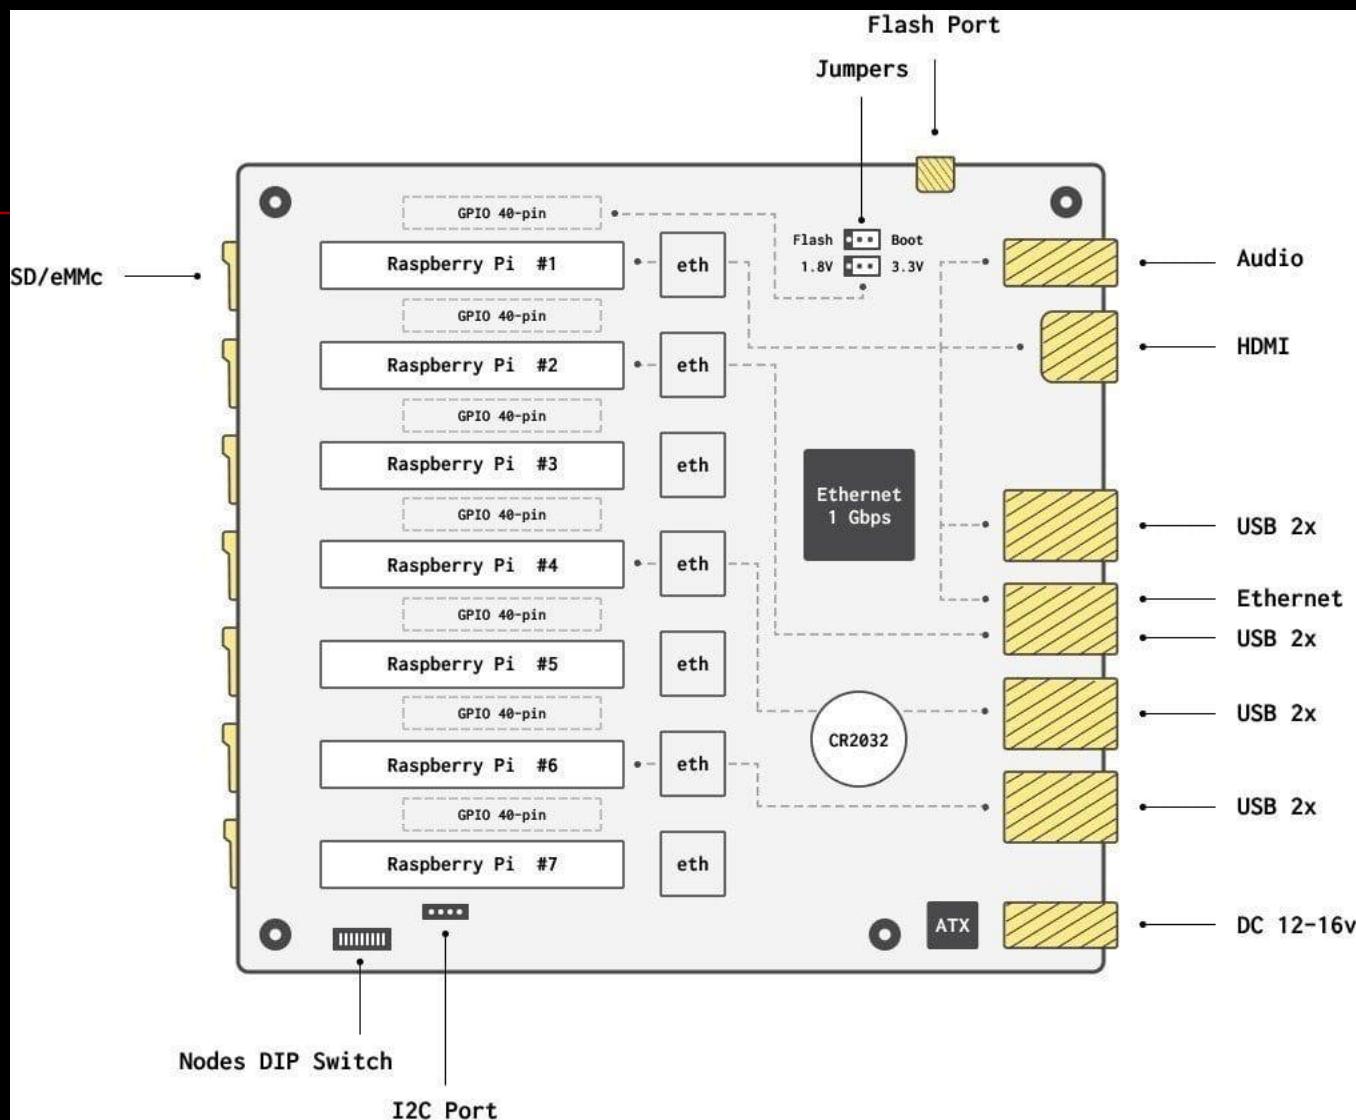
## Turing Pi 1

- <https://turingpi.com/>
- <https://docs.turingpi.com/specs>
- **Support Raspberry Pi Compute Module 1, 3, 3+**



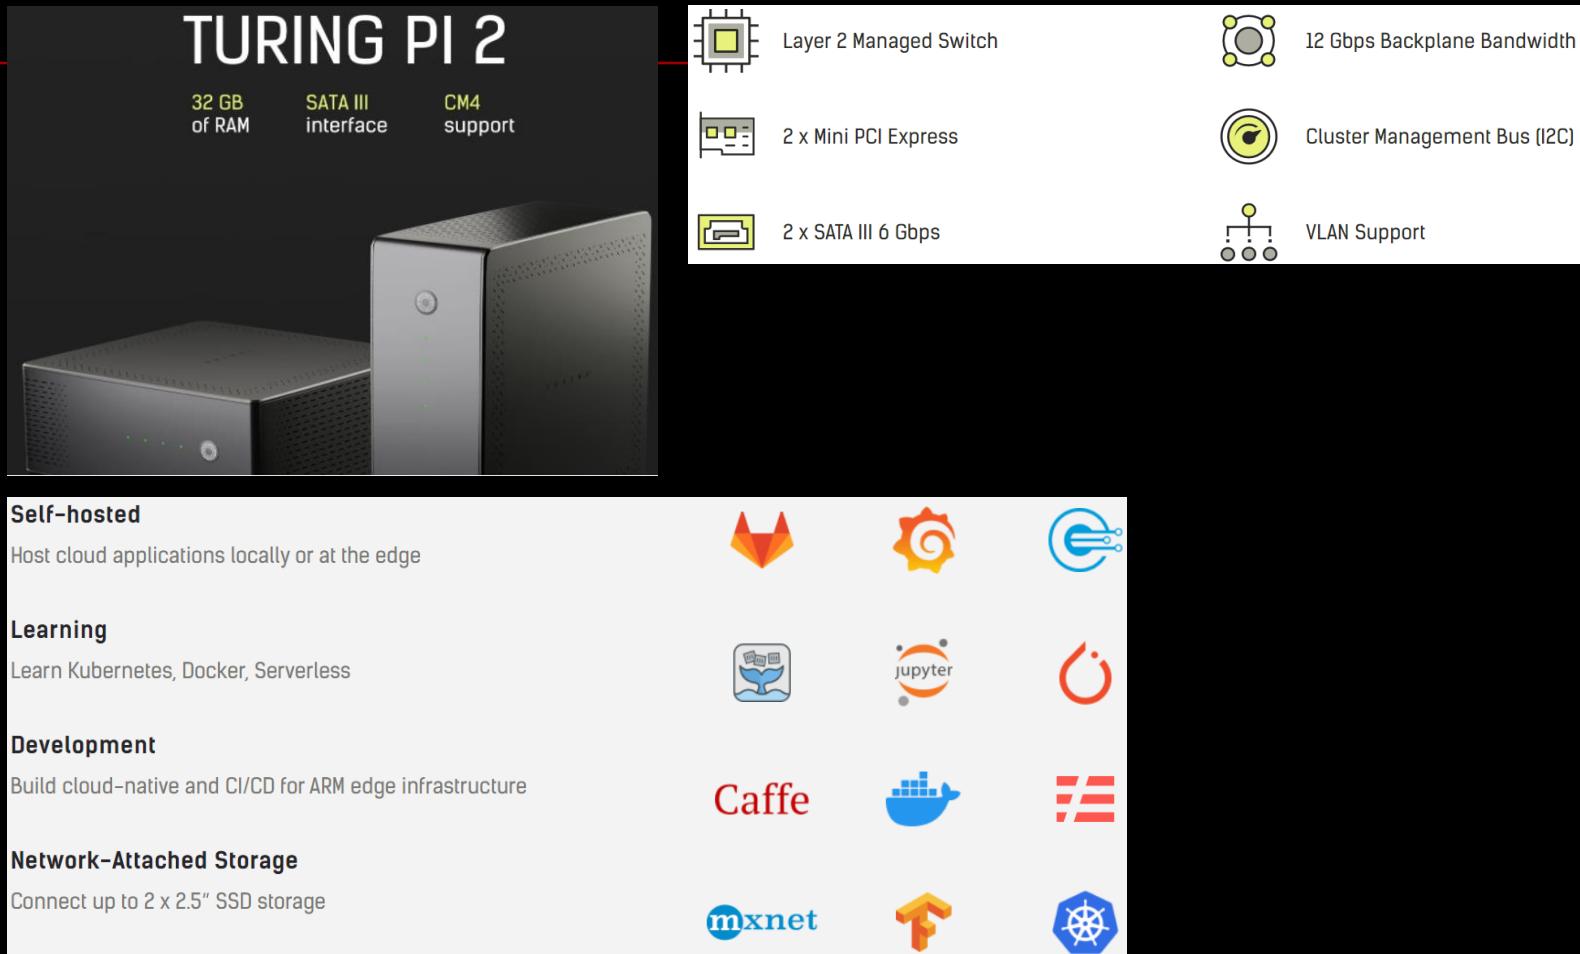
- **Host Kubernetes, Containers and Cloud Native Apps locally**

## ■ Scheme



## Upcoming Turing Pi 2

- <https://turingpi.com/v2/>
- <https://turingpi.com/turing-pi-2-announcement/>



The diagram illustrates the Turing Pi 2 hardware and its software ecosystem. On the left, the Turing Pi 2 hardware is shown with three main components: a small form-factor computer, a network switch, and a storage unit. Key specifications listed are 32 GB of RAM, SATA III interface, and CM4 support. To the right, the system's internal architecture is detailed with icons: a Layer 2 Managed Switch (represented by a square with a yellow center), 2 x Mini PCI Express slots (represented by two rectangles with yellow centers), and 2 x SATA III 6 Gbps drives (represented by two yellow rectangular ports). A separate box on the right lists four external features with corresponding icons: 12 Gbps Backplane Bandwidth (a circular icon with green and yellow segments), Cluster Management Bus (I2C) (a circular icon with a green center and a yellow ring), and VLAN Support (a tree-like icon with three nodes).

**TURING PI 2**

32 GB of RAM    SATA III interface    CM4 support

Layer 2 Managed Switch

12 Gbps Backplane Bandwidth

2 x Mini PCI Express

Cluster Management Bus (I2C)

2 x SATA III 6 Gbps

VLAN Support

**Self-hosted**  
Host cloud applications locally or at the edge

**Learning**  
Learn Kubernetes, Docker, Serverless

**Development**  
Build cloud-native and CI/CD for ARM edge infrastructure

**Network-Attached Storage**  
Connect up to 2 x 2.5" SSD storage

**Caffe**

**mxnet**

**Jupyter**

**Docker**

**Hadoop**

**Flink**

**TensorFlow**

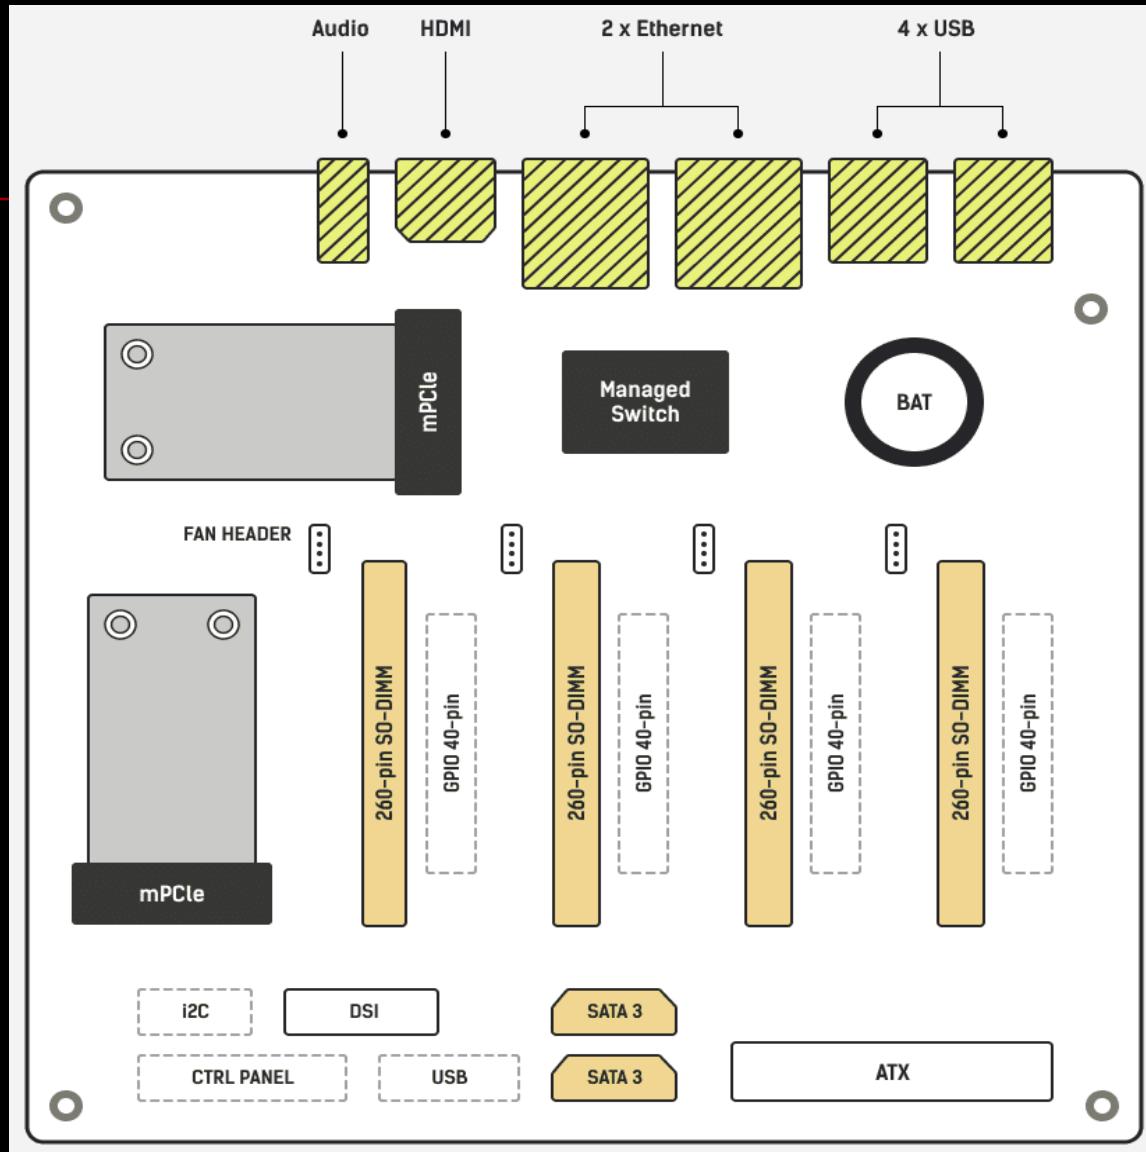
**Apache HDFS**

**Apache Spark**

**Apache Flink**

- **Reliable, Scalable, Cloud-Native Infrastructure for the Edge**

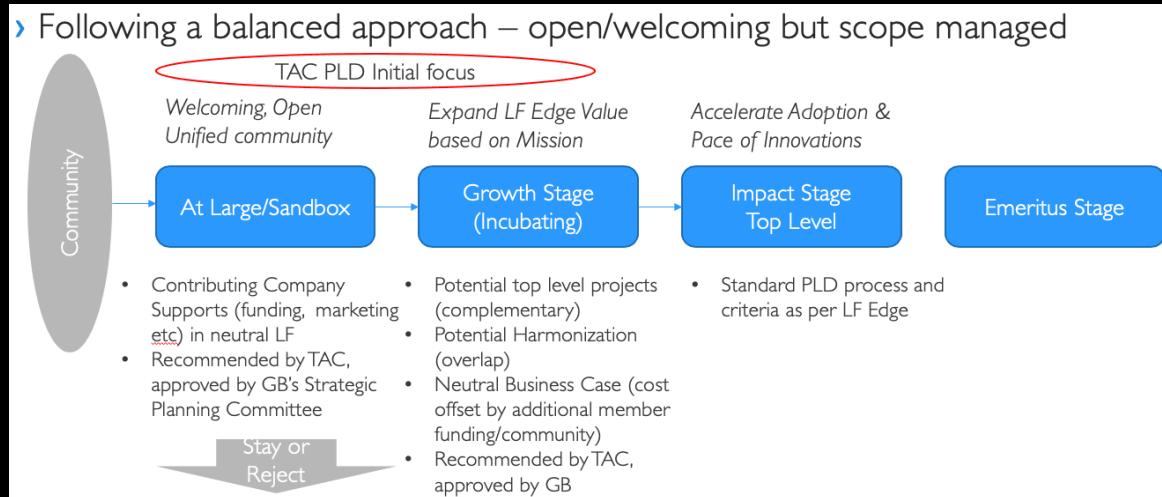
## ■ Scheme



# II. LF Edge

## 1) Overview

- <https://www.lfedge.org/>
- **IOT + Telecom + Cloud + Enterprise + Industrial**
- <https://www.linuxfoundation.org/en/press-release/the-linux-foundation-launches-new-lf-edge-to-establish-a-unified-open-source-framework-for-the-edge/>
- <https://www.lfedge.org/resources/publications/>
- <https://landscape.lfedge.org/>
- <https://www.lfedge.org/projects/>



IMPACT STAGE



GROWTH STAGE

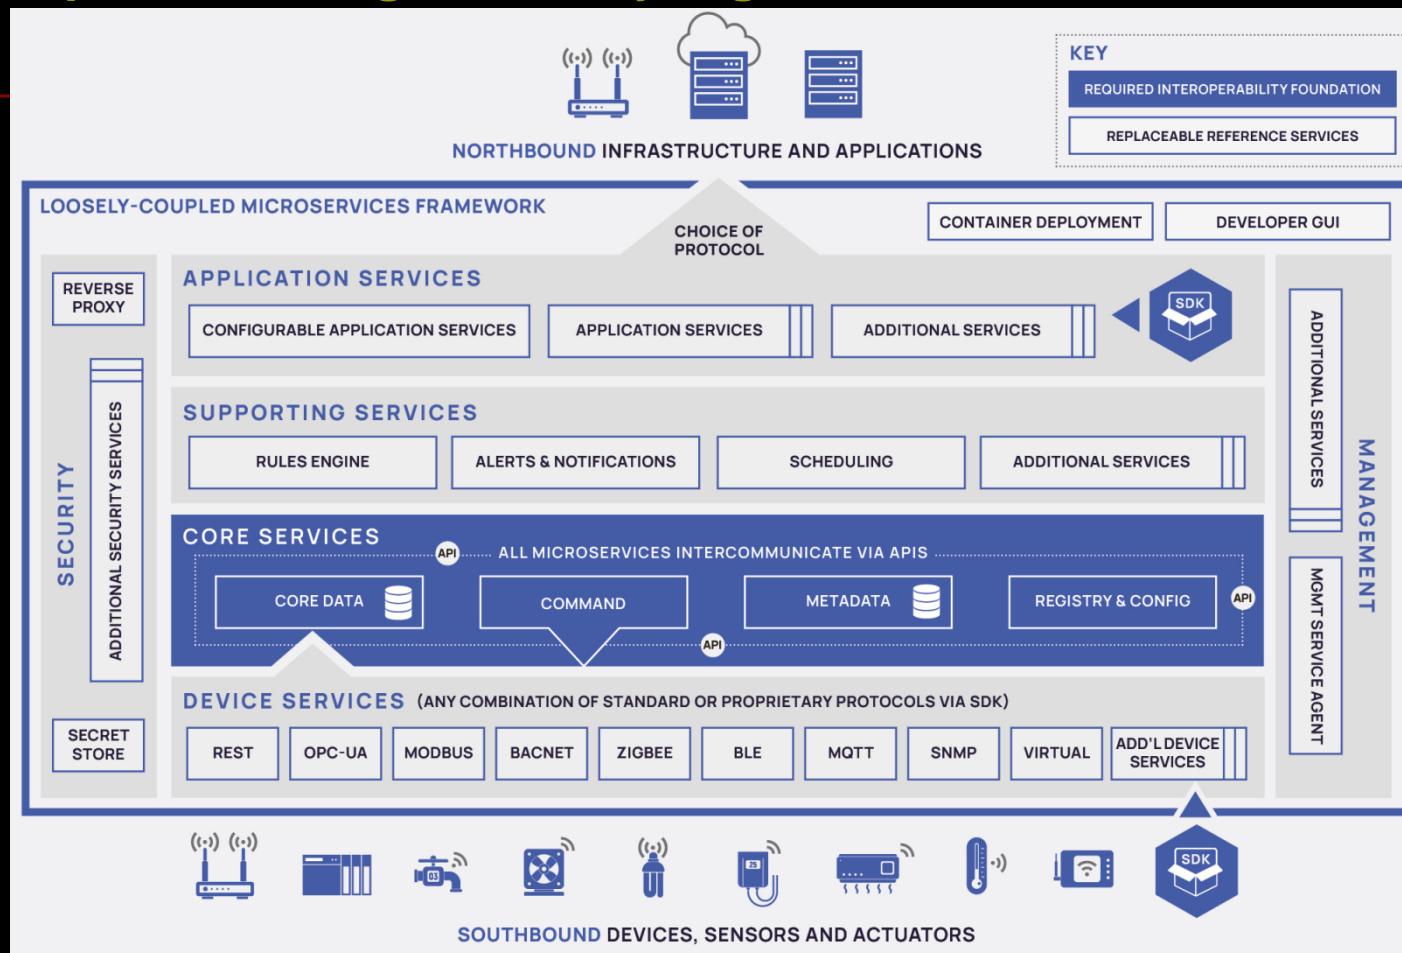


AT LARGE STAGE



## EdgeX

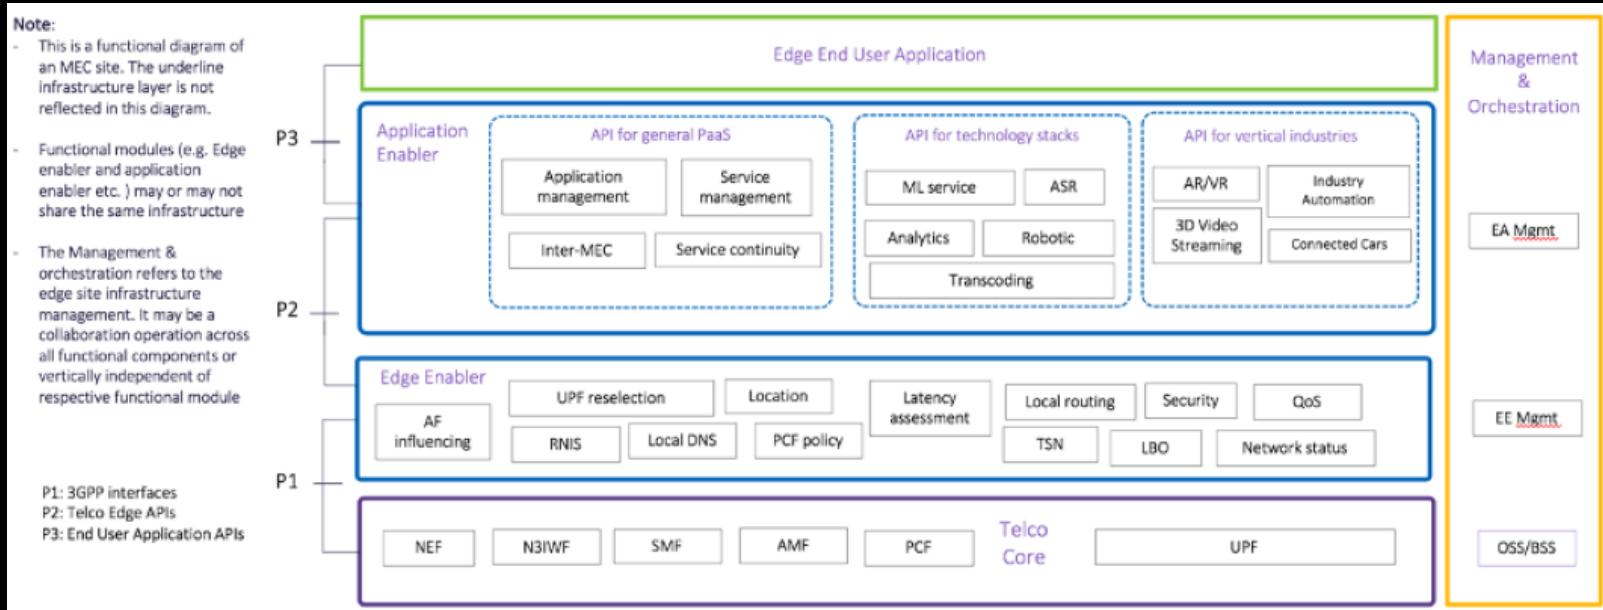
- <https://www.lfedge.org/projects/edgexfoundry/>
- <https://www.edgexfoundry.org/>



- **Migrating from Java to Go since 1.0 release**

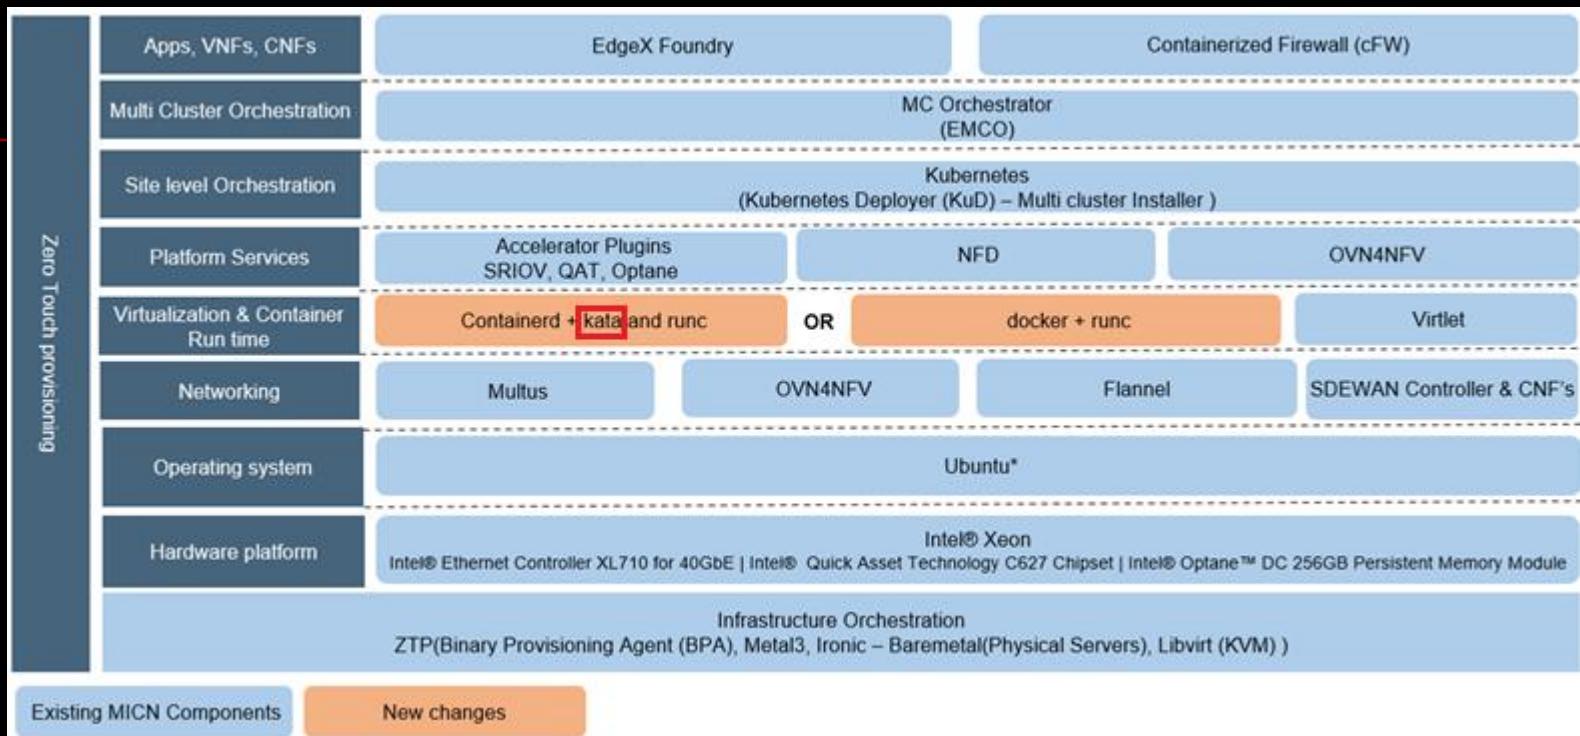
## Akraino

- <https://www.lfedge.org/projects/akraino/>
- **Akraino Edge Stack**  
<https://wiki.akraino.org/display/AK/Akraino+Edge+Stack>
- **Edge Stack Functional Layers**



- <https://www.lfedge.org/projects/akraino/release-4-2/>

■ <https://wiki.akraino.org/display/AK/Multi-tenant+Secure+Cloud+Native+Platform+Architecture>



## 2) EVE

- <https://www.lfedge.org/projects/eve/>
- **Key Capabilities**

The goal of Project EVE is to enable IoT edge computing deployments with the following capabilities:

- Access to hardware root of trust (e.g. TPM) when deployed on bare metal, supporting functions such as crypto-based ID (no device usernames and passwords), measured boot, remote attestation, signed updates, encryption, etc.
- “Secure by default” deployment profile
- High efficiency and usage of device resources including remote control of CPU, memory, networking and edge device I/O ports
- Hosting of any combination of apps in virtual machines, containers and Kubernetes clusters
- Hosting of any guest operating system deployable in a virtual machine
- Ability to assign CPU cores and co-processing (e.g. GPU) to specific apps
- Ability to block unused I/O ports to prevent physical tampering
- Remote updates of entire software stack with rollback capability to prevent bricking
- Automated patching for security updates
- Automated connectivity to one or more backends (cloud or on premises)
- Distributed firewall to securely route data over networks per policy

Project EVE is building EVE-OS, a universal, open Linux-based operating system for distributed edge computing. EVE-OS aims to do for the distributed edge what Android did for mobile by creating an open foundation that simplifies development, orchestration and security of edge computing nodes deployed on-prem and in the field. Supporting Docker containers, Kubernetes clusters and virtual machines, EVE-OS provides a flexible foundation for distributed edge deployments with choice of any hardware, application and cloud.

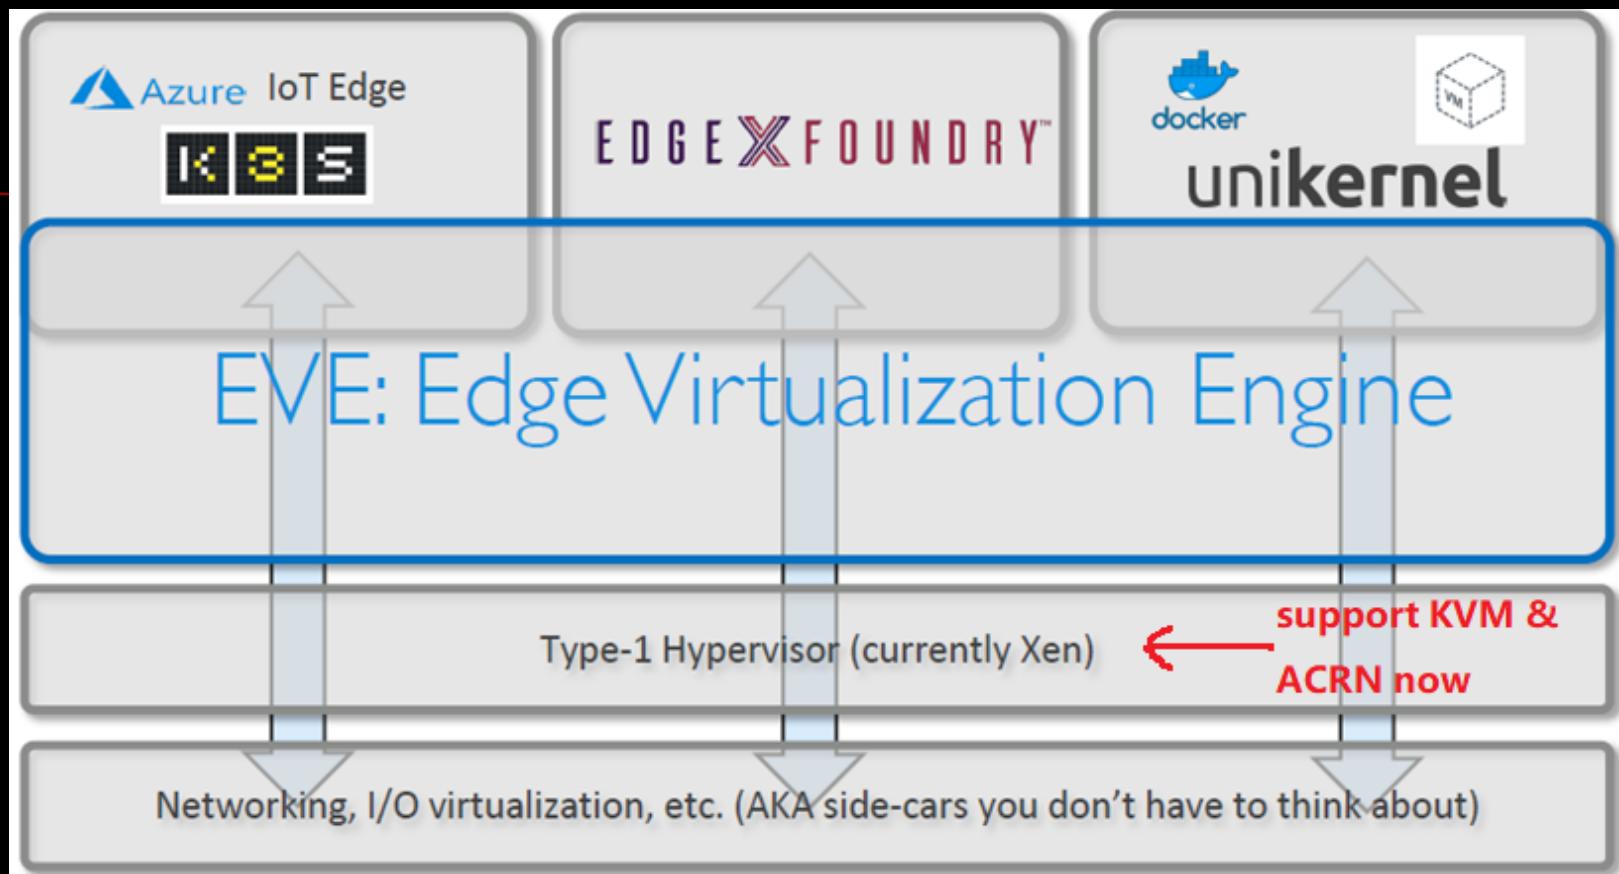
EVE-OS can be deployed on any bare metal hardware (e.g. x86, Arm, GPU) or within a VM to provide consistent system and orchestration services and provides the ability to run applications in a variety of formats. Support for VMs enables users to continue to use existing software investments while building new containerized innovations in parallel. Compared to agent-based edge management solutions, the bare metal EVE-OS eliminates the possibility of bricking a device in the field during an update, requiring an expensive truck roll.

- <https://github.com/lf-edge/eve>

EVE supports both ARM and Intel architectures and requires hardware-assisted virtualization. While EVE can run on a board as small as a \$20 Orange Pi, the sweet spot for its deployment are IoT Gateways and Industrial PCs.

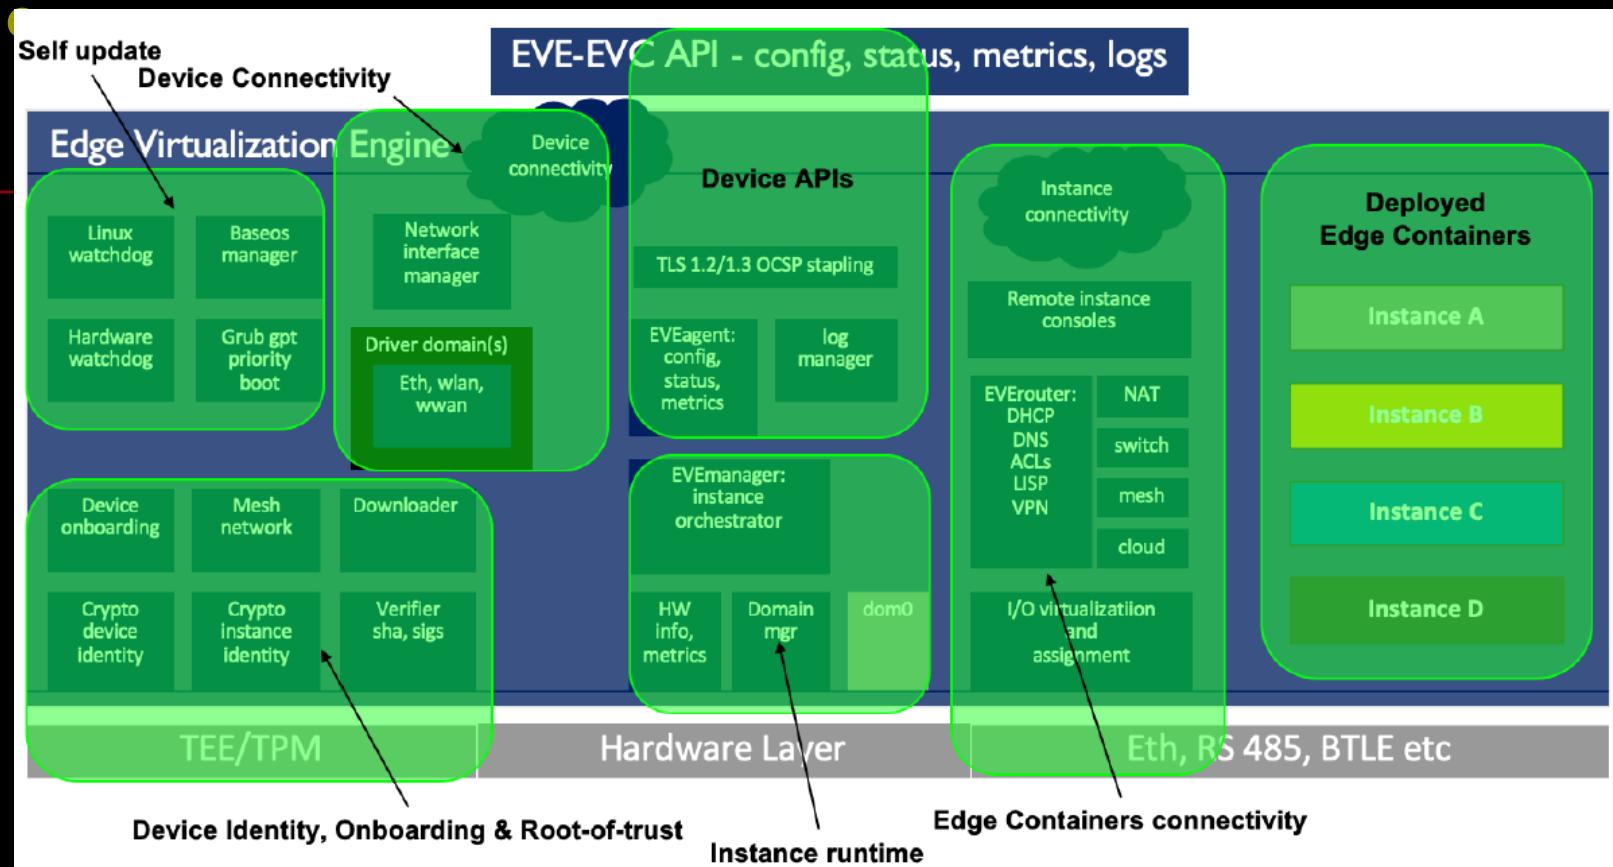
To get its job done, EVE leverages a lot of great open source projects: [Xen Project](#), [Linuxkit](#) and [Alpine Linux](#) just to name a few. All of that functionality is being orchestrated by the Go microservices available under [pkg/pillar](#). Why pillar? Well, because pillar is the kind of a monolith we need to break out into true, individual microservices under [pkg/](#).

## Architecture



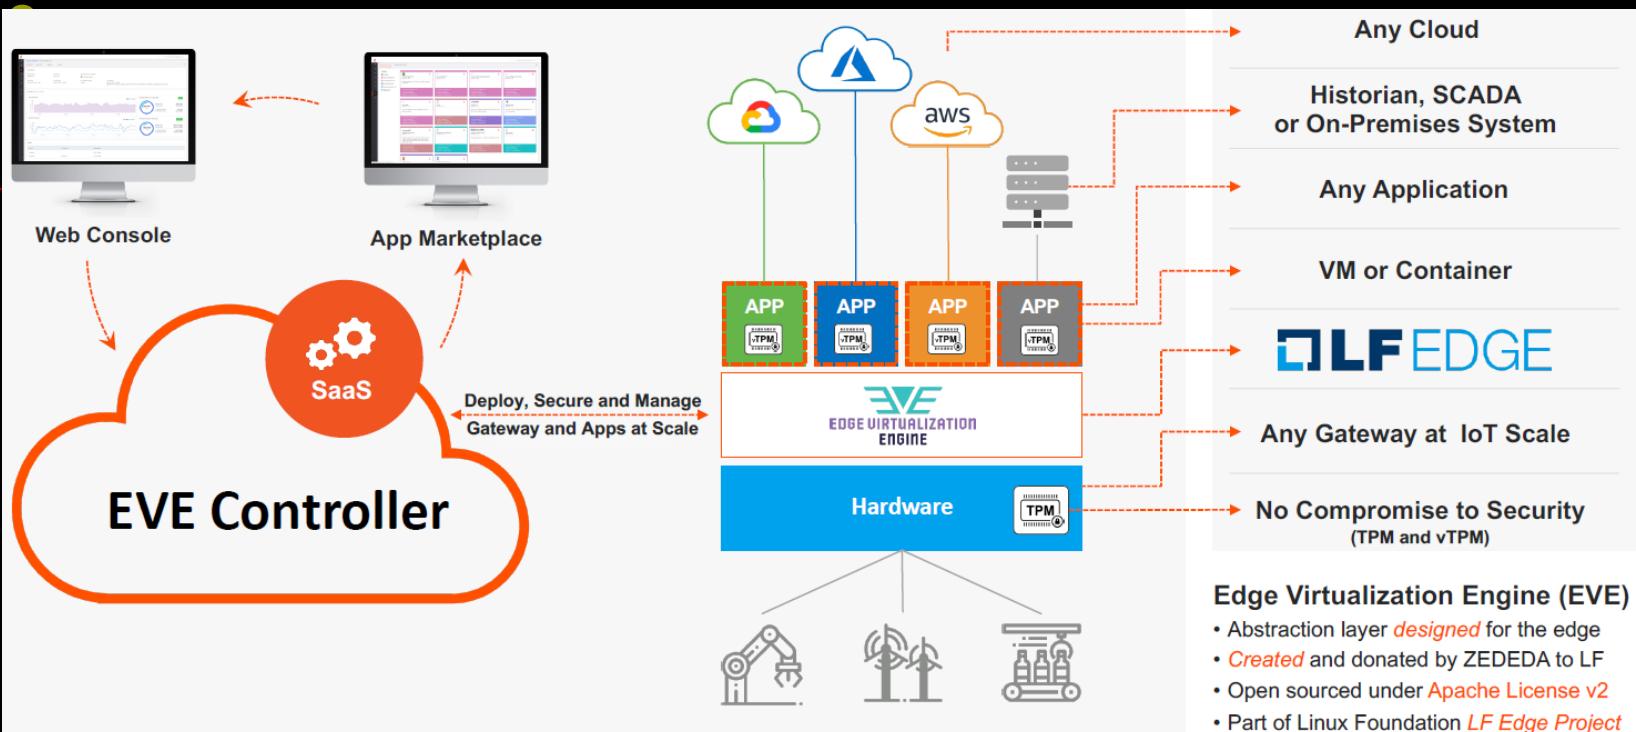
Source: “Linux Foundation’s Project EVE: a Cloud-Native Edge Computing Platform”, Roman Shaposhnik, QConSF2019

## Internals



Source: “Linux Foundation’s Project EVE: a Cloud-Native Edge Computing Platform”, Roman Shaposhnik, QConSF2019

## Edge Virtualization



Source: “Linux Foundation’s Project EVE: a Cloud-Native Edge Computing Platform”, Roman Shaposhnik, QConSF2019

### Edge Virtualization Engine (EVE)

- Abstraction layer *designed* for the edge
- *Created* and donated by ZEDEDATA to LF
- Open sourced under Apache License v2
- Part of Linux Foundation *LF Edge Project*

## 2.1 Edge Containers

■ <https://github.com/lf-edge/edge-containers>

### Edge Containers

- A true extension to the OCI specification
  - Image specification (not much of a change)
  - Runtime specification
  - Registry Support (via OCI Artifacts Initiative)
- Related initiatives
  - Kata Containers, Singularity Containers, etc.
  - Weave.works's Project Ignite (Firecracker MicroVMs)
  - Rancher's K3S + K3OS
- Top 3 goals:
  - Filesystem-level composition (aka OCI layers)
  - Block-level composition (VMs and Unikernels)
  - Hardware mapping
- Registry as a "nexus of Liquid Software"

Source: “Linux Foundation’s Project EVE: a Cloud-Native Edge Computing Platform”, Roman Shaposhnik, QConSF2019



K3S



K3OS



Unikernel

## 2.2 EVE for RPi

- <https://github.com/lf-edge/eve/blob/master/Makefile>

- How to use on a Raspberry Pi 4 ARM board

Raspberry Pi 4 is a tiny, but capable enough ARM board that allows EVE to run with either Xen or KVM hypervisors. While EVE would run in the lowest memory configuration (1GB) if you plan to use it for actual EVE development we strongly recommend buying a 4GB RAM option.

Since a full Raspberry Pi 4 support is only available in upstream Linux kernels starting from 5.6.0, you'll have to use that bleeding edge kernel for your build. Another peculiar aspect of this board is that it doesn't use a standard bootloader (e.g. u-boot or UEFI) so we need to trick it into using our own u-boot as UEFI environment. Thankfully, our Makefile logic tries to automate as much of it as possible. Thus, putting it all together, here are the steps to run EVE on Raspberry Pi 4:

1. Make sure you have a clean build directory (since this is a non-standard build) `rm -rf dist/arm64`
2. Build a live image `make ZARCH=arm64 HV=rpi live-raw` (or `make ZARCH=arm64 HV=rpi-kvm live-raw` if you want KVM by default)
3. Flash the `dist/arm64/live.raw` live EVE image onto your SD card by [following these instructions](#)

Once your Raspberry Pi 4 is happily running an EVE image you can start using EVE controller for further updates (so that you don't ever have to take an SD card out of your board). Build your rootfs by running `make ZARCH=arm64 rootfs-rpi` (or `make ZARCH=arm64 rootfs-kvm-rpi` if you want KVM by default) and give resulting `dist/arm64/installer/rootfs.img` to the controller.

One final note about Raspberry Pi 4 GPU support is that since we are running EVE in 64bit (aarch64) mode we are still waiting for the proper [VC4 DRM BCM2711 drivers](#) to be upstreamed. Currently it is expected that Kernel 5.7 may actually ship the fully functional driver.

- <https://github.com/lf-edge/eve/blob/master/Makefile>
- <https://hub.docker.com/r/lfedge/eve>

## No prebuilt image for RPi

- <https://github.com/lf-edge/eve/releases>

Release 6.6.0

github-actions released this 4 days ago

---

Assets 16

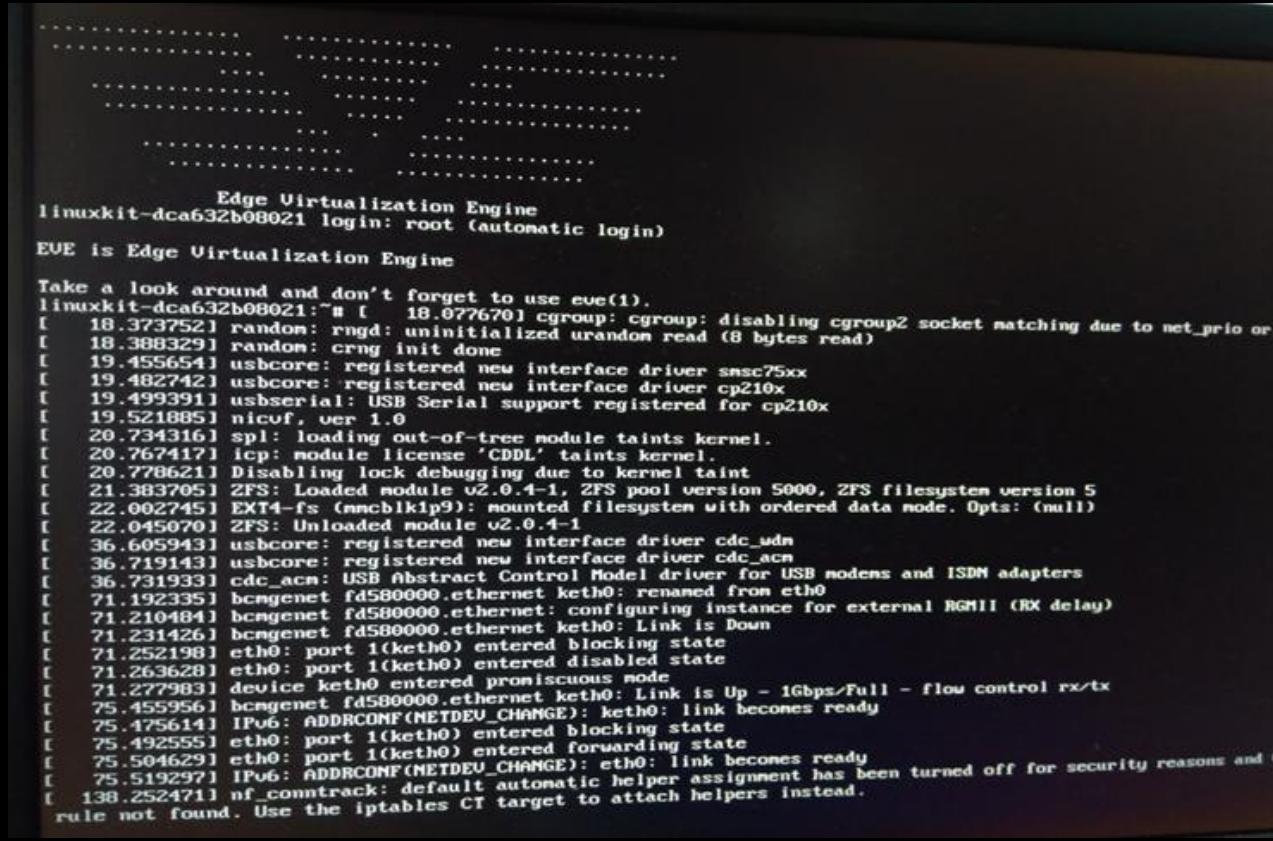
<a href="#">amd64.initrd.bits</a>	118 KB
<a href="#">amd64.initrd.img</a>	3.55 MB
<a href="#">amd64.installer.img</a>	8.3 MB
<a href="#">amd64.ipxe.efi.cfg</a>	1.24 KB
<a href="#">amd64.ipxe.efi.ip.cfg</a>	1.23 KB
<a href="#">amd64.kernel</a>	9.31 MB
<a href="#">amd64.rootfs.img</a>	231 MB
<a href="#">arm64.initrd.bits</a>	118 KB
<a href="#">arm64.initrd.img</a>	3.42 MB
<a href="#">arm64.installer.img</a>	8.13 MB
<a href="#">arm64.ipxe.efi.cfg</a>	1.24 KB
<a href="#">arm64.ipxe.efi.ip.cfg</a>	1.23 KB
<a href="#">arm64.kernel</a>	28.7 MB
<a href="#">arm64.rootfs.img</a> 	198 MB
<a href="#">Source code (zip)</a>	
<a href="#">Source code (tar.gz)</a>	

## Retrieve live image from container

- <https://www.lfedge.org/2020/10/21/xen-on-raspberry-pi-4-with-project-eve/>  
e.g.:

**docker pull lfedge/eve:6.6.0-rpi-xen-arm64**

**docker run lfedge/eve:6.6.0-rpi-xen-arm64 live > live.raw**



The screenshot shows a terminal window with a black background and white text. At the top, it displays the system information: "Edge Virtualization Engine", "linuxkit-dca632b08021 login: root (automatic login)", and "EVE is Edge Virtualization Engine". Below this, there is a large amount of log output from the kernel and various drivers. The log includes messages about random number generation, USB drivers (smc75xx, cp210x), file systems (ZFS, EXT4), and network interfaces (eth0, keth0). It also shows the configuration of bcangenet and bcntrack modules, along with IPv6 and flow control information. The log ends with a note about nf\_conntrack and iptables.

```
Edge Virtualization Engine
linuxkit-dca632b08021 login: root (automatic login)
EVE is Edge Virtualization Engine

Take a look around and don't forget to use eve(1).
linuxkit-dca632b08021:~# [ 18.077670] cgroup: disabling cgroup2 socket matching due to net_prio or
[ 18.373752] random: rngd: uninitialized urandom read (8 bytes read)
[ 18.388329] random: crng init done
[ 19.455654] usbcore: registered new interface driver smc75xx
[ 19.482742] usbcore: registered new interface driver cp210x
[ 19.499391] usbserial: USB Serial support registered for cp210x
[ 19.521885] nicuf, ver 1.0
[ 20.734316] spl: loading out-of-tree module taints kernel.
[ 20.767417] icp: module license 'CDDL' taints kernel.
[ 20.778621] Disabling lock debugging due to kernel taint
[ 21.383705] ZFS: Loaded module v2.0.4-1, ZFS pool version 5000, ZFS filesystem version 5
[ 22.002745] EXT4-fs (mmcblk1p9): mounted filesystem with ordered data mode. Opts: (null)
[ 22.045070] ZFS: Unloaded module v2.0.4-1
[ 36.605943] usbcore: registered new interface driver cdc_udn
[ 36.719143] usbcore: registered new interface driver cdc_acm
[ 36.731933] cdc_acm: USB Abstract Control Model driver for USB modems and ISDN adapters
[ 71.192335] bcangenet fd580000.ethernet keth0: renamed from eth0
[ 71.210484] bcangenet fd580000.ethernet: configuring instance for external RGMII (RX delay)
[ 71.231426] bcangenet fd580000.ethernet keth0: Link is Down
[ 71.252198] eth0: port 1(keth0) entered blocking state
[ 71.263628] eth0: port 1(keth0) entered disabled state
[ 71.277983] device keth0 entered promiscuous mode
[ 75.455956] bcangenet fd580000.ethernet keth0: Link is Up - 1Gbps/Full - flow control rx/tx
[ 75.475614] IPv6: ADDRCONF(NETDEV_CHANGE): keth0: link becomes ready
[ 75.492555] eth0: port 1(keth0) entered blocking state
[ 75.504629] eth0: port 1(keth0) entered forwarding state
[ 75.519297] IPv6: ADDRCONF(NETDEV_CHANGE): eth0: link becomes ready
[ 138.252471] nf_conntrack: default automatic helper assignment has been turned off for security reasons and a
rule not found. Use the iptables CT target to attach helpers instead.
```

## 2.3 RunX

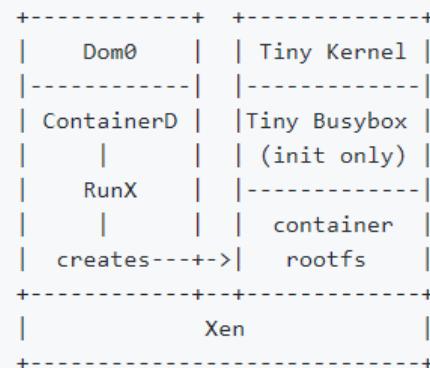
- <https://github.com/lf-edge/runx>
- an OCI Runtime Spec compliant containers runtime that runs containers as VMs

### RunX and KataContainers

Both KataContainers and RunX are containers runtimes that use hypervisors to start containers as virtual machines. However, there are a few key differences.

KataContainers focuses on KVM-based virtual machines. RunX focuses on Xen virtual machines. KataContainers uses an agent running inside each VM, while RunX does not do that by design. RunV (KataContainers' parent) uses libxl to create Xen VMs; thus, it has a build dependency on the Xen Dom0 libraries. RunX doesn't have any build or runtime dependencies on libraries as it invokes the command-line tool `xl`.

### Architecture



# III. Lightweight Kubernetes

## 1) K3S

- <https://k3s.io/>
- Why is it

### Perfect for Edge

K3s is a highly available, certified Kubernetes distribution designed for production workloads in unattended, resource-constrained, remote locations or inside IoT appliances.

### Simplified & Secure

K3s is packaged as a single <40MB binary that reduces the dependencies and steps needed to install, run and auto-update a production Kubernetes cluster.

### Optimized for ARM

Both ARM64 and ARMv7 are supported with binaries and multiarch images available for both. K3s works great from something as small as a Raspberry Pi to an AWS a1.4xlarge 32GiB server.

Additionally, K3s simplifies Kubernetes operations by maintaining functionality for:

- Managing the TLS certificates of Kubernetes components
- Managing the connection between worker and server nodes
- Auto-deploying Kubernetes resources from local manifests, in realtime as they are changed.
- Managing an embedded etcd cluster (work in progress)

## Features

### ■ <https://github.com/k3s-io/k3s/blob/master/README.md>

K3s is a [fully conformant](#) production-ready Kubernetes distribution with the following changes:

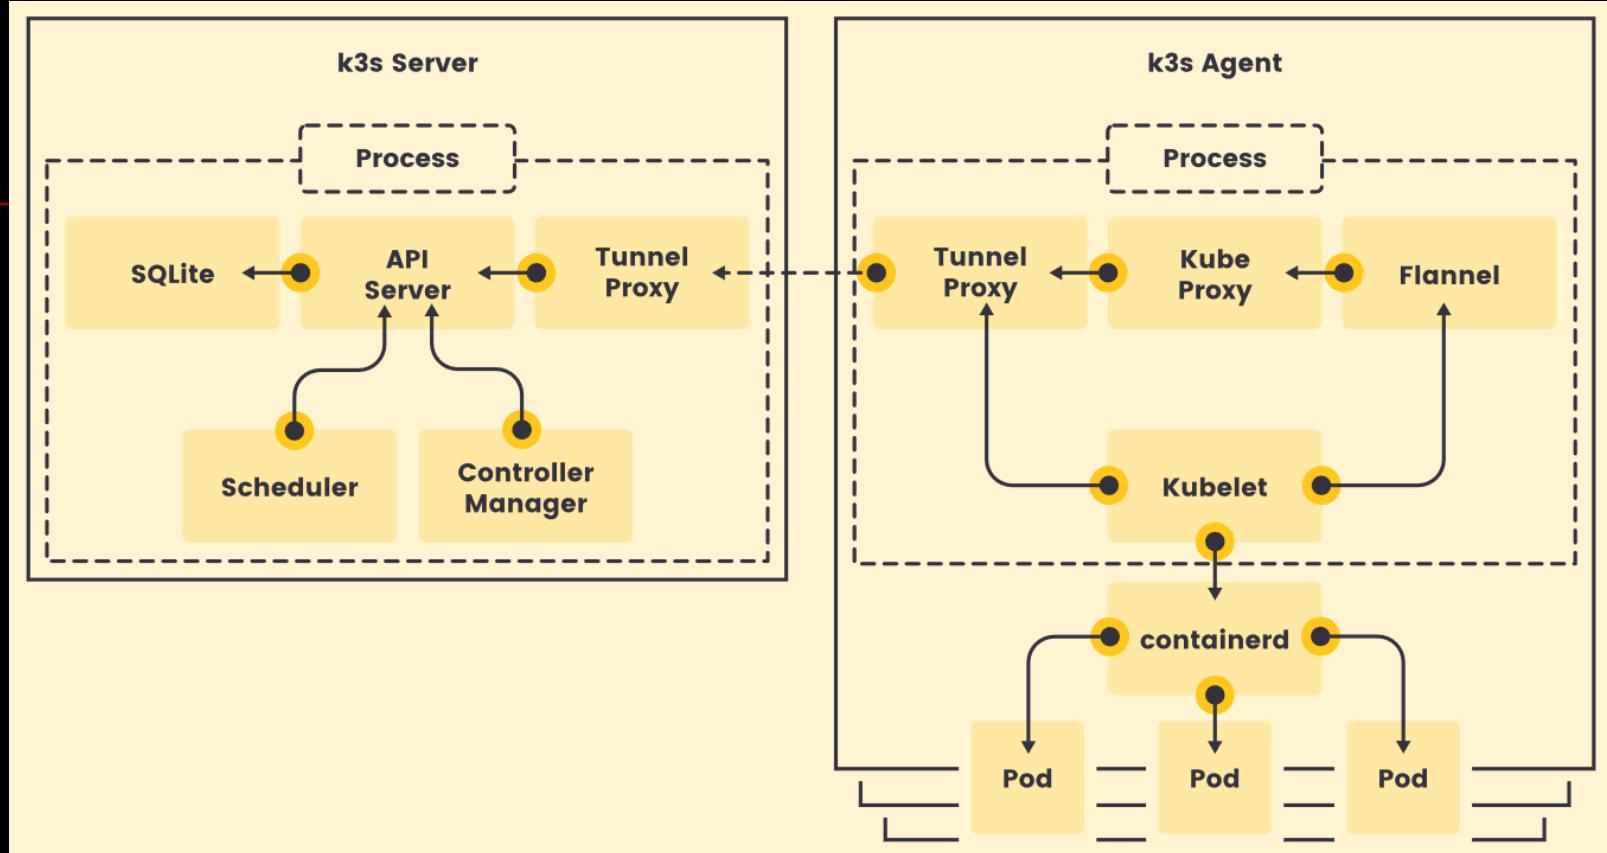
1. It is packaged as a single binary.
2. It adds support for sqlite3 as the default storage backend. Etcd3, MySQL, and Postgres are also supported.
3. It wraps Kubernetes and other components in a single, simple launcher.
4. It is secure by default with reasonable defaults for lightweight environments.
5. It has minimal to no OS dependencies (just a sane kernel and cgroup mounts needed).
6. It eliminates the need to expose a port on Kubernetes worker nodes for the kubelet API by exposing this API to the Kubernetes control plane nodes over a websocket tunnel.

K3s bundles the following technologies together into a single cohesive distribution:

- [Containerd](#) & [runc](#)
- [Flannel](#) for CNI
- [CoreDNS](#)
- [Metrics Server](#)
- [Traefik](#) for ingress
- [Klipper-lb](#) as an embedded service loadbalancer provider
- [Kube-router](#) for network policy
- [Helm-controller](#) to allow for CRD-driven deployment of helm manifests
- [Kine](#) as a datastore shim that allows etcd to be replaced with other databases
- [Local-path-provisioner](#) for provisioning volumes using local storage
- [Host utilities](#) such as iptables/nftables, ebtables, ethtool, & socat

These technologies can be disabled or swapped out for technologies of your choice.

## How it works



## 2) Arkade

- <https://github.com/alexellis/arkade>
- **The Open Source Kubernetes Marketplace**
- <https://kccncna20.sched.com/event/ekAF/the-past-present-and-future-of-kubernetes-on-raspberry-pi-alex-ellis-openfaas-ltd>

### The Past, Present, and Future of Kubernetes on Raspberry Pi

KubeCon, North America

Alex Ellis - OpenFaaS Ltd - CNCF Ambassador

### **3) Sandboxed Containers Operator**

- <https://github.com/openshift/sandboxed-containers-operator>
  - **An operator to enhance an Openshift/Kubernetes cluster to support running sandboxed containers**
-

# IV. Project ACRN

## 1) Overview

- <https://projectacrn.org/>

ACRN™ is a flexible, lightweight reference hypervisor, built with real-time and safety-criticality in mind, and optimized to streamline embedded development through an open source platform. ACRN defines a device hypervisor reference stack and an architecture for running multiple software subsystems, managed securely, on a consolidated system using a virtual machine manager (VMM). It also defines a reference framework implementation for virtual device emulation, called the "ACRN Device Model".

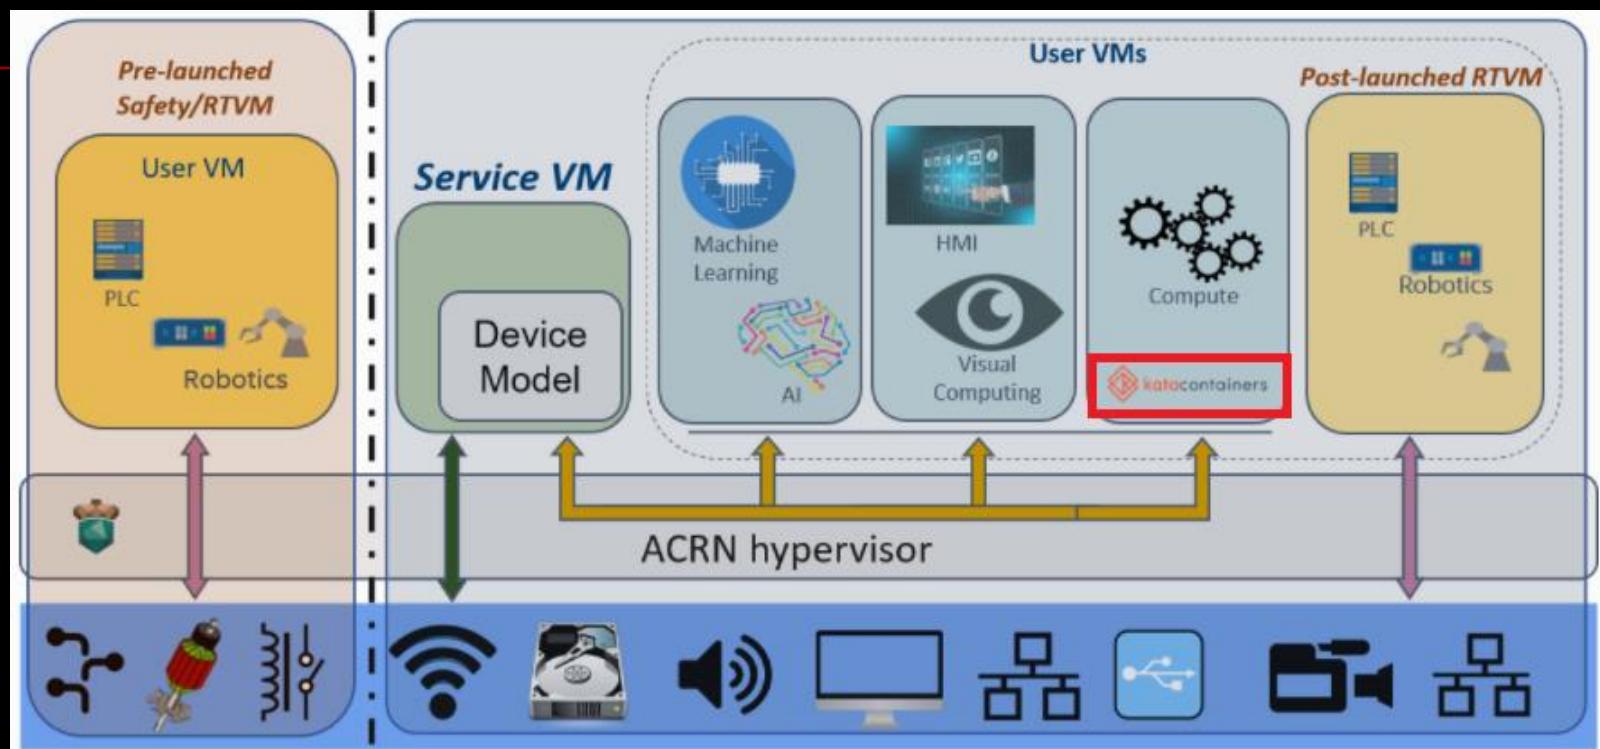
The ACRN Hypervisor is a Type 1 reference hypervisor stack, running directly on the bare-metal hardware, and is suitable for a variety of IoT and embedded device solutions. The ACRN hypervisor addresses the gap that currently exists between datacenter hypervisors, and hard partitioning hypervisors. The ACRN hypervisor architecture partitions the system into different functional domains, with carefully selected user VM sharing optimizations for IoT and embedded devices.

- [https://projectacrn.github.io/latest/release\\_notes/release\\_notes\\_2.4.html](https://projectacrn.github.io/latest/release_notes/release_notes_2.4.html)
- <https://projectacrn.org/acrn-hypervisor-service-module-upstreamed-to-linux-kernel/>  
<https://lwn.net/Articles/845372/> //HSM driver for ACRN hypervisor  
<https://www.kernel.org/doc/html/latest/virt/acrn/index.html>

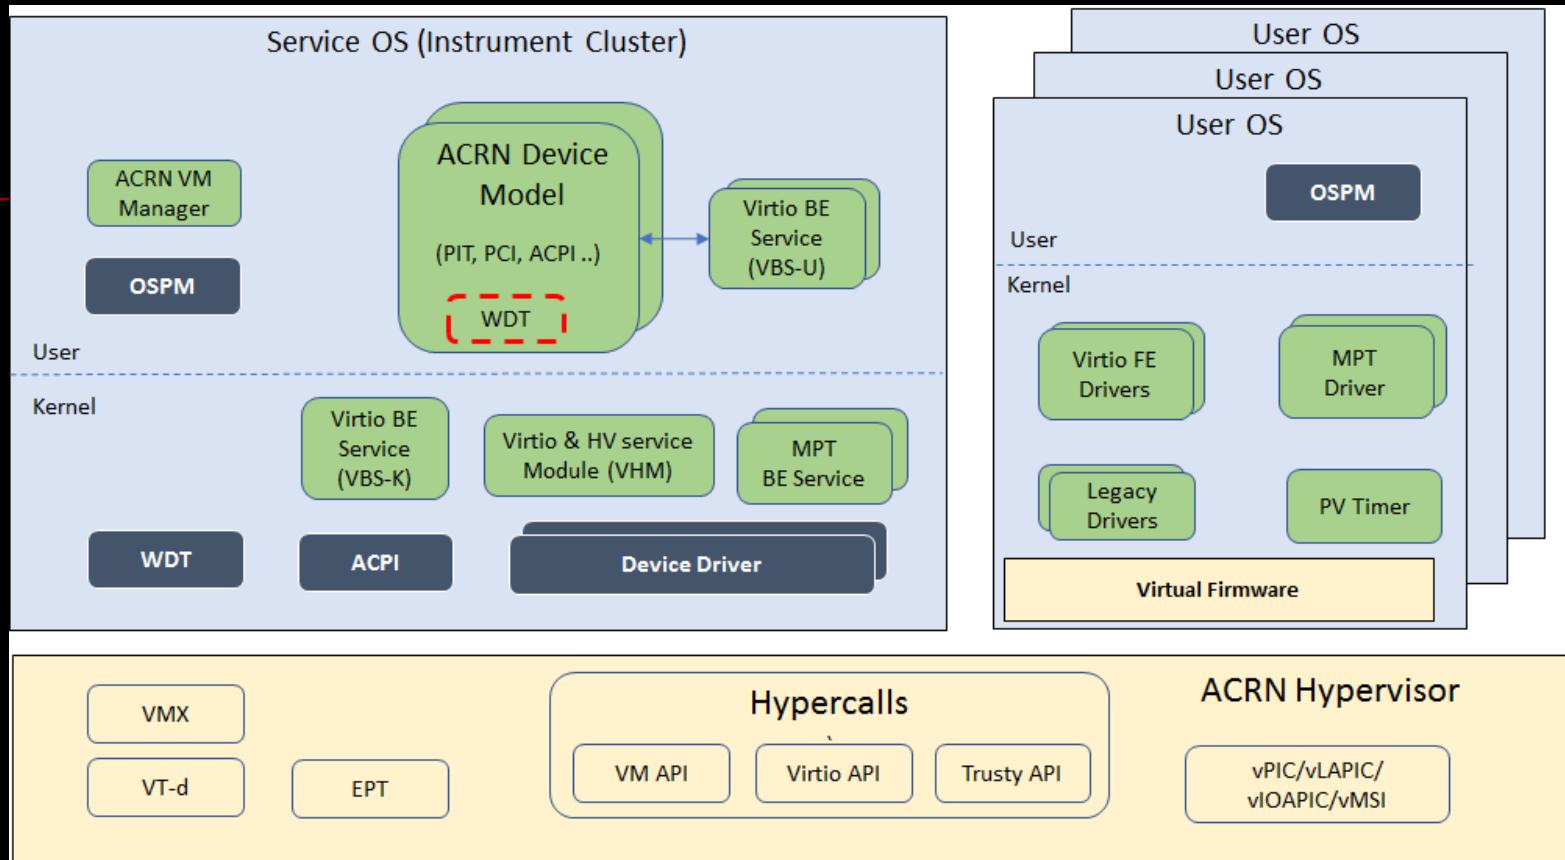
[https://kernelnewbies.org/Linux\\_5.12](https://kernelnewbies.org/Linux_5.12)

## Architecture

- <https://projectacrn.github.io/latest/introduction/index.html>
- High-level Architecture



## Hypervisor Architecture



## 2) X86 only

- <https://projectacrn.github.io/latest/reference/hardware.html>

### Minimum System Requirements for Installing ACRN

Hardware	Minimum Requirements	Recommended
Processor	Compatible x86 64-bit processor	2 core with Intel Hyper-threading Technology enabled in the BIOS or more cores
System memory	4GB RAM	8GB or more (< 32G)
Storage capabilities	20GB	120GB or more

### Minimum Requirements for Processor

1 GB Large pages

### Known Limitations

Platforms with multiple PCI segments

ACRN assumes the following conditions are satisfied from the Platform BIOS

- All the PCI device BARs should be assigned resources, including SR-IOV VF BARs if a device supports.
- Bridge windows for PCI bridge devices and the resources for root bus, should be programmed with values that enclose resources used by all the downstream devices.
- There should be no conflict in resources among the PCI devices and also between PCI devices and other platform devices.

- <https://github.com/kata-containers/kata-containers/blob/main/src/runtime/arch/arm64-options.mk>

## ***Run Kata Containers on a Service VM***

- [https://projectacrn.github.io/latest/tutorials/run\\_kata\\_containers.html](https://projectacrn.github.io/latest/tutorials/run_kata_containers.html)

### **Prerequisites**

1. Refer to the [ACRN supported hardware](#).
2. For a default prebuilt ACRN binary in the end-to-end (E2E) package, you must have 4 CPU cores or enable “CPU Hyper-threading” in order to have 4 CPU threads for 2 CPU cores.
3. Follow the [Getting Started Guide for ACRN Industry Scenario With Ubuntu Service VM](#) to set up the ACRN Service VM based on Ubuntu.
4. This tutorial is validated on the following configurations:
  - ACRN v2.0 (branch: `release_2.0`)
  - Ubuntu 20.04
5. Kata Containers are supported for ACRN hypervisors configured for the Industry or SDC scenarios.

- <https://github.com/kata-containers/packaging/tree/master/kata-deploy>
- <https://github.com/kata-containers/kata-containers/blob/main/docs/how-to/how-to-use-kata-containers-with-acrn.md>

## so many Kata projects archived recently...

**osbuilder** Archived  
Kata Containers version 1.x guest OS building scripts (for version 2.x see <https://github.com/kata-containers/kata-containers>).  
Shell Apache-2.0 84 ⭐ 105 0 Updated 10 days ago

---

**runtime** Archived  
Kata Containers version 1.x runtime (for version 2.x see <https://github.com/kata-containers/kata-containers>).  
docker kubernetes security containers virtual-machine  
virtualization container  
Go Apache-2.0 381 ⭐ 2,095 0 Updated 10 days ago

---

**community**  
Python Apache-2.0 55 ⭐ 193 5 2 Updated 11 days ago

---

**shim** Archived  
Kata Containers version 1.x shim (for version 2.x see <https://github.com/kata-containers/kata-containers>).  
Go Apache-2.0 52 ⭐ 67 0 0 Updated 12 days ago

---

**proxy** Archived  
Kata Containers version 1.x proxy (for version 2.x see <https://github.com/kata-containers/kata-containers>).  
Go Apache-2.0 56 ⭐ 56 0 0 Updated 12 days ago

---

**packaging** Archived  
Kata Containers version 1.x packaging (for version 2.x see <https://github.com/kata-containers/kata-containers>).  
Logos Apache-2.0 89 ⭐ 117 0 0 Updated 12 days ago

---

**ksm-throttler** Archived  
Kata Containers KSM throttling daemon  
Go Apache-2.0 21 ⭐ 24 0 0 Updated 12 days ago

---

**documentation** Archived  
Kata Containers version 1.x documentation (for version 2.x see <https://github.com/kata-containers/kata-containers>).  
Shell Apache-2.0 271 ⭐ 454 0 0 Updated 12 days ago

---

**agent** Archived  
Kata Containers version 1.x agent (for version 2.x see <https://github.com/kata-containers/kata-containers>). Virtual Machine agent for hardware virtualized containers  
Go Apache-2.0 118 ⭐ 233 0 0 Updated 12 days ago

## Implementation

### ■ Kata adaptations to support ACRN

#### ○ Kata-runtime:

- Added ACRN hypervisor as a supported hypervisor, so that kata config could pickup ACRN instead of QEMU when launching VM

```
[hypervisor.acrn]
path = "/usr/bin/acrn-dm"
kernel = "/usr/share/kata-containers/vmlinuz.container"
# initrd = "/usr/share/kata-containers/kata-containers-initrd.img"
image = "/usr/share/kata-containers/kata-containers.img"
```
- Implemented Sandbox Management APIs such as CreateSandBox, FetchSandBox..
  - <https://github.com/kata-containers/documentation/blob/master/design/kata-api-design.md#sandbox-management-api>
- Implemented Sandbox Operation APIs such as StartSandbox, StopSandbox, PauseSandbox..
  - <https://github.com/kata-containers/documentation/blob/master/design/kata-api-design.md#sandbox-operation-api>
- Implemented hot-plug API (currently have a WA. Working on adding PCI device hot-plug support in ACRN-DM)
  - <https://github.com/kata-containers/documentation/blob/master/design/kata-api-design.md#sandbox-hotplug-api>
- Prime the devices, image and kernel as parameters for ACRN-DM when launching the VM.

Source: <https://www.slideshare.net/ProjectACRN/acrn-kata-container-on-acrn>

## ■ ACRN-DM features to support Kata

- Socket Backend support:
  - Kata uses socket communication to talk between kata-runtime <->Kata-proxy, Kata-shim<->Kata-Proxy and Kata-Proxy<->Kata-agent running inside VM.
  - Implemented socket backend for virtio-console device in acrn-dm. (This feature is already merged upstream)
    - -s x,virtio-console,socket:"socket name"="socket path" where x is the PCI slot number.
- PCI device Hot-Plug support:
  - Kata containers does hot plugging of container roofs for both Docker and Kubernetes.
    - Kubernetes when creating a pod, first creates a \*pause container and then subsequently creates the application container(s). During the launch of the VM, the roofs for the application container is not known, it needs to be hot-plugged.
  - Looking at both ACPI based hot-plug and PCI/PCIe hot-plug support for PCI device (container roots is passed as virtio-blk device). Scoped out initial design and code changes that would be needed in acrn-dm. (WIP)
- Graceful shutdown of VMs:
  - When kata containers finishes its job, expectation is that VM associated with container will be shutdown gracefully. (WIP)

Source: <https://www.slideshare.net/ProjectACRN/acrn-kata-container-on-acrn>

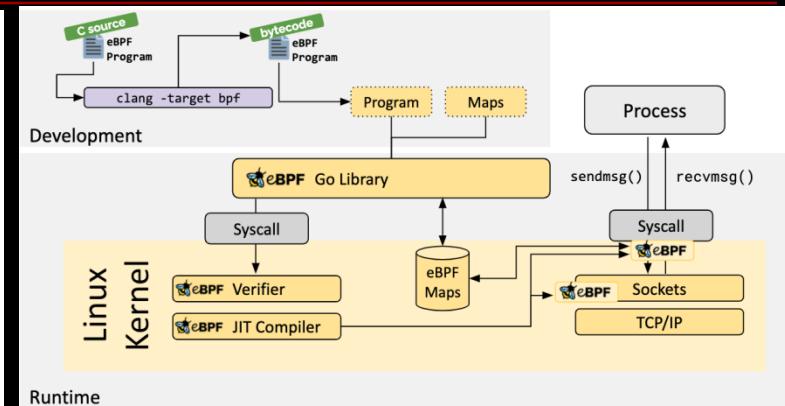
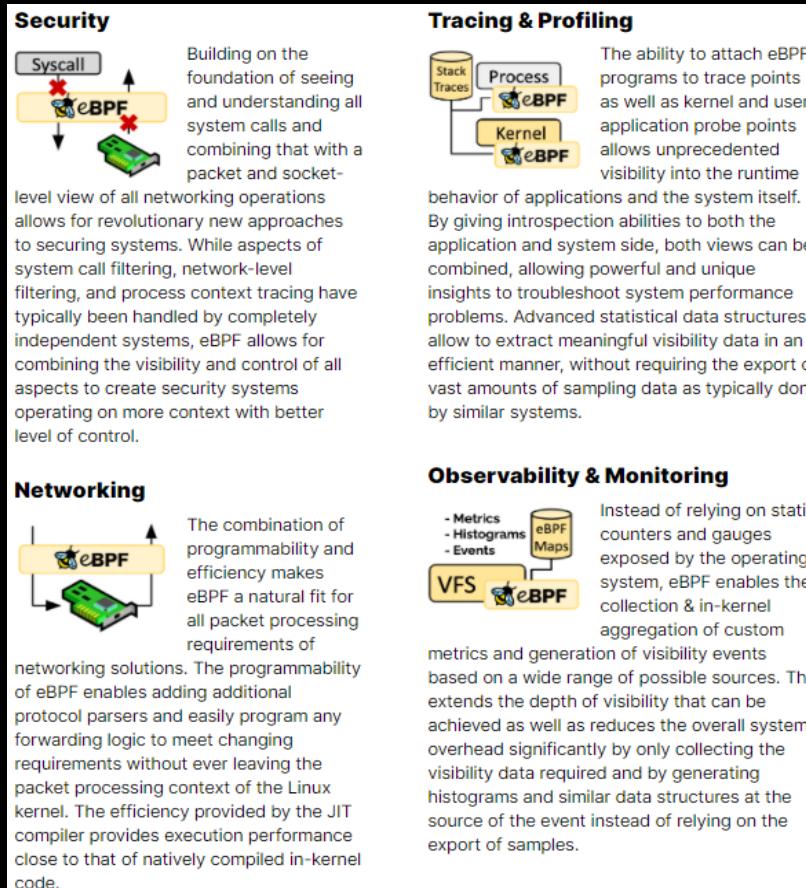
## V. Rethinking Infrastructure for Edge Cloud

### 1) Why Linux is critical to Edge Computing

- <https://opensource.com/article/21/2/linux-edge-computing>
-

## 2) eBPF-centric

- [https://en.wikipedia.org/wiki/Berkeley\\_Packet\\_Filter](https://en.wikipedia.org/wiki/Berkeley_Packet_Filter)
- <https://www.kernel.org/doc/html/latest/bpf/index.html>
- <https://ebpf.io/>



### *List of BPF features per kernel version*

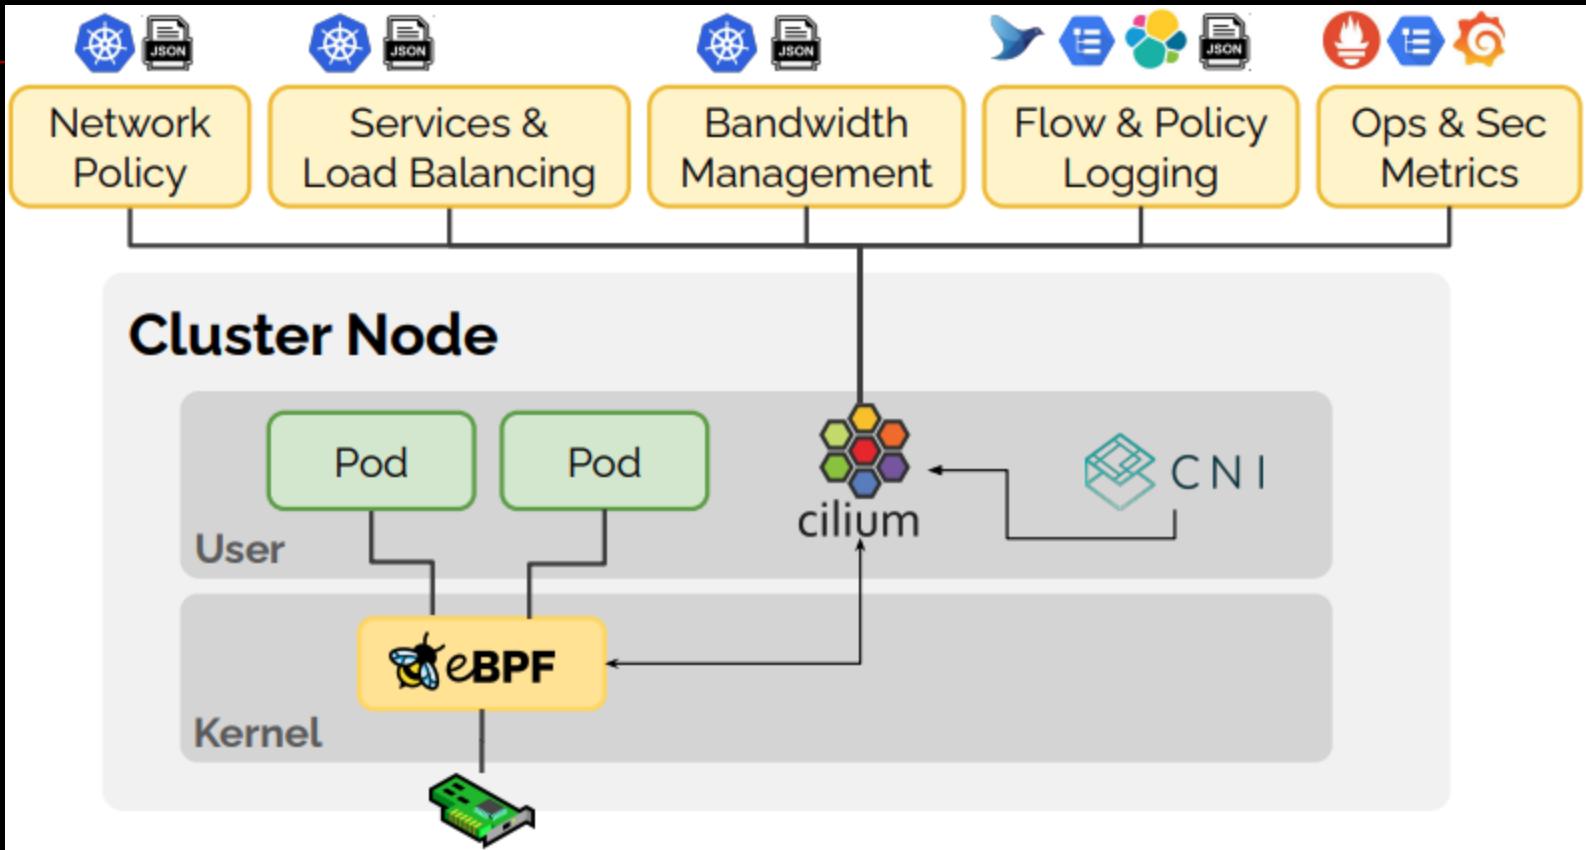
- <https://github.com/iovisor/bcc/blob/master/docs/kernel-versions.md>
  - <https://kernelnewbies.org/>
-

## 2.1 eBPF in Kata Containers

- <https://github.com/cilium/ebpf>
  - <https://github.com/opencontainers/runc/tree/master/libcontainer/cgroups/ebpf>
  - <https://github.com/vishvananda/netlink>
  - ...
-

## Cilium

- <https://cilium.io/>
- **eBPF-based Networking, Security, and Observability**



- <https://cilium.io/blog/2021/05/11/cni-benchmark>

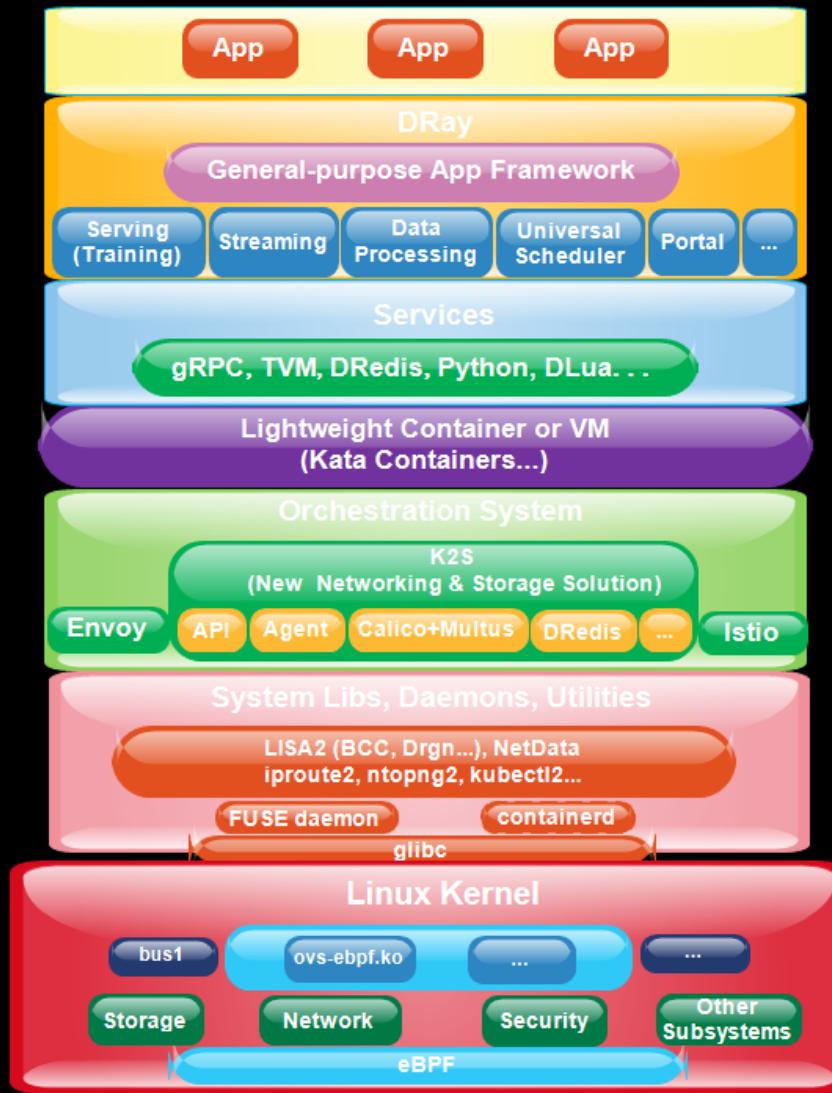
## 2.2 Our New Design

- an eBPF-centric new lightweight for Edge Computing
- Stage I



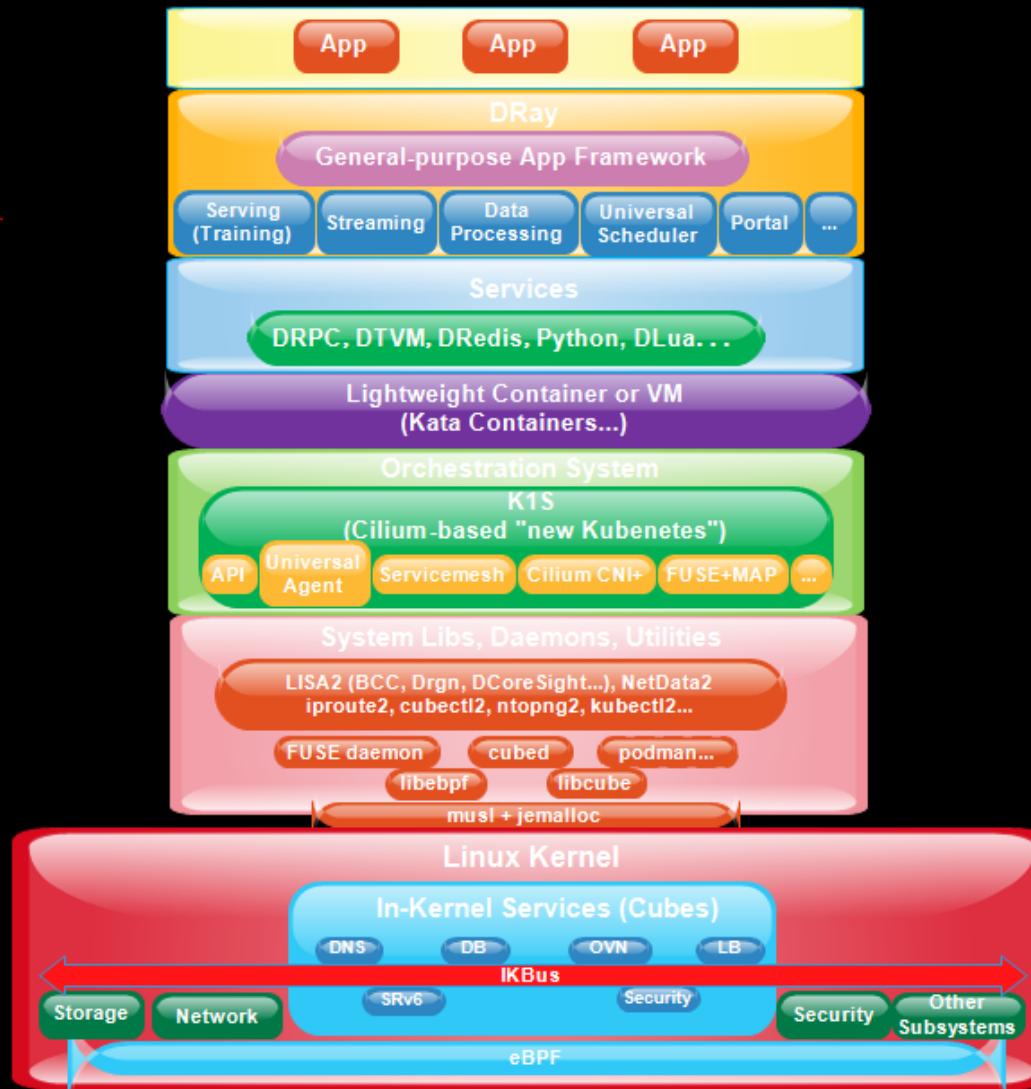
Source: “Rethinking Hyper-Converged Infrastructure for Edge Computing”, Feng Li, OpenInfra Days 2019(Shanghai)

## ■ Stage II



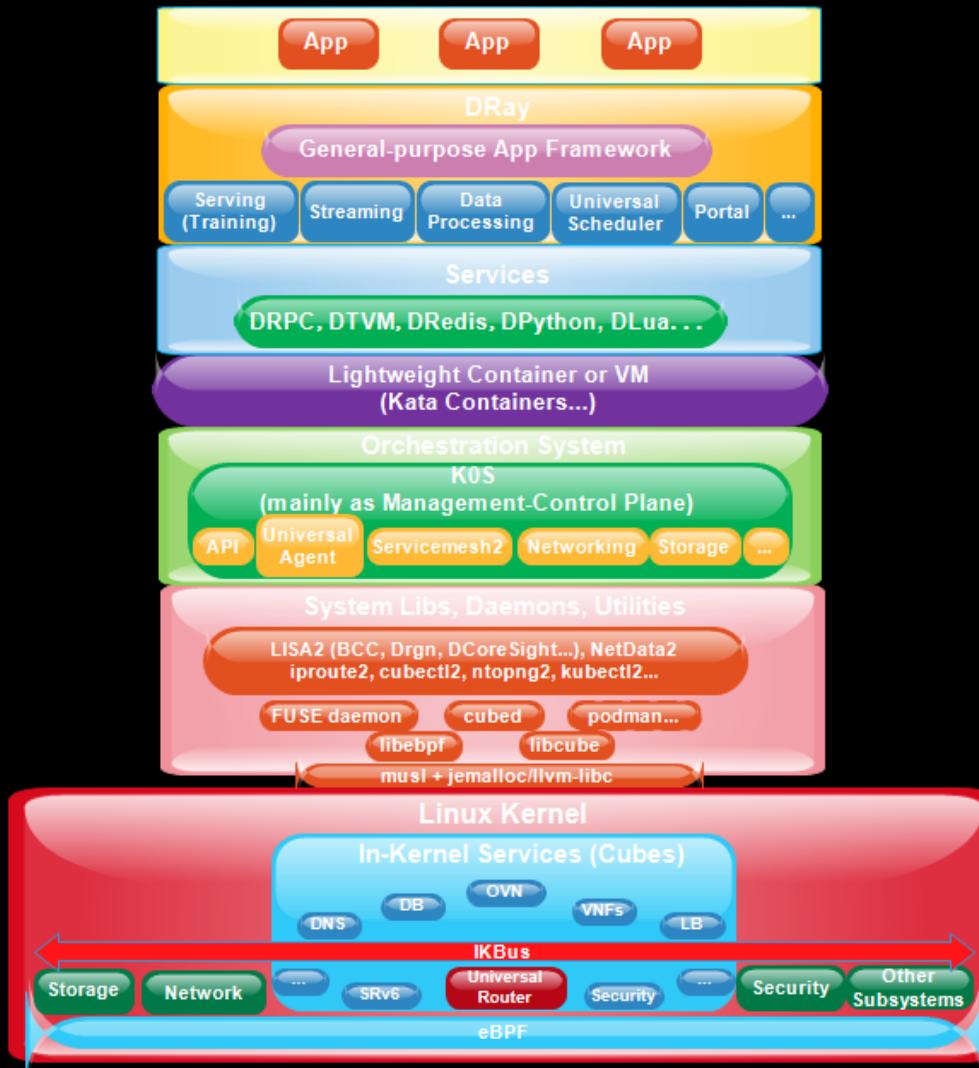
Source: “Rethinking Hyper-Converged Infrastructure for Edge Computing”, Feng Li, OpenInfra Days 2019(Shanghai)

## ■ Stage III



Source: “Rethinking Hyper-Converged Infrastructure for Edge Computing”, Feng Li, OpenInfra Days 2019(Shanghai)

## ■ Stage IV



Source: “Rethinking Hyper-Converged Infrastructure for Edge Computing”, Feng Li, OpenInfra Days 2019(Shanghai)

### 3) Rust for Linux Kernel Development

- <https://lwn.net/Articles/853423/> //Rust heads into the kernel?
- <https://lwn.net/Articles/852704/> //Rust in the Linux kernel
- <https://lwn.net/Articles/849849/> //Rust support hits linux-next
- <https://lwn.net/Articles/829858/> //Supporting Linux kernel development in Rust
- <https://blogs.gartner.com/manjunath-bhat/2021/01/03/why-2021-will-be-a-rusty-year-for-system-programmers/>
- <https://developers.slashdot.org/story/21/04/17/009241/linus-torvalds-says-rust-closer-for-linux-kernel-development-calls-c-a-crap-language>
- <https://www.infoq.com/news/2021/04/rust-linux-kernel-development/>
- <https://itwire.com/open-source/rust-support-in-linux-may-be-possible-by-5-14-release-torvalds.html>
- <https://github.com/libbpf/libbpf-rs> //Idiomatic rust wrapper around libbpf
- <https://github.com/foniod/redbpf> //A Rust eBPF toolchain

## Cloud Hypervisor

### ■ <https://github.com/cloud-hypervisor/cloud-hypervisor>

Cloud Hypervisor is an open source Virtual Machine Monitor (VMM) that runs on top of [KVM](#) and the [MSHV](#) hypervisors .

The project focuses on exclusively running modern, cloud workloads, on top of a limited set of hardware architectures and platforms. Cloud workloads refers to those that are usually run by customers inside a cloud provider. For our purposes this means modern operating systems with most I/O handled by paravirtualised devices (i.e. virtio), no requirement for legacy devices, and 64-bit CPUs.

Cloud Hypervisor is implemented in [Rust](#) and is based on the [rust-vmm](#) crates.

## ■ Objectives

### High Level

- Runs on KVM or MSHV
- Minimal emulation
- Low latency
- Low memory footprint
- Low complexity
- High performance
- Small attack surface
- 64-bit support only
- CPU, memory, PCI hotplug
- Machine to machine migration

### Architectures

Cloud Hypervisor supports the `x86-64` and `AArch64` architectures. There are some small differences in functionality between the two architectures (see [#1125](#)).

### Guest OS

Cloud Hypervisor supports `64-bit Linux` and `Windows 10/Windows Server 2019`.

## 4) WebAssembly

- <https://en.wikipedia.org/wiki/WebAssembly>
  - <https://www.rust-lang.org/what/wasm>
  - <https://rustwasm.github.io/book>
  - <https://github.com/rustwasm>
- 

### In-kernel WASM

- <https://github.com/wasmerio/kernel-wasm>  
*//sandboxed kernel mode WebAssembly runtime*

- WASI support (incomplete; work in progress)
- Asynchronous networking extension with `epoll` support
- Modular host API provider interface
- Fully sandboxed execution environment with software fault isolation
- Faster than native (partially achieved)
- Device drivers in WASM
- "eBPF" in WASM

- <https://github.com/kenny-ngo/wasmjit>

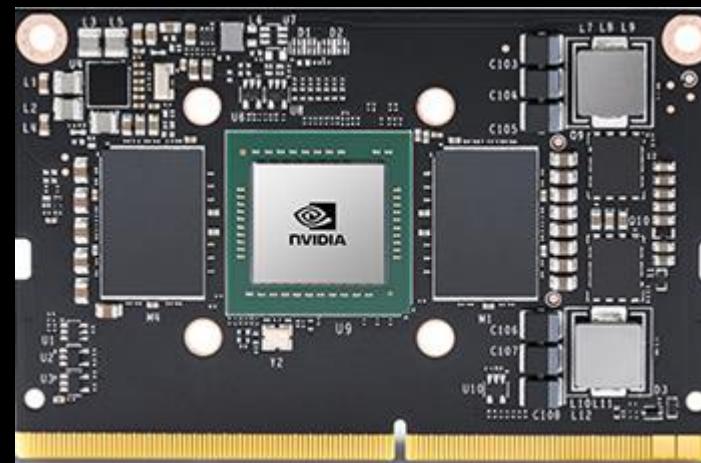
Its primary target is a Linux kernel module that can host Emscripten-generated WebAssembly modules. In this configuration it runs WebAssembly modules in kernel-space (ring 0) and provides access to system calls as normal function calls. This configuration avoids user-kernel transition overhead, as well as scheduling overheads from swapping page tables. This results in a significant performance increase for syscall-bound programs like web servers or FUSE file systems. The performance increase is more pronounced in a post-Spectre/Meltdown world due to PTI overhead. Check it out:

## 5) Edge AI

- **Definition?**
- <https://towardsdatascience.com/edge-ai-is-the-next-wave-of-ai-a3e98b77c2d7>
- <https://wiki.akraino.org/display/AK/The+AI+Edge+Blueprint+Family>

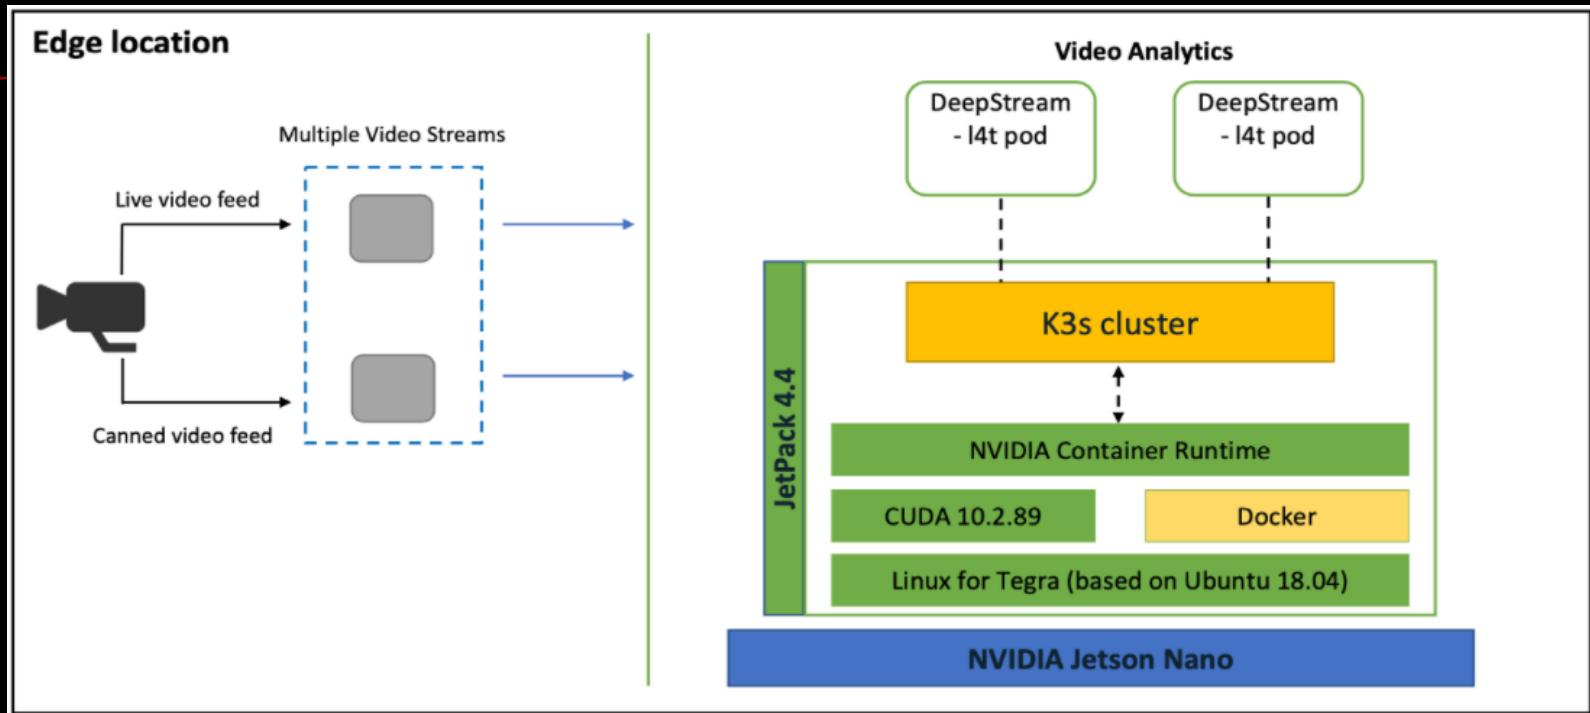
## 5.1 Nvidia Jetson

- <https://www.nvidia.com/en-us/autonomous-machines/jetson-store/>  
**Jetson Nano; Jetson Xavier NX; Jetson AGX Xavier; Jetson TX2...**



## AI at the Edge with K3s and NVIDIA Jetson Nano

- <https://www.suse.com/c/ai-at-the-edge-with-k3s-nvidia-jetson-nano-object-detection-real-time-video-analytics-src/>



## NVidia Jetson Lab Environment - exploring k3s on ARM

- <https://github.com/cloudxabide/jetsons.lab/>



# Running Kata Containers on Jetson series?



## ■ Ubuntu 18.04.5 LTS

```
mydev@mydev-jetson1:~$ uname -a
Linux mydev-jetson1 4.9.201-tegra #1 SMP PREEMPT Fri Feb 19 08:40:32 PST 2021 aarch64 aarch64 aarch64 GNU/Linux
mydev@mydev-jetson1:~$
mydev@mydev-jetson1:~$ dmesg |grep -i kvm
mydev@mydev-jetson1:~$
mydev@mydev-jetson1:~$ sudo zcat /proc/config.gz |grep -i kvm
[sudo] password for mydev:
# CONFIG_KVM is not set
mydev@mydev-jetson1:~$
mydev@mydev-jetson1:~$ sudo zcat /proc/config.gz |grep -i virtual
# CONFIG_REGULATOR_VIRTUAL_CONSUMER is not set
# CONFIG_FB_VIRTUAL is not set
CONFIG_DMA_VIRTUAL_CHANNELS=y
# CONFIG_TEGRA_VIRTUALIZATION is not set
CONFIG_VIRTUALIZATION=y
# CONFIG_CRYPTO_DEV_TEGRA_VIRTUAL_SE_INTERFACE is not set
mydev@mydev-jetson1:~$
mydev@mydev-jetson1:~$ sudo zcat /proc/config.gz |grep -i cgroup
CONFIG_CGROUPS=y
CONFIG_CGROUP_DEBUG=y
CONFIG_CGROUP_FREEZER=y
CONFIG_CGROUP_PIDS=y
CONFIG_CGROUP_DEVICE=y
CONFIG_CGROUP_CPUACCT=y
CONFIG_BLK_CGROUP=y
# CONFIG_DEBUG_BLK_CGROUP is not set
CONFIG_CGROUP_WRITEBACK=y
CONFIG_CGROUP_SCHED=y
CONFIG_CGROUP_HUGETLB=y
CONFIG_CGROUP_PERF=y
CONFIG_SOCK_CGROUP_DATA=y
# CONFIG_NETFILTER_XT_MATCH_CGROUP is not set
CONFIG_NET_CLS_CGROUP=y
CONFIG_CGROUP_NET_PRIO=y
CONFIG_CGROUP_NET_CLASSID=y
```

## ■ <https://developer.nvidia.com/embedded/linux-tegra>

NVIDIA L4T is the board support package for Jetson. It includes Linux Kernel 4.9, bootloader, NVIDIA drivers, flashing utilities, sample filesystem based on Ubuntu 18.04, and more for the Jetson platform.

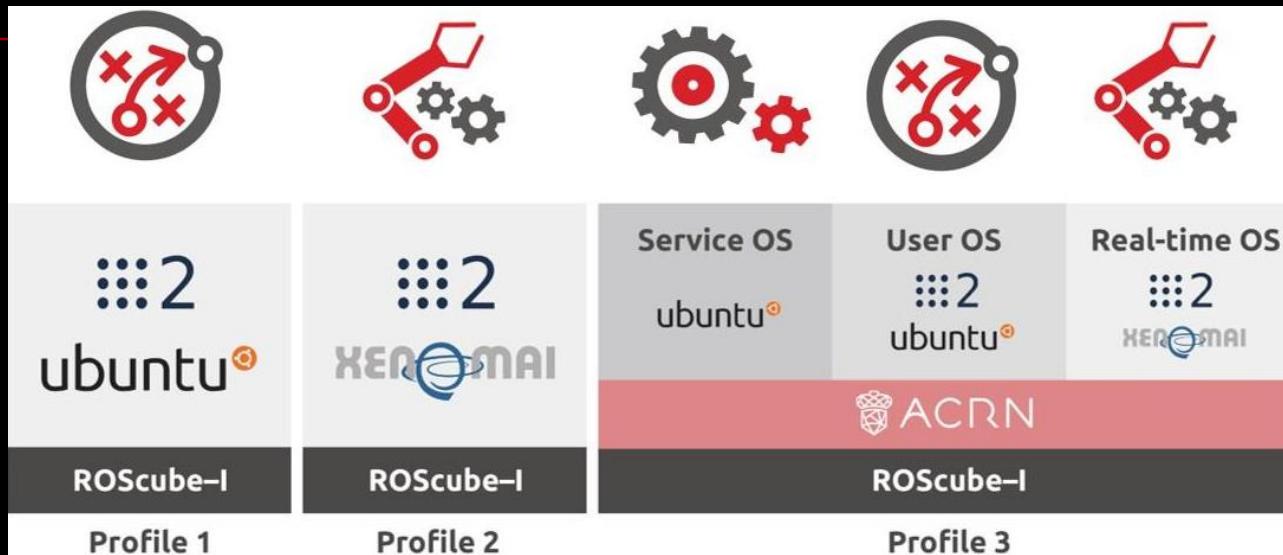
NVIDIA L4T 32.5.1 supports all Jetson modules: Jetson AGX Xavier series, Jetson Xavier NX, Jetson TX2 series, Jetson TX1, and Jetson Nano. All Jetson developer kits are also supported.

L4T 32.5.1 is included as part of [JetPack 4.5.1](#)

## 5.2 ROS

### ROScube-I

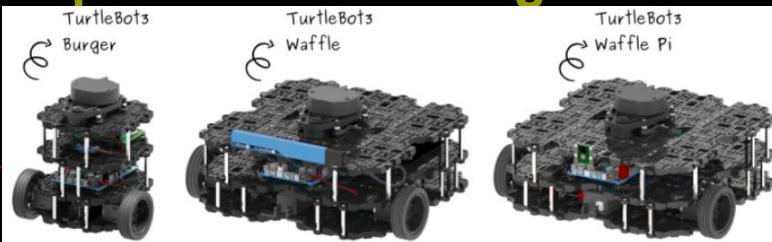
- <https://www.adlinktech.com/en/News/ADLINK-ROScube-I-ROS-2-controller-now-with-ACRN-real-time-hypervisor>



- <https://projectacrn.github.io/latest/getting-started/roscube/roscube-gsg.html>

## Turtlebot3

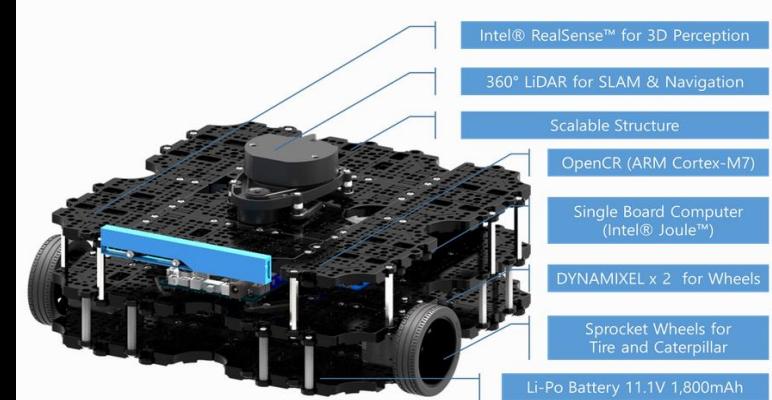
- <https://robots.ros.org/turtlebot3/>



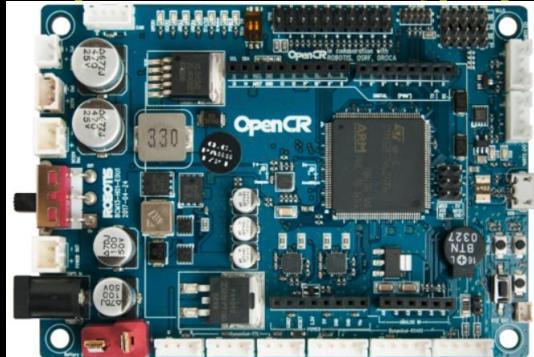
Burger



Waffle



- <http://wiki.ros.org/opencr>



## **Containerization of ROS2**

- <https://docs.ros.org/en/foxy/Guides/Run-2-nodes-in-single-or-separate-docker-containers.html>
- <https://developer.nvidia.com/blog/implementing-robotics-applications-with-ros-2-and-ai-on-jetson-platform-2/>
- <https://ubuntu.com/blog/ros-2-on-kubernetes-a-simple-talker-and-listener-setup>
- <https://discourse.ros.org/t/ros-2-on-kubernetes/17182>

## 6) Serverless

- [https://en.wikipedia.org/wiki/Serverless\\_computing](https://en.wikipedia.org/wiki/Serverless_computing)
- Open Source Serverless Platforms



## 6.1 vAccel

■ <https://vaccel.org/>

■ Give your **Serverless** deployments the power of hardware accelerators with vAccel

■ **Problem: hardware acceleration in the Cloud & at the Edge**

### Hardware partitioning

- bound to hardware device/vendor support
- inflexible sharing of diverse accelerator resources

### API remoting

- still either device/vendor API specific, or
- can incur significant performance overhead
- not fit for infrastructures with resource or performance (i.e. latency) constraints

### Paravirtualization

- users have to program the hardware directly
- multiple schedulers doing the same job (VM, VMM, runtime system)
- software stack duplication

Source: “ML inference acceleration for lightweight VMMs”, Anastassios Nanos, Babis Chalios, Kostis Papazafeiropoulos, Fosdem 2021

■ **Design Goals**

- programmability / simplicity (device/vendor-agnostic)
- performance (minimal overhead)
- portability / interoperability (run anywhere)
- security / isolation (virtualization support)

Source: “ML inference acceleration for lightweight VMMs”, Anastassios Nanos, Babis Chalios, Kostis Papazafeiropoulos, Fosdem 2021

# Architecture & Design



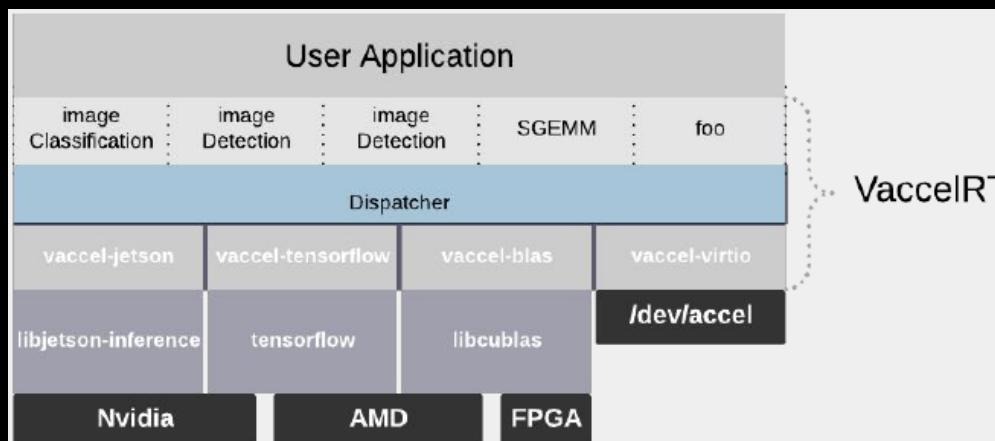
core component: vAccelRT (vAccel runtime system)

user-facing API: function prototypes

- abstracted by the underlying frameworks or
- defined by the system as a superset / subset of individual acceleration functions

hardware abstraction layer: **acceleration frameworks, transport layer**

- low-level APIs (openCL, CUDA, openACC etc.)
- higher-level frameworks (TensorRT, tensorflow, pytorch etc.)
- user-facing APIs (jetson-inference, libBLAS etc.)
- virtio-accel

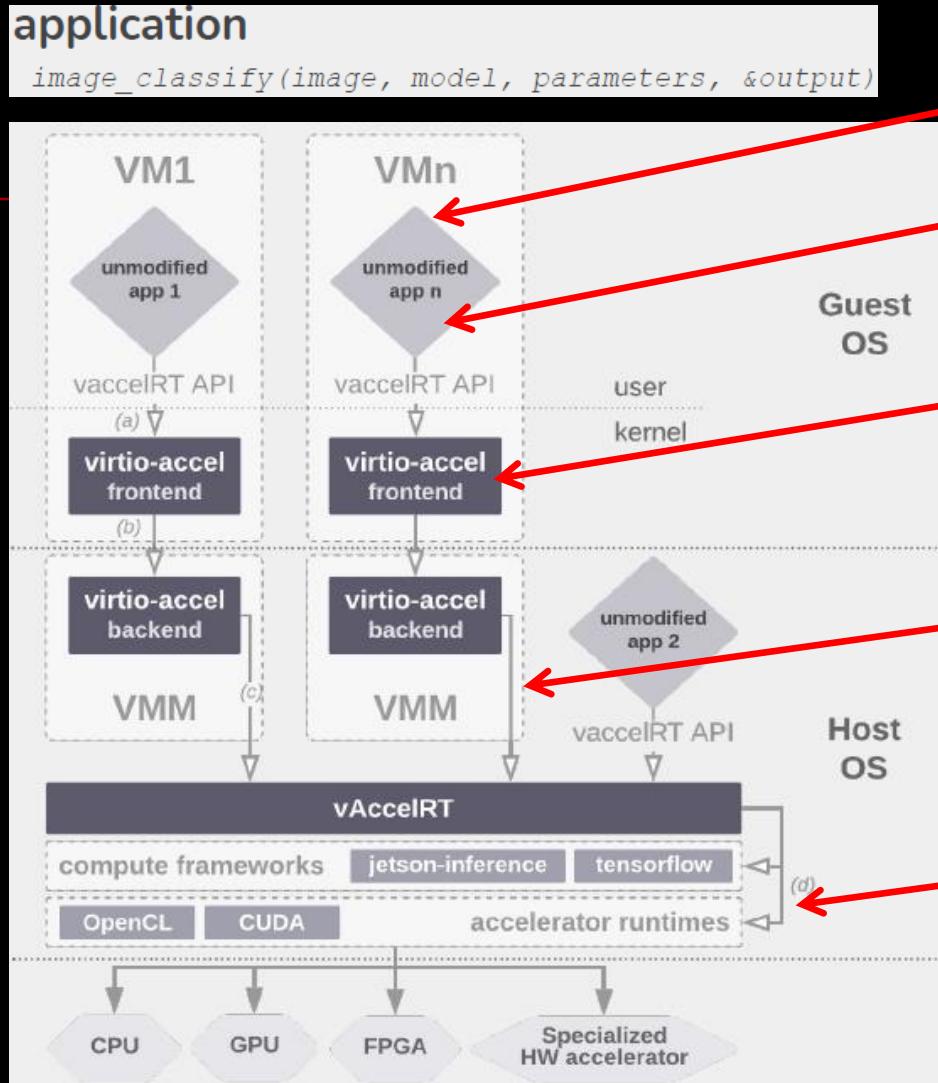


Source: “ML inference acceleration for lightweight VMMs”, Anastassios Nanos, Babis Chalios, Kostis Papazafeiopoulos, Fosdem 2021

## Example execution flow for a VM

### application

```
image_classify(image, model, parameters, &output)
```



#### VaccelRT (inside the guest)

1. Look for backend plugin that implements image classification
2. If supported, call the plugin code

#### vaccel-virtio plugin

1. Check that the virtio device is present
2. Prepare arguments and issue `ioctl` command to the vaccel device

#### vaccel-virtio front-end driver

1. Create virtio request
2. Insert user arguments inside the request
3. Insert the request inside the virtqueue
4. Kick the virtqueue

#### VMM

1. The guest VM exits in the VMM
2. The virtio-backend inside the VMM
  - a. Parses the request
  - b. Validates the request header and corresponding arguments
  - c. `image_classify(image, model, parameters, &output)`

#### VaccelRT (in the host)

1. Look for the backend plugin that implements image classification
2. If supported, call the plugin code

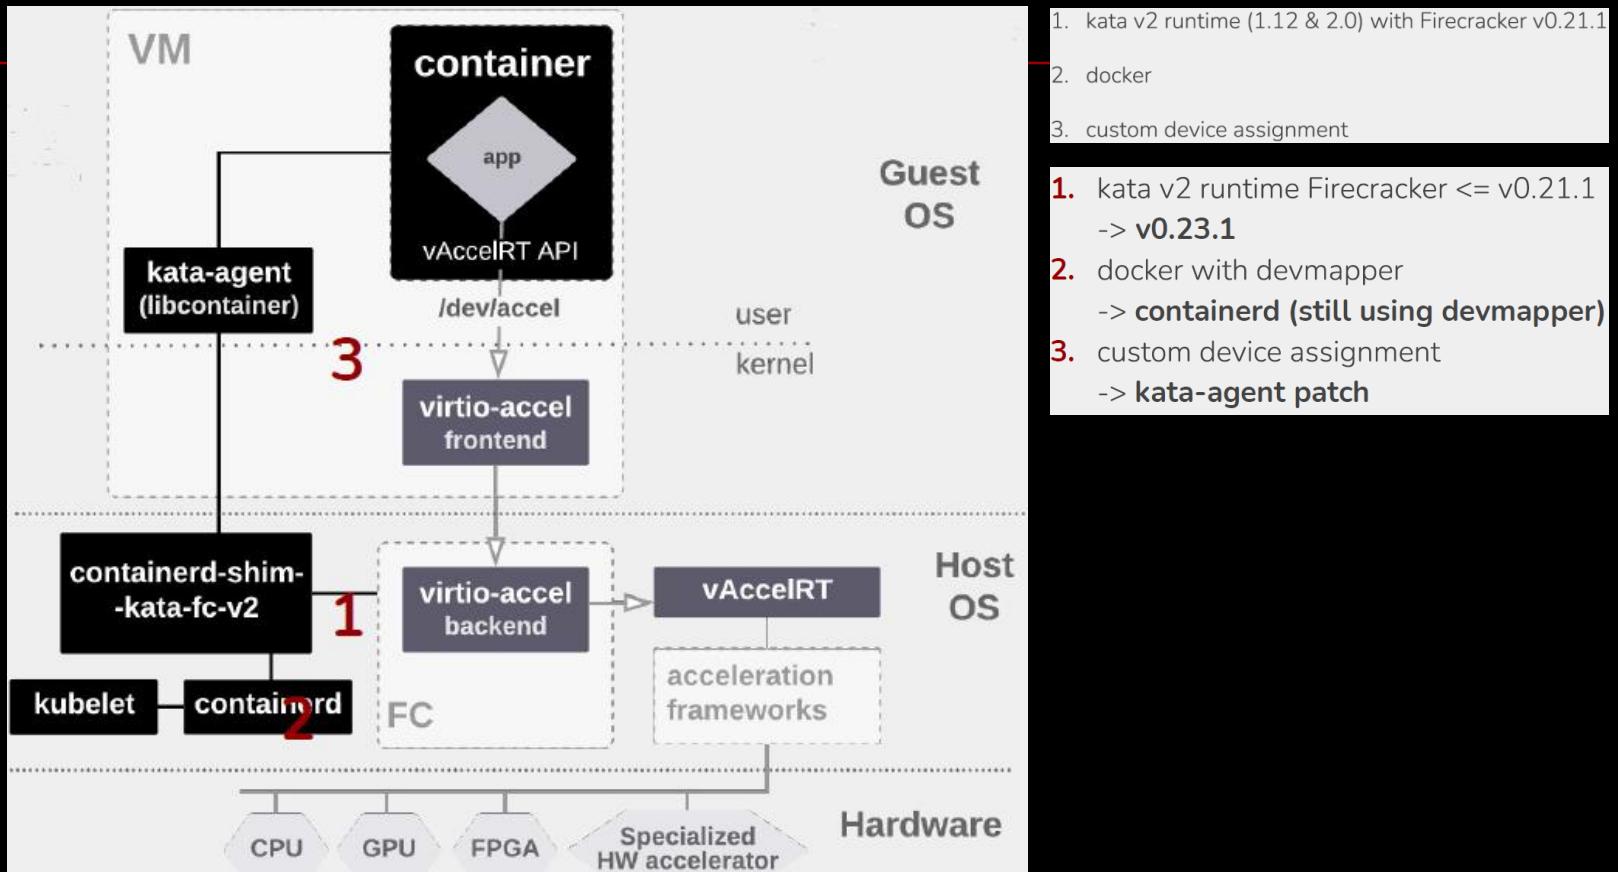
#### vaccel-jetson plugin

1. Perform operation using jetson-inference framework
2. Offload computation to GPU

Source: “ML inference acceleration for lightweight VMMs”, Anastassios Nanos, Babis Chalios, Kostis Papazafeiropoulos, Fosdem 2021

## Orchestration

- <https://docs.vaccel.org/kata-vaccel/>
- AWS Firecracker with vAccel on K8S(using Kata Containers)



Source: “ML inference acceleration on K8s using kata containers & AWS Firecracker”, Orestis Lagkas Nikolos, Babis Chalios, Anastassios Nanos, Fosdem 2021

- <https://docs.vaccel.org/unikernels/unikraft/>

## 7) Unified Runtime for eBPF & WASM

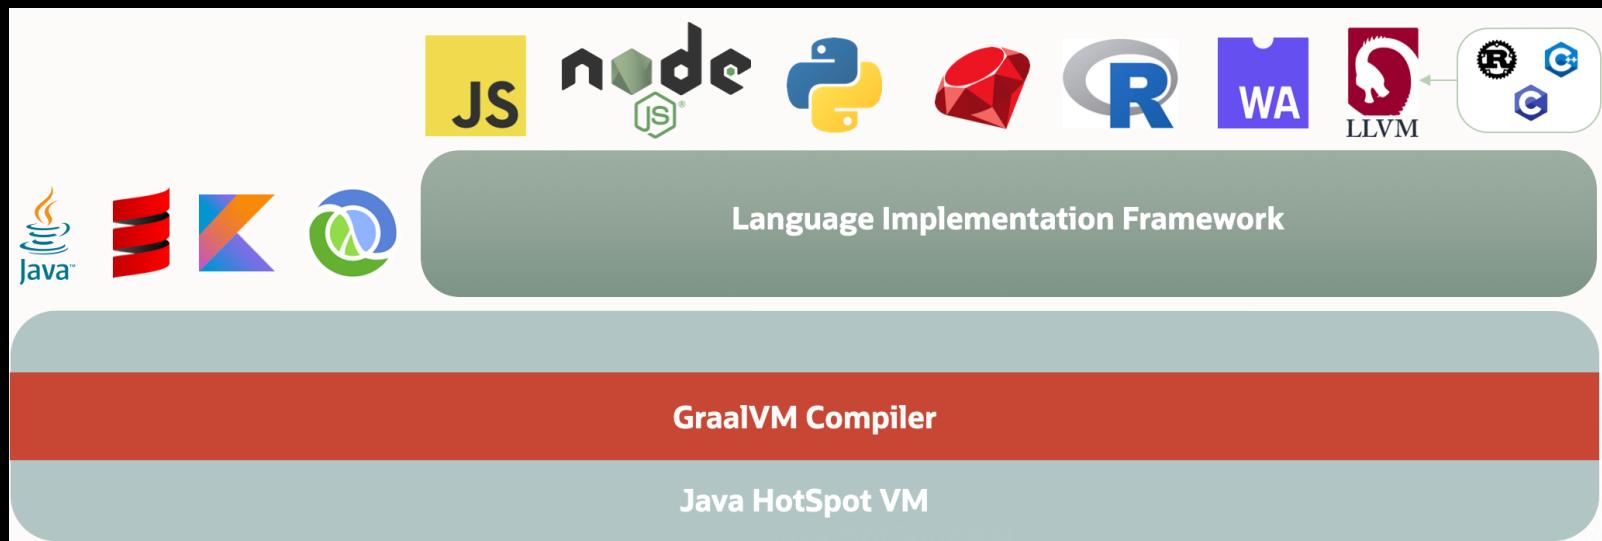
### In-Kernel

- AIO runtime in **Rust**?

### User space

- Implement uBPF in **GraalVM**

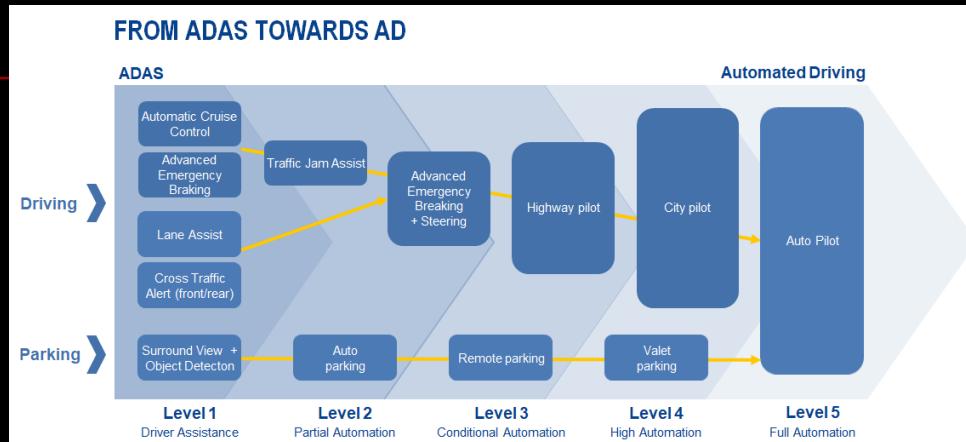
<https://github.com/oracle/graal/blob/master/wasm/README.md>



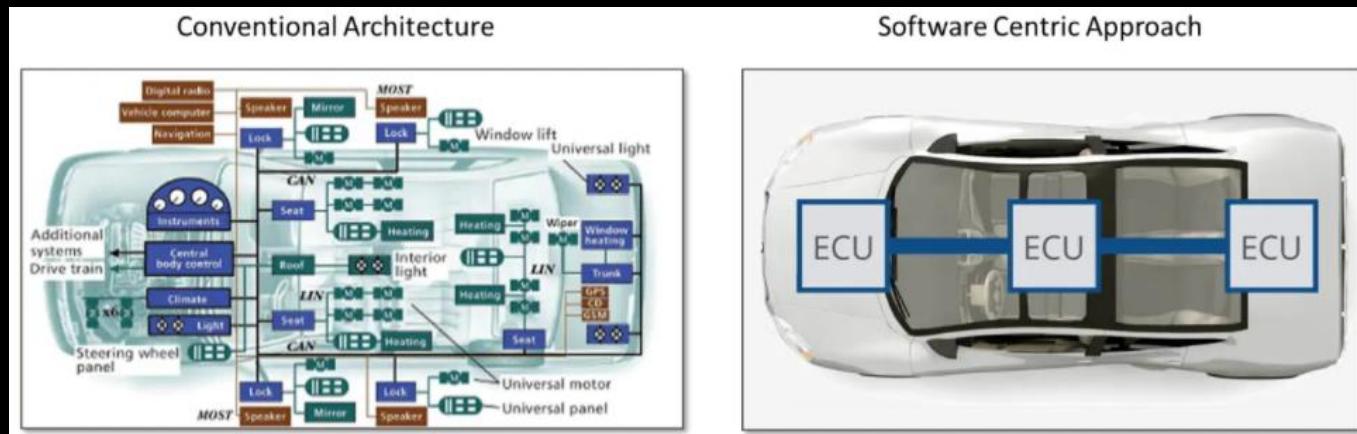
## 8) Autonomous Vehicles

[https://en.wikipedia.org/wiki/Self-driving\\_car](https://en.wikipedia.org/wiki/Self-driving_car)

[https://en.wikipedia.org/wiki/Advanced\\_driver-assistance\\_systems](https://en.wikipedia.org/wiki/Advanced_driver-assistance_systems)



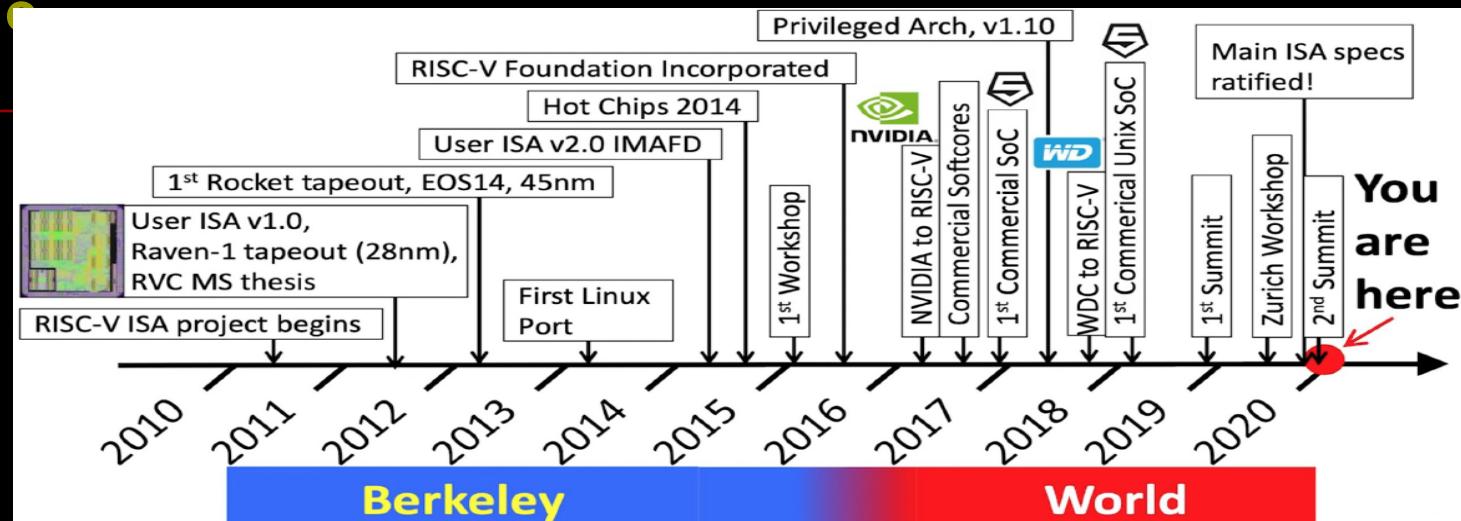
<https://www.autovision-news.com/hmi/vehicle-control/changes-vehicle-electrical-architecture/>



[https://en.wikipedia.org/wiki/Time-Sensitive\\_Networking](https://en.wikipedia.org/wiki/Time-Sensitive_Networking)

## 9) RISC-V

<https://en.wikipedia.org/wiki/RISC-V>



Source: "Linux on RISC-V", D. Fustimi, ELC2020

## A New Golden Age for Computer Architecture

- <https://cacm.acm.org/magazines/2019/2/234352-a-new-golden-age-for-computer-architecture/fulltext>



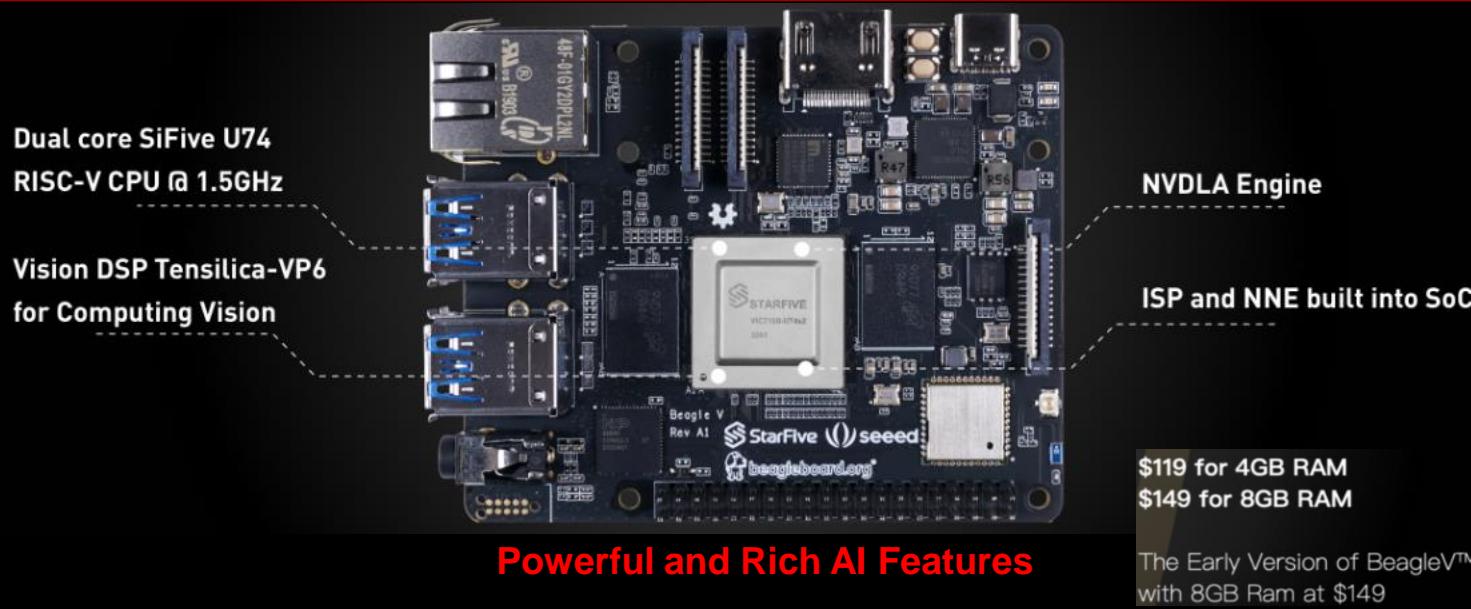
**“the Linux of the chip world”**

- <https://www.zdnet.com/article/risc-v-the-linux-of-the-chip-world-is-starting-to-produce-technological-breakthroughs/>



## BeagleV

- <https://beagleboard.org/beaglev>
- <https://beaglev.seeed.cc/>
- **The First Affordable RISC-V Computer Designed to Run Linux**



## TRULY OPEN SOURCE



Open Source Software

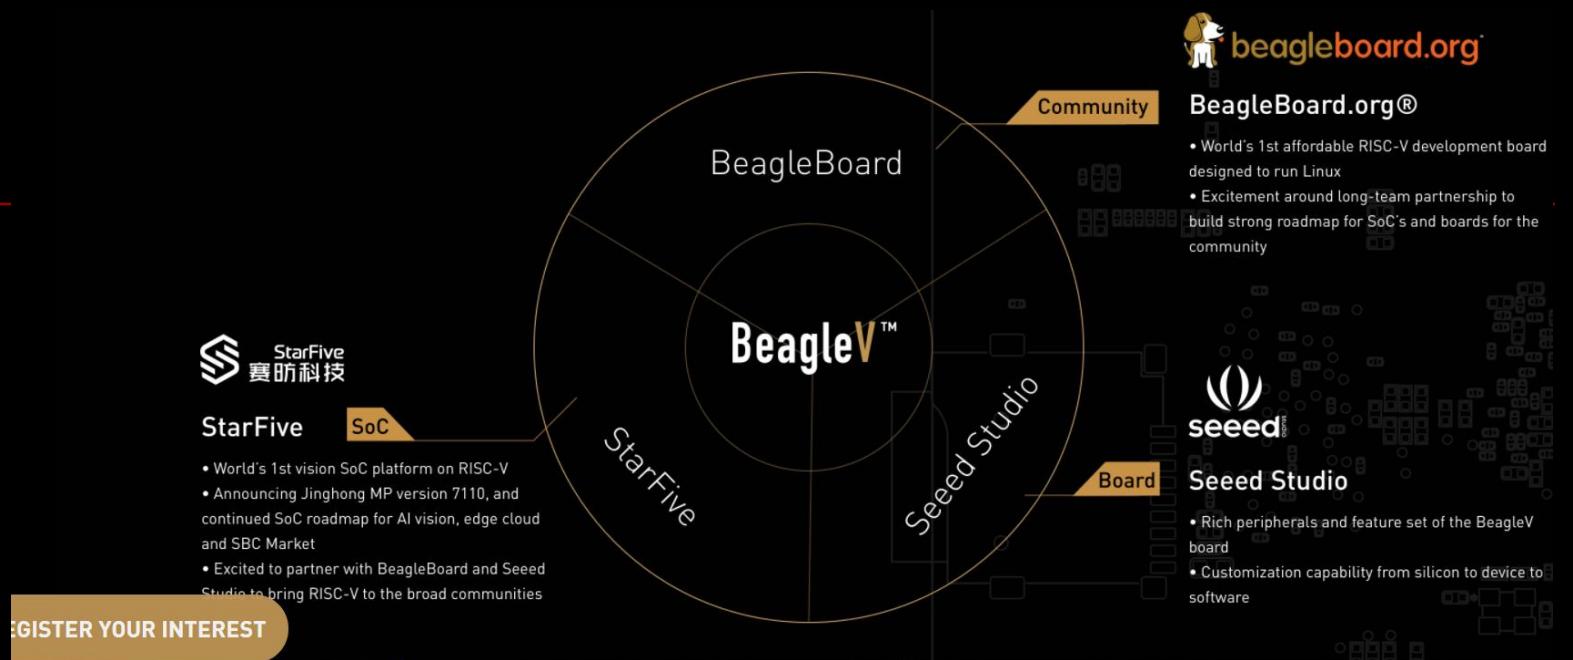


Based on RISC-V Open Architecture



Open Hardware Design

## ■ Strategy Partners



<https://www.imaginationtech.com/news/press-release/imagination-gpu-selected-by-starfive-to-create-high-performance-small-and-low-cost-beaglev-risc-v-ai-single-board-computer/>

### IMG B-SERIES MULTI-CORE

Unprecedented performance for GPU IP



### TILE REGION PROTECTION

Enabling safety-critical rendering for automotive applications

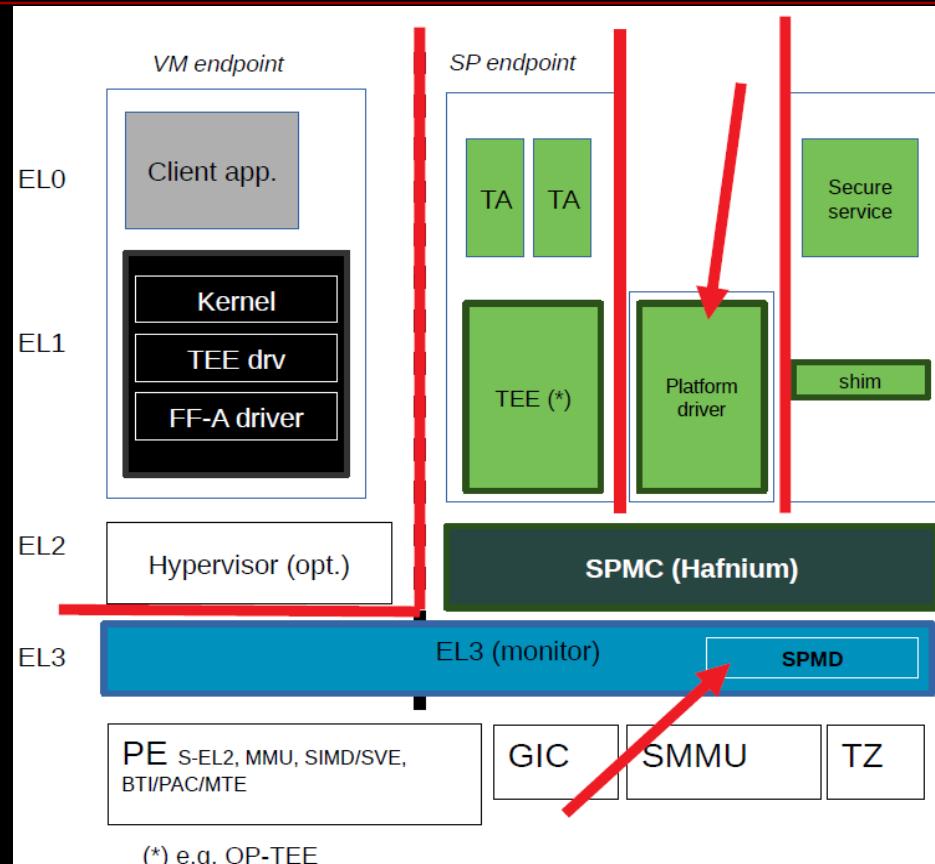


## **10) HW-SW Co-designed System Security**

- **Hardware-Software Co-designed System and System Security is the trend in the real world**
  - [https://en.wikipedia.org/wiki/Trusted\\_execution\\_environment](https://en.wikipedia.org/wiki/Trusted_execution_environment)
-

## 10.1 Security on ARM

- <https://developer.arm.com/architectures/architecture-security-features>
- ARMv8.4 Secure EL2



Source: “Secure Partition Manager(Armv8.4 Secure EL2)”, Olivier Deprez & Madhukar Pappireddy, LVC21-303

## Project Cassini for a Cloud-Native Edge

- <https://www.arm.com/solutions/infrastructure/edge-computing/project-cassini>
- extends PSA to the infrastructure Edge

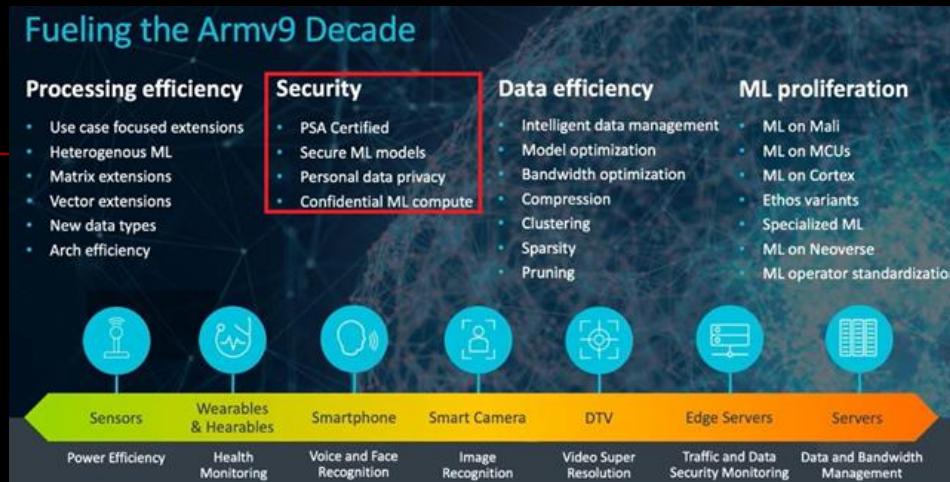
PSA for M-profile (constrained IoT) deliverable	PSA for Infrastructure Edge deliverable
Threat models	New Threat Model and Security Analysis (TMSA) for the Edge
Firmware Framework-M	Firmware Framework-A
Trusted Boot and Firmware Update	Same
PSA APIs (Crypto, Attestation, Secure Storage)	Same A new open-source security microservice: <a href="#">PARSEC</a>
Trusted Firmware-M	Trusted Firmware-A New features being added to support the needs of infrastructure
Certification scheme	In progress – Working to understand market need
-	New SBSG – Server Base Security Guide

Source: [https://www.arm.com/-/media/global/solutions/artificial-intelligence/Project\\_Cassini.pdf](https://www.arm.com/-/media/global/solutions/artificial-intelligence/Project_Cassini.pdf)

- **PSA (Platform Security Architecture)**  
<https://www.arm.com/why-arm/architecture/platform-security-architecture>

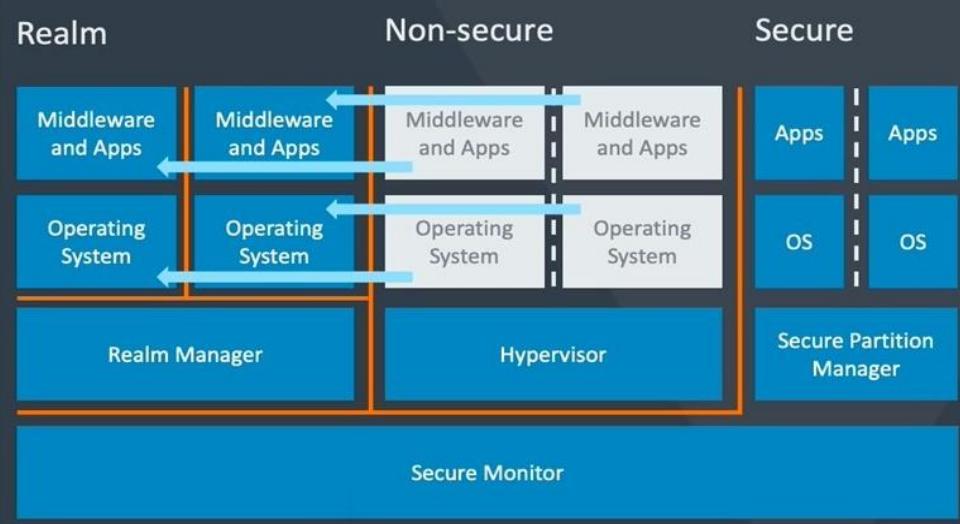
## ARM v9

- <https://www.embedded.com/new-armv9-takes-on-security-and-ai/>



- **CCA (Confidential Compute Architecture)**

- **Realms** enable you to run an application or service in such a way that your data is protected from:
  - The host
  - The hardware manufacturer or service supplier
  - Other software running on the device or sharing the machine



# VI. Wrap-up

- Lightweight VM like **Kata Containers** is well suited for Edge- and Serverless- computing
- Linux is a key player in the edge-computing revolution
- Unified runtime for **eBPF & WASM?**
- Rethinking App Runtime – User space and Kernel space  
**Repartitioning or Unifying**



---

**Q & A**

# Thanks!

---



# Reference

Slides/materials from many and varied sources:

- <http://en.wikipedia.org/wiki/>
- <http://www.slideshare.net/>
- <https://katacontainers.io/collateral/kata-containers-onboarding-deck.pptx>
- <https://www.openstack.org/videos/summits/denver-2019/kata-containers-story-of-a-container-runtime-1>
- <https://www.lfedge.org/resources/publications/>
- <https://medium.com/wasmer/running-webassembly-on-the-kernel-8e04761f1d8e>
- <https://www.arm.com/products/silicon-ip-security>
- <https://developer.arm.com/documentation/102142/latest/Secure-virtualization>
- [https://community.arm.com/developer/ip-products/processors/b/processors-ip-blog/posts/project-cassini-ensuring-a-cloud-native-experience-for-secure-edge?\\_ga=2.45148257.414898915.1620708938-1735275506.1620708938](https://community.arm.com/developer/ip-products/processors/b/processors-ip-blog/posts/project-cassini-ensuring-a-cloud-native-experience-for-secure-edge?_ga=2.45148257.414898915.1620708938-1735275506.1620708938)
- <https://github.com/firecracker-microvm/firecracker-containerd/blob/main/docs/architecture.md>
- <https://github.com/cloud-hypervisor/>

- <https://github.com/firecracker-microvm>
  - <https://github.com/edgexfoundry>
  - <https://gerrit.akraino.org/r/q/status:open+-is:wip>
  - <https://containerd.io/>
  - <https://github.com/opencontainers>
  - <https://www.graalvm.org/docs/introduction/>
  - ...
-